

The ACM Multimedia 2022 Deep Video Understanding Grand Challenge

Keith Curtis

National Institute of Standards and Technology
Gaithersburg, Maryland, USA
keith.curtis@nist.gov

Shahzad Rajput*

National Institute of Standards and Technology
Gaithersburg, Maryland, USA
shahzad.rajput@nist.gov

George Awad

National Institute of Standards and Technology
Gaithersburg, Maryland, USA
george.awad@nist.gov

Ian Soboroff

National Institute of Standards and Technology
Gaithersburg, Maryland, USA
ian.soboroff@nist.gov

ABSTRACT

This is the overview paper for the Deep Video Understanding (DVU) Grand Challenge. In recent years, a growing trend towards working on understanding videos (in particular movies) to a deeper level started to motivate researchers working in multimedia and computer vision to present new approaches and datasets to tackle this problem. This is a challenging research area which aims to develop a deep understanding of the relations which exist between different individuals and entities in movies using all available modalities such as video, audio, text and metadata. The aim of this grand challenge is to foster innovative research in this new direction and to provide benchmarking evaluations to advance technologies in the deep video understanding community.

CCS CONCEPTS

• **Information systems** → *Multimedia content creation; Multimedia and multimodal retrieval.*

KEYWORDS

video understanding; multimedia; multimodal interaction; information retrieval; video ontology

ACM Reference Format:

Keith Curtis, George Awad, Shahzad Rajput, and Ian Soboroff. 2022. The ACM Multimedia 2022 Deep Video Understanding Grand Challenge. In *Proceedings of the 30th ACM International Conference on Multimedia (MM '22)*, October 10–14, 2022, Lisboa, Portugal. ACM, New York, NY, USA, 4 pages. <https://doi.org/10.1145/3503161.3551582>

1 INTRODUCTION

Deep video understanding is a difficult task which requires systems to develop a deep analysis and understanding of the relationships between different entities in video, to use known information to

reason about other, more hidden information, and to populate a knowledge graph (KG) representation with all acquired information [2]. To work on this task, a system should take into consideration all available modalities (speech, image/video, and in some cases text). The aim of this challenge series is to push the limits of multimodal extraction, fusion, and analysis techniques to address the problem of analyzing long duration videos holistically and extracting useful knowledge to utilize it in solving different types of queries. The target knowledge includes both visual and non-visual elements. As videos and multimedia data are getting more and more popular and usable by users in different domains and contexts, the research, approaches and techniques we aim to be applied in this Grand Challenge will be very relevant in the coming years and near future.

This is the third in the series of Deep Video Understanding Grand Challenges at ACM Multimedia. In our first challenge, we focused on video understanding at the whole movie level, such as the relations between different characters and entities in the movie. In the second of these Grand Challenges we extended video understanding to a deeper level and introduced queries at the scene-level. These scene-level queries focused on finding the unique scene given a set of interactions, finding the next or previous interaction between two characters, matching the scene with a textual description, and scene sentiment classification. For this third year of the DVU challenge, we are introducing an Explainable AI component, in which participating teams are asked to submit evidence for their answers to scene-level queries. This evidence will be in the form of a start and end times of a short video clip from the scene in question in which evidence of the queried interaction will exist, or in the movie-level queries by providing a scene number for evidence of the queried relation.

2 RELATED WORK

There has been efforts to encourage research in high level video understanding such as the "MovieQA" [6] and "The Large Scale Movie Description Challenge" [5]. However these tasks revolve around isolated visual concepts retrieval and not about testing systems for their overall understanding of entities, relations and events within the video/movie. Early visions of this kind of work [3] proposed to use visual and audio descriptors, in addition to employing semantic analysis and linking with external knowledge sources in order to populate a knowledge graph. A large-scale dataset of corresponding

*Georgetown University

This paper is authored by an employee(s) of the United States Government and is in the public domain. Non-exclusive copying or redistribution is allowed, provided that the article citation is given and the authors and agency are clearly identified as its source.

MM '22, October 10–14, 2022, Lisboa, Portugal

© 2022 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-1-4503-9203-7/22/10...\$15.00

<https://doi.org/10.1145/3503161.3551582>

movie trailers, plots, posters, and metadata was developed by [1] who study the effectiveness of visual, audio, text, and metadata-based features for predicting high-level information about movies such as their genre or estimated budget.

3 DATASET

The Deep Video Understanding Training Set described in Table 1 consists of 14 Creative Commons (CC) license movies with total duration of about 17 hours¹. This training set has been annotated by human assessors and final ground truth, both at the overall movie level (Ontology of relations, entities, actions & events, Knowledge Graph, and names and images of all main characters), and the individual scene level (Ontology of locations, people/entities, interactions and their order between people, sentiments, and text summary) has been provided for the training set to participating researchers for training and development of their systems. In summary, we hired 5 annotators in addition to a summer student and on average each movie took about 20 hours of work to annotate both movie and scenes. A sample from a scene-level knowledge graph annotation can be seen in figure 2. For more detailed information about the annotation framework please refer to our paper [4].

The DVU Test Set described in Table 2 contains 6 movies licensed from KinoLorberEdu² platform with total duration of about 8.5 hours. Participants were required to complete a data access form in order to access these movies. The testing set was fully annotated by human annotators to the same degree as the training set. A set of queries, described in more details in Section 3, were then automatically extracted from human annotations and released to participants, along with the set of movies and annotated images of the movie characters identified during annotation.

More information about movies' genres and duration are provided.

4 ANNOTATION

Human assessors annotated each movie of the full DVU dataset. Full movies were annotated to a Knowledge Graph (KG) indicating the relationships and connections between every major character, entity, and concept in the movie. Images of each character and entity was also provided. Figure 1 shows an example of a movie-level KG. Following this, every scene within each movie was also annotated to a scene-level KG indicating the locations and characters within each scene, at least one sentiment label for that scene, non-neutral mental states of characters, the interactions between characters, and the ordering of the interactions as they happened in that scene. Figure 2 shows an example of a scene-level KG.

5 QUERIES

5.1 Movie-Level

5.1.1 Relations between characters. How is character X related to character Y? This query type question asks participants about all routes through the KG from one person to another. The main

Movie	Genre	Duration
Honey	Romance	86 min
Let's Bring Back Sophie	Drama	50 min
Nuclear Family	Drama	28 min
Shooters	Drama	41 min
Spiritual Contact		
The Movie	Fantasy	66 min
Super Hero	Fantasy	18 min
The Adventures of Huckleberry Finn	Adventure	106 min
The Big Something	Comedy	101 min
Time Expired	Comedy / Drama	92 min
Valkaama	Adventure	93 min
Bagman	Drama / Thriller	107 min
Manos	Horror	73 min
Road to Bali	Comedy / Musical	90 min
The Illusionist	Adventure / Drama	109 min

Table 1: The Full DVU Training Set

Movie	Genre	Duration
Calloused Hands	Drama	92 min
Chained For Life	Comedy / Drama	88 min
Liberty Kid	Drama	88 min
Like Me	Horror / Thriller	79 min
Littlerock	Drama	82 min
Losing Ground	Comedy / Drama	81 min

Table 2: The Full DVU Testing Set




Figure 1: Movie-level KG Sample

¹<https://www-nlpir.nist.gov/projects/trecvid/dvu/dvu.development.dataset/>

²<https://www.kinolorber.com/>




Figure 2: Scene-level KG Sample

objective of this query type is to test the quality of the established KG. If the system managed to build a representative KG reflecting the real story line of the movie, then it should be able to return back all valid paths, including the shortest, between characters (i.e. how they are related to each other).

5.1.2 Fill in the graph space. Fill in spaces in the Knowledge Graph (KG). Given the listed relationships, events or actions for certain nodes, where some nodes are replaced by variables X, Y, etc., solve for X, Y etc. Example of The Simpsons: X Married To Marge. X Friend Of Lenny. Y Volunteers at Church. Y Neighbor Of X. Solution for X and Y in that case would be: X = Homer, Y = Ned Flanders.

5.1.3 Question Answering. This query type represents questions on the resulting KG, including actions and events, of the movies in the described dataset. For example, we may ask 'How many children does Person A have?', in which case participating researchers should count the 'Parent Of' relationships Person A has in the Knowledge Graph. These are multiple choice questions.

5.1.4 Explainability. A new Explainable AI component has been added to relation based movie-level queries. For this component, teams are asked to submit a scene number along with their answer to the relationship question. This number should indicate a scene in which evidence exists of the submitted relationship.

5.2 Scene-Level

5.2.1 Find the Unique Scene. Given a full, inclusive list of interactions, unique to a specific scene in the movie, teams should find which scene this is.

5.2.2 Fill in the graph space. Find the person in a specific scene with the following attributes and interactions with others. Participating teams will be given a scene number, a list of person attributes, and a list of interactions to and from other people. Teams should find the only person in that scene with those attributes and interactions.

5.2.3 Find next or previous interaction. Given a specific scene and a specific interaction between person X and person Y, participants will be asked to return either the previous interaction or the next interaction, in either direction, between person X and Person Y. This can be specifically the next or previous interaction within the

same scene, or over the entire movie. These will be multiple choice questions selected from a list of possible interactions, only one of which will be correct.

5.2.4 Find the 1-to-1 relationship between scenes and natural language descriptions. Given a set of scenes, and a set of natural language descriptions of movie scenes, match the correct natural language description for each scene.

5.2.5 Classify scene sentiment from a given scene. Given a specific movie scene and a set of possible sentiments, classify the correct sentiment label for each given scene.

5.2.6 Explainability. A new Explainable AI component has been added to certain next or previous interaction scene-level queries. For this component, teams are asked to submit a clip start time and end time along with their chosen interaction. This start and end time should correspond to a video clip from within the specified scene in which evidence exists of the interaction in question.

6 METRICS

6.1 Movie-Level

6.1.1 Relations between characters. In this query type, systems are asked to submit all valid paths from a source node to another target node with the goal of maximizing recall and precision. Each path will be evaluated if it is a valid path (i.e the submitted order of nodes and edges leads to a path from the source person to the target person of the query) and finally the recall, precision and F1 measures will be reported.

6.1.2 Fill in the graph space. Results will be treated as ranked list of result items per each unknown variable and the Reciprocal Rank score will be calculated per unknown variable and Mean Reciprocal Rank (MRR) per query.

6.1.3 Question Answering. Scores for this query will be calculated by the number of Correct Answers / number of Total Questions.

6.2 Scene-Level

6.2.1 Find the Unique Scene. Results will be treated as ranked list of result items per each unknown variable and the Reciprocal Rank score will be calculated per unknown variable and Mean Reciprocal Rank (MRR) per query.

6.2.2 Fill in the graph space. Results will be treated as ranked list of result items per each unknown variable and the Reciprocal Rank score will be calculated per unknown variable and Mean Reciprocal Rank (MRR) per query.

6.2.3 Find next or previous interaction. Scores for this query will be calculated by the number of Correct Answers / number of Total Questions.

6.2.4 Find the 1-to-1 relationship between scenes and natural language descriptions. Scores for this query will be calculated by the number of Correct Answers / number of Total Questions.

6.2.5 Classify scene sentiment from a given scene. Scores for this query will be calculated by the number of Correct Answers / number of Total Questions.

Group	Rank 1	Rank 2	Rank 3
Group 1	UZH	N/A	N/A
Group 2	Nanjing Univ.	HERO TVQA	UZH

Table 3: Movie-Level Rank Table

Group	Rank 1	Rank 2	Rank 3	Rank 4	Rank 5
Group 1	E-VG	Graphen	Nanjing Univ.	HERO TVQA	UZH
Group 2	DVU-SQL	HERO TVQA	Nanjing Univ.	Graphen	N/A

Table 4: Scene-Level Rank Table

7 RUN SUBMISSION TYPES

Teams are allowed submit to either the Movie-level queries only, the scene-level queries only, or both the movie and scene queries. For movie-level queries, teams can choose to submit results to either or both groups of:

- Group 1 - Question 1 (relation between characters)
- Group 2 - Question 2 & 3 (fill in the graph and question answering)

For scene-level queries, teams can choose to submit results to either or both groups of:

- Group 1 - Questions 1, 2, & 3 (all related to interactions between characters)
- Group 2 - Questions 4 & 5 (matching of scenes and their descriptions in text, and sentiment classification)

Submissions for both movie-level and scene-level queries were in XML format. This allowed for the automatic scoring of solutions to the challenge. Below is an example XML solution to a scene-level previous interaction question from a previous challenge. This indicates the submitted answer of "watches" for the type "Interaction".

```
<DeepVideoUnderstandingTopicResult id="2" question="4">
  <item answer="watches" type="Interaction" />
</DeepVideoUnderstandingTopicResult>
```

8 PARTICIPATING TEAMS AND RANKINGS

A total of 6 teams participated in this years challenge, representing a 200% increase on the previous year. We believe that providing the 4 different group options so that participating teams had the option of which query types for which to submit answers encouraged more participation this year. In this years challenge we had 4 participating teams from Asia, 1 team from Europe, and 1 from North America.

In scene-level queries, 5 teams participated in Group 1 and 4 teams in Group 2. For the movie-level queries, 3 teams participated in Group 1 and 3 teams in Group 2. Tables 3 and 4 show final ranks for movie and scene level queries.

9 CONCLUSIONS

Running the DVU grand challenge during the last 3 years has been very exciting and challenging at the same time. It is exciting to see new research and techniques trying to address the various research questions posed by the organizers, and challenging due to the labor and cost involved in identifying reliable data collection sources, human annotation, and reviewing generated queries. We hope that attendees and readers of systems' papers and sessions find a good diversity of research presented at this challenge, and we hope that we have encouraged and inspired further work in this new direction.

ACKNOWLEDGMENTS

The authors would like to thank the Information Technology Laboratory at NIST for sponsoring all annotation work done on the DVU dataset and sponsoring the first three organizers of this grand challenge. Also we thank the ACM MM GC chairs and main conference for continuing this challenge this year, and all paper authors for their research efforts and contributions.

Disclaimer: Certain commercial entities, equipment, or materials may be identified in this document in order to describe an experimental procedure or concept adequately. Such identification is not intended to imply recommendation or endorsement by the National Institute of Standards and Technology, nor is it intended to imply that the entities, materials, or equipment are necessarily the best available for the purpose. The views and conclusions contained herein are those of the authors and should not be interpreted as necessarily representing the official policies or endorsements, either expressed or implied, of NIST, or the U.S. Government.

REFERENCES

- [1] Paola Cascante-Bonilla, Kalpathy Sitaraman, Mengjia Luo, and Vicente Ordonez. 2019. Moviescope: Large-scale Analysis of Movies using Multiple Modalities. *arXiv preprint arXiv:1908.03180* (2019).
- [2] Keith Curtis, George Awad, Shahzad Rajput, and Ian Soboroff. 2020. HLUU: A New Challenge to Test Deep Understanding of Movies the Way Humans do. In *Proceedings of the 2020 International Conference on Multimedia Retrieval*. 355–361.
- [3] Jeremy Debattista, Fahim A Salim, Fasih Haider, Clare Conran, Owen Conlan, Keith Curtis, Wang Wei, Ademar Crotti Junior, and Declan O'Sullivan. 2018. Expressing Multimedia Content Using Semantics—A Vision. In *2018 IEEE 12th International Conference on Semantic Computing (ICSC)*. IEEE, 302–303.
- [4] Erika Loc, Keith Curtis, George Awad, Shahzad Rajput, and Ian Soboroff. 2022. Proceedings of LREC2022 Workshop "People in language, vision and the mind" (P-VLAM2022). In *Proceedings of LREC2022 Workshop "People in language, vision and the mind" (P-VLAM2022)*.
- [5] Anna Rohrbach and Jae Sung Park. 2019. Large Scale Movie Description Challenge (LSMDC) 2019. <https://sites.google.com/site/describingmovies/lsmdc-2019>, Last accessed on 2019-11-06.
- [6] Makarand Tapaswi, Yukun Zhu, Rainer Stiefelhofen, Antonio Torralba, Raquel Urtasun, and Sanja Fidler. 2016. Movieqa: Understanding stories in movies through question-answering. In *Proceedings of the IEEE conference on computer vision and pattern recognition*. 4631–4640.