

AI Trust Score: A User-Centered Approach to Building, Designing, and Measuring the Success of Intelligent Workplace Features

Jennifer Wang
Microsoft, USA
jennifer.wang@microsoft.com

Angela Moulden
Microsoft, USA
angelam@microsoft.com

ABSTRACT

This paper describes how to help product teams move increasingly towards responsible AI experiences by examining user trust within business contexts. We discuss the development, application, and validation of the AI Trust Score, which allows teams to measure the success of an AI feature from enterprise users' perspectives. The Trust Score is a multi-dimensional metric which consists of several statements users respond to that inform their trust, including whether an AI feature helps with job efficiency and effectiveness and understanding how and when to use the feature for their job role. The metric, along with usability feedback, can be adapted to encourage UX-related conversations and actions within product or feature development processes.

CCS CONCEPTS

• **Human-centered computing** → Human computer interaction (HCI).

KEYWORDS

Artificial intelligence, user experience, methods, measurement

ACM Reference Format:

Jennifer Wang and Angela Moulden. 2021. AI Trust Score: A User-Centered Approach to Building, Designing, and Measuring the Success of Intelligent Workplace Features. In *CHI Conference on Human Factors in Computing Systems Extended Abstracts (CHI '21 Extended Abstracts)*, May 08–13, 2021, Yokohama, Japan. ACM, New York, NY, USA, 7 pages. <https://doi.org/10.1145/3411763.3443452>

1 INTRODUCTION

In the workplace, people are increasingly relying on AI to do their jobs more efficiently and effectively. From sales to retail to financial services, AI is changing businesses in fundamental ways. Spending on AI technologies will grow to \$97.9 billion USD in 2023—more than two and a half times the spending level of \$37.5 billion USD in 2019 [1].

Due to the unique challenges of AI, it is important for product and feature teams to exercise care and responsibility for AI experiences – that is, building trustworthy experiences for users. In this paper we discuss the development, application, and validation of the Trust

Score, which allows teams to measure the success of an AI feature from the perspectives of enterprise users. The Trust Score is a multi-dimensional metric consisting of several statements that capture users' initial experiences with an AI feature, which is pivotal in whether users will use or return to that feature and continue to trust the product and company.

1.1 AI and trust

The ability of AI systems to affect the input and output of data can make results less predictable. Because of this, understanding user trust in the results or insights is vital yet complex, due to its multi-dimensional and multi-phased nature.

Companies such as Google [2] and IBM [3] have developed guidebooks to help UX professionals engage in human-centered approaches to AI experiences. In these public-facing resources, many companies have taken a *multi-dimensional approach* on how trust is defined, ranging from understanding how technology works to protecting user data and privacy to providing feedback when predictions are unhelpful or incorrect.

In addition to multi-dimensionality, user trust can evolve over time and *initial experiences* with a new AI feature is critical as users may scrutinize the experience more. Interactions often start with predictability, then dependability or reliance, followed by more habitual faith in the interaction [4]. With predictability, user understanding is key. The initial experiences of interacting with new AI features are critical because trust tends to be fragile and based on the perceived risks weighed against potential benefits [5]. Confusing, frustrating, or overwhelming AI results and experiences can negatively impact a user's desire to use and return to an AI feature, and potentially erode trust from the company brand.

1.2 Enterprise considerations

"For an intelligent feature to help me to prioritize and plan my day, I need it to be relevant, personalized, and efficient, so that I can do my job well." – Seller, Study Participant

AI has impacted consumer and enterprise applications, yet less research has focused on the unique considerations for AI enterprise experiences for users. For example, these enterprise experiences encompass the configuration or setup of AI, building AI-related apps or tools (e.g., building models for conversational capabilities or forms processing), and consumption of AI (e.g., being provided with insights or recommendations, forecasting trends).

There are at least three major distinctions with enterprise AI UX: the emphasis on users' job efficiency and effectiveness, data protection concerns, and need for understanding and control in AI features.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

CHI '21 Extended Abstracts, May 08–13, 2021, Yokohama, Japan

© 2021 Copyright held by the owner/author(s).

ACM ISBN 978-1-4503-8095-9/21/05.

<https://doi.org/10.1145/3411763.3443452>




Figure 1: User trust over time with emphasis on how trust can be easily lost in the initial experiences interacting with a new AI feature. Used with permission from Penny Collisson, Microsoft. [Private domain]

First, trust in AI is tied to users' need for *job efficiency and effectiveness* in completing work tasks. Within our own business applications organization, we have found that enterprise users often focus on taking the next best action for their job role. In a recent review of our AI features, many of these features fell into two broad categories to help our users: providing insight-based needs and task-based needs. Insight-based needs focus on providing information or content necessary for one's job roles, such as providing more nuanced or more in-depth insights and helping users determine relationships between data points. Task-based needs are based on helping users solve complex tasks and assisting to reduce time on repetitive or menial tasks. Both categories aim to provide users help with job efficiency and effectiveness.

Second, trust in AI is often tied to *data protection* concerns (e.g., privacy and security) across many industries, such as in healthcare, finance, and banking. For example, serious ramifications can occur when a fraud protection predictor inaccurately predicts compromised accounts. Furthermore, data breaches can have a profound impact on a company's perceived privacy and security and overall brand perception [6]. In our research, we have found that data protection is a top priority for enterprise users across industries because of compliance and usage requirements (e.g., GDPR).

Third, user *understanding* and *control* are critical because inconsistent and unpredictable behaviors can confuse users, erode trust, and lead to abandonment of AI technology [5]. However, within an enterprise setting, business end users may not be able to abandon their use of an AI feature or product entirely at work even if they feel that it is not relevant or personalized to their needs. For example, many sellers rely on a multitude of factors to determine leads

or opportunities for a sales deal. If an AI feature makes a prediction that runs counter to a seller's knowledge of a potential customer or against their gut instinct, they may be less likely to depend or rely on that feature over time. This example illustrates the need to focus on how user trust relates to user satisfaction in enterprise settings. In addition, users need to understand how and when AI features are relevant to their job role and to have a sense of control – such as in the form of user feedback when results are incongruent with expected outcomes.

1.3 Measuring success and scorecards

More companies have begun to provide guidance on building usable, useful, and trustworthy AI experiences. For example, Microsoft [7] and Google [2] have provided applicable design guidelines for human-AI interaction. Furthermore, there are an increasing number of frameworks and guidance on how to approach UX research on AI experiences [8, 9]. With an accruing emphasis to build, design, and research AI UX, how do we begin to measure the trustworthiness of AI features, particularly within the enterprise space?

Measuring user trust consistently across features can be challenging. Additionally, strong organizational incentives for fast-paced development and deployment lifecycles can be at odds with practicing cultural change towards responsible AI [10]. To ensure we were able to work within a fast-paced environment to deliver results, we worked across disciplines to build and scale our approach in two major ways:

- We initially conducted an “audit” of our AI features across 30+ product areas. Because many AI features have similar underlying technologies and algorithms, we worked with

our cross-discipline stakeholders (i.e., data science, project management, and engineering) to build on a results-driven culture and to align with understanding of these features.

- We adapted a scorecard view to measure success. We have a precedent in our organization with UX benchmarking which is a familiar process to drive engineering culture toward change. In this effort, scorecards employ a consistent tool to assess UX quality across product pillars, help improve customer experience over time, and inform plans to address key blockers.

These approaches set up our organization's Design Return-on-Investment (ROI) by connecting Design key performance indicators (KPIs) (i.e., measurement of performance activities) and critical business metrics. By understanding how trustworthy our AI features are in conjunction with standard business metrics such as user satisfaction, we can support a continual process that provides an incentive for teams to incorporate UX quality into the fast-paced development and deployment lifecycle. To do so, we needed to develop a user-centered metric on trust in AI.

2 CASE STUDY

In the following section, we detail how we developed a user-centered, multi-dimensional AI Trust Score aimed to measure the success of our enterprise AI features. We then applied it in seven UX research studies with enterprise AI features to validate the metric. We discuss learnings and next steps, along with how the AI Trust Score can evolve for product teams in the enterprise space.

Across 30+ product groups in our organization, there are at least 90 user profiles; however, these user profiles can be categorized into four larger persona categories in which users configure, create, or consume AI experiences: admins (technical workers who manage their companies' IT infrastructure), makers (semi-technical workers who build apps for their companies or clients), analysts (data and statistics experts who analyze data and share it with end users), and business end users (less technical business people who use data to make decisions).

2.1 Trust Score development

2.1.1 Item development and feedback. Our goal was to start broadly and identify a larger range of factors that may impact trust in AI feature interactions among users in enterprise job roles. To do so, approximately 50 items were developed from a review of internal research and interviews with product teams. All items began with "I trust an intelligent feature for work..." to help situate participants in the context of thinking about trust within a work setting. An example of the item included: "I trust an intelligent feature for work when I know that my data is protected."

These 50+ items covered categories appropriate within an enterprise context and relevant across different persona profiles. Furthermore, these items aimed to capture a range of feelings, thoughts, and actions that can reflect trust [11] in an enterprise space: job efficiency and effectiveness, data protection, understanding, and control along with feature comprehension, absence of concerning elements, and not infringing on human rights (e.g., not stereotyping or discriminating users).

2.1.2 Survey with qualitative and quantitative feedback. Participants ($N=166$ end users, analysts, makers, and admins) were presented with scenarios containing a brief description of an AI feature relevant for their job role. For example, sales end users were presented with the following:

You want to save time and be more organized. Your organization has a product with an intelligent feature that helps you prioritize and plan your day, including automatically showing you changes to your meetings, tasks, and the next best action to close a deal.

After being presented with a scenario, each participant was prompted to respond to an open-ended question: "What would make you trust [the intelligent feature] more or less?"

Following this, each participant was randomly presented with each of the 50+ items ("I trust an intelligent feature when...") and rated their agreement to each item on a 1 (strongly disagree) to 7 (strongly agree), including a N/A option as well. The intent of presenting each item one at a time was to encourage participants to respond with their initial, gut reaction of tying each item to trustworthiness. Most participants moved swiftly, only thinking aloud if they wanted to clarify their responses; however, many explained their responses for each item presented.

2.1.3 Item reduction and inferential statistical analyses. Each of the 50+ items were ranked from the most agreed upon to the least agreed upon per persona profile. To capture the items most relevant to trust, items in the top third were analyzed (all items in this range were rated on average between "agree" to "strongly agree"). From this, the most agreed upon items across all persona profiles were analyzed via exploratory factor analysis, an item reduction technique that identifies underlying latent groupings between item responses.

Factor analyses revealed four factors or themes that were consistent across all persona profiles: data protection for self and customers, job efficiency, understanding/control (due to a lower internal consistency and qualitative evidence, this theme was broken into two themes), and safety (based on qualitative evidence, this theme was combined with data protection).

In addition, other trust themes were identified for each persona profile. For example, among sales end users, being able to take the next best action with a feature, knowing that a feature is accurate, and having confidence using a feature were also important for trust. For this report, however, we focus on the trust items most agreed upon across personas.

2.2 Trust Score themes

The trust themes (i.e., job efficiency, understanding, control, data protection) were supported by the open-ended responses participants provided prior to being exposed to each of the trust items tested. Furthermore, many participants explained their responses to items aloud, allowing us to analyze their perceptions of the items related to each theme.

Job efficiency and effectiveness. Job efficiency items were correlated with doing one's job well and there were two essential elements – one that is time-based (job efficiency) and another that is performance-based (job effectiveness). This theme was especially critical for end users, to help them take the next best action.

Table 1: Items analyzed in the meta-analysis

Categories	Items
Trust Score items	[The AI feature] will help me do my job more efficiently and effectively. I understand how and when to use [the AI feature]. I have control using [the AI feature]. I know my data and my customers' data will be protected with [the AI feature].
Additional items	I am confident in the results made by [the AI feature]. I trust the results made by [the AI feature].
User satisfaction	Overall, how satisfied are you with using [the AI feature]?

"If the [AI feature] works accurately, if it actually saves me time." – Seller, Study Participant

Understanding. The range of understanding responses varied from wanting to know the ins-and-outs of an AI feature (e.g., among admins) to wanting to understand how it impacts their job, their goals, and what they can do (e.g., among end users). Some users, particularly analysts and makers, also wanted to be able to check on the results or insights provided by a feature to verify accuracy.

"Understanding the data, how it learns, interprets, and generates the results. . . and how it was trained and how it can be corrected." – Analyst, Study Participant

Control. Having control over AI features was critical for admins and makers and was tied to understanding, as they needed to know more about the inner workings of features. Control ranged from needing to configure for customization to providing feedback. For example, among end users, control fell into two categories: 1) making sure that the results or insights could be of value to their tasks and 2) having control when something "goes wrong" in the "1 percent scenario."

"The ability to control security features and update or stop the program in its tracks if necessary. And ultimately roll out any adverse changes that may occur."
– Admin, Study Participant

Data protection for self and customers. In general, data protection and safety were important themes across all persona profiles, which include protecting privacy and security. Correlational analyses and qualitative responses suggested relationships between data protection and safety items. Admins tended to be more specific about their data protection needs (e.g., anonymized data) whereas there was variability among end users - some just wanted to know the bare minimum around data protection and privacy for a feature.

"Data being gathered is not misused and if any human is seeing that data, then it is appropriately masked so that the identity of the person is not revealed."
– Admin, Study Participant

2.3 Trust Score application

After the Trust Score was created, it was applied to seven enterprise AI features in usability studies within our product group. These features were creation and consumption-based, ranging from predicting values and providing recommendations for business end users to helping analysts and makers set up models. Each of these studies had participant sample sizes ranging from 6 to 12, which are

greater than typical sample sizes used to uncover major usability problems [12].

All seven UX studies incorporated standard usability study procedures, including prompting participants to complete core tasks relevant to each AI feature. Each researcher captured attitudinal and behavioral observations along with response ratings to the Trust Score items and additional UX metrics. For the purposes of this report, we focused on Trust Score items, confidence and trust in the results made by an AI feature, and user satisfaction in an AI feature.

2.4 Trust Score validation

A meta-analysis was completed to examine patterns across the seven UX research studies that applied the AI Trust Score. For the purposes of this meta-analysis, two major questions were addressed.

"Are we doing the right thing?" (i.e., construct & content validity)

Based on inter-item consistency, most of the Trust Score items "hung together" to represent the construct of trust/confidence. Internal consistency was good ($\alpha = 0.7$ when the data protection item was dropped). We also examined whether the Trust Score items were associated with asking about user trust and confidence directly in the results of the AI feature. When examining the Trust Score items as a composite score (i.e., job efficiency and effectiveness, understanding, control, and data protection items), the Trust Score was positively correlated with feeling confident, having trust, and user satisfaction in a feature ($r's = 0.5$ to 0.7 , $p < 0.05$).

"Does the metric matter?" (i.e., predictive validity)

As noted earlier, in our organization the Trust Score is viewed as a bridge metric between Design KPIs and standardized business metrics such as user satisfaction. When running simple linear and hierarchical regressions controlling for related predictors of trust and confidence items, a higher Trust Score predicts user satisfaction in an AI feature ($p's < 0.01$). Notably, the Trust Score had predictive value beyond asking trust and confidence directly. Furthermore, all three predictors (Trust Score and the two items on confidence and trust) explained nearly 60% of variance in user satisfaction for an AI feature.

A question that might arise is whether user satisfaction could be explained by directly asking users if they have confidence or trust in the results of an AI feature. We believe that these items should still be included to examine their associations with other UX metrics. However, a major advantage of the Trust Score is that




Figure 2: Trust Score as a bridge metric between Design KPIs (i.e., measurement of performance activities) and business metrics.




Figure 3: Trust Score includes a standardized set of items and was found to be predictive of user satisfaction in an AI feature in the meta-analysis.

it helps us understand the specific elements of trust that may lead to user satisfaction.

Specifically, these items are asked the same way every time so that there is an understanding of what is relevant to trust, and product teams can compare the items across features and measure improvement over time. Essentially, it is like a heatmap that helps teams understand how to up-level the user trustworthiness of AI features.

In our analysis, the data protection item was less strongly correlated with the other Trust Score items. There are at least two possible explanations. First, previous research has found that the item is associated with privacy, security, and safety and compliance with regulations. Some UX studies from the analysis found that

participants examined this item from the lens of company reputation in these areas. Second, data protection may be more of a salient concern in specific contexts compared to the other three items. Previous research and additional external perspectives [2, 3] suggest that user concern may be escalated and trust eroded when, for example, there are data breaches or awareness when system is using data they did not know it had access to.

3 DISCUSSION

3.1 Learnings

Within our organization, we emphasize a learning-based culture and we anticipate that the Trust Score will evolve over time as we




Figure 4: Scorecard view of the Trust Score to help product teams compare items across features and measure improvement over time.

gain more understanding in how this metric is correlated with other key user and business metrics aside from user satisfaction.

Because trust is a complex social process [13–15], we are also examining additional factors that may influence trust beyond a user's interaction with an AI feature, such as brand perception (i.e., opinion of company and tools) and AI/ML proficiencies (e.g., understanding of probability and statistics, experience with data and models, current use of AI/ML tools, etc.). We are beginning to pilot metrics related to these factors and over time anticipate building frameworks that include the role of dispositional, design, and system-level factors. For now, we have found that the four items in the Trust Score have spurred conversations and action among product teams to improve on building user trustworthiness.

Finally, how do we help users avoid over-indexing on trust in AI? We view the Trust Score as an overall reflection of user trust in an AI feature. Specifically, the user is providing a reflection on multi-dimensional elements that are predictive of user satisfaction – “getting a green scorecard” on these items represents a general sense of their view on the AI feature for the distinct elements of trust. From many product teams' perspectives, enhancing trust in AI is critical to improve adoption [16]. On the other hand, it is critical to reduce the tendency for some users to blindly trust probabilistic outcomes that are generated in major decision-making contexts which may unfairly allocate opportunities, resources, or information to people [17] or when model performance or prediction confidence is less than ideal [16, 18]. It will also be important to continually examine the interaction between factors such as brand perception and AI/ML proficiencies in this discussion of blind trust.

As AI experiences become more embedded in daily lives and we have more data on how the Trust Score is used, we will continue to evolve what the metric looks like.

3.2 Expected insights and next steps

Within our organization, we are actively applying the Trust Score across our AI features and evolving our approach with additional learnings. Aside from providing more standardized ways to ask about trust which reduces subjective interpretations from users and teams, we believe that the distinct advantages of the Trust Score include the following:

- It helps bring multi-disciplinary teams together (UX, Project Management, Data Science, Engineering) throughout the development and deployment cycle.
- It encourages UX-related conversations early within AI development processes and reviews by having a standardized set of trust elements to operationalize on, which can help drive continual organizational change.
- It helps build collaboration based on understanding and aligning on areas for improvement after consolidating user feedback from scorecard metrics along with attitudinal and behavioral observations.

In addition, the Trust Score development approach can be adapted to different product teams' needs and can be viewed as a toolkit. The methods taken in this report can be modified for various needs – this is, to start with an internal synthesis and interviews, and a brainstorm of elements of trust that would be relevant

The Trust Score can also evolve to be used with experienced users of AI features and products via within-product intercept surveys. Furthermore, the Trust Score items can be modified beyond the initial experience, to encompass dependence and even habitual faith in a feature over time [4]. We look forward to a discussion on the Trust Score evolution and learning more about additional frameworks in understanding user trust with AI features.

We thank the Business Applications Group at Microsoft for the opportunity to develop the Trust Score and conduct research, along with many researchers, designers, data scientists, product/program managers, engineers, and leaders across the company for their time, expertise, and support. We thank Kelley Rand for feedback throughout the Trust Score development process. We are also appreciative of all users who contributed their perspectives on trust in AI.

- [1] IDC. 2019. Worldwide Spending on Artificial Intelligence Systems Will Be Nearly \$98 Billion in 2023, According to New IDC Spending Guide. <https://www.idc.com/getdoc.jsp?containerId=prUS45481219>
- [2] Google People and AI Guidebook. 2020. <https://pair.withgoogle.com/guidebook>
- [3] IBM Trusting AI. 2020. <https://www.research.ibm.com/artificial-intelligence/trusted-ai/>
- [4] John Rempel, John Holmes, & Mark Zanna (1985). "Trust in close relationships." *Journal of Personality and Social Psychology*, 49:1, pp. 311-331. <https://psycnet.apa.org/record/1985-30794-001>
- [5] Penny Collisson. 2015. Earning Trust in Intelligent Assistance. Internal Microsoft report.
- [6] The Economist Intelligence Unit. 2020. Staying ahead of the curve: The business case for responsible AI. [EIUSTayingAheadOfTheCurve.pdf](#)
- [7] Saleema Amershi, Dan Weld, Mihaela Vorvoreanu, Adam Fournay, Besmira Nushi, Penny Collisson, Jina Suh, Shamsi Iqbal, Paul N. Bennett, Kori Inkpen, Jaime Teevan, Ruth Kikin-Gil, and Eric Horvitz. 2019. Guidelines for Human-AI Interaction. In Proceedings of CHI Conference on Human Factors in Computing Systems (CHI '19), May 4–9, 2019, Glasgow, Scotland UK. ACM, New York, NY,

- USA 14 Pages. <https://doi.org/10.1145/3296055.3300233>
- [8] Clara Kliman-Silver, Oliver Siy, Kira Awadalla, Alison Lentz, Gregorio Convertino and Elizabeth Churchill. 2020. Adapting User Experience Research Methods for AI-Driven Experiences. In Proceedings of CHI '20: CHI Conference on Human Factors in Computing Systems (CHI'20 Extended Abstracts), April 25-30, 2020, Honolulu, HI, USA. ACM, New York, NY, USA, <https://doi.org/10.1145/3334480.3375231>
- [9] Qian Yang, Aaron Steinfeld, Carolyn Rosé, and John Zimmerman. 2020. Re-examining Whether, Why, and How Human-AI Interaction Is Uniquely Difficult to Design. In CHI Conference on Human Factors in Computing Systems (CHI '20), April 25-30, 2020, Honolulu, HI, USA. ACM, New York, NY, USA 17 Pages. <https://doi.org/10.1145/3313831.3376301>
- [10] Michael A. Madaio, Luke Stark, Jennifer Wortman Vaughan, and Hanna Wallach. 2020. Co-Designing Checklists to Understand Organizational Challenges and Opportunities around Fairness in AI. In CHI Conference on Human Factors in Computing Systems (CHI '20), April 25-30, 2020, Honolulu, HI, USA. ACM, New York, NY, USA 15 Pages. <https://doi.org/10.1145/3313831.3376445>
- [11] Penny Collisn, Gwen Hardiman, and Trish Miner. 2020. Why people may not trust your AI, and how to fix it. <https://www.microsoft.com/en-us/research/group/customer-insights-research/articles/why-people-may-not-trust-your-ai-and-how-to-fix-it/>
- [12] Jakob Nielsen. 2012. How many Test Users in a Usability Study? <https://www.nngroup.com/articles/how-many-test-users/#:~:text=Summary%3A%20The%20answer%20is%205,b%20bigger%20and%20some%20smaller.>
- [13] Arathi Sethumadhavan (2018). "Trust in artificial intelligence." *Ergonomics in Design*, 27:2, pp.34-34. <https://journals.sagepub.com/doi/pdf/10.1177/1064804618818592>
- [14] Kevin Hoff and Masooda Bashir (2015). "Trust in automation: Integrating empirical evidence on factors that influence trust." *Human Factors*, 57:3, pp. 407-434. https://journals.sagepub.com/doi/10.1177/0018720814547570?url_ver=Z39.88-2003&rft_id=ori:rid:crossref.org&rft_dat=cr_pub%20%20pubmed
- [15] John Lee and Katrina See (2004). "Trust in automation: Designing for Appropriate Reliance." *Human Factors*, 46:1, pp.50-80. https://journals.sagepub.com/doi/pdf/10.1518/hfes.46.1.50_30392
- [16] Q. Vera Liao, Daniel Gruen, and Sarah Miller. 2020. Questioning the AI: Informing Design Practices for Explainable AI User Experiences. In Proceedings of CHI Conference on Human Factors in Computing Systems (CHI '20), April 25-30, 2020, Honolulu, HI USA. ACM, New York, NY, USA 15 Pages <http://dx.doi.org/10.1145/3313831.3376590>
- [17] Microsoft Aether Fairness & Inclusiveness. 2020. <https://www.microsoft.com/en-us/ai/responsible-ai/activetab=pivot1%3Aprimaryr6>
- [18] Yunfeng Zhang, Q. Vera Liao, and Rachel K. E. Bellamy. 2020. Effect of Confidence and Explanation on Accuracy and Trust Calibration in AI-Assisted Decision Making. In Conference on Fairness, Accountability, and Transparency (FAT* '20), January 27-30, 2020, Barcelona, Spain. ACM, New York, NY, USA, 11 pages. <https://doi.org/10.1145/3351095.3372852>