

GenAICHI 2023: Generative AI and HCI at CHI 2023

Michael Muller
IBM Research
Cambridge MA, USA
michael_muller@us.ibm.com

Lydia B. Chilton
Columbia University
New York, NY, USA
chilton@cs.columbia.edu

Anna Kantosalo
University of Helsinki and Siili
Solutions
Helsinki, Finland
anna.kantosalo@helsinki.fi

Q. Vera Liao
Microsoft Research
Montreal QU, Canada
veraliao@microsoft.com

Mary Lou Maher
University of North Carolina at
Charlotte
Charlotte NC, USA
m.maher@unc.edu

Charles Patrick Martin
The Australian National University
Canberra, Australia
charles.martin@anu.edu.au

Greg Walsh
University of Baltimore
Baltimore MD, USA
gwalsh@ubalt.edu

ABSTRACT

This workshop applies human centered themes to a new and powerful technology, generative artificial intelligence (AI). Unlike AI systems that produce decisions or descriptions, generative AI systems can produce new and creative content that can include images, texts, music, video, code, and other forms of design. The results are often similar to results produced by humans. However, it is not yet clear how humans make sense of generative AI algorithms or their outcomes. It is also not yet clear how humans can control and more generally, interact with, these powerful capabilities in ethical ways. Finally, it is not clear what kinds of collaboration patterns will emerge when creative humans and creative technologies work together.

Following a successful workshop in 2022, we convene the interdisciplinary research domain of generative AI and HCI. Participation in this invitational workshop is open to seasoned scholars and early career researchers. We solicit descriptions of completed projects, works-in-progress, and provocations. Together we will develop theories and practices in this intriguing new domain.

CCS CONCEPTS

• **Human-centered computing** → **Human computer interaction (HCI)**; • **Computing methodologies** → **Artificial intelligence**.

KEYWORDS

Generative AI; Design; Bias; Uncertainty.

ACM Reference Format:

Michael Muller, Lydia B. Chilton, Anna Kantosalo, Q. Vera Liao, Mary Lou Maher, Charles Patrick Martin, and Greg Walsh. 2023. GenAICHI 2023: Generative AI and HCI at CHI 2023. In *CHI '23: CHI '23: ACM Symposium on Human Factors in Computer Systems, Apr 23–28, 2023, Hamburg, Germany or Hybrid or Virtual*. ACM, New York, NY, USA, 7 pages. <https://doi.org/10.1145/3544549.3573794>

1 INTRODUCTION

In the past year, we have seen or made powerful tools that can create images from textual descriptions [1, 48] or conduct reasonably coherent conversations [32, 62], make writing suggestions for creative writers [33], and write code as a pair programmer [23]. We have also seen claims of what an historical person “really looked like” [4], and of a “completed” version of a musical compositions left unfinished by their composer’s untimely death [47]. What all of these examples have in common is that the AI does not simply categorize data and interpret text as determined by models, but instead creates something new—e.g., in images [35, 56], molecules [45, 65], or designs [46]. This work moves the potential of AI systems from problem solving to problem finding, and it tends to change the “role” of the AI from decision-maker to human-supporter [57, 71]. Following a successful CHI workshop in 2022 [40], we focus on various aspects of generative AI and its interactions with humans, including

- new sociotechnical opportunities for work and recreation that are afforded by powerful new interactive capabilities
- novel design challenges of systems that produce a *different* outcome after each invocation
- ethical issues related to their design and use; and
- useful patterns for collaboration between humans and generative AI in different domains.

Generative AI can be defined as an AI system that uses existing media to create new, plausible media [45, 46, 56, 65]. This scope is broad, and the generative potential of AI systems varies greatly.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

CHI '23, Apr 23–28, 2023, Hamburg, Germany or Hybrid or Virtual

© 2023 Copyright held by the owner/author(s).

ACM ISBN 978-1-4503-9422-2/23/04.

<https://doi.org/10.1145/3544549.3573794>

Over the last decade, we have seen a shift in methodology moving from expert systems based on patterns and heavy human curating [e.g., 21, 28, 49] towards stochastic and generative models such as Generative Adversarial Networks (GANs) that use big data to produce convincingly human-like results in various domains [18, 56], and Large Language Models (LLMs) that can generate text [62], source code [55], and images [1, 53] from simple instructions (“prompts”) [32].

1.1 HCI Challenges of Generative AI

In this workshop, we focus on the unique challenges that emerging generative AI methods pose to designers in various fields. These include, for example, the limited ability that designers and users have in gaining an understanding of the inner workings of these models: We cannot inspect the code or make sense of the latent vectors used in them and the generated results may vary. In addition the creative process is inherently fuzzy and involves elements of surprise and problem invention rather than achieving specific goals. Therefore we posit that generative AI proposes new challenges for human-centered design due to the serendipitous and uncertain nature of the design space and especially for use cases where users work with the AI.

Drawing from the long tradition of computer supported collaborative work and creativity support systems, the idea of humans and AI agents working together to achieve creative results is becoming more and more commonplace, and a research paradigm of human-computer co-creativity and mixed-initiative creative interfaces is emerging. In human-computer co-creativity, creative partners accept each others’ suggestions into the space of possible creative products [60] and both the human and the computer are influenced by each others’ contributions [12] culminating in sharing creative responsibility over the resulting product [28]. The new paradigm is characterized by the mixing of computer and human initiative [70] and human-computer co-creativity can be seen to flourish along a continuum between human creativity and autonomous computational creativity [15]. Emerging generative AI methods offer opportunities for designing creative AI agents that can generate useful artifacts with which they can contribute to the creative process of humans.

Despite powerful generative AI methods becoming more and more accessible for designers of creative systems, we still know relatively little about how to make generative AI methods interactive. There is a firm basis for designing UIs to AI-based systems that produce reliable and replicable outcomes [2]. By contrast, generative AI provokes us with questions of how to design good user experiences to deal with *generative variability* - i.e., systems whose purpose is *not* to produce the same outcome to invariant inputs - and which can sometimes provide incomplete or even incorrect outcomes [69]. Variable outcomes seem well-suited to the open-endedness of creative work (e.g., [17]), and yet even less is known about the long term effects of the new technology for the creative practice of artists, designers and laypersons; the role generative AI based interactive systems will eventually take in society; and what kind of regulations will eventually govern the space of design in this area. We propose this workshop to continue to unite

a disparate community around generative AI to investigate questions related to the design, evaluation, deployment and ethics of interactive generative AI.

1.2 Previous Workshops

The general topic of AI has led to multiple workshops in the SIGCHI traditions [e.g., 3, 16, 31, 36, 39, 41]; however, the theme of *generative* AI and user experience [22, 44, 61] has only been addressed via a workshop at last year’s CHI conference [40]. With this submission, we propose to build on the rich HCI possibilities of this emerging technology. Below, we describe some projects that have inspired our workshop.

2 BACKGROUND

2.1 Generative AI

Researchers have applied generative AI methods to diverse media, such as images [20, 51, 51], text [50, 62], music [18, 26], physical products [37, 38], source code [5, 8, 55], and movement [66]. Whereas previous eras of generative AI were rule-based or built on relatively small models, new models are trained on massive Internet-scale data and have untold abilities that we are only beginning to discover (see Figure 1). Moreover, these models (or versions of them) are open to the public without having to retrain them yourself [1, 9, 14, 64]. This gives many more people the ability to access powerful models. For example, there are thriving internet communities of artists collectively discovering how text-to-image synthesis can be used in the artistic process [1, 10, 27, 52, 53, 63]. However, there are many unanswered questions about how and what AI can generate and how can and should people be involved in the process (e.g., [11, 19, 24, 25]).

2.2 Designing for Creativity

Novelty and utility or value [54], sometimes accompanied by surprise [35], form three basic criteria for evaluating creativity. The products of emerging generative AI methods can be said to be creative since they can be novel, valuable and surprising in a context of use. This means emerging generative AI methods have great potential for facilitating creative work in some form or another. We propose that creative work in various domains, such as music, art, writing, and design, offers important application areas for generative AI methods.

There are domain independent considerations for facilitating interaction design for creativity. The open-ended nature of creative work challenges expectations that designers have about designing software for traditional, productivity centered domains; the requirements of creative work are vague, the measures of success unclear and the behaviors of users can be unorthodox [58]. This open-ended nature of creative work raises questions about the roles that Generative AI can play in the creative process [30] and how we evaluate and compare the human-AI systems that are emerging from the design for creativity [29]. Users of creative software are not necessarily even working towards a specific goal, but the process of creation itself can be the goal of autotelic creativity [13]. Combining powerful generative AI methods with users working and playing in creative domains opens up the design space to unforeseen dimensions.




Figure 1: Images generated from text prompts in the VQGAN+CLIP model. Images for multiple subjects and various styles are shown to demonstrate the wide ability of image generation models trained with internet-scale data.

2.3 Imperfection in AI Outcomes for Human Use

These human-oriented and human-directed systems are just now becoming usable and useful. Initial findings suggest that, while perfect operation and outcomes would be preferable, for human creativity support even imperfect results can be better than no support at all, as studied in humans’ use of generated documentation from source code [43, 67] and in translated source code from one programming language to another [69]. Software engineers both adopt and adapt (and occasionally replace) the outcomes of generative AI systems. The principle of AI-imperfection has not yet been parameterized: How bad does a generated outcome have to be, for a human to abandon it and re-do the work from scratch? How would domain experts negotiate the trade-offs in quality vs. productive-cost of generated outputs?

Many generative systems have been *productive* of new instances and new combinations learned from a class of exemplars or other strategies that make use of known histories of work [35, 46, 56, 65]. However, systems that create novel instances that humans would consider to be *creative* (in terms of novelty, utility and surprise) are less understood. Spoto and Oleynik [59] created a *Library of Mixed Initiative Creative Interfaces* based on [15]. However, only a few of these experiments used specifically generative algorithms. Systems capable of producing novel instances are only recently documented [e.g., 7, 17], and tooling that allows domain experts to control those novel or creative outcomes is even rarer [e.g., 34]. There are few formal human-centric, empirical evaluations of the quality of generated outcomes.

2.4 Responsible GenAI

HCI research has played an imperative role in developing AI technologies responsibly by producing design solutions, guidelines, and

methods that align AI technology design with stakeholder values and mitigate potential harms. Understanding and tackling the ethical issues and harms of generative AI is an emerging area that draws much interest in the HCI and AI communities [6, 25, 42, 68]. Besides creating new challenges to the core ethical AI principles including fairness, transparency, accountability, privacy, and so on, generative AI can lead to other types of risk and potential harms due to the generative nature and unique characteristics of application domains, many of which are yet to be understood. For example, text generation models can produce unintended toxic content that harms certain groups and individuals or spreads misinformation [68]. Many genAI-powered productivity support tools, whether for programming, writing, or other creative work, can risk having over-reliance and threatening to user agency and ownership [33]. Large Language Model’s societal harms including environmental and socioeconomic risks are also being actively researched [68].

3 TOPICS AND THEMES

Our workshop is open to diverse interpretations of interactive generative AI, characterized by the AI systems’ abilities to make new things, learn new things and foster serendipity and emergence. We are interested in several dimensions of generative AI, including mixed initiative, human–computer collaboration, or human–computer competition, with the main focus on interaction between humans and generative AI agents. We welcome researchers from various disciplines, inviting researchers from different creative domains including, but not limited to art, images, music, text, style transfer, text-to-image, programming, architecture, design, fashion and movement.

Because of the far-reaching implications of Generative AI, we propose the following list of non-exhaustive, thematic questions to guide our discussions at the workshop:

- What is generative AI and how can we leverage diverse definitions of it? Does generative AI go beyond intelligent interaction? What distinguishes generative AI?
- How do you design in this characteristically uncertain space? What design patterns do we need to think about? How does uncertainty play into this and how to we help people set expectations to designing *with* AI?
- Do generative AI algorithms contribute needed serendipity to the design process—or simply randomness—or worse, chaos?
- Is presenting AI as a desirable and “objective” method appropriate for generative AI?

We encourage people to write and answer their own questions as well. We hope that the workshop leads to new ways-of-thinking.

These themes can be addressed within the following topics:

- The emerging capabilities of generative AI.
- Generative AI existence in different domains including (but not limited to) images, music, text, design, and motion.
- The role of generative AI in assisting, replacing, and regimentering human work.
- Human-AI collaboration and co-creative systems.
- Ethical issues including misuses and abuses, provenance, copyright, bias, and diversity.
- The uncanny valley in Human-AI interactions.
- Speculative futures of generative AI and its implications for human-AI utopias and dystopias.

As above, we encourage people to add new topics and domains.

4 WEBSITE

Our workshop website is: <https://generativeaiandhci.github.io> (currently online).

5 PRE-WORKSHOP PLANS

We will publicize the workshop via ACM, HCI, and AI distribution lists, plus social media [i.e., similarly to 3, 16, 31, 36, 39, 41]. We will also reach out to selected scholars to encourage them to submit their work.

We will request participants to submit completed work, works-in-progress, and provocations. Submissions should use the single-column ACM small format (see <https://chi2022.acm.org/for-authors/presenting/papers/chi-publication-formats/>). While the CHI 2022 instructions for *papers* say that there is no page-limit, we request that each submission be limited to four pages of content; references may appear on subsequent pages, and are not counted toward the four-page limit. The four-page submissions may take the form of reports, essays, or pictorials (for more information about pictorials, please see the section “What are Papers and Pictorials?” at <https://dis.acm.org/2021/papers-and-pictorials/>).

The co-organizers will select submissions for inclusion in the workshop. We anticipate a workshop size of 20-30 participants. Selection will be based on uniqueness and/or provocativeness of content, engagement with the themes and topics in the workshop call, and potential for contribution to this new community. All submissions will be subjected to anonymous peer-review by at least two experts. If we receive a large number of submissions, we will recruit a program committee from known researchers and

scholars who have worked at the intersection of HCI and generative AI.

6 WORKSHOP STRUCTURE

The goals of the workshop are to build community, exchange knowledge and to co-create new understandings in this emerging field. As we note below in Section 8, the workshop will run remote and it may be configured in multiple sessions to provide more inclusive access for participants in different timezones. We postpone setting up the exact schedules until we have a tentative list of participants, so that we can adjust the timing to fit their timezones.

The following list of topics describes the general content of our workshop sessions designed to reflect our goals to build a community, to exchange knowledge, and to co-create new understandings in this emerging field:

- (1) We will reserve significant time for the participants’ and organizers’ self-introductions and descriptions of research interests.
- (2) We will invite a subset of the participants to present the most influential or promising work submitted to the workshop - either as individual talks or in a panel format (see Section 8).
- (3) We will preserve time for small-group discussions around themes of mutual interest, or generative HCI design challenges. These themes and challenges may be identified ahead of time, or they may be emergent from the presentations.

All participants’ submissions will be made available on our website before the workshop, and communications platforms will be arranged to enable asynchronous engagement before, during, and after our sessions. Online tools (e.g., Padlet, Miro) will be used to engage our community in collaborative ideation and critical discussion.

7 POST-WORKSHOP PLANS

We hope to build a persistent hybrid research area among HCI and generative AI. Toward that goal, we plan to submit an *Interactions* article or post for the general HCI audience. In addition, we want to create resources for teaching and researching in this emergent domain. Therefore, we hope to publish selected papers from the workshop, as an invitational journal issue or a book of readings.

8 REMOTE/ONSITE PLANS

We assume that this workshop will be entirely virtual, so as to include people from around the world, and so as to avoid visa issues at international borders. To make a successful workshop under these circumstances, we propose to schedule several partial sessions, held so as to accommodate timezones of multiple continents. For example, if we receive submissions from Asia, Europe and Africa, and the Americas, then we would schedule three partial sessions (perhaps four hours per session) so as to include people from each of those regions. We will need to work adaptively and inclusively, depending on the submissions.

We realize that participants may not be able to attend all sessions, and we will therefore request that presenters prepare a video version of their presentation, and/or a slide-deck and a “script” that one of the organizers can read during the slide-deck. Each presentation will be followed by discussion. If authors can attend the discussion,

that will make the discussion richer. However, we know from many reading groups that it is possible to conduct a rich and productive discussion even if the author is not present.

9 ORGANIZERS

Lydia B. Chilton is an Assistant Professor in the Computer Science Department at Columbia University. She is an early pioneer in decomposing complex tasks so that crowds and computers can solve them together. Her current research is in computational design - how computation and AI can help people with design, innovation and creative problem solving. Applications include: conveying a message within an image for journalism and advertising, developing technology for public libraries, improving risk communication during hurricanes, and helping scientists explain their work on Twitter.

Anna Kantosalo is a Postdoctoral Researcher at the University of Helsinki and a Service Designer at Siili Solutions. Anna focuses on the design and evaluation of AI solutions with a creativity twist. Anna has co-chaired the PC committee for the 13th International Conference on Computational Creativity and organized several workshops related to interactions with creative systems.

Q. Vera Liao is a Principal Researcher at Microsoft Research Montréal, where she is part of the FATE (Fairness, Accountability, Transparency, and Ethics of AI) group. Her current research interests are in human-AI interaction and responsible AI. She serves as the Co-Editor-in-Chief for Springer HCI Book Series, in the Editors team for CSCW, and on the Editorial Board of ACM Transactions on Interactive Intelligent Systems (TiiS).

Mary Lou Maher is a Professor in the Software and Information Systems Department at the University of North Carolina at Charlotte. Her early research in AI-based generative design has led to a human centered approach to computational creativity and co-creative systems. She has Chaired the Creativity and Cognition Conference (2019) and the International Conference on Computational Creativity (2019) as well as organized several workshops on AI-based design and creativity.

Charles Martin is a Senior Lecturer in Computer Science at the Australian National University. Charles works at the intersection of music, AI/ML and HCI. He studies how humans can interact creatively with intelligent computing systems and how such systems might fit in the real world. Charles has organized multiple generative-AI-focused workshops at the New Interfaces for Musical Expression conference.

Michael Muller works as a Research Scientist at IBM Research in Cambridge MA USA, and is co-author of *Human Centered Data Science: An Introduction*. He has analyzed how domain experts make use of generative AI outcomes, and how humans intervene between "the data" and "the model" in responsible and accountable data science work. He has co-organized workshops on human centered data science at CHI, CSCW, GROUP, and NeurIPS.

Greg Walsh is an associate professor at the University of Baltimore where he teaches courses in Design. He is an interaction design researcher who focuses on user-centered, inclusive design for children and adults. His work seeks to include more voices in the design process and has been a recipient of a prestigious Google

Faculty Research Award. His work has included participatory design sessions in Baltimore City libraries and is now exploring the use of generative AI as a co-design partner.

10 CALL FOR PARTICIPATION (250 WORDS)

This workshop applies human centered themes to a new and powerful technology, generative artificial intelligence (AI). Unlike AI systems that produce decisions or descriptions, generative AI produces new instances of types of data that can include images, texts, music, design, and motion.

However, it is not yet clear how humans can make sense of generative algorithms and outcomes. We have yet to understand what user interface technologies will enable humans to control, and more generally to interact with these powerful, variable capabilities. These human-like capabilities put into question our current paradigms for mixed initiative user interfaces. Further, the unpredictability of "creative" algorithms raises new questions about how, when, and how much control humans may wish to share with these algorithms, and what types of societal governance may be needed for such powerful capabilities. Finally, it is not clear what kinds of collaboration patterns will emerge when creative humans and creative technologies work together.

In this one-day workshop, we will convene the interdisciplinary research domain of generative AI and HCI. Participation in this invitational workshop is open to seasoned scholars and early career researchers from diverse disciplines. We solicit descriptions of completed projects, works-in-progress, and provocations. Together we will develop theories and practices in this intriguing new domain. Please visit <https://generativeaiandhci.github.io/> to learn more about the workshop, and about how you can participate. Please send any questions to Michael Muller, michael_muller@us.ibm.com.

REFERENCES

- [1] Lisa C Adams, Felix Busch, Daniel Truhn, Marcus R Makowski, Hugo JWL Aerts, and Keno K Bressem. 2022. What Does DALL-E 2 Know About Radiology? *arXiv preprint arXiv:2209.13696* (2022).
- [2] Saleema Amershi, Dan Weld, Mihaela Vorvoreanu, Adam Fourney, Besmira Nushi, Penny Collisson, Jina Suh, Shamsi Iqbal, Paul N Bennett, Kori Inkpen, et al. 2019. Guidelines for human-AI interaction. In *Proceedings of the 2019 chi conference on human factors in computing systems*. 1–13.
- [3] Cecilia Aragon, Clayton Hutto, Andy Echenique, Brittany Fiore-Gartland, Yun Huang, Jinyoung Kim, Gina Neff, Wanli Xing, and Joseph Bayer. 2016. Developing a research agenda for human-centered data science. In *Proceedings of the 19th ACM Conference on Computer Supported Cooperative Work and Social Computing Companion*. 529–535.
- [4] Anon. at DeMilked.com. 2020. *Here's How 20 Famous Historical And Fictional Figures 'Really' Looked Like*. Retrieved Oct 4, 2022 from <https://www.demilked.com/historical-figures-recreated-bas-uterwijk/>
- [5] Shraddha Barke, Michael B James, and Nadia Polikarpova. 2022. Grounded Copilot: How Programmers Interact with Code-Generating Models. *arXiv preprint arXiv:2206.15000* (2022).
- [6] Su Lin Blodgett, Q Vera Liao, Alexandra Olteanu, Rada Mihalcea, Michael Muller, Morgan Klaus Scheuerman, Chenhao Tan, and Qian Yang. 2022. Responsible Language Technologies: Foreseeing and Mitigating Harms. In *CHI Conference on Human Factors in Computing Systems Extended Abstracts*. 1–3.
- [7] Terence Broad, Sebastian Berns, Simon Colton, and Mick Grierson. 2021. Active Divergence with Generative Deep Learning—A Survey and Taxonomy. *arXiv preprint arXiv:2107.05599* (2021).
- [8] Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. 2021. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374* (2021).
- [9] Katherine Crowson. 2021. Introduction to VQGAN CLIP. Google Docs. <https://docs.google.com/document/d/>

- 1Lu7XPRKINHBJQjK8k8qRzUzbW7kzxb5VU72GMRn2E/edit
- [10] Katherine Crowson. 2021. Rivers Have Wings. <https://twitter.com/RiversHaveWings>
 - [11] Jay L Cunningham, Gabrielle Benabdallah, Daniela K Rosner, and Alex S Taylor. 2022. On the Grounds of Solutionism: Ontologies of Blackness and HCI. *ACM Transactions on Computer-Human Interaction* (2022).
 - [12] Nicholas Davis. 2013. Human-computer co-creativity: Blending human and computational creativity. In *Doctoral Consortium of the Ninth AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment* (Boston, Massachusetts, USA, October 14-18.). AAAI, 9–12.
 - [13] Nicholas Davis, Chih-Pin Hsiao, Kunwar Yashraj Singh, Lisa Li, and Brian Magerko. 2016. Empirically studying participatory sense-making in abstract drawing with a co-creative cognitive agent. In *Proceedings of the 21st International Conference on Intelligent User Interfaces*. 196–207.
 - [14] Bestiario del Hypogripho. 2021. Ayuda:Generar imágenes con VQGAN CLIP/English. https://tuscriaturas.miraheze.org/w/index.php?title=Ayuda:Generar_imágenes_con_VQGANCLIP/English
 - [15] Sebastian Deterding, Jonathan Hook, Rebecca Fiebrink, Marco Gillies, Jeremy Gow, Memo Akten, Gillian Smith, Antonios Liapis, and Kate Compton. 2017. Mixed-initiative creative interfaces. In *Proceedings of the 2017 CHI Conference Extended Abstracts on Human Factors in Computing Systems*. 628–635.
 - [16] Upol Ehsan, Philipp Wintersberger, Q Vera Liao, Martina Mara, Marc Streit, Sandra Wachter, Andreas Riemer, and Mark O Riedl. 2021. Operationalizing Human-Centered Perspectives in Explainable AI. In *Extended Abstracts of the 2021 CHI Conference on Human Factors in Computing Systems*. 1–6.
 - [17] Ahmed Elgammal, Bingchen Liu, and Kunpeng Song. 2020. Sketch-to-Art: synthesizing stylized art images from hand-drawn sketches with no semantic labeling. In *ACM SIGGRAPH 2020 Real-Time Live!* 1–1.
 - [18] Jesse Engel, Kumar Krishna Agrawal, Shuo Chen, Ishaan Gulrajani, Chris Donahue, and Adam Roberts. 2019. GANSynth: Adversarial Neural Audio Synthesis. In *Proceedings of the International Conference on Learning Representations*. <https://openreview.net/pdf?id=H1xQVn09FX>
 - [19] Patricia Garcia, Tonia Sutherland, Marika Cifor, Anita Say Chan, Lauren Klein, Catherine D'Ignazio, and Niloufar Salehi. 2020. No: critical refusal as feminist data practice. In *Conference Companion Publication of the 2020 on Computer Supported Cooperative Work and Social Computing*. 199–202.
 - [20] Leon A. Gatys, Alexander S. Ecker, and Matthias Bethge. 2015. A Neural Algorithm of Artistic Style. *arXiv:1508.06576 [cs.CV]*
 - [21] Pablo Gervás. 2001. An Expert System for the Composition of Formal Spanish Poetry. In *Applications and Innovations in Intelligent Systems VIII*, Ann Macintosh, Mike Moulton, and Frans Coenen (Eds.). Springer London, London, 19–32.
 - [22] Werner Geyer, Lydia B Chilton, Justin D Weisz, and Mary Lou Maher. 2021. HAI-GEN 2021: 2nd Workshop on Human-AI Co-Creation with Generative Models. In *26th International Conference on Intelligent User Interfaces*. 15–17.
 - [23] GitHub. 2022. *GitHub Copilot - Your AI pair programmer*. Retrieved September, 2022 from <https://github.com/features/copilot/>
 - [24] Anna Lauren Hoffmann. 2021. Even when you are a solution you are a problem: An uncomfortable reflection on feminist data ethics. *Global Perspectives* 2, 1 (2021).
 - [25] Stephanie Houde, Vera Liao, Jacquelyn Martino, Michael Muller, David Piorkowski, John Richards, Justin Weisz, and Yunfeng Zhang. 2020. Business (mis) use cases of generative AI. *arXiv preprint arXiv:2003.07679* (2020).
 - [26] Cheng-Zhi Anna Huang, Ashish Vaswani, Jakob Uszkoreit, Noam Shazeer, Ian Simon, Curtis Hawthorne, Andrew Dai, Matt Hoffman, Monica Dinulescu, and Douglas Eck. 2019. Music Transformer: Generating Music with Long-Term Structure. <https://arxiv.org/abs/1809.04281>
 - [27] Justin. 2021. Somewhere Systems Twitter. <https://twitter.com/somewheresy>
 - [28] Anna Kantosalo, Jukka M. Toivanen, Ping Xiao, and Hannu Toivonen. 2014. From Isolation to Involvement: Adapting Machine Creativity Software to Support Human-Computer Co-Creation. In *Proceedings of the Fifth International Conference on Computational Creativity, Ljubljana, Slovenia*. 1–7. http://computationalcreativity.net/icc2014/wp-content/uploads/2014/06/1.1_Kantosalo.pdf
 - [29] Pegah Karimi, Kaz Grace, Mary Lou Maher, and Nick Davis. 2018. Evaluating Creativity in Computational Co-Creative Systems. In *Proceedings of the 2018 international conference on computational creativity*, Vol. 147. Citeseer.
 - [30] Pegah Karimi, Mary Lou Maher, Nick Davis, and Kaz Grace. 2019. Deep Learning in a Computational Model for Conceptual Shifts in a Co-Creative Design System. In *Proceedings of the 2019 international conference on computational creativity*, Vol. 147. Citeseer.
 - [31] Marina Kogan, Aaron Halfaker, Shion Guha, Cecilia Aragon, Michael Muller, and Stuart Geiger. 2020. Mapping Out Human-Centered Data Science: Methods, Approaches, and Best Practices. In *Companion of the 2020 ACM International Conference on Supporting Group Work*. 151–156.
 - [32] Harsh Kumar, Ilya Musabirov, Jiakai Shi, Adele Lauzon, Kwan Kiu Choy, Ofek Gross, Dana Kulzhabayeva, and Joseph Jay Williams. 2022. Exploring The Design of Prompts For Applying GPT-3 based Chatbots: A Mental Wellbeing Case Study on Mechanical Turk. *arXiv preprint arXiv:2209.11344* (2022).
 - [33] Mina Lee, Percy Liang, and Qian Yang. 2022. Coauthor: Designing a human-ai collaborative writing dataset for exploring language model capabilities. In *CHI Conference on Human Factors in Computing Systems*. 1–19.
 - [34] Ryan Louie, Andy Coenen, Cheng Zhi Huang, Michael Terry, and Carrie J Cai. 2020. Novice-AI music co-creation via AI-steering tools for deep generative models. In *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems*. 1–13.
 - [35] Mary Lou Maher, Katherine Brady, and Douglas H Fisher. 2013. Computational models of surprise in evaluating creative design. In *Proceedings of the fourth international conference on computational creativity*, Vol. 147. Citeseer.
 - [36] Charles Patrick Martin, Fabio Morreale, Benedikte Wallace, and Hugo Scurto. 2021. Critical Perspectives on AI/ML in Musical Interfaces. Workshop at NIME 2021. <https://critical-ml-music-interfaces.github.io>
 - [37] Justin Matejka, Michael Glueck, Erin Bradner, Ali Hashemi, Tovi Grossman, and George Fitzmaurice. 2018. Dream Lens : Exploration and Visualization of Large-Scale Generative Design Datasets. (2018), 1–12.
 - [38] Justin Matejka, Michael Glueck, Erin Bradner, Ali Hashemi, Tovi Grossman, and George Fitzmaurice. 2018. *Dream Lens: Exploration and Visualization of Large-Scale Generative Design Datasets*. Association for Computing Machinery, New York, NY, USA, 1–12. <https://doi.org/10.1145/3173574.3173943>
 - [39] Michael Muller, Cecilia Aragon, Shion Guha, Marina Kogan, Gina Neff, Cathrine Seidelin, Katie Shilton, and Anissa Tanweer. 2020. Interrogating Data Science. In *Conference Companion Publication of the 2020 on Computer Supported Cooperative Work and Social Computing*. 467–473.
 - [40] Michael Muller, Lydia Chilton, Anna Kantosalo, Mary Lou Maher, Charles Patrick Martin, and Greg Walsh. 2022. *GenAICHI: Generative AI and Computer Human Interaction*. Retrieved Oct 4, 2022 from <https://sites.google.com/view/genaichi2022/home>
 - [41] Michael Muller, Melanie Feinberg, Timothy George, Steven J Jackson, Bonnie E John, Mary Beth Kery, and Samir Passi. 2019. Human-centered study of data science work practices. In *Extended Abstracts of the 2019 CHI Conference on Human Factors in Computing Systems*. 1–8.
 - [42] Michael Muller, Steven Ross, Stephanie Houde, Mayank Agarwal, Fernando Martinez, John Richards, Kartik Talamadupula, and Justin D Weisz. 2022. Drinking Chai with Your (AI) Programming Partner: A Design Fiction about Generative AI for Software Engineering. *HAI-GEN Workshop at IUI 2022: 3rd Workshop on Human-AI Co-Creation with Generative Models* (2022). <https://hai-gen.github.io/2022/>
 - [43] Michael Muller, April Y. Wang, Steven I. Ross, Justin D. Weisz, Mayank Agarwal, Kartik Talamadupula, Stephanie Houde, Fernando Martinez, John Richards, Jaimie Drozdal, Xie Lui, David Piorkowski, and Dakuo Wang. 2021. *How data scientists improve generated code documentation in Jupyter notebooks*. Retrieved October 5, 2021 from <https://hai-gen2021.github.io/program/>
 - [44] NeurIPS. 2019. Machine Learning for Creativity and Design. Retrieved Oct 10, 2021 from <https://neurips2019creativity.github.io/>
 - [45] Akshat Kumar Nigam, Robert Pollice, Mario Krenn, Gabriel dos Passos Gomes, and Alan Aspuru-Guzik. 2021. Beyond generative models: superfast traversal, optimization, novelty, exploration and discovery (STONED) algorithm for molecules using SELFIES. *Chemical science* (2021).
 - [46] Amin Heyrani Nobari, Muhammad Fathy Rashad, and Faez Ahmed. 2021. Creativegan: editing generative adversarial networks for creative design synthesis. *arXiv preprint arXiv:2103.06242* (2021).
 - [47] NPR. 2021. Team uses AI to complete Beethoven's unfinished masterpiece. *NPR* (Oct 2021). <https://www.npr.org/2021/10/02/1042742330/team-uses-ai-to-complete-beethovens-unfinished-masterpiece>
 - [48] Jonas Oppenlaender. 2022. The Creativity of Text-based Generative Art. *arXiv preprint arXiv:2206.02904* (2022).
 - [49] Rafael Pérez Y Pérez and Mike Sharples. 2001. MEXICA: A computer model of a cognitive account of creative writing. *Journal of Experimental & Theoretical Artificial Intelligence* 13, 2 (2001), 119–139. <https://doi.org/10.1080/09528130010029820> arXiv:<https://doi.org/10.1080/09528130010029820>
 - [50] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. 2021. Learning Transferable Visual Models From Natural Language Supervision. *arXiv:2103.00020 [cs.CV]*
 - [51] Aditya Ramesh, Mikhail Pavlov, Gabriel Goh, Scott Gray, Chelsea Voss, Alec Radford, Mark Chen, and Ilya Sutskever. 2021. Zero-Shot Text-to-Image Generation. *arXiv:2102.12092 [cs.CV]*
 - [52] Reddit. 2021. *r/bigsleep: A subreddit for AI imagery generated from text descriptions*. <https://www.reddit.com/r/bigsleep/>
 - [53] Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. 2022. High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 10684–10695.
 - [54] Mark A. Runco and Garrett J. Jaeger. 2012. The Standard Definition of Creativity. *Creativity Research Journal* 24, 1 (2012), 92–96. <https://doi.org/10.1080/10400419.2012.650092> arXiv:<https://doi.org/10.1080/10400419.2012.650092>

- [55] Advait Sarkar, Andrew D. Gordon, Carina Negreanu, Christian Poelitz, Sruti Srinivasa Ragavan, and Ben Zorn. 2022. What is it like to program with Artificial Intelligence? <https://arxiv.org/abs/2208.06213v1>
- [56] Othman Sbai, Mohamed Elhoseiny, Antoine Bordes, Yann LeCun, and Camille Couprie. 2018. Design: Design inspiration from generative networks. In *Proceedings of the European Conference on Computer Vision (ECCV) Workshops*. 0–0.
- [57] Isabella Seeber, Eva Bittner, Robert O Briggs, Triparna De Vreede, Gert-Jan De Vreede, Aaron Elkins, Ronald Maier, Alexander B Merz, Sarah Oeste-Reiß, Nils Randrup, et al. 2020. Machines as teammates: A research agenda on AI in team collaboration. *Information & management* 57, 2 (2020), 103174.
- [58] Ben Shneiderman. 2007. Creativity support tools: Accelerating discovery and innovation. *Commun. ACM* 50, 12 (2007), 20–32.
- [59] Angie Spoto and Natalia Oleynik. [n. d.]. *Library of Mixed Initiative Creative Interfaces*. Retrieved Oct 10, 2021 from <http://mici.codingconduct.cc/>
- [60] Ivo Swartjes and Mariët Theune. 2009. Iterative authoring using story generation feedback: debugging or co-creation?. In *Joint International Conference on Interactive Digital Storytelling*. Springer, 62–73.
- [61] ICC21 The Twelfth International Conference on Computational Creativity. 2021. Second Workshop on the Future of Co-Creative Systems. Retrieved Oct 11, 2021 from <https://computationalcreativity.net/icc21/wfcs/>
- [62] M Onat Topal, Anil Bas, and Imke van Heerden. 2021. Exploring transformers in natural language generation: Gpt, bert, and xlnet. *arXiv preprint arXiv:2102.08036* (2021).
- [63] Twitter 2021. VQGANCLIP Hashtag. https://twitter.com/hashtag/vqganclip?src=hashtag_click
- [64] @someheresy Twitter. 2021. VQGAN CLIP Colab Notebook. https://colab.research.google.com/drive/1_4Jl0a7WlJeqy5LTjPjFZOwMZopG5C-W?usp=sharing#scrollTo=ZdlpRFL8UAIW
- [65] Quentin Vanhaelen, Yen-Chu Lin, and Alex Zhavoronkov. 2020. The advent of generative chemistry. *ACS Medicinal Chemistry Letters* 11, 8 (2020), 1496–1505.
- [66] Benedikte Wallace, Charles Patrick Martin, Jim Tørresen, and Kristian Nymoen. 2021. Exploring the Effect of Sampling Strategy on Movement Generation with Generative Neural Networks. In *Artificial Intelligence in Music, Sound, Art and Design: 10th International Conference, EvoMUSART 2021*. Springer International Publishing, 344–359. https://doi.org/10.1007/978-3-030-72914-1_23
- [67] April Yi Wang, Dakuo Wang, Jaimie Drozdal, Michael Muller, Soya Park, Justin D Weisz, Xuye Liu, Lingfei Wu, and Casey Dugan. 2021. Themisto: Towards Automated Documentation Generation in Computational Notebooks. *arXiv preprint arXiv:2102.12592* (2021).
- [68] Laura Weidinger, Jonathan Uesato, Maribeth Rauh, Conor Griffin, Po-Sen Huang, John Mellor, Amelia Glaese, Myra Cheng, Borja Balle, Atoosa Kasirzadeh, et al. 2022. Taxonomy of risks posed by language models. In *2022 ACM Conference on Fairness, Accountability, and Transparency*. 214–229.
- [69] Justin D Weisz, Michael Muller, Stephanie Houde, John Richards, Steven I Ross, Fernando Martinez, Mayank Agarwal, and Kartik Talamadupula. 2021. Perfection Not Required? Human-AI Partnerships in Code Translation. In *26th International Conference on Intelligent User Interfaces*. 402–412.
- [70] Georgios N Yannakakis, Antonios Liapis, and Constantine Alexopoulos. 2014. Mixed-initiative co-creativity. In *Proceedings of the 9th International Conference on the Foundations of Digital Games, FDG 2014* (Liberty of the Seas, Caribbean, April 3-7).
- [71] Yijun Zhou, Yuki Koyama, Masataka Goto, and Takeo Igarashi. 2020. Generative Melody Composition with Human-in-the-Loop Bayesian Optimization. *arXiv preprint arXiv:2010.03190* (2020).