

I Was Never in the AI.R.A: an Experimental Study on the Use of Manipulated Audio for the Censorship of Political Dissidents

John Twomey
School of Applied Psychology
University College Cork
Cork, Ireland
LERO, the Science Foundation Ireland
Research Centre for Software
Limerick, Ireland
118447782@umail.ucc.ie

Didier Ching
School of Applied Psychology
University College Cork
Cork, Ireland
LERO, the Science Foundation Ireland
Research Centre for Software
Limerick, Ireland
116378223@umail.ucc.ie

Cian O'Mahony
School of Applied Psychology
University College Cork
Cork, Munster, Ireland
LERO, the Science Foundation Ireland
Research Centre for Software
Limerick, Ireland
115526303@umail.ucc.ie

Conor Linehan
School of Applied Psychology
University College Cork
Cork, Ireland
LERO, the Science Foundation Ireland
Research Centre for Software
Limerick, Ireland
conor.linehan@ucc.ie

Gillian Murphy
School of Applied Psychology
University College Cork
Cork, Ireland
LERO, the Science Foundation Ireland
Research Centre for Software
Limerick, Ireland
gillian.murphy@ucc.ie

Abstract

During the conflict in Northern Ireland known as “the Troubles”, audio recordings of 10 controversial political groups were banned by the British government from being broadcast. Interestingly, video recordings of these people were instead overdubbed by paid actors, implying that the political figures’ original voices were somehow more convincing, legitimate or sexy than the actors. This study uses AI-generated audio to examine whether this censorship method affects opinions of political dissidents. We carried out a between-subjects experimental study ($n = 636$), where we compared three different types of AI-generated vocal manipulations with unedited footage of two well-known Northern Irish politicians. We found a significant negative effect of voice manipulation on perceived charisma, legitimization and credibility. Participants also voiced strong disapproval and concern at the idea of using AI to revoice historical footage. We discuss implications of AI assisted audio and video translation of politically sensitive media.

CCS Concepts

• Empirical studies in HCI; • Social aspects of security and privacy; • Psychology;

Keywords

Censorship, GerAI Adams, LondonPsychology,

ACM Reference Format:

John Twomey, Didier Ching, Cian O'Mahony, Conor Linehan, and Gillian Murphy. 2025. I Was Never in the AI.R.A: an Experimental Study on the Use of Manipulated Audio for the Censorship of Political Dissidents. In *Extended Abstracts of the CHI Conference on Human Factors in Computing Systems (CHI EA '25)*, April 26–May 01, 2025, Yokohama, Japan. ACM, New York, NY, USA, 8 pages. <https://doi.org/10.1145/3706599.3716238>

1 Introduction

Recent advances in the quality and accessibility of audio and video manipulation technologies have made the regular creation and dissemination of edited “fake” videos by many actors a reality. Academics and media commentators have pointed out the democratizing potential of AI powered media editing tools. High quality media editing was previously only available to skilled professionals with expensive hardware and tools, whereas now can be carried out using free to download smartphone applications [25]. The majority of academic discourse has focused on identifying the potential dangers of such democratization. For example, concerns have been raised that such videos could be used for spreading effective misinformation regarding elections and referenda and generally carrying out smear campaigns of politicians [26]. However, history tells us that new developments in information technology, in the long term, typically serve to centralize power, and support the interests of corporations and capital more than those of dissidents, separatists, freedom fighters or political agitators. In this paper we focus on an application of generative AI technology that has been previously unexplored in the HCI literature; state censorship of political dissidents.

Media editing has always been a part of the process of political censorship within domestic and international conflict [19]. In this paper, we draw on a bizarre historical case of media censorship in our own country, Ireland, as a model of how traditional and AI fueled media editing technologies can be used for novel and creative

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

CHI EA '25, Yokohama, Japan

© 2025 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-1395-8/2025/04

<https://doi.org/10.1145/3706599.3716238>

state censorship practices. We carry out and report the findings of a psychological experiment that asks whether subtly manipulating the voices of political speakers can delegitimize their message and undermine their natural magnetism. In this paper we also explore how the same technology could be used in a democratizing manner; in restoring archival footage that has previously been edited and censored. The democratizing of dubbing technologies may facilitate activists who seek to challenge censorship and erasure of the voices of victims of active ethnic cleansing and genocides.

2 Related work

The Troubles” in Northern Ireland describe a conflict that took place from 1969 until 1998, between one group of people (the Unionists) who believe that Northern Ireland should remain in the United Kingdom, and another group (the Republicans) who believe that the island of Ireland should consist of a single nation. The Republican cause is defined by struggles for the voting rights and freedom to demonstrate of the marginalized Catholic communities, as well as the suppression of the native Irish language and culture. On the other side, the Unionist community is concerned with the potential erasure of their freedoms under a Catholic Ireland and were further fueled by the Republic of Ireland’s strict and repressive Catholic government at the time of the Troubles.

The conflict has roots in Ireland’s history of colonization by the British Empire, in various stages, over a period of 800 years, and the subsequent War of Independence that allowed for the creation of an independent Irish State. However, the Troubles itself emerged from the shutting down of a Civil Rights movement in Northern Ireland by UK authorities and a subsequent descent into inter-group violence [5]. The violence in the North of Ireland was characterized by numerous high-profile acts of political violence by Unionist and Republican paramilitary groups, as well as by the British police force and military. In particular, the bombings carried out by the Irish Republican Army in British cities often characterize popular perception of the conflict. Much academic research has considered the morality, history and ethics of this conflict and the lives it destroyed. [5] [15]. However, little work has considered its implications for human-computer interaction.

Our work draws inspiration from a peculiar law enacted by the UK government during the Troubles that banned audio of members of 10 political groups from TV and radio (mostly focusing on Sinn Féin, an Irish political party concerned with achieving Irish independence from the United Kingdom), notably targeting the president of Sinn Féin, Gerry Adams. The ban was put in place to deprive violent paramilitary groups and political dissidents of air-time. The motivation of media censorship as stated by British prime minister Margaret Thatcher was to deprive political dissidents of the “oxygen of publicity” [18]. While the ban was intended to censor footage entirely, the wording allowed for dubbed versions of original footage to be broadcast [13]. There was an equivalent but much stricter ban in the Republic of Ireland which ran for almost 30 years and included the state broadcaster RTE claiming it would not allow a Sinn Féin member on TV to advertise a bar of soap [12] [21].

For reasons that we do not fully understand, the prevailing method through which British broadcasters chose to comply with



Figure 1: The Day Today- comedy sketch about the IRA revocation

this censorship law consisted of over-dubbing Irish republican figures with different voices, while the text of the message was kept the same. It has been suggested by some commentators that this method of censorship (like similar “security” measures [20]) counterintuitively increased support for the speakers. Instead of hearing a nervous and defensive representative being pushed on challenging questions, audiences were treated to the calm authoritative voice of a professional actor smoothly and charmingly explaining the political motivations behind dissident and loyalist violence [3]. The practice was famously lampooned in “The Day Today” (fig. 1), a British parody news show, which suggested members of Sinn Féin be forced to inhale helium before speaking to reduce the credibility of their voices [2]. In other popular culture around the conflict, it has been suggested that Gerry Adams’s voice was dubbed over because his voice was “too sexy” and thus was unfairly persuasive [11]. Surprisingly, there has been a lack of empirical work testing this hypothesis and demonstrating the effectiveness of traditional and AI overdubbing for political censorship of dissidents.

Our work tests the effectiveness of AI powered voice overdubbing by; 1) dubbing original videos of northern Irish political figures Gerry Adams and Ian Paisely, 2) testing how the original is perceived versus the dubbed voice, 3) testing the affect other vocal manipulations/qualities have on its persuasiveness, and 4) accounting for whether prior familiarity with the politicians’ effects this. We have carried out a small demonstrative study where we show users interview footage with either an original voice, a deepfake dub, or a mocking/humorous voice (e.g. squeaky like a mouse or demon). We measure the agreement of participants towards the statement made by the speaker, and how persuasive they find the voice. We also control for participants’ pre-existing opinions on the politicians. Our Null hypothesis is that there is no difference in persuasiveness between the three conditions.

Of course, there is some previous research that explores the effect of different voice qualities on perception of the speaker. For example, prosody has been linked with credibility in political speech [6]. Specific aspects of speech such as (prosody and volume for example) are associated with perceptions of charisma [17] [22]. Little work has been done approaching how these findings apply to censorship practices which use media editing technology to

manipulate voices. Lots of work focuses on specific characteristics of voices such as pitch, prosody and tone – our work situates this in a real-world example drawing from the Troubles as a real-life application of this work.

This paper proudly follows a rich tradition of HCI research that inserts generative AI into places it does not belong. We speculatively ask whether it is “good” to use AI in restoring historical footage of controversial historical conflicts. Specifically, existing reviews of generative AI technologies such as deepfakes have highlighted that many of the supposedly positive uses of these technologies in historical education materials have not been demonstrated empirically [25]. We explore the potential implications of restoring the voices of separatist leaders in previously censored archival footage.

Our research questions are as follows:

Do manipulated vocal characteristics affect the perceived credibility, legitimacy and charisma of historical politicians and their statements?

How do viewers feel about the potential of using generative AI to revoice existing censored footage?

2.1 A special note to readers

With regard to the broader application of this work; we must acknowledge that there is none. Our research is focused specifically on Irish and British conflicts and should be judged solely within that sphere. We make no claims around the effects of political censorship and counterterrorism. The Irish experience has little relevance to broader global political affairs, as our history is quite specific and unlike any other violence in the world today- a former colony of the British Empire divided along religious grounds between the settler-colonizer and indigenous populations - the establishment of the resulting two state system contributing to a vicious cycle of terrible political violence and unrest on both sides- defined by many powerful figures of the British Empire including William Churchill, David Lloyd George, Lord Mountbatten and Arthur Balfour – a conflict defined by unnecessary death and violence, organized acts of revenge, ethnic cleansing, land seizures and the suppression and destruction of cultural identities.

3 Methods

3.1 Ethics and preregistration

This project received ethical clearance from our institution’s ethical board as part of a series of projects on the potentials of deepfake media in political disinformation. We obtained consent for participants prior to the study and gave them the option of withdrawing this consent before submitting their data. This study was also preregistered (available online here: <https://aspredicted.org/m6gt-g4d4.pdf>).

3.2 Participants

We recruited 636 valid participants from Britain and Ireland using the website Prolific. We did not include participants who did not fully complete the survey or who failed the attention checks. For the attention checks we asked participants a simple question about the content of the video and removed those who answered incorrectly for either of the two politicians. We gathered participants from the United Kingdom and Ireland. While we didn’t specifically collect

location as demographic data – the majority of participants on Prolific were British. Three participants withdrew their consent at the end of the study and their data was not analyzed. We regathered participants to make up for those who had failed one of the attention checks. In total 25 people failed the Ian Paisley attention checks and 36 failed the Gerry Adams attention checks. We also regathered participants who reported being unable to hear one or more of the conditions. We had 702 distinct participants in total in the dataset. The majority of the participants described themselves as female ($n = 384$), followed by male ($n = 312$). Four participants selected non-binary/third gender. Another self-described as bi-gender and one participant chose not to say. Participants had an average age of 43 ($SD = 13.91$).

3.3 Materials

3.3.1 Videos. For both of our target politicians, we created four versions of the same brief interview clip – original video, AI generated dubbed voice with similar Northern Irish accent, AI generated dubbed voice with an accent approximating received pronunciation (RP – an accent associated with upper-class England notable for its historical use in television news broadcasts) and a high-pitched voice similar to an individual who had inhaled helium. The motivation for the dubbed voice with a similar accent was to emulate the methods used during the historical censorship. The received pronunciation (RP) condition was meant to simulate newsreader voices, the type which translated political speech is often translated for viewers using. The high-pitch conditions were selected as ways that broadcasters might seek to delegitimize speakers – by casting them as unserious (high-pitch). The two target politicians were Gerry Adams and Ian Paisley. Gerry Adams was chosen as he is the most well-known target of the broadcasting ban and one of the figureheads of the Republican side of the Troubles. Ian Paisley was chosen as he could be considered a similarly high profile and controversial Unionist figure. The specific clips we chose both speak to the differing interpretations of the two politicians on the idea of peace in the North and how they saw themselves as advocating for self-defense and non-violence (fig. 2).

The AI dubbed voice was generated using play.ht, an online voice cloning software. For Gerry Adams, the voice was generated based on a recording of the voice actor who was used to overdub Gerry Adams during the ban. However, Ian Paisley was never overdubbed – for his AI generated dubbed replacement we picked Jim Alister, a unionist politician from the same part of Northern Ireland as Ian Paisley. For the received pronunciation voice, we used narakeet, an online AI voice generator and we chose a British accent which most closely resembled received pronunciation. For the high-pitched condition, we raised the pitches by 35 per cent in both politicians as we found this value was a good balance between being noticeable but not unintelligible. All materials are available on our OSF repository (https://osf.io/jvgx2/?view_only=1bcf196ad3094e498a3dc98a6e76fc1c).

3.3.2 Outcome Measures. While we set out to measure the role of audio dubbing in the perceived sexiness of political dissidents, there are in fact few robust measures for “sexiness” and it is hard to account for the impact of sexuality and sexual preference. As a result, we chose to measure charisma as a way of measuring the




Figure 2: Transcripts of the two authentic videos used in this study, Participants saw a version of either video, with either unaltered audio, or video that was doctored in one of three ways.

magnetism of these politicians. The Generalised Charisma Inventory [24] is a 6-item psychometric scale with 3 items measuring influence (e.g. knows how to lead a group) and 3 items measuring affability (e.g. “smiles at people often”). The scale is measured on a five-point Likert scale. We also used measures of political legitimacy and credibility drawing from existing deepfake literature [4]. We measured political legitimacy by asking “Based on the video you just watched, please indicate how you feel about this politician by indicating to what extent you agree with the following statements.” (1 = completely disagree, 7 = completely agree)”, with items such as “This politician sells out their principles”. Participants were asked to rate the credibility of the video by their agreement with certain descriptors on the same scale. (E.g. this video is: Informative (1 = completely disagree, 7 = completely agree)).

3.4 Procedure.

After providing informed consent and demographic information, participants were shown three videos in a random order: one of the four versions of the Gerry Adams video, one of the four versions of the Ian Paisley video, and a control video of the former British Prime Minister Margaret Thatcher. After each video, participants completed an attention check and then rated the charisma, legitimacy and credibility of the video. After they were debriefed, participants were then told about the purpose of the study and shown an original dubbed video of Gerry Adams that was broadcast during the ban. We also showed them a redubbed version we

made using deepfake technology, and asked participants to reflect on how they felt about this use of AI technology.

For the revoiced Gerry Adams we used the same voice cloning software on the original audio that we used for the control condition. This generated our Gerry Adams audio clone. We manually matched this to the original video with the intent to match the lips of the original video as we felt using lip-syncing technology would defeat the purpose of “restoring” the original video. This process was carried out iteratively with multiple versions of each sentence trialed to capture the cadence and also his accent. We were limited in our choice for the dubbed audio by the few recordings of dubbed Gerry Adams that exist online (this was the most significant effect of the broadcasting ban: it reduced coverage of Sinn Féin altogether due to extra hassle of revoicing footage) [13]. This task was perhaps the most labor intensive of the materials to produce, and inevitably also the poorest quality.

3.5 Positionality.

As an Irish research team, we acknowledge the sensitive issue of political violence in Ireland and Britain. The authors do not make any claims on being pro-terrorism or censorship but simply encourage HCI to realize that political agitators are a unique stakeholder group who often get ignored by traditional user-focused HCI literature. Though we as HCI researchers may ignore how people will use technology for political violence across the globe [1], it won't stop this from happening. Similarly, while HCI has made some




Figure 3: Perceived charisma, legitimacy and credibility of the target politicians, as a function of audio condition.

significant and laudable contributions to the field of censorship by media giants and authoritarian states, we urge readers to further consider how generative AI may revolutionize political censorship, especially for social media companies who are already party to genocides worldwide [14]. As part of this positionality statement, we also must disclose that our university was built by the British government during the state-sponsored famine of the 1840's which killed a quarter of the Irish population and forced another quarter to emigrate. One university building is also built on the bones of 13 IRA prisoners of war who were interned and executed during the Irish war of Independence. While we plan to visit the mass grave to pay respects, we cannot make any claims as to what they would have thought about the use of deepfake technology in political conflicts.

4 Results

We first checked whether there was any evidence that Gerry Adams was particularly magnetic, relative to our other politician, Ian Paisley, as has been claimed [11]. We found that in fact, in every condition of our experiment, Gerry Adams was rated as marginally but significantly less charismatic ($M = 18.44$, $SD = 4.33$) than Ian Paisley ($M = 19.50$, $SD = 3.88$), $t(1347) = 4.74$, $p < .001$, at least in the eyes of our (mostly British) sample.

4.1 Do manipulated vocal characteristics affect the credibility, legitimacy, and perceived charisma of politicians and their statements?

As per our preregistration, we conducted 3 one-way between groups ANOVAs to find out whether the perceived charisma, legitimization and credibility of political statements differed between our four conditions. The perceived charisma differed significantly across all four conditions $F(3, 1345) = 4.579$, $p = .003$. The Northern Irish dub ($M = 19.42$, $SD = 4.08$) and the original video ($M = 19.35$, $SD = 3.93$) were viewed as more charismatic than the received pronunciation

($M = 18.57$, $SD = 4.07$) and high pitched ($M = 18.97$, $SD = 4.41$) conditions (see fig. 3). Post-hoc Tukey tests indicated that the latter two conditions were significantly different from the first two.

We also found significant differences in perceived legitimacy between the conditions, $F(3, 1345) = 4.08$, $p = .009$. The AI dubbed video with the Northern Ireland voice was rated the highest ($M = 36.77$, $SD = 7.69$), followed by the original condition ($M = 35.99$, $SD = 7.35$) and received pronunciation dub ($M = 35.77$, $SD = 6.76$). The high-pitched voice was seen as the least legitimate ($M = 34.81$, $SD = 7.57$). Post-hoc Tukey tests indicated that there were significant differences between the Northern Irish dub and high-pitched conditions.

There were also significant results for the credibility ratings $F(3, 1345) = 16.275$, $p < .001$. Unsurprisingly the unedited footage was seen as the most credible ($M = 28.80$, $SD = 6.32$), followed by the dubbed footage with a similar accent ($M = 27.33$, $SD = 7.92$) and the high pitched voice ($M = 26.46$, $SD = 7.58$). For this condition the received pronunciation dub scored lowest on average ($M = 24.88$, $SD = 7.78$) perhaps reflecting how it is a more obvious manipulation. Post-hoc Tukey tests indicated that the unedited condition was significantly more credible than the high pitched and received pronunciation dubs.

We carried out a robustness check by using an ANCOVA which showed that these results were stronger when controlling for participants pre-existing support for the target politicians (further details on this are available on our OSF repository).

4.2 How do viewers feel about the potential of using generative AI to revoice existing censored footage?

For our second research question, we asked the participants for their opinion of our completely novel and groundbreaking AI revoiced Gerry Adams and compared their responses and found an overwhelmingly negative response to the potentials of generative AI (fig. 4). We discuss these findings in further detail and




Figure 4: Opinions on the use of AI revoiced audio of Gerry Adams

also present the qualitative data which often contextualizes their responses.

When asked as to whether they considered the use of AI revoicing a good thing on a scale of one to 5, participants had mixed feelings. Even participants who saw the revoicing manipulation as a good thing ($M = 2.36$, $SD = 1.25$) usually added caveats around the potential misuse of the technology and the controversial nature of the footage/conflict: *“I think it’s good to help recreate the past for educational purposes but only if it’s made clear it is AI. Also, it should not be used to manipulate, misinform or misrepresent.”* Some participants noted that the restored footage matched the original lip movements better. *“The AI generated voice gives more credence as it fits with speech pattern better. We should be able to look back at history - whether we agree with the politics of the time, it is still important to learn from it, and to give a voice to those involved.”*

“It was very hard to spot the difference I wouldn’t approve in these circumstances due to the person being involved and his reputation. There are times when AI dubbed interviews is a good thing to bring back memories of good times and there are other times when it shouldn’t happen like in the above.”

On the whole, participants tended to view the revoiced footage quite negatively. When asked to rate from 1-5 whether they agree this use of AI was disrespectful, participants tended to agree ($M = 3.19$, $SD = 1.28$).

“It is scary, it’s messing with history. I remember those times clearly, lived through the troubles & this manipulation of history for future generations will distort the truth & could be used to insight trouble again. It’s totally pointless.”

There were many different reasons given for this including the potential offensive connotations of modified accents and their values *“Dubbing is inaccurate and highlights negative connotations associated with certain accents”*. There were many who viewed this use of AI as being particularly controversial due to the nature of the conflict and the various controversies surrounding Gerry Adams as a politician. *“I think this was a different time and to return back to it is negative hopefully Ireland has moved on and realise the pen is mightier than the sword. While I doubt very much Irelands troubles will be gone in my lifetime we cannot return to a time when people lived in fear blaming each other.”* Many focused specifically on Gerry Adams as a particularly controversial politician whose inclusion cast shade on the legitimacy of the entire idea: *“It is irrelevant as Gerry Adams is trying to legitimise himself through politics when all he is, is a blatant TERRORIST.”*

On average, most participants disagreed that AI could be useful for Historical records ($M = 2.82$, $SD = 1.35$).

“I just don’t like this idea. It makes me feel very uncomfortable. Further strengthens my doubts and worries surrounding AI.”

Some people agreed that the AI would be useful for recreating historical records *“I think that AI has the ability to add context to these old pieces of dialogue helping to bring greater clarity to what was attempting to be conveyed at the time, thus helping to provide a more*

accurate historical record” / “I think, although it is a simulation of the original recording, it is closer to the person’s voice and intonation than the actor – so in a way, you are getting a more accurate snapshot of the interview as it actually was”. Many more expressed the opinion that it was “rewriting history.

“Whatever AI is used requires some human input in training which particularly in a hostile political sphere may introduce biases”

Participants felt that a better approach would involve releasing the original recordings. *“The AI restoration is clumsy, and actually uninspiring because of that. The original dubbed footage is more convincing! (AI can, surely, do better!) But regardless of that shoddy AI the concept is just odd – the ‘ban’ was how things were, it’s better to have the footage as it was then viewed – not least as the politicians involved knew they would be dubbed, so we cannot judge how they spoke in reality, and they never meant us to know. By the way, if you want to ‘restore’ these interviews someone – the BBC etc – must have the actual real recordings, so why faff about with shoddy AI?”.* They often complained about the quality of the AI revoiced video – which the primary researcher took very personally. *“It didn’t sound like a real person far less Gerry Adams-unless he’s suddenly developed a pseudo American accent. It actually felt as if I was watching some bizarre Planet of the Apes re-enactment”.*

Many agreed with the idea that historical revoicing of AI is dangerous ($M = 3.66$, $SD = 1.223$) Most criticism more broadly criticized AI as being inherently biased and hostile towards truth. *“I am hesitant to support any form of AI at all, let alone when it comes to representing historical events.”* Perhaps the most interesting takes on the historical context of this redubbing were from those who remembered and experienced the Troubles, the broadcasting ban and its consequences. *“I’m old enough to have heard Gerry Adams speak, so I know what he sounds like. I think the content of what he says is more important than the way it sounds. I think that from an historical point of view hearing his voice is important.”.* Participants described the importance of learning about censorship and political conflict without the need for gaudy “restored” content.

“I was born and raised in Belfast and lived there until 1999. I remember Sinn Féin not being allowed to have their voice on any broadcast. I think it would better to leave these videos the way that they were originally shown, but also explain exactly why they were shown like this and how things came to be like this.”

Overall, like most use cases of AI, the use of AI in historical records perhaps brings more challenges than benefits, whether it be the inherent falsity of the content or the questions it raises about how easy it may be to revoice historical footage.

5 Implications

We believe that our work is the first to experimentally test the use of AI supported censorship of political voices. We find that the voice used to dub a politician affects how they are perceived, highlighting the potentially significant power that lies in the hands of those who manipulate the voices of political speakers. This finding has implications beyond the absurdly specific context of the ban enforced during the Troubles. Most significantly, it underscores how international media can sway public opinion with their choice of translator. Historically too, the current study provokes important questions about how this ban was conceived of and strictly enforced for so many years, without any evidence to support the idea that simply revoicing figures like Gerry Adams could reign in their apparently dangerous magnetism. The condition which most closely approximated the manipulation of voices used in the Troubles was rated as significantly more credible, legitimate and charismatic than either the received pronunciation or high-pitched dubs. Perhaps the trick to censoring dissident Irish republican voices this whole time was, ironically, to make them sound English. This work recontextualizes the censorship employed during the Troubles by asking if and how such censorship works.

5.1 Implications for HCI non-work in historical records

This work provides strong critiques of the use of generative AI in editing historical records. We hope that this work is off-putting for researchers who seek to design technologies which use AI to “restore” existing footage, particularly where such footage is already controversial or edited. We have shown empirically that the use of AI in historical education and record keeping is tainted by overwhelming fears of the technology being used to rewrite history – adding to calls for more work interrogating the positive uses of media editing technology [25]. Our work highlights how footage edited by AI is not seen by viewers as historical anymore and is perhaps tantamount to automating the job of Winston Smith in 1984. Other forms of automated historical recreation such as the restoring or recoloring of footage should also be critically considered. Our work further highlights once again how every well-intentioned use of generative AI is tainted by the well of negative public opinion on the technology. This is consistent with work done on the supposedly positive, but mostly creepy use of deepfakes in recasting movies [16]. Faith in this technology was hard to find in the responses. The sentiment of responses indicates that the marketing of hype of AI is dangerous and disruptive may be turning people away from the technology itself.

Outside of the implications for generative AI, there is a broader implication here for how translated political speech is dubbed over using technology. As international conflict is often subject to translation – there should be an increased focus on how the voices we use for translated political speech affect people’s attitudes. As previous literature has explored the potential of “resurrecting Margaret Thatcher” using deepfakes [7], republican activists may attempt to resurrect dissident republican leaders to challenge her. There are also broader implications for the ethics of overdubbing translated footage that need to be encouraged. If a hypothetical French TV station ran a clip of, for example, a buffoonish former/current U.S.

president, would the French language dub affect the persuasiveness of the content? The attractiveness of the politician? How do the voices we apply affect the impact of political statements? There are numerous design implications for the applications of AI voices.

5.2 Future HCI work in political censorship

There are two main takeaways of this paper for HCI researchers: different voices have different values, and neutrality does not exist. While technology is often used by those intending political violence, there has been a generalized handwaving regarding the use of technology by state and corporate actors, and the role we as HCI researchers have to play in this area. We strongly believe that it is time to take a stance on whether or not we approve of the wholesale use of technology in murdering and censoring political radicals. Certainly, a large number of industry researchers at CHI may already think this is an admirable goal based on their companies' track records. Instead of assuming our work won't fuel acts of violence and acting surprised when it does [23] - we urge researchers to take a more active position on how media editing technology begets censorship.

An alternate reading of this paper is that censorship works and governments looking to manipulate audio should particularly consider making enemy politicians sound like chipmunks. We situate this potential in an approach we think of as unprincipled, dishonourable, uncaring, sellout HCI inspired by existing alt.chi practices [9] [10]. Our paper is submitted at a time when alt.chi is discontinued, as CHI embraces the move towards a conference where it becomes impossible to tell which papers are intentionally disgusting alt.chi papers and which papers are just chi papers with poor taste. Building on established ideas such as anarchi and punk HCI [8], we acknowledge the "you know I was punk once" HCI or "conservatives are the new punks" HCI approaches. This is a radical move away from utopian futures and instead imagines a dystopian past/present. Is it time to embrace CHI's move away from overly provocative and speculative work. We should instead move toward "non-speculative fact" - everything is fine as it is, and our work aims to keep it that way.

Acknowledgments

This publication is the result of research conducted with the financial support of Taighde Éireann – Research Ireland under Grant number 13/RC/2094_2.

References

- [1] Aniruddha Kumar Baro. 2019. Information and Communication Technology and Ethnic Riots: Riots in Goalpara District of Assam, India. *Int. J. Innov. Technol. Explor. Eng.* 8(12), 2072-2075.
- [2] James Brassett. 2021. Irony and the Liminality of Resistance. In *The Ironic State* (pp. 77-96). Bristol University Press.
- [3] Tony Hall. 1994. A gag that chokes freedom: 3 edition. Independent (London, England : 1986)
- [4] Michael Hamelers, Toni GLA. van der Meer, and Tom Dobber. 2024. Distorting the truth versus blatant lies: The effects of different degrees of deception in domestic and foreign political deepfakes. *Computers in Human Behavior*, 152, 108096.
- [5] Caroline Kennedy-Pipe. 2014. The origins of the present troubles in Northern Ireland. Routledge.
- [6] Gabrijela Kišček. 2018. Persuasive power of prosodic features. *Argumentation and advocacy*, 54(4), 345-350.
- [7] Dominic Lees, Tom Bashford-Rogers, and Marcus Keppel-Palmer (2021). The digital resurrection of Margaret Thatcher: Creative, technological and legal dilemmas in the use of deepfakes in screen drama. *Convergence*, 27(4), 954-973.
- [8] Conor Linehan and Ben Kirman. 2014. Never mind the bollocks, i wanna be anarCHI: a manifesto for punk HCI. In *CHI'14 Extended Abstracts on Human Factors in Computing Systems* (pp. 741-748).
- [9] Conor Linehan, and Ben Kirman. "Mc hammer presents: the hammer of transformative nostalgia-designing for engagement at scale." In *Proceedings of the 2017 CHI conference extended abstracts on human factors in computing systems*, pp. 735-746. 2017.
- [10] Conor Linehan, Sabine Harrer, Ben Kirman, Shaun Lawson, and Marcus Carter. 2015. "Games against health: a player-centered design philosophy." In *Proceedings of the 33rd Annual ACM Conference Extended Abstracts on Human Factors in Computing Systems*, pp. 589-600.
- [11] Lisa McGee. 2018. Derry Girls (TV series). Hat Trick Productions.
- [12] Niall Meehan. 2003. How RTE censored its censorship. *Sunday Business Post*.
- [13] Eoin Moloney. 1991. Closing down the airwaves: The story of the broadcasting ban. In *The media and Northern Ireland: covering the troubles* (pp. 8-50). London: Palgrave Macmillan UK.
- [14] Paul Mozur. 2018. A genocide incited on Facebook, with posts from Myanmar's military. *The New York Times*, 15(10), 2018.
- [15] Orla T. Muldoon. 2004. Children of the Troubles: The impact of political violence in Northern Ireland. *Journal of Social Issues*, 60(3), 453-468.
- [16] Gillian Murphy, Didier Ching, John Twomey and Conor Linehan. 2023. Face/Off: Changing the face of movies with deepfakes. *Plos one*, 18(7), e0287503.
- [17] Oliver Niebuhr and Radek Skarnitzl. 2021. Those who shout the loudest—Do they sound more charismatic?. In *Proc. 1st International Conference on Tone and Intonation (TAD)* (pp. 1-5).
- [18] Max Pettigrew. 2016. The 'oxygen of publicity' and the suffocation of censorship: National newspaper representations of the British broadcasting ban (1988–94). In *The Northern Ireland Troubles in Britain* (pp. 227-242). Manchester University Press.
- [19] Julie Posetti, and Alice Matthews. 2018. "A short guide to the history of 'fake news' and disinformation." *International Center for Journalists* 7.
- [20] Nathan Pluquet. 2023. The Dangers of Emergency Powers and the Downfall of Democracy: British Security Response in Northern Ireland during the "Troubles".
- [21] Betty Purcell. 1991. The silence in Irish broadcasting. In *The Media and Northern Ireland: covering the troubles* (pp. 51-68). London: Palgrave Macmillan UK.
- [22] Andrew Rosenberg and Julia Hirschberg. 2009. Charisma perception from text and speech. *Speech communication*, 51(7), 640-655.
- [23] Miriam Sturdee, Joseph Lindley, Conor Linehan, Chris Elsdon, Neha Kumar, Tawanna Dillahun, Regan Mandryk, and John Vines. 2021. "Consequences, Schmonsequences! Considering the Future as Part of Publication and Peer Review in Computing Research." In *Extended Abstracts of the 2021 CHI Conference on Human Factors in Computing Systems*, pp. 1-4. 2021.
- [24] Konstantin O. Tskhay., Rebecca Zhu, Christopher Zou, and Nicholas O. Rule. 2018. Charisma in everyday life: Conceptualization and validation of the General Charisma Inventory. *Journal of personality and social psychology*, 114(1), 131.
- [25] John Twomey, Didier Ching, Matthew Peter Aylett, Michael Quayle, Conor Linehan, and Gillian Murphy. 2024. What Is So Deep About Deepfakes? A Multi-Disciplinary Thematic Analysis of Academic Narratives About Deepfake Technology. *IEEE Transactions on Technology and Society*.
- [26] John Twomey, Didier Ching, Matthew Peter Aylett, Michael Quayle, Conor Linehan, and Gillian Murphy. 2023. "Do deepfake videos undermine our epistemic trust? A thematic analysis of tweets that discuss deepfakes in the Russian invasion of Ukraine." *Plos one* 18, no. 10: e0291668.