

Data Paradigms in the Era of LLMs: On the Opportunities and Challenges of Qualitative Data in the WILD

Shengqi Zhu
Information Science
Cornell University
Ithaca, New York, USA
sz595@cornell.edu

Jeffrey M. Rzeszutarski
Information Science
Cornell University
Ithaca, New York, USA
jeffrz@cornell.edu

David Mimno
Information Science
Cornell University
Ithaca, New York, USA
mimno@cornell.edu

Abstract

This work considers current and future affordances for a prospective data paradigm for Human-AI Interaction studies: real-World Interactions as Large-scale Data sources (**WILD**). We envision how the records of natural interactions with LLMs, a common yet uncontrolled product of contemporary chatbots, may serve as a comprehensive, sustainable data source for qualitative user studies. We embed the opportunities in detailed discussions of the background and current form of such data, and enumerate the practical directions for improvement and fundamental challenges. Our work seeks to provide tangible starting points for the incorporation of the new data paradigm, and we call for more future work with the joint efforts of the HCI and NLP communities.

CCS Concepts

• **Human-centered computing** → **User studies**; **Natural language interfaces**; • **Computing methodologies** → **Language resources**.

Keywords

data, interaction, conversation, Large Language Models, user-LLM interaction, data paradigms

ACM Reference Format:

Shengqi Zhu, Jeffrey M. Rzeszutarski, and David Mimno. 2025. Data Paradigms in the Era of LLMs: On the Opportunities and Challenges of Qualitative Data in the WILD. In *Extended Abstracts of the CHI Conference on Human Factors in Computing Systems (CHI EA '25)*, April 26–May 01, 2025, Yokohama, Japan. ACM, New York, NY, USA, 8 pages. <https://doi.org/10.1145/3706599.3720285>

1 Introduction

Large Language Models (LLMs) have defined a ubiquitous mode of interaction [23, 26, 48, 52]. The transformational capabilities of LLMs created a wide range of heterogeneous applications and interaction modes, and inspired many works to probe the features, capabilities, and possibilities of these AI products [35, 36, 42, 44], evoke the perceptions of users [15, 25, 43, 45], and reveal the implications on broader groups [22, 23, 51, 60].

However, eliciting adequate, high-quality data and user feedback, which has long been essential to qualitative research [27], is especially complicated. This is essentially because the goals of model training are often in conflict with the needs of HCI experiments. LLMs are (pre-)trained with vast amounts of unstructured data, in order to provide a wide range of capabilities to a variety of users [26, 46, 47]. Even though human feedback at scale can then be used to further fine-tune these models to be more effective for specific tasks and interactions [31, 33, 56, 59], such data are generally still in the form of minimalist “gold labels” or “preferences” for standardized evaluations like accuracy.

On the other hand, conventional HAI studies pursue fine-grained, focused input and output, in order to distinguish individuals’ perceptions of interfaces [21, 54]. As these two paradigms collide, practical considerations emerge as key limitations of validity. The unprecedented *scale* leads to synchronous interactions spanning diverse languages and cultural contexts. It creates challenges for common qualitative toolkits such as interviews or self-paced cognitive records to accommodate the user base of billions with broad socio-cultural diversity. In practical data collection, this also puts forward the question of *representation*: it is challenging, if not impossible, to extrapolate findings from a handful of participants to the diverse and multifaceted ways people engage with AI, and more broadly, describe their views and lives. This is further complicated as LLMs are also redefining the *boundaries of functionality* with their versatility. Whereas user studies conventionally seek to create precise and well-grounded profiles of the target groups and use scenarios, LLMs as a jack-of-all-trades service cast uncertainty on the models and goals of user data gathering.

More importantly, this inconsistency reinforces the bifurcation of the two correlated discourses. On one hand, it is less clear how detailed qualitative insights can influence the iteration of model parameters; on the other hand, the ample data resources surrounding LLMs have not been able to benefit downstream user-centered research as much as how they assist parameter fine-tuning. As such, the back- and front-ends of LLMs increasingly fall into their respective closed loops. This is a risky situation; for instance, if one seeks to develop a downstream interface, it is necessary to obtain an accurate evaluation and understanding of the base model, such that any practical design choices correspond to real needs [17, 38]. To this end, we require sustainable paradigms of collecting and analyzing diverse, unpredictable user experiences and motivations at scale within an acceptable amount of effort and intervention.

In this work, we envisage the possibilities and remaining challenges of incorporating the real-World Interactions as Large-scale Data sources (**WILD**) in the qualitative HAI context. Specifically,

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

CHI EA '25, Yokohama, Japan

© 2025 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-1395-8/25/04

<https://doi.org/10.1145/3706599.3720285>




Figure 1: A brief overview of two respective “loops” on the “front-” and “back-end” of the same chatbot (and its interface) as core. In Natural Language Processing (left), real-world chat logs, together with other human feedback data, are utilized at scale to further fine-tune the model parameters. For qualitative user studies (right) with the current model as a version (or prototype) of the product, user feedback is carefully elicited for designing the next generation of applications. The scope of this work is generally represented by the grey arrow on top: Can we build a reliable connection between the two schemes at data level? How might the real-world interactions be properly handled to contribute to qualitative insights?

we explore whether and how unstructured chat logs, a common by-product of LLM deployment, can foster new data paradigms that enrich the scope of current data practices. The main goal of this work is to evaluate the current affordance from a comprehensive lens. We discuss three prospective directions of incorporating data in the WILD: as a giant pool of data instances in general, as materials for longitudinal studies, and as a reflexive form of data at the model level. More importantly, we draw attention to two major challenges before the new data paradigm can be applied: the balance between control and naturality (and lack thereof), and the ethical and consent considerations. In response, we provide initial thoughts on how they may be addressed. Overall, we believe that with joint efforts across related communities, this data paradigm can be prospective in the future data practices.

Positionality. This paper connects to recent directions in HCI and NLP utilizing user feedback for model fine-tuning and linguistic studies. In conducting this work, we seek to leverage our understanding of engineering practices for the data sources in discussion from a technical lens. We acknowledge that at this moment the intersection of humans and AI, especially with respect to qualitative research, is in tension. We aim to build bridges rather than assert preference over any particular research method, but we acknowledge that presenting a technical perspective in and of itself does promote a specific line of investigation which we hope will encourage community discussion. We hope to foster such systematic discourse regarding a supplementary path connecting user-LLM data at scale with qualitative studies, and thus encourage future work and cross-disciplinary data resource collaborations.

2 Background

2.1 Real-World Human-AI Interaction Data: What is data in this paper about?

LLM-based conversational agents are tailored towards completing user requests. Current practice starts with “base models” that are trained on gigantic amounts of unstructured text purely for text prediction, with no specific user interaction pattern. Base models are then fine-tuned (further trained) to support specific and fine-grained capabilities to follow input instructions, using pairs of task instructions and desired responses. This paradigm, termed “instruction tuning”, has been a key contributor to the success of prompt-based LLMs [31, 40, 56]. It also leads to the implementation of many basic yet important built-in feature chatbots (e.g., formatting an answer as a polite and effective response) as well as a constant need for reliable instruction and response data.

OpenAI’s ChatGPT [32] introduced LLM products to a mainstream audience, meeting a combination of enthusiasm and frustration from users. ShareGPT¹, a web extension that allows users to export the current chat log as a separate webpage for easy URL sharing, quickly rose to the attention of many as a supplement for socialization needs. While the tool emerged as a simple gadget, it soon became a de facto communal medium for a solid fan base. While the details of user consent in the process are somewhat unclear, its “Explore” function allows the public to browse user-selected conversations. The browsability of ShareGPT led quickly to shared datasets created by scraping chat logs, such as [1], and these early practices were soon brought to public attention. Model developers found they could use ShareGPT data to fine-tune new models

¹<https://sharegpt.com/>

with impressive chat performance, such as Vicuna [6], which was trained on 70k logs. Public discourses also surged as the creator of the tool claimed² that Google had been inappropriately scraping the data for their own LLMs. Although the original ShareGPT datasets are less frequently employed now due to limited quality and questionable provenance, a lay-user-based data paradigm it represents has demonstrated its potential. Importantly, if the deployment of language models can create useful data in return, and the scale of these grows in proportion to the user base, this creates a sustainable and self-contained data practice. In some sense, the user experience and use feedback are unitary in this loop, which were incorporated into the next generation update.

More recently, the field of NLP has seen more formal setups for gathering user-LLM chat logs while preserving minimal intervention. The most representative real-world user-LLM datasets are LMSYS-CHAT-1M [58] and WildChat [57], both at the scale of ~1M instances. The creation of both datasets follows an intuitive and voluntary scheme: interested participants give consent to share their conversation logs, in exchange for the free and/or unlimited language model services covered by reliable platforms like HuggingFace. With the more systematic setup, these task-specific interaction data were successfully amassed to the level of millions, in a natural time span of months to years. These data at the million scale not only serve as good input text for LLMs, but may also ripple in various other directions for fundamentally new paradigms and use cases.

The concept of WILD in this work originates from these handy resources. By WILD, we refer to data of the same format: the real and naturalistic user-LLM conversations in their original form, collected (with consent) at a large scale (at least the level of thousands). Note, however, that the vision here is not confined to what the existing data has to offer; In fact, we discuss in Section 4 why the existing data may *not* be the ideal input for a practical user-centered study. Our work highlights fundamental challenges in the protocol and design of this type of data resource, and we call to both the NLP and HCI communities for collaborative efforts towards a reliable data paradigm.

2.2 Related Work: LLMs as contributors in HCI and User Research

The impressive text generation abilities and ease of use of LLMs have inspired various related areas, especially ones with the largest demand for evaluation or text generation resources. For instance, a large body of work explores *LLMs as (crowd) workers* that seek reduced effort from crowdsourcing [18, 20, 49]. In the context most related to this paper, existing work has discussed the prospects and concerns to employ LLMs as pseudo users and generate synthetic survey questions and user research data [10, 16].

Another major application is in Deductive Coding/Thematic Analysis: As the heavy labor of text processing and the open-ended nature match well with the strengths of LLMs, various works proposed to prompt LLMs to code in place of humans [8, 50, 55]. More recently, there has been rising awareness of the more complicated dynamics and more ethical practices in reforming research. Instead of letting models take over certain onerous duties, automation is more often applied to enable better preparation, implementation,

and collaboration in the academic context. For instance, [34] sets up a Beginner-Friendly Human-AI Platform with a trained AI as assistant, and [14] uses LLMs to mediate Collaborative Qualitative Coding between human coders.

We'd also like to bring to attention the recent meta-study [41] interviewing the exact qualitative researchers in this field about LLMs in their workflow. The work finds the interviewees "largely open to or interested in the conscientious use of LLMs across the research pipeline"; meanwhile, they are more comfortable and familiar with the features and functionality pertaining to the "chatbot" side, compared with the norms and values at the core of qualitative analysis.

Our work is in line with the question of LLMs' roles: with the involvement of AI as a direct replacement potentially questionable in the current paradigms, what are some different ways to enhance HAI research? We discuss a new data-guided LLM-research(er) model here: LLMs do not involve with or contribute to the user study process per se; rather, we seek to start from the existing interaction data, harvesting the valuable information in the WILD and applying meaningful analyses with rigorous methods.

3 Opportunities of Real-World Interactions as Large-Scale Data Sources

In this section, we start with the new opportunities presented by large-scale, real-world user-AI interactions. We describe how WILD may serve as a systematic data resource for qualitative research in HCI, and exemplify prospective directions at finer grain: the case of longitudinal methodologies, and the models as reflexive data snapshots. These two schemes respectively feature the analysis of individual users and models supported by the time axis.

3.1 Documenting Real-World User-AI Interactions as Qualitative Databases

3.1.1 Scope. Real-world user interactions provide a naturally occurring ocean of data that captures authentic user behavior. With a potentially worldwide scope, WILD data captures the voices, contexts, and behaviors at scale across diverse demographic, cultural, and social backgrounds with little additional cost. The scale can enable researchers to uncover patterns that would be invisible in smaller datasets, e.g., the nuanced variations in usage trends (e.g. across geographic regions) essential for designing equitable AI systems well-tuned to the diverse realities of their user base. Leveraging WILD could also significantly reduce the costs and logistical challenges of data practices at similar scales, e.g., the substantial effort to recruit worldwide participants for an online survey.

3.1.2 Natural Experiments. Meanwhile, such data is inherently naturalistic, reflecting "in-the-wild" behavior without the artificial constraints or interventions imposed [5]. This mitigates the influence of observation, i.e., the presence of researchers and knowledge of the topic inadvertently shaping participant behaviors [28]. As WILD data is collected passively and unobtrusively, it can be well-suited for topics that might be sensitive or hard to address in a face-to-face setup with researchers. This can offer rich understanding of how users navigate complex, nuanced realities via conversations with an AI assistant.

²<https://x.com/steventey/status/1641267979399704576>

While observational data do not naturally map to experimental design, we note the similarities here with the ample threads of work via *natural experiments* [11, 39], a now-common methodology in Social and Medical Sciences. Unlike traditional lab experiments or structured observational studies with a specifically designed intervention, natural experiments leverage real-world events — such as policy changes, historical incidents, or natural disasters — to analyze (causal) effects and/or extend research to the general public. We highlight this discourse because the trajectory of natural experiments gaining prevalence as a research method provides referential insights for the discussion of the WILD data paradigm. This involves not only the proven strengths of such methods (enabling research at a large scale and in non-experimental real-world conditions) but more importantly the explorations for the key challenges [7, 9, 19] — for instance, whether and how a meaningful randomization can be captured in the uncontrolled public environment. Our point is that the WILD paradigm, in some sense an adaptation of the natural experiment notion in the Human-AI Interaction context, is to the very least not impossible; much can be observed from the study of human behaviors in neighboring fields where conventional research design are successfully transitioned to natural, unpredictable interventions at scale [24, 37].

3.1.3 Incorporation with existing paradigms. It is worth noting that WILD can also resonate with existing data paradigms. For instance, consider a hypothetical study on students’ interaction patterns with AI tools. Parallel with well-established methods like focus groups and interviews, we can also gather real-world interactions between students and LLMs. Much as the former reveal broad perceptions, the latter could enable more granular insights and traceable observations, e.g., the types of errors they repeatedly ask the system to correct, or how they craft their prompts. The direct access of use cases helps to remove the layer of ambiguity, dishonesty, and information loss in communication. A further exciting possibility is the meta-analysis of the data paradigms by comparing the data collected from different scenarios for comparative research. For example, researchers could juxtapose findings from traditional ethnographic studies with insights derived from large-scale interaction data to validate or challenge existing theories. Additionally, this can also enhance effectiveness when the issue in question is hard to describe or abstract by users but frequently be found as a pattern in natural interactions.

3.2 Longitudinal and Time-Lapse Research

Longitudinal research like diary studies and ethnographies offer invaluable and irreplaceable insights into how a target group’s behaviors, practices, and perceptions evolve over time; meanwhile, they are also most constrained by the practical challenges as discussed: limited sample sizes, high costs, and the difficulty of sustaining participant engagement. We propose to give attention to the WILD as a relevant component in understanding temporal dimensions, because of not only its scale and richness but also the unique time-lapse user-LLM conversational setup it offers.

A particularly relevant and natural approach for exploring WILD data is user-level modeling, i.e., tracking the behaviors and interactions of specific users over time. This closely aligns with established

methodologies, such as diary studies and narrative inquiry, by capturing the evolution of individual experiences. Meanwhile, WILD possesses a unique level of scalability and naturality. For a basic instance, researchers can observe how users adapt their interaction strategies with an LLM as they become more familiar with its capabilities, or how their prompts evolve as their needs and expectations shift. Meanwhile, there are relatively few additional logistical burdens of recruitment or management as the interactions proceed., enabling the documentation of insights on a topic parallelly, extending both depth and breadth in understanding longitudinal trends. Note also that this paradigm is not limited to the studies of LLMs per se; with the appropriate post-processing (e.g., retrieval and filter) and reasonable protocols (Section 4), WILD may serve many broader social contexts via user conversations.

Beyond individual user-level tracking, WILD also enables longitudinal analyses at a macro level, especially insights into sociocultural trends and collective behaviors in the AI context. For example, it might be interesting to explore how the collective attitudes toward AI-assisted communication change over time, particularly in response to key external factors and time stamps such as regulatory developments or technological advancements. Likewise, elicited remarks, debates, and (dis)agreements surrounding key topics (which can be either about LLMs themselves or existing societal discussions) will be an interesting source of public opinions and values, which has the potential to supplement data collected from social media or surveys.

3.3 Model as Data

So far we have discussed the prospective use cases of WILD as the documentation to better understand the behaviors and demands of *users*. In this point, we conceptualize a distinct data paradigm that considers *models* per se from a “data view”, in addition to the role as a product. Language models, as complex data artifacts trained on the real-world data, reflect the cultural, social, and linguistic patterns embedded in the datasets they are trained on. At the same time, they are also a natural-language interface for such gigantic collections of human activities. This inspires new paths to inquire societal narratives, biases, and norms, in the ways that they are embedded in and mediated by computational systems as “living data complexes” of human societal data and knowledge. Consider the example of an insufficient or incorrect output when a prompt is adapted to a different culture/language. While classic methods can inquire into the practical usability and performance, there can be more layers of strategies. For instance, the observed insufficiency and inaccuracy of the outputs might draw attention to systematic epistemic gaps and mismatches in the accessible field data, e.g., the sociocultural groups or narratives that are underrepresented in terms of a target scenario (as hence in the training data). Likewise, analyzing how biases manifest in LLM outputs could provide insights into underlying inequities within the source datasets and the deeper cultural formation, e.g., gender biases in professions, or cultural stereotypes in long-established idiomatic expressions.

Adding the time dimension to this direction similarly, LLMs can be positioned as “time capsules” of collective memory, which preserve how the world was represented in public discourse, media, and literature at specific times. As language models are essentially

a snapshot of collective human knowledge at the time of training, researchers can compare outputs from models that synthesize data from different periods, in order to examine how societal values, norms, and language use evolve. For example, the ways models discuss complex topics such as climate change, gender roles, or political polarization can be viewed as a certain type and source of data. In this case, it is arguably at least as significant to systematically understand how societal trends influence computational knowledge and their presentation, as compared to how this may be described as an unwanted quirk of the model itself, or as a training objective where the models are taught not to comment on the topic.

4 Key Remaining Challenges for Practical Data Work

4.1 Control vs. Naturality

The major challenge for incorporating any WILD paradigms into HCI practices is the intrinsic, subtle balance between *naturality* and *control*. As discussed, real-world interactions offers a novel path, capturing authentic interactions at scale as they occur naturally in the world without the artificial constraints. Nonetheless, this very naturality introduces an evident and indispensable trade-off. Much as these interactions can reflect genuine and varied user needs, they could be heterogeneous for the same reasons. In the absence of well-structured setups and predefined tasks, WILD data might more often be found to be noisy, unpredictable, and difficult to analyze (or even see any entrance for analysis).

Much like the collision of existing data paradigms, there exists an intrinsic contrast between question-oriented research schemas (top-down) and documentation of the real world (bottom-up). In conventional qualitative studies, researchers define the tasks, scope, and participant pool, and then collect data to implement the experimental design in well-constrained conditions. Conversely, current practices to obtain data in the WILD involve a vast, uncontrolled user base interacting with LLMs in diverse and unpredictable contexts. Moreover, users' goals and contexts in the WILD are often implicit or ambiguous. What a dialog is about and whether it's relevant might need much context reading to determine, or even remain unclear after the whole chat, as users do not need to describe everything they were thinking. Therefore, the key question to be answered is *how to connect these two worlds*, either by introducing more/less desired variations in the corresponding end(s), or by applying suitable data processing and analysis techniques for the scenario. Here, we briefly introduce several prospective directions, with strategies spanning different parts of a study:

Targeted Interventions and Semi-structured protocols. The problem with existing datasets is that they don't have any actual design or constraints within the data collection. Mediating and combining the paradigm with a well-planned research goal will drastically increase its value and feasibility. For instance, if we seek to study to co-creation of a story, it is possible to add a requirement when providing free access in exchange for data (with everything else not impacted): users will need to complete a writing task with AI at the end of the contract. Another simple yet effective protocol is to constrain access accordingly, e.g., providing language model service for college students on educational usage (instead of anyone), if

that is the research topic. With these examples, we intend to show that WILD *can* be controlled to serve as a self-contained source of user data, but rather it's limited by its current implementations and we call for attempts with more rigorous interventions.

Hybrid data schemes. While the open-ended nature of WILD creates challenges, we highlight the possible application of *hybrid* data paradigms, i.e., strategically combining large-scale, uncontrolled datasets with smaller, well-curated subsets. Starting from a few high-quality data instances as anchors, larger datasets can be effectively "calibrated" based on the attributes of interest, e.g., by retrieving conversations with matching keywords, topics, or time of conversation. Another notable perspective is how such resources can be used to enrich available data in small batches. Real-world datasets mediate between human-centered user studies and data bootstrapping techniques: they provide the middle ground where the candidate pool of additional data indeed rises from real user demands. On the other hand, larger datasets with potentially more variations and branches can serve as a prior for more precise data collection. Researchers could use WILD data to identify broad behavioral patterns and trends, and then design fine-grained experiments accordingly with conventional qualitative methods to validate or refine these findings in more rigorous setups.

Making sense of real-world data. Aside from improving the source quality, it is also possible to benefit from proper downstream processing, such that resources like WILD can be effectively analyzed and filtered. The post-processing can be further paired with well-established means of automation — for instance, a clustering [2] or topic modeling [3, 53] module may be constructive for narrowing down to a focused data subset of interest. Like most other data sources, WILD can also benefit from corpus annotation, e.g., crowdsourcing fine-grained data for intentions or dialog state [4]. If appropriate in the context, more advanced language technologies may also be a major contributor, e.g., training a deep-learning based topic classifier, or fine-tuning a retrieval tool to extract the most relevant instances.

4.2 Ethical Considerations

Next, we identify the pressing challenges to ethically implementing the proposed data paradigms. We review the practices from existing datasets, discuss their limitations, and propose our suggestions on the adaptation of current NLP artifacts.

4.2.1 Consent and User Rights. Major existing resources like LM-Sys and WildChat are based upon the consent from self sign-up users to share their chat logs and metadata. In exchange, the research group provides free access to various latest LLMs. This is implemented as a Terms of Agreement notice that users must confirm before accessing the prompting interface. Nonetheless, these user consent practices still require significant improvement and clarification. The terms of ShareGPT and LMSys are especially limited and do not specifically tackle the challenges of user conversations. WildChat, as a more recent artifact, includes a more standardized set of terms that explicitly indicates the collection and possible use

of chat logs and meta-data³. Nonetheless, the ambiguous and overarching language still forms obstacles to a constructive consensus of consent. Users deserve a necessary level of transparency, for instance, which entities will have access to the data in “sharing with third parties” or “reporting purposes”, and why these third parties are trusted. More broadly, the practice of covering the key elements with underspecified umbrella terms essentially deprives users of their rights of choice. For instance, one might make different decisions on whether to participate depending on where the data will appear and be used. Besides, more rigorous terms shall be set in terms of data ownership or whether users can request access to or deletion of their data.

The earlier attempts represented by ShareGPT follow a different scheme, where users choose on their own which of their conversation threads might be shared. This is a viable approach with foreseeable positive use cases when adapted to a research schema, e.g., for a study where participants are asked to write things down when and only when under a specific condition. However, the original ShareGPT platform lacks the fundamental constraints of public access. Participants, especially the earlier users when the community hasn’t formed, likely possess the reasonable expectation that the major reason they participated is for personal purposes, not for public promotion of the tool. It is necessary to incorporate more systematic control over who can access the data in what way, which, as the community has witnessed, created a large chaos for ShareGPT. Yet, this protocol indeed forms a more controlled possibility of WILD usage, and we encourage future work to explore more rigorous implementations.

4.2.2 The conversation content and post-processing: where are the boundaries? The content of less-controlled user-LLM interaction needs to be handled in an especially rigorous manner. While the nature of one-to-one human-computer conversations is bound with a reasonable expectation of privacy, these borders are blurred as they form publicly available datasets. While datasets like WildChat apply desensitizing strategies to “fuzz out” the logs, these are not well informed to users, and the combination of metadata might still reveal more information than expected. Even when the users agree to the terms and conditions, practical issues can occur in the publication of the data. In fact, recent studies report an astonishing amount of Personal Identifying Information in WildChat [30]. These concerns of information leakage are further complicated by the unclear user intent. For instance, it is possible that a participant maliciously or unconsciously put *another* individual’s personal information in the chat logs. At a high level, these facts are a major thought for the promotion of the qualitative, user-centered usage of these data: We believe that it may be risky to directly feed these user logs to models as the training objective, and mix-method explorations with more careful interpretations are required.

Even when the data practices fully abide by the user-agreed terms and conditions, additional considerations should apply for topics that are intrinsically risky for the data paradigm. For instance, AI consulting records for issues like mental health from the real world can largely benefit related studies, but it requires careful ethical examination, e.g., confirming that users had the conversation with the full awareness that the data will be collected, and that the

data piece is entirely desensitized when adding it to the dataset. This is also similar to the natural experiments method in terms of the common scenarios that may “exploit” the vulnerabilities of relevant groups, e.g., in the post-disaster setup where a tragic incident is marked as a natural source of intervention [12, 13, 19, 29]. To mediate between the goals of research and ethical practice in the HAI context, we underscore again the importance of a more complete data collection protocol design, especially in terms of users’ agency and right of choice at the level of each dialog. For instance, instead of automatically documenting all conversations, researchers can give out subtle reminders, via the extra step of a confirmation inquiry or a self-paced review (e.g. let users create diaries of their usage).

5 Conclusions

This short opinion piece briefly walks readers through the 5W1H of the unstructured user-LLM interaction from the by-product of language technology as a data resource. We show that, while the idea of making use of user data out there may seem intuitive, the implementation is non-trivial. Yet, we highlight that the reward of this supplementary data source is especially meaningful as our research and daily life are impacted via these interactions. We hope our work can be a prequel to more systematic work on the data paradigm for user studies in the contemporary context, including but not limited to our WILD imaginations here in this work.

References

- [1] Ryoko AI. 2023. ShareGPT52K. <https://huggingface.co/datasets/RyokoAI/ShareGPT52K>. <https://huggingface.co/datasets/RyokoAI/ShareGPT52K>
- [2] Nicholas O Andrews and Edward A Fox. 2007. Recent developments in document clustering. (2007).
- [3] Jordan Boyd-Graber, Yuening Hu, David Mimno, et al. 2017. Applications of topic models. *Foundations and Trends® in Information Retrieval* 11, 2-3 (2017), 143–296.
- [4] Paweł Budzianowski, Tsung-Hsien Wen, Bo-Hsiang Tseng, Inigo Casanueva, Stefan Ultes, Osman Ramadan, and Milica Gasic. 2018. MultiWOZ-A Large-Scale Multi-Domain Wizard-of-Oz Dataset for Task-Oriented Dialogue Modelling. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*. 5016–5026.
- [5] James W Chesebro and Deborah J Borisoff. 2007. What makes qualitative research qualitative? *Qualitative research reports in communication* 8, 1 (2007), 3–14.
- [6] Wei-Lin Chiang, Zhuohan Li, Zi Lin, Ying Sheng, Zhanghao Wu, Hao Zhang, Lianmin Zheng, Siyuan Zhuang, Yonghao Zhuang, Joseph E. Gonzalez, Ion Stoica, and Eric P. Xing. 2023. Vicuna: An Open-Source Chatbot Impressing GPT-4 with 90%* ChatGPT Quality. <https://lmsys.org/blog/2023-03-30-vicuna/>
- [7] Peter Craig, Cyrus Cooper, David Gunnell, Sally Haw, Kenny Lawson, Sally Macintyre, David Ogilvie, Mark Petticrew, Barney Reeves, Matt Sutton, et al. 2011. Using natural experiments to evaluate population health interventions: guidance for producers and users of evidence. *Medical Research Council* (2011).
- [8] Shih-Chieh Dai, Aiping Xiong, and Lun-Wei Ku. 2023. LLM-in-the-loop: Leveraging Large Language Model for Thematic Analysis. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, Houda Bouamor, Juan Pino, and Kalika Bali (Eds.). Association for Computational Linguistics, Singapore, 9993–10001. doi:10.18653/v1/2023.findings-emnlp.669
- [9] Angus Dawson and Julius Sim. 2015. The nature and ethics of natural experiments. *Journal of Medical Ethics* 41, 10 (2015), 848–853.
- [10] Peitong Duan, Jeremy Warner, Yang Li, and Bjoern Hartmann. 2024. Generating Automatic Feedback on UI Mockups with Large Language Models. In *Proceedings of the CHI Conference on Human Factors in Computing Systems*. 1–20.
- [11] Thad Dunning. 2012. *Natural experiments in the social sciences: A design-based approach*. Cambridge University Press.
- [12] Thomas C Erren, Philip Lewis, and David M Shaw. 2020. The COVID-19 pandemic: ethical and scientific imperatives for “natural” experiments. *Circulation* 142, 4 (2020), 309–311.
- [13] Leslie P Francis. 2021. Understanding the ethics of natural experiments in a pandemic. *American Journal of Public Health* 111, 2 (2021), 212–214.

³<https://huggingface.co/spaces/yuntian-deng/ChatGPT4>

- [14] Jie Gao, Yuchen Guo, Gionnieve Lim, Tianqin Zhang, Zheng Zhang, Toby Jia-Jun Li, and Simon Tangi Perrault. 2024. CollabCoder: a lower-barrier, rigorous workflow for inductive collaborative qualitative analysis with large language models. In *Proceedings of the CHI Conference on Human Factors in Computing Systems*. 1–29.
- [15] Xiao Ge, Chunchen Xu, Daigo Misaki, Hazel Rose Markus, and Jeanne L Tsai. 2024. How Culture Shapes What People Want From AI. In *Proceedings of the CHI Conference on Human Factors in Computing Systems*. 1–15.
- [16] Perttu Hämäläinen, Mikke Tavast, and Anton Kunnari. 2023. Evaluating large language models in generating synthetic HCI research data: a case study. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*. 1–19.
- [17] AKM Bahalul Haque, AKM Najmul Islam, and Patrick Mikalef. 2023. Explainable Artificial Intelligence (XAI) from a user perspective: A synthesis of prior literature and problematizing avenues for future research. *Technological Forecasting and Social Change* 186 (2023), 122120.
- [18] Zeyu He, Chieh-Yang Huang, Chien-Kuang Cornelia Ding, Shaurya Rohatgi, and Ting-Hao Kenneth Huang. 2024. If in a Crowdsourced Data Annotation Pipeline, a GPT-4. In *Proceedings of the CHI Conference on Human Factors in Computing Systems*. 1–25.
- [19] Clare Herrick and Kirsten Bell. 2024. The social life of natural experiments in epidemiology and public health. *Sociology of Health & Illness* 46, 2 (2024), 276–294.
- [20] Xinyu Hu, Mingqi Gao, Sen Hu, Yang Zhang, Yicheng Chen, Teng Xu, and Xiaojun Wan. 2024. Are LLM-based Evaluators Confusing NLG Quality Criteria?. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Lun-Wei Ku, Andre Martins, and Vivek Srikumar (Eds.). Association for Computational Linguistics, Bangkok, Thailand, 9530–9570. doi:10.18653/v1/2024.acl-long.516
- [21] Felicity Agwu Kalu and Jack C Bwalya. 2017. What makes qualitative research good research? An exploratory analysis of critical elements. *International Journal of Social Science Research* 5, 2 (2017), 43–56.
- [22] Enkelejd Kasneci, Kathrin Seßler, Stefan Küchemann, Maria Bannert, Daryna Dementieva, Frank Fischer, Urs Gasser, Georg Groh, Stephan Günemann, Eyke Hüllermeier, et al. 2023. ChatGPT for good? On opportunities and challenges of large language models for education. *Learning and individual differences* 103 (2023), 102274.
- [23] Callie Y Kim, Christine P Lee, and Bilge Mutlu. 2024. Understanding large-language model (llm)-powered human-robot interaction. In *Proceedings of the 2024 ACM/IEEE International Conference on Human-Robot Interaction*. 371–380.
- [24] Scott T Leatherdale. 2019. Natural experiment methodology for research: a review of how different methods can support real-world research. *International Journal of Social Research Methodology* 22, 1 (2019), 19–35.
- [25] Michael Xieyang Liu, Advait Sarkar, Carina Negreanu, Benjamin Zorn, Jack Williams, Neil Toronto, and Andrew D Gordon. 2023. “What it wants me to say”: Bridging the abstraction gap between end-user programmers and code-generating large language models. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*. 1–31.
- [26] Pengfei Liu, Weizhe Yuan, Jinlan Fu, Zhengbao Jiang, Hiroaki Hayashi, and Graham Neubig. 2023. Pre-train, prompt, and predict: A systematic survey of prompting methods in natural language processing. *Comput. Surveys* 55, 9 (2023), 1–35.
- [27] Nicholas Mays and Catherine Pope. 2000. Assessing quality in qualitative research. *Bmj* 320, 7226 (2000), 50–52.
- [28] Jim McCambridge, John Witton, and Diana R Elbourne. 2014. Systematic review of the Hawthorne effect: new concepts are needed to study research participation effects. *Journal of clinical epidemiology* 67, 3 (2014), 267–277.
- [29] Signe Mezinska, Péter Kakuk, Goran Mijaljica, Marcin Waligóra, and Dónal P O’Mathúna. 2016. Research in disaster settings: a systematic qualitative review of ethical guidelines. *BMC medical ethics* 17 (2016), 1–11.
- [30] Niloofer Mirehshgallah, Maria Antoniak, Yash More, Yejin Choi, and Golnoosh Farnadi. 2024. Trust No Bot: Discovering Personal Disclosures in Human-LLM Conversations in the Wild. In *First Conference on Language Modeling*.
- [31] Swaroop Mishra, Daniel Khashabi, Chitta Baral, and Hannaneh Hajishirzi. 2022. Cross-Task Generalization via Natural Language Crowdsourcing Instructions. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 3470–3487.
- [32] OpenAI. 2022. Introducing ChatGPT. <https://openai.com/index/chatgpt/>
- [33] Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. 2022. Training language models to follow instructions with human feedback. *Advances in neural information processing systems* 35 (2022), 27730–27744.
- [34] Cassandra Overney, Belén Saldías, Dimitra Dimitrakopoulou, and Deb Roy. 2024. SenseMate: An Accessible and Beginner-Friendly Human-AI Platform for Qualitative Data Analysis. In *Proceedings of the 29th International Conference on Intelligent User Interfaces*. 922–939.
- [35] Akhil Padmanabha, Jessie Yuan, Janavi Gupta, Zulekha Karachiwalla, Carmel Majidi, Henny Admoni, and Zackory Erickson. 2024. Voicepilot: Harnessing LLMs as speech interfaces for physically assistive robots. In *Proceedings of the 37th Annual ACM Symposium on User Interface Software and Technology*. 1–18.
- [36] Hyanghee Park and Daehwan Ahn. 2024. The Promise and Peril of ChatGPT in Higher Education: Opportunities, Challenges, and Design Implications. In *Proceedings of the CHI Conference on Human Factors in Computing Systems*. 1–21.
- [37] Mark Petticrew, Steven Cummins, Catherine Ferrell, Anne Findlay, Cassie Higgins, Caroline Hoy, Adrian Kearns, and Leigh Sparks. 2005. Natural experiments: an underused tool for public health? *Public health* 119, 9 (2005), 751–757.
- [38] Mark O Riedl. 2019. Human-centered artificial intelligence and machine learning. *Human behavior and emerging technologies* 1, 1 (2019), 33–36.
- [39] Mark R Rosenzweig and Kenneth I Wolpin. 2000. Natural “natural experiments” in economics. *Journal of Economic Literature* 38, 4 (2000), 827–874.
- [40] Victor Sanh, Albert Webson, Colin Raffel, Stephen Bach, Lintang Sutawika, Zaid Alyafeai, Antoine Chaffin, Arnaud Stiegler, Arun Raja, Manan Dey, et al. 2022. Multitask Prompted Training Enables Zero-Shot Task Generalization. In *International Conference on Learning Representations*.
- [41] Hope Schroeder, Marianne Aubin Le Quéré, Casey Randazzo, David Mimno, and Sarita Schoenebeck. 2024. Large language models in qualitative research: Can we do the data justice? *arXiv preprint arXiv:2410.07362* (2024).
- [42] Orit Shaer, Angelora Cooper, Osnat Mokryn, Andrew L Kun, and Hagit Ben Shoshan. 2024. AI-Augmented Brainwriting: Investigating the use of LLMs in group ideation. In *Proceedings of the CHI Conference on Human Factors in Computing Systems*. 1–17.
- [43] Shreya Shankar, JD Zamfirescu-Pereira, Björn Hartmann, Aditya Parameswaran, and Ian Arawjo. 2024. Who validates the validators? aligning llm-assisted evaluation of llm outputs with human preferences. In *Proceedings of the 37th Annual ACM Symposium on User Interface Software and Technology*. 1–14.
- [44] Donghoon Shin, Lucy Lu Wang, and Gary Hsieh. 2024. From Paper to Card: Transforming Design Implications with Generative AI. In *Proceedings of the CHI Conference on Human Factors in Computing Systems*. 1–15.
- [45] Chandan Kumar Tiwari, Mohd Abass Bhat, Shagufta Tariq Khan, Rajaswaminathan Subramaniam, and Mohammad Atif Irshad Khan. 2024. What drives students toward ChatGPT? An investigation of the factors influencing adoption and usage of ChatGPT. *Interactive Technology and Smart Education* 21, 3 (2024), 333–355.
- [46] Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. 2023. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971* (2023).
- [47] Yizhong Wang, Swaroop Mishra, Pegah Alipoormolabashi, Yeganeh Kordi, Amirreza Mirzaei, Atharva Naik, Arjun Ashok, Arut Selvan Dhanasekaran, Anjana Arunkumar, David Stap, Eshaan Pathak, Giannis Karamanolakis, Haizhi Lai, Ishan Purohit, Ishani Mondal, Jacob Anderson, Kirby Kuznia, Krma Doshi, Kuntal Kumar Pal, Maitreya Patel, Mehrad Moradshahi, Mihir Parmar, Mirali Purohit, Neeraj Varshney, Phani Rohitha Kaza, Pulkit Verma, Ravsehaj Singh Puri, Rushang Karia, Savan Doshi, Shailaja Keyur Sampat, Siddhartha Mishra, Sujay Reddy A, Sumanta Patro, Tanay Dixit, and Xudong Shen. 2022. Super-NaturalInstructions: Generalization via Declarative Instructions on 1600+ NLP Tasks. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, Yoav Goldberg, Zornitsa Kozareva, and Yue Zhang (Eds.). Association for Computational Linguistics, Abu Dhabi, United Arab Emirates, 5085–5109. doi:10.18653/v1/2022.emnlp-main.340
- [48] Tongshuang Wu, Michael Terry, and Carrie Jun Cai. 2022. Ai chains: Transparent and controllable human-ai interaction by chaining large language model prompts. In *Proceedings of the 2022 CHI conference on human factors in computing systems*. 1–22.
- [49] Tongshuang Wu, Haiyi Zhu, Maya Albayrak, Alexis Axon, Amanda Bertsch, Wenxing Deng, Ziqi Ding, Bill Guo, Sireesh Gururaja, Tzu-Sheng Kuo, et al. 2023. LLMs as workers in human-computational algorithms? replicating crowdsourcing pipelines with llms. *arXiv preprint arXiv:2307.10168* (2023).
- [50] Ziang Xiao, Xingdi Yuan, Q Vera Liao, Rania Abdelghani, and Pierre-Yves Oudeyer. 2023. Supporting qualitative analysis with large language models: Combining codebook with GPT-3 for deductive coding. In *Companion proceedings of the 28th international conference on intelligent user interfaces*. 75–78.
- [51] Lixiang Yan, Vanessa Echeverria, Gloria Milena Fernandez-Nieto, Yueqiao Jin, Zachari Swiecki, Linxuan Zhao, Dragan Gašević, and Roberto Martínez-Maldonado. 2024. Human-AI Collaboration in Thematic Analysis using ChatGPT: A User Study and Design Recommendations. In *Extended Abstracts of the CHI Conference on Human Factors in Computing Systems*. 1–7.
- [52] Diyi Yang. 2024. Human-AI Interaction in the Age of Large Language Models. In *Proceedings of the AAAI Symposium Series*, Vol. 3. 66–67.
- [53] Limin Yao, David Mimno, and Andrew McCallum. 2009. Efficient methods for topic model inference on streaming document collections. In *Proceedings of the 15th ACM SIGKDD international conference on Knowledge discovery and data mining*. 937–946.
- [54] Nur Yildirim, Mahima Pushkarna, Nitesh Goyal, Martin Wattenberg, and Fernanda Viégas. 2023. Investigating how practitioners use human-ai guidelines: A case study on the people+ ai guidebook. In *Proceedings of the 2023 CHI Conference*

- on *Human Factors in Computing Systems*. 1–13.
- [55] He Zhang, Chuhao Wu, Jingyi Xie, ChanMin Kim, and John M Carroll. 2023. QualiGPT: GPT as an easy-to-use tool for qualitative coding. *arXiv preprint arXiv:2310.07061* (2023).
- [56] Shengyu Zhang, Linfeng Dong, Xiaoya Li, Sen Zhang, Xiaofei Sun, Shuhe Wang, Jiwei Li, Runyi Hu, Tianwei Zhang, Fei Wu, et al. 2023. Instruction tuning for large language models: A survey. *arXiv preprint arXiv:2308.10792* (2023).
- [57] Wenting Zhao, Xiang Ren, Jack Hessel, Claire Cardie, Yejin Choi, and Yuntian Deng. 2024. WildChat: 1M ChatGPT Interaction Logs in the Wild. In *The Twelfth International Conference on Learning Representations*.
- [58] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Tianle Li, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zhuohan Li, Zi Lin, Eric Xing, et al. 2024. LMSYS-Chat-1M: A Large-Scale Real-World LLM Conversation Dataset. In *The Twelfth International Conference on Learning Representations*.
- [59] Daniel M Ziegler, Nisan Stiennon, Jeffrey Wu, Tom B Brown, Alec Radford, Dario Amodei, Paul Christiano, and Geoffrey Irving. 2019. Fine-tuning language models from human preferences. *arXiv preprint arXiv:1909.08593* (2019).
- [60] Caleb Ziems, William Held, Omar Shaikh, Jiaao Chen, Zhehao Zhang, and Diyi Yang. 2024. Can large language models transform computational social science? *Computational Linguistics* 50, 1 (2024), 237–291.