

Mitigating Toxicity in Social Media: Redesign Guidelines for Cultivating Positive User Interactions in the Instagram Threads App

Hamza Ahmad
Chisel Lab
Lahore University of Management
Sciences
Lahore, Pakistan
hamzaahmad277@gmail.com

Noverah Khan
Lahore University of Management
Sciences
Lahore, Punjab, Pakistan
noverah.khan@lums.edu.pk

Suleman Shahid
Computer Science Department
Lahore University of Management
Sciences
Lahore, Punjab, Pakistan
suleman.shahid@lums.edu.pk

Abstract

This paper explores the systemic challenges contributing to toxicity in social media platforms, focusing on the case study of Threads, a direct competitor of X (formerly known as Twitter). With its launch in July 2023, Threads gained popularity among X users who cited concerns surrounding increased toxicity on the platform. In this paper, we propose interventions after thoroughly reviewing catalysts such as the spread of misinformation, cyberbullying, and problems in online governance. The nascent platform, Threads, is not yet afflicted by these pitfalls but could be in the future. Our interventions include a collaborative fact-checking mechanism, AI bots for responding to violations against community guidelines, a mood selector board to curate content, and the option of anonymous replies. Through usability testing, we aimed to prove the effectiveness of these measures in reducing toxicity in online communities.

CCS Concepts

• **Human-centered computing** → Collaborative and social computing; Collaborative and social computing systems and tools; Social networking sites; Human computer interaction (HCI); Interaction design.

Keywords

Online Communities, Toxicity, Threads: An App by Instagram

ACM Reference Format:

Hamza Ahmad, Noverah Khan, and Suleman Shahid. 2025. Mitigating Toxicity in Social Media: Redesign Guidelines for Cultivating Positive User Interactions in the Instagram Threads App. In *Extended Abstracts of the CHI Conference on Human Factors in Computing Systems (CHI EA '25)*, April 26–May 01, 2025, Yokohama, Japan. ACM, New York, NY, USA, 7 pages. <https://doi.org/10.1145/3706599.3719924>

1 INTRODUCTION

The digital age has been marked by the rise of online communities that foster global connections and offer spaces for niche interests.

The internet has provided users unparalleled opportunities to interact, communicate, and exchange ideas, from subreddits to social media platforms like Facebook and X (formerly Twitter). Yet, as with any vast territory, there are darker corners where negativity can fester. Meta's Threads, launched in July 2023 as a direct competitor to X, has rapidly emerged as a significant platform, gaining over 320 million monthly active users by January 2025 [22]. This growth, coupled with its appeal to users fleeing X due to rising toxicity—such as hate speech and polarization [17, 23]—positions Threads at a critical juncture. Unlike X, which has struggled with content moderation and governance since its transformation under Elon Musk [24], or Reddit, where community-specific norms can both foster and challenge healthy interactions [15], Threads remains in its formative years. This infancy offers a unique opportunity to proactively design a platform that mitigates toxicity before it takes root.

Online communities have burgeoned as pivotal spaces for discourse, entertainment, and social networking. However, maintaining decorum and positive interactions within these platforms has been a challenge [19], given the vast array of user backgrounds, beliefs, and intentions. The intrinsic nature of online anonymity and detachment can sometimes lead to harmful behaviors like trolling, cyberbullying, and hate speech [19]. X, for instance, has been widely criticized for its toxic environment [20], with studies linking its use to increased outrage and diminished well-being [25], prompting user migration to alternatives like Threads [26]. Similarly, Reddit's reliance on community-driven moderation reveals both strengths and vulnerabilities, as rapid influxes of new users can disrupt established norms [15]. Threads, while built on Instagram's community guidelines, inherits a design akin to X, suggesting it could wield comparable influence over public perceptions and opinions in the future [25]. Focusing on its design now, at just over two years old, is thus imperative to harness its potential as a leader in positive digital interaction rather than a successor to X's pitfalls.

The redesign of Threads provides a unique chance to reassess the structure and functionalities of the platform. By studying user needs, behaviors, and feedback, and contrasting them with the well-documented pain points of X [17] and Reddit [15], Threads can leverage its early stage to implement innovative features—such as collaborative fact-checking and AI-driven moderation—to foster inclusivity and safety. This proactive approach is essential not only to meet the expectations of its growing user base but also to set a precedent for how social media platforms can evolve to prioritize

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

CHI EA '25, Yokohama, Japan

© 2025 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-1395-8/2025/04

<https://doi.org/10.1145/3706599.3719924>

positive engagement over unchecked toxicity [2]. Threads currently faces challenges in various domains:

- Navigational difficulties like having to switch to the Instagram app to view privacy settings or make profile changes.
- Lack of clarity and prominence of terms and conditions during user onboarding can lead to uninformed users inadvertently breaching platform guidelines.
- Absence of community-specific guidelines, relying heavily on Instagram’s overarching regulations.
- Lack of sensitive content warnings and tools to hide potentially offensive material, leaving users vulnerable to unsolicited harmful content.

The core problem lies in ensuring that Threads is both user-friendly and a safe environment. It is imperative to strike a balance between user autonomy and protective regulations. Addressing these challenges can transform Threads from being just another social media platform to a benchmark in digital community interaction.

2 RELATED WORKS

While online platforms, like Threads, offer users the opportunity to network globally, there are systemic challenges embedded into these platforms that compromise the quality of experience for the users. Though their popularity only continues to increase, having become an everyday essential for billions of users already, the governance of these platforms for their given scale of activities is insufficient [1]. Taylor Owen [2] in his comprehensive report about the need for platform governance discusses the potential outcomes of the phenomenon coined as “techlash”. Techlash refers to a large-scale backlash towards the tech industry and its products, primarily fueled by concerns from consumers [2, 18]. In order to create healthier online communities, the approach to avoiding techlash is through a holistic and global response. Atkinson et al. (2023) advocate for “tech-realism”—the acceptance of advancements in technology as the driving force for innovation while devising smart solutions to combat the problems they introduce [18].

Online communities tend to develop patterns of interactions in the form of direct and indirect reciprocity that remain consistent over time [3]. Considering that these patterns form the nature of the environment that is created in these digital platforms, they should be taken into account when implementing positive changes in these platforms. This is further elevated by the power law distributions observed in online communities through a study suggesting that there is a combination of preferential attachment along with direct and indirect reciprocity that explains the social mechanisms behind online communities [4]. Further implying that the dynamics that form in online communities are reflective of real-life social patterns and behaviors.

Online platforms consist of all kinds of users, the ones who actively participate and engage with the content and the ones who passively scroll. These users align themselves to the values of the online communities they are engaged in, and people who do not adhere to the core values of the community are generally seen as an out-group, and are, among others, targets of internet bullying [5]. Users with similar communities also tend to share information primarily with other users from the same communities [6]. This

gives online communities a more exclusive environment, despite the openness of the internet.

However, as the platforms continue to expand, concerns about how they are moderated are raised. Langdon Winner [7] had opened the discussion whether tech is going to have a utopic or dystopic future; using the term “cyberlibertarianism”, which refers to a system where you can be both digital and free. While such spaces can give you freedom of expression, moderation is vital in ensuring misinformation or malice is not spread through these platforms. With the growth of Artificial Intelligence, the idea of using AI models to improve online communities is gaining attention. However, while implementing moderation in online platforms can ensure an improved quality of content [8], using AI models for moderation can result in a lot of false positives because their technology is not advanced enough to understand the nuances of content [9]. Only a few, larger platforms can position themselves to invest in operations that can ensure the users a secure interaction in these communities, using data analytics and AI [10]. This means that many unmoderated platforms could be vulnerable to the spread of misinformation, malice, violations of community guidelines, and overall create a hostile or unsafe environment for users.

Based on our literature review and the challenges outlined, this study aims to answer the following research questions:

RQ1: What are the key challenges and dynamics of online community interaction on platforms like Threads, and how do these dynamics influence user behavior and content quality?

RQ2: How effective are current moderation techniques, including AI and human moderation, in managing toxic behaviors like misinformation, cyberbullying and hate speech on platforms like Threads, and what improvements can be made to these systems?

RQ3: How do community guidelines and their enforcement impact users’ perceptions of platform safety and their willingness to engage in online communities, particularly on platforms like Threads?

We now proceed to analyze the platform and community dynamics of Threads’ foremost competitor, X.

2.1 Competitor Analysis: X (formerly Twitter)

Threads, launched in July 2023 as a direct competitor to X, shares a text-based, real-time interaction design, making X a critical case study for understanding toxicity and informing Threads’ redesign. Threads’ design—featuring short-form posts, public feeds, and rapid reply mechanics—mirrors X’s structure, positioning it as a benchmark for both potential influence and pitfalls like misinformation and hate speech [17]. Analyzing X is thus essential to address our research questions on community dynamics (RQ1), moderation effectiveness (RQ2), and guideline impact (RQ3), while guiding design solutions to mitigate toxicity preemptively, such as AI moderation and collaborative fact-checking. To illustrate these challenges and inform Threads’ approach, we examine notable instances of X’s struggles with toxicity and moderation efforts.

X’s history reveals significant toxicity challenges, exemplified by the 2014 Gamergate controversy, marked by harassment, misogyny, hate speech, and cyberbullying directed primarily at feminist

women in the gaming industry [11]. Surprisingly, Gamergate supporters had considerable social connections and often met in person, challenging the notion that toxicity stems solely from isolated, antisocial behavior. X suspends accounts for spam, compromised security, or abusive content that violates platform rules [12]. Research revealed that Gamergate supporters faced higher suspension rates than baseline users, likely due to their heightened popularity and activity [13]. In an effort to combat harassment, X introduced a safe search feature, which restricts potentially harmful tweets in search results but keeps them accessible if users deliberately seek them. Enforcement operates at three levels—Tweet, Direct Message, and Account—and users can block, mute, or report accounts. This case underscores the need for Threads to develop robust moderation tools, like AI bots, to address harassment proactively.

Another critical example is X's handling of misinformation during the 2020 US Presidential elections, where it employed soft moderation by labeling Donald Trump's tweets containing unverified claims of election fraud [14]. While these labels did not significantly alter user engagement, they still served as cautionary markers regarding the accuracy of content shared on the platform. This approach highlights the limitations of passive moderation, reinforcing Threads' need for a collaborative fact-checking mechanism to tackle misinformation effectively, aligning with RQ2 and informing our design to balance user autonomy with platform safety.

2.2 Comparative Analysis: Reddit

While Reddit differs structurally from Threads with its subreddit-based model, its community-driven dynamics and user-generated content challenges make it a valuable comparison for studying toxicity mitigation. Threads, as a nascent platform with a centralized feed akin to X, may evolve to host diverse user groups with varying norms, similar to Reddit's subreddits [16]. This relevance justifies analyzing Reddit to inform our research questions—particularly RQ1 and RQ3—and to brainstorm design solutions like community-specific guidelines and mood selectors that adapt to Threads' growth. Reddit's experiences with norm adaptation and policy impacts offer critical lessons for these objectives, which we explore below.

A study by Rajadesingan et al. (2020) underscores that the survival of online communities relies on attracting newcomers who adapt to community norms through self-selection, pre-entry learning, post-entry learning, and retention [15]. Analysis of political subreddits found these adaptations are community-specific rather than broadly transformative—individuals less toxic in one subreddit may be highly toxic in another, especially when a surge of new users disrupts established norms. This variability in dynamics (RQ1) suggests Threads could face similar challenges as it scales, informing our design of community-specific guidelines to foster consistent, positive interactions across diverse user groups.

In July 2023, Reddit's decision to increase fees for third-party apps (e.g., Apollo) prompted widespread protests, with many subreddits switching from public to private visibility [16], highlighting how platform policies can erode community health and user loyalty, outcomes central to RQ3. Reddit employs upvotes and downvotes to signal content quality, interpreted differently across

subreddits, while volunteer moderators, unpaid and reliant on platform tools, shape user experiences through grassroots governance. These examples emphasize the need for Threads to ensure transparent, user-centric guidelines to maintain trust and engagement, avoiding Reddit's policy-driven pitfalls.

3 RESEARCH METHODOLOGY

To assess the nature and scope of toxicity in online communities, particularly in the nascent platform "Threads," a two-pronged methodology was employed: secondary research and primary research. We commenced by studying existing literature on online community dynamics, regulations, and interventions. The resources provided insights into existing frameworks and mechanisms platforms use to regulate content, including the challenges and benefits they offer. We conducted in-depth, semi-structured interviews with 30 participants, aged 19 to 26, including undergraduate and Master's students from three universities in Lahore, Pakistan—two private and one public. Participants were recruited through a combination of random and snowball sampling techniques, ensuring the gender ratio was maintained throughout the sample. All participants were active users of at least one social media platform, such as Instagram, X, Facebook, WhatsApp, Reddit, or TikTok.

Participants' interactions with the platforms were crucial in understanding the nuances of user experiences and expectations from platforms like Threads. Thus, our interviews were aimed at exploring participants' experiences with online community dynamics, the impact of moderation techniques on toxic behaviors, and how community guidelines influence users' perceptions of platform safety and engagement. We analyzed our interview findings through thematic analysis using inductive coding and identified four major themes regarding online community dynamics, user concerns, and platform expectations, which are detailed in the Findings section.

3.1 Interview Findings

3.1.1 Nature of Online Communities. Our interviewees claimed to actively participate in at least one or more online communities focused on gaming, academics, technology, online dating, K-pop, freelancing, anime, and e-commerce. The nature of interactions within online communities is community-specific and depends on the level of moderation within each community, but can be described as generally positive and helpful with occasional toxicity, i.e., unpleasant or harmful behavior with the intent of malice. A participant mentioned, *"It's like a global village. There's no nationality. Somebody could be portraying themselves as someone they're not. And there's a lot of randomness in the interactions that you get"* [P4, 22 years]. This reflects the diverse and unpredictable nature of online interactions, where anonymity can lead to both positive and negative experiences.

3.1.2 Experiences with Toxicity. When asked to document their experiences with toxicity, participants referred to observing, and at times being on the receiving end of, instances such as cyberbullying (which in extreme cases escalated to racial slurs), doxing, verbal harassment, and often found themselves in the middle of general conflicts resulting from disagreements over political, religious, and fandom-related topics. Different social media platforms witnessed

varying manifestations of toxic behaviors. X, for example, bore the brunt of religious and political disputes, whereas Discord and Facebook's gaming and programming communities encountered a higher degree of systematic ridicule and exclusion, especially for new users. According to one participant, *"I think a lot of the time toxic behavior arises when there's an aspect of vulnerability in the original post and this is the part where people double down on it for whatever reason. I think it's kind of a power trip to be able to attack someone's vulnerabilities"* [P7, 19 years]. This observation underscores how toxic behaviors often exploit personal vulnerabilities, contributing to a hostile online environment.

3.1.3 Impact on Mental Health. The impact on mental health varied among users as some participants reported being unaffected while others found it particularly difficult to navigate the budding toxicity in certain communities, often forcing them to leave these communities as a final resort. However, participants also reported numerous instances of positive behavior such as receiving encouragement, support, and acceptance, often leading them to form lasting friendships, and connections, and foster a sense of belonging. While the availability of an anonymity feature in online communities was not acknowledged, its role as a conduit for increased hostility and unregulated behavior was. One participant mentioned, *"On Subtle Curry Dating, I came across this case where there was a guy who was divorced, and he was looking for a partner and he was also open to other divorcees who were, you know, open to marrying and everybody was very encouraging towards him"* [P3, 22 years]. This example illustrates how online communities can provide support and positive interactions, fostering empathy and a sense of belonging among users.

3.1.4 Awareness and Effectiveness of Community Guidelines. Regarding their familiarity with community guidelines, participants concurred that they were not well-informed and tended to only skim through them but also pointed out that guidelines were tailored as inaccessible in nature. The consensus was that current guidelines, while present, were not enforced with consistency, or in certain cases, at all. Experiences with reporting toxic behavior yielded different results, and while some found the reporting process to be quite effective, others remained dissatisfied and felt as though their concerns were not adequately addressed. Participants appreciated AI's role in moderation but wanted human involvement to remain an integral part of the process. A participant mentioned, *"To be honest, I just tend to be as polite as possible, but I never really read the guidelines"* [P10, 22 years]. This indicates a general lack of engagement with community guidelines among users. On average, participants reported a satisfaction rating of 7 out of 10 with online community dynamics and social media platforms. Suggestions were targeted around community guidelines, emphasizing accessibility via constant reminders and notifications and a focus on diversity to protect different groups. According to another participant, *"Stronger guidelines can include a more detailed explanation of what content may be considered eligible for reporting"* [P1, 23 years]. This highlights the participants' desire for clearer and more comprehensive guidelines to improve the effectiveness of moderation.

4 DESIGN SOLUTION

We began by brainstorming solutions to the challenges highlighted in the Related Works and Findings sections. Our focus was on finding effective solutions to these problems by making community guidelines more accessible, reviewing fact-checking mechanisms employed by social media platforms and the protocols that govern them, the role of artificial intelligence (AI) and AI-powered moderation, interpersonal dynamics within groups, and methods that could prove beneficial to foster positive discussions. Some initial ideas included Temporary Identity Masking (offering anonymized user interactions), Collaborative Fact-Checking (community-driven validation of information), and Virtual Meditation Spaces (calming visuals with the option to mute regular feeds). We then created low-fidelity wireframes and storyboards (see Figure 1) to illustrate core user challenges and demonstrate how these features mitigate toxic behavior.

In the first scenario, a *"positive"* user named Jack, tired of toxicity on X joins Threads. When another user starts using profane language, Jack sees these messages blurred with a sensitivity label. After repeated offenses, an AI bot publicly alerts the offender to respect community guidelines. Persistent disregard leads to intervention by a moderator, who removes the toxic user from the community. In contrast, Sam, a *"negative"* user, also joins Threads but quickly resorts to offensive language. He receives private AI alerts to be respectful, then public alerts in the community. When he continues ignoring these warnings, a moderator ultimately removes him from the group.

We conducted concept testing with undergraduate students who were frequent social media users and had created accounts on Threads. Their feedback prompted a more holistic approach to our design, leading us to develop high-fidelity prototypes in Figma (rather than paper prototypes) for seamless integration with Threads' existing UI.

Our prototypes (as shown in Figure 2) include a revamped home page which now boasts a **community feature** and a feature to **flag misinformation**. Administrators can create and manage communities, each with its guidelines and an AI bot to assist with moderation. If enabled, the AI bot issues private alerts to users violating guidelines, starting with positive reinforcement (reminders to remain respectful). On repeated violations, public alerts focusing on potential consequences are displayed. After a set number of violations, a human moderator intervenes. This workflow addresses toxic behavior by merging automated detection with human oversight.

A **mood selector** on the explore page allows users to choose from moods such as happiness, sadness, frustration, anger, curiosity, and love. The platform then tailors content accordingly, hiding or labeling potentially triggering material, thus fostering a more personalized and supportive user experience.

To mitigate the fear of backlash and encourage honest opinions, we introduced **temporary identity masking**. By selecting the masked icon before replying, a user's name and profile picture are replaced with an alias and an avatar. However, to prevent misuse, thread owners can choose whether to permit masked replies in their posts.




Figure 1: Storyboards illustrating the positive vs negative user experience.




Figure 2: Final Prototypes (1) Fact Checking, (2) Mood Selection, (3) Communities, (4) AI Moderation

4.1 Evaluation

safer online environments. However, they also suggested improvements to better address key challenges like cyberbullying, misuse of offensive language, misinformation, weak moderation, and unclear community guidelines.

With regards to AI moderation, the general feedback was that it encourages healthy discourse but there were concerns with the accuracy of the AI bot in flagging messages that might contain content that goes against community guidelines. One respondent claimed that the feature relies too much on the “goodness” of people with the initial alerts focused on positive reinforcement and that a total of five alerts seemed exhausting when you already have human moderators in the community who can review misdemeanors and act accordingly. Further concerns revolved around the freedom of speech for individuals who regularly utilize offensive words in their language to better express themselves, how the bot would

deter cyberbullying given the evolving nature of meme culture, and whether it would cater to the use of profanity in languages such as Roman Urdu. To quote a participant, *"International figures like Gary Vee use swearing as a form of expression, which creates an ethical dilemma for me—does this feature limit freedom of speech?"*

Regarding the mood selector board, one of the primary concerns was that people might be indecisive if they don't know what they're feeling which might lead to the board becoming a pain point, and it was suggested that we add a 'skip' option there to prevent that from happening. From a UI perspective, it was suggested that instead of having a separate screen for the moods, we should consider adding them to the categories already available on the explore page to avoid having an additional layer between the user and their goal, i.e., to explore content. Another respondent claimed that it would be better to filter based on user activity instead of moods and even if something triggering offsets you on the feed, that would just be part of the social media experience. Overall feedback for this feature was overwhelmingly positive and users pointed out that this could be a healthy alternative to the algorithmic curation already employed by many social media platforms. One participant mentioned, *"Constantly selecting 'not interested' to remove potentially triggering posts from your feed is frustrating. Having a mood selector board to curate content from the start seems like a better solution."*

Moving towards two-tier verification, it was suggested that an AI bot review the descriptions attached with reports to prevent flooding community fact-checkers with excessive reports on a day-to-day basis. It was also suggested to introduce a flagging limit for misinformation which necessitates that a certain number of users flag a thread for the same thing before it is forwarded to the community fact-checkers. Furthermore, it was suggested that once information has been fact-checked and found to be inaccurate, the correct news, along with a list of sources, be shown to the users along with the appropriate label. If community fact-checkers largely share the same perspective, they could reinforce echo chambers and suppress opposing viewpoints. It was also suggested that if new facts emerge after fact-checkers have verified a situation, there should be an option to reflag the content. To quote a participant, *"The misinformation flag could be misused as a 'dislike' button alongside the 'like' feature."*

When examining temporary identity masking, it was suggested that the platform only allow anonymous replies instead of both anonymous replies and posts, and introduce some real-time monitoring for this mode similar to Reddit where comments are immediately removed if they violate community guidelines. Since we gave thread owners the option to enable or disable masked replies, a potential parent-child inheritance issue was identified. This issue could arise when someone replies using their original profile and subsequently receives unwanted anonymous sub-replies. There were some general concerns about anonymity being misused which were alleviated when users went through the specific community guidelines for this mode. According to one participant, *"Anonymous options are helpful, and while anonymous posting is common, anonymous replies are not. The latter could encourage more participation in discussions about personal topics like family or relationships, as people may feel more comfortable sharing their thoughts without revealing their identity."*

Lastly, most respondents preferred having community guidelines included in the placeholder text when starting a new thread, although some felt the guidelines should be made even more accessible. One respondent claimed, *"Since nobody actually reads the community guidelines, placing them in the placeholder text where people are more likely to see them is a good idea."*

5 DISCUSSION

Our redesign of Threads took into consideration the pressing need to strike a balance between free speech and regulation on the platform, while also creating safer communities focused keenly on fostering inclusivity and participation, directly responding to concerns about platform governance as highlighted in the literature [1]. The features we sought to introduce for users, e.g., creating a two-tier system for flagging misinformation, in which the first tier was solely focused on community fact-checkers, aligned with Taylor Owen's emphasis on combating techlash through a holistic response [2] and was a step towards ensuring better quality of content [8]. The use of AI bots within community groups to detect offensive language and reinforce guidelines through public and private reminders to members is inspired by discussions on secure interactions in online platforms enabled by AI [10]. The preventative role of the bots, along with our decision to explicitly provide *human moderators* with the power of removing/banning users from the community after their messages were repeatedly flagged by the bots, is aimed at minimizing the number of false positive events [9] and counteract the limitations of AI.

AI systems trained predominantly on data from the Global North are likely to embed ideological and political biases inherent in that data [21]. These biases can shape how AI models interpret and respond to issues related to both the North and the South, making the systems neither apolitical nor neutral. When deployed, these models not only reflect these biases but also risk perpetuating harm, particularly to marginalized communities, instead of promoting safety and fairness.

Offering users the option of masked or anonymous responses was to encourage them to participate in public forums on Threads without being held back or hamstrung by the fear of being shunned. This was a response to concerns about internet bullying and the exclusion of out-group members, both of which are significant issues highlighted in the literature [5]. This served as a fair attempt at promoting a platform that encouraged cross-community participation and a safe space for expression on public forums, reducing the risk of negative interactions and protecting users from harm and exclusion [6].

The selective but comprehensive redesign of Threads—emphasizing community-specific rules, improved accessibility of guidelines, increased transparency in their enforcement, sensitive content controls, and a balanced approach to both positive and negative reinforcement—reflects a deliberate effort to create a safer, more inclusive, and more tolerant environment. This approach aims to avoid known pitfalls that often lead to toxicity. These changes are rooted in the systemic analysis of interaction patterns in the digital landscape and a review of current guidelines, including an assessment of their shortcomings. They also reflect the importance of direct and indirect reciprocity—encompassing both active and

passive user engagement on a platform—and the role of power dynamics in fostering healthy online communities [3, 4].

Future work will extend beyond the initial focus of undergraduate and graduate students, outlining the rapid expansion of platforms and their future constantly being shaped by a diverse set of user interactions [7]. This broadening of research participants will ensure that Threads continues to evolve in response to the multifaceted nature of social media platforms and their users, further solidifying its place as a leader in promoting positive online community experiences.

6 CONCLUSION

Overall, we aimed to find ways to implement the core principles of a cyber-collectivist approach to make online communities a safer place. However, our goal was never to eliminate toxicity in online communities because we believed that the inherent toxicity in any individual is something that they, themselves, must work on. So instead, our solutions were focused on reducing toxicity in online communities, particularly Threads, by a significant degree. We wanted to focus our efforts on Threads because of its standing as a direct competitor to X, and a platform to which the user base of X was migrating. Furthermore, our solutions were focused on creating a balanced online ecosystem, where an equilibrium between free speech and regulation could be attained.

References

- [1] Fay, R., & Centre for International Governance Innovation. (2019). Digital Platforms Require a Global Governance Framework. In *Models for Platform Governance* (pp. 27–31). Centre for International Governance Innovation. <http://www.jstor.org/stable/resrep26127.7>
- [2] Owen, T., & Centre for International Governance Innovation. (2019). Introduction: Why Platform Governance? In *Models for Platform Governance* (pp. 3–6). Centre for International Governance Innovation. <http://www.jstor.org/stable/resrep26127.3>
- [3] Faraj, S., & Johnson, S. L. (2011). Network Exchange Patterns in Online Communities. *Organization Science*, 22(6), 1464–1480. <http://www.jstor.org/stable/41303137>
- [4] Johnson, S. L., Faraj, S., & Kudaravalli, S. (2014). Emergence of Power Laws in Online Communities: The Role of Social Mechanisms and Preferential Attachment. *MIS Quarterly*, 38(3), 795–A13. <https://www.jstor.org/stable/26634997>
- [5] ESCARTIN, M. C. P. D. (2015). Rogue Cops Among Rogues: Trolls and Trolling in Social Networking Sites. *Philippine Sociological Review*, 63, 169–190. <http://www.jstor.org/stable/24717164>
- [6] Hwang, E. H., Singh, P. V., & Argote, L. (2015). Knowledge Sharing in Online Communities: Learning to Cross Geographic and Hierarchical Boundaries. *Organization Science*, 26(6), 1593–1611. <http://www.jstor.org/stable/43663672>
- [7] WINNER, L. (1997). Technology Today: Utopia or Dystopia? *Social Research*, 64(3), 989–1017. <http://www.jstor.org/stable/40971195>
- [8] CHEN, J., XU, H., & WHINSTON, A. B. (2011). Moderated Online Communities and Quality of User-Generated Content. *Journal of Management Information Systems*, 28(2), 237–268. <http://www.jstor.org/stable/41304626>
- [9] Etlinger, S., & Centre for International Governance Innovation. (2019). What's So Difficult about Social Media Platform Governance? In *Models for Platform Governance* (pp. 20–26). Centre for International Governance Innovation. <http://www.jstor.org/stable/resrep26127.6>
- [10] Troup, E. G. (2017). Growing Role of Platforms in Cybersecurity. *The Cyber Defense Review*, 2(1), 61–70. <http://www.jstor.org/stable/26267401>
- [11] Romano, A. (2021, January 7). What was Gamergate? The lessons we still haven't learned. Vox. <https://www.vox.com/culture/2020/1/20/20808875/gamergate-lessons-cultural-impact-changes-harassment-laws>
- [12] Help on your suspended Twitter account. (2022). help.twitter.com. <https://help.twitter.com/en/managing-your-account/suspended-twitter-accounts>
- [13] Chatzakou, D., Kourtellis, N., Blackburn, J., De Cristofaro, E., Stringhini, G., & Vakali, A. (2017). Hate is not Binary. <https://doi.org/10.1145/3078714.3078721>
- [14] Papakyriakopoulos, O., & Goodman, E. (2022). The Impact of Twitter Labels on Misinformation Spread and User Engagement: Lessons from Trump's Election Tweets. In *Proceedings of the ACM Web Conference 2022*. <https://doi.org/10.1145/3485447.3512126>
- [15] Rajadesingan, A., Resnick, P., & Budak, C. (2020). Quick, Community-Specific Learning: How Distinctive Toxicity Norms Are Maintained in Political Subreddits. In *Proceedings of the International AAAI Conference on Web and Social Media*, 14, 557–568. <https://doi.org/10.1609/icwsm.v14i1.7323>
- [16] Hern, A. (2023, June 11). Reddit communities to “go dark” in protest over third-party app charges. The Guardian. <https://www.theguardian.com/technology/2023/jun/11/reddit-communities-to-go-dark-in-protest-over-third-party-app-charges>
- [17] Konikoff, D. (2021). Gatekeepers of toxicity: Reconceptualizing Twitter's abuse and hate speech policies. Policy & Internet. <https://doi.org/10.1002/poi3.265>
- [18] Atkinson, R. D., Brake, D., Castro, D., Cunliff, C., Kennedy, J., McLaughlin, M., ... & New, J. (2019). A policymaker's guide to the “techlash”—what it is and why it's a threat to growth and progress. Information Technology and Innovation Foundation.
- [19] Sheth, A., Shalin, V. L., & Kursuncu, U. (2021). Defining and detecting toxicity on social media: context and knowledge are key. *Neurocomputing*, 490, 312–318. <https://doi.org/10.1016/j.neucom.2021.11.095>
- [20] Kender, K. (2022). Tumblr is queer and Twitter is toxic: Speculating about the vibe of social media spaces. In *Proceedings of the 2022 Nordic Human-Computer Interaction Conference (NordCHI '22)* (pp. 1–8). Association for Computing Machinery. <https://doi.org/10.1145/3546155.3547279>
- [21] Ray, P. P. (2023). CHATGPT: A comprehensive review on background, applications, key challenges, bias, ethics, limitations and future scope. *Internet of Things and Cyber-Physical Systems*, 3, 121–154. <https://doi.org/10.1016/j.iotcps.2023.04.003>
- [22] Backlinko Team. (2025, February 3). Number of Threads Users. <https://backlinko.com/threads-users>
- [23] Lyons, K. (2024, April 27). The demise of Twitter: How a “utopian vision” for social media became a “toxic mess.” The Guardian. <https://www.theguardian.com/technology/2024/apr/28/the-demise-of-twitter-how-a-utopian-vision-for-social-media-became-a-toxic-mess>
- [24] Hurst, L. (2023, October 27). “More toxic, less valuable, less useful”: Twitter's year of transformation to X under Elon Musk. Euronews. <https://www.euronews.com/next/2023/10/27/more-toxic-less-valuable-less-useful-twitters-year-of-transformation-to-x-under-elon-musk>
- [25] Oldemburgo de Mello, V., Cheung, F., & Inzlicht, M. (2024). Twitter (X) use predicts substantial changes in well-being, polarization, sense of belonging, and outrage. *Communications Psychology*, 2(1), 15.
- [26] MVMEMNT Staff. (2024, September 13). Why are people leaving X (twitter) for threads? here's the real tea. <https://www.mvemnt.com/why-are-people-leaving-x-twitter-for-threads-heres-the-real-tea/>