

Understanding Journalistic Practices Surrounding AI: Towards the Design of Support Resources for Journalists to Foster Public AI Literacy

Charlotte Li

Department of Communication Studies
Northwestern University
Evanston, Illinois, USA
charlotte.li@u.northwestern.edu

Duri Long

Department of Communication Studies
Northwestern University
Evanston, Illinois, USA
duri@northwestern.edu

Abstract

AI literacy refers to an important set of skills for thinking critically about AI for the general public, and media plays a significant role in shaping public understanding and perception of AI. Research has called for the design of socio-technical systems for journalism practitioners to foster AI literacy in media. In this paper, we conduct an interview study investigating how journalists produce news articles about AI based on a qualitative text analysis of recently published news articles about AI. Our analyses synthesize the current landscape of journalistic practices for producing AI news and surface journalism-specific considerations for communicating AI to the public. We recommend opportunities for supporting journalists with design interventions at different stages of the news production cycle. Our findings have implications for both designers of support tools for journalists, and the computational journalism community.

CCS Concepts

• **Human-centered computing** → **User studies**; • **Computing methodologies** → **Artificial intelligence**; • **Social and professional topics** → **Computing literacy**.

Keywords

AI Literacy, Computational Journalism, Design

ACM Reference Format:

Charlotte Li and Duri Long. 2025. Understanding Journalistic Practices Surrounding AI: Towards the Design of Support Resources for Journalists to Foster Public AI Literacy. In *Extended Abstracts of the CHI Conference on Human Factors in Computing Systems (CHI EA '25)*, April 26–May 01, 2025, Yokohama, Japan. ACM, New York, NY, USA, 9 pages. <https://doi.org/10.1145/3706599.3719953>

1 Introduction

AI technologies are becoming an integral part of many people's day-to-day lives, despite skepticism and confusion surrounding the capabilities and implications of these technologies. While AI is a heavily studied topic in academia, at least half of the respondents of a 2019 survey reported they get information about AI via media

outlets, which extends beyond educational institutions [15]. Anecdotally, there are many misrepresentations of AI in the news, which makes it ever more urgent to promote AI literacy [2], a set of skills for thinking critically about how to use AI, among the general public so people can make informed decisions in their lives and work. Much work in the HCI research community has investigated how to promote AI literacy through educational interventions [17, 18, 32] and system design [3, 30]. Computational journalism researchers similarly have turned their attention to the representation of AI in English media [5, 7] and point to socio-technical interventions for journalists to improve AI literacy in communication about AI technologies [14]. Nevertheless, little work has been done on fostering AI literacy through media and designing tools and resources to support journalists in producing news articles that foster AI literacy. In this paper, we address the need for socio-technical solutions to improve AI literacy in media by embarking on an investigation of how AI has been written about in the news, how journalists produce news pieces about AI, and how we can design interventions to support journalists in writing about AI. We pose the following research questions for this exploration: RQ1: *What are some existing AI literacy concerns in news published about AI technologies?*, RQ2: *What are the existing practices in journalism for producing journalistic pieces around the topic of AI, among different newsrooms and journalist roles?*, and RQ3: *What are considerations for designing an intervention to support journalistic practices that foster public AI literacy?*

To answer these research questions, we first piloted a qualitative text analysis of recently published news articles related to AI, surveying a list of commonly seen practices that could lead to misunderstandings of AI through the news. Using findings from the text analysis, we conducted semi-structured interviews with 10 news practitioners at different career stages about their workflow when producing news related to AI to investigate existing practices and difficulties in the production of AI news stories. Participants also interacted with a low-fidelity prototype of an editing support tool to investigate acceptable and usable interventions that help with the process of reporting and writing about AI in journalism.

Our qualitative analysis of interview data identifies several journalism specific factors and considerations for communicating about AI. We additionally recommend opportunities for supporting journalists with design at different stages of the news production cycle. We contribute a qualitative analysis of common AI literacy concerns in news writing and design implications for supporting journalists to foster AI literacy in news production, which are applicable for

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

CHI EA '25, Yokohama, Japan

© 2025 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-1395-8/25/04

<https://doi.org/10.1145/3706599.3719953>

both designers of support tools for journalists and the computational journalism community.

2 Related Work

An area of related research investigates communications of AI in mass media. Researchers began examining how AI was portrayed in media around 2017 when AI entered public discourse, and they applied both quantitative and qualitative approaches to gauge media representations of AI. Fast and Horvitz used natural language processing and crowdsourcing to understand attitudes towards AI in New York Times articles over a 30-year time span [11]. This study was complemented by an array of qualitative studies that either expanded the scope of investigation to other major newsrooms in the U.S. [7], explored representations of AI in media in the U.K. [5], or honed in on representations of ethical issues related to AI in media [23]. Additional quantitative research verifying prior findings was also conducted using more recent data [9]. Nevertheless, all of this research was produced prior to the popularization of generative AI technologies in public in 2022. Given the spike in discussions and opinions surrounding generative AI since then, the lack of recent literature on how AI is represented in media invites an analysis of how current news media covers AI technologies. *Our work analyzing AI literacy concerns in existing news coverage (RQ1) updates this area of research in the post gen-AI era and establishes a baseline for further systematic investigation.*

Many computational support tools have been designed for use in journalism, from background research for stories to understanding readers' feedback [4, 13, 20–22, 24, 29]. Many studies focused on the design of technologies to assist with tasks in early stages of a news production cycle such as a computational news discovery tool for science journalism [22] and a tool for supporting ideation of news story angles [24]. While some research has explored how to use LLMs to assist with science writing [12] and legal writing [1], little work has investigated how to design tools to support journalists in news creation. *In this paper, we build on research on computational support tools for journalists by examining possibilities for assistance in the news production process to promote more accurate representations of AI.*

Lastly, we respond to an emergent area of research to define AI literacy in journalism, as initiated by work promoting AI intelligibility in journalism, which recommended a socio-technical approach to the issue of AI technology illiteracy in newsrooms [14]. Similar solution was called for in a study on generative AI in newsrooms [8]. Research also introduced literacy curricula and spaces for experimentation as potential ways to improve AI literacy among journalists[6], laying the groundwork for the development of socio-technical interventions to foster AI literacy in journalism and the general public. *Our work extends these research by investigating journalists' needs and reactions to potential interventions (RQ3).*

3 Methodology

3.1 Text Analysis

We conducted a qualitative analysis on published news articles. This preliminary analysis allowed us to search for and identify a list of common AI literacy concerns in news writings to address

RQ1, and serve as basis for further investigation – a baseline of topics to discuss in our interviews. Our findings are exploratory and descriptive and are not intended as an exhaustive summation of AI literacy concerns in the U.S. media.

After initial skimming of articles on AI to verify anecdotal accounts of misrepresentation of AI in media, we searched for news articles published from July 2023 to June 2024 in U.S. Newsstream [26], a database that hosts U.S. national and regional newspaper articles. Specifically, we searched the full text of articles in the database using “AI OR “Artificial Intelligence” OR A.I.” as search terms, and applied several filters such as source types (newspaper), document types (news), language (English), and location (U.S.) to limit the search results to 10778 news articles from a diverse range of newsrooms. Since the goal of this analysis was to establish a baseline of commonly seen issues surrounding AI literacy in news to inform our interviews, we hoped to quickly and meaningfully conduct a qualitative analysis of articles without losing diversity in newsrooms represented. Thus, we divided the search result by month and selected the first two most relevant articles ranked by the database that were produced by two distinct news organizations. This returned 24 documents that were published after the popularization of generative AI technologies and spanned a full year to avoid bias towards a particular news cycle.

The first author coded the documents based on the following sub-questions to RQ1, informed by her background working in journalism and conducting computational journalism research: RQ1.1) What kinds of AI and opinions about AI are represented in news?; RQ1.2) What kinds of attitudes or emotions about AI do the articles elicit?; and RQ1.3) How is AI defined and portrayed, both in terms of its capabilities and its role in human-AI interactions?

Memos on codes and potential themes were created and discussed between the authors of the paper while coding. Once all documents were coded, codes and corresponding quotes were organized into themes related to each of the research questions. For RQ1.1, we examined occurrences of different AI products and experts in news articles. For RQ1.2, we examined uses of sensationalist language. For RQ1.3, based on definitions of AI literacy in related literature [2], we analyzed how news articles addressed questions of “What can AI do?”; “How does AI work?”; and “What should AI do?” Lastly, we consolidated themes into a list of common AI literacy concerns, based on our team's expertise in AI literacy and journalism practice.

3.2 Interviews with Journalists

To address RQ2 and RQ3, we conducted semi-structured interviews with 10 journalists in various career stages. To recruit participants for interviews, we employed social media posting, direct emailing of published journalists, and snowball sampling. For social media posting, we used a colleague with a network of journalists and both authors' social media accounts to reach a large group of journalists. We also directly emailed journalists who published on AI to diversify the types of newsrooms, career stages, and job titles represented in the sample. We also reached out to graduate and undergraduate journalism students at both authors' home institution. All participants were required to have some experience in editorial

roles at U.S.-based newsrooms. While most of our interviewees had experiences writing technology related articles, it was not required.

Each interview had three parts and ran for 1-2 hours. In part 1, we asked participants about the process of producing content related to AI. In part 2, we asked participants to reflect on resources they use to understand AI. In part 3, participants were presented with a list of common concerns in news writing about AI (developed via textual analysis). We then asked participants to edit three passages of news articles based on the list, using a web-based text editor where highlights and comments designed by the first author provided differing levels of suggestions for edits on each of the passages (See Appendix C). One passage was generated using GPT-4 using the prompt “Please write a news paragraph that describes an AI technology with sensationalist language,” and the other two were collected from the internet for incorporating extreme languages in AI news. We conducted this editing exercise with a low-fidelity prototype because we wanted to investigate journalists’ needs, work practices, and preferences before building a tool. Participants were not given specific time limits, but on average took 15-20 minutes to complete the edits. We concluded the interview by asking the participants to provide feedback on their experience with the edits.

We qualitatively analyzed the transcript of each interview, with the first author inductively coding them based on RQ2 and RQ3 and generating memos along the process of key observations. Both authors reviewed codes and memos, resolved conflicts through synchronous and asynchronous discussion, and grouped similar codes to develop themes and a diagram of AI in news production cycles.

4 Results

4.1 Textual Analysis of News Articles

Themes identified in the textual analysis include: *over-representation of well-known AI technologies and related entities*, *writing about AI with sensationalized language*, and *misrepresentations of AI’s capabilities with human-like abilities*. While all of the themes point to potential AI literacy issues in the representation of AI in news, we recognize that many positive journalistic practices exist in media, such as those discussed in Cools et al.[9].

4.1.1 Representation in AI News (RQ1.1). We observed frequent mentions (6 articles) of OpenAI, Google and their corresponding products in articles. Some other tech companies were also mentioned (although less often), including Microsoft, Amazon, Anthropic, etc. Most of the companies mentioned were based either in the U.S. or Europe. Correspondingly, stakeholders at these tech companies were strongly represented in articles, including Elon Musk, Sam Altman, Peter Thiel, and etc. When including voices outside of the tech industry, 2 articles only included one source, often academic, on their opinions about AI. The academics represented in news usually came from the field of CS and were affiliated with institutions located near the newsroom. However, some academics quoted came from fields like law or psychology, presenting AI as a social science discussion and a technological one. Occasionally (1 article), the opinions of policymakers and state officials were also included.

4.1.2 Attitudes and Emotions Towards AI in News (RQ1.2). We observed a broad range of attitudes towards AI in news, from praise to criticism, and uncertainties. 10 news articles represented AI technologies in either a positive or negative light—seldom did articles include both extreme positive and negative views. Extreme positive perspectives described AI as an unprecedented innovation. These articles emphasized AI’s importance for the future of various sectors of work, and they appealed to ethos such as the inevitability of incorporation of AI in work (“We’re at the cusp of using AI for probably the biggest positive transformation that education has ever seen” –A11), fear about humans’ inability to perform as well as AI systems, and fear of being left behind in technological advances (“AI is probably the single most important thing in the future of health care...If you don’t join in, you will be left behind” –A9). Articles with positive attitudes promoted the creation and use of AI for businesses and other work, and they fell short at highlighting potential harms and misuses of AI and users’ agency in AI usage. Conversely, extreme negative representations of AI often appealed to exaggerated ways AI might lead to harms. These articles employed militant language like “kill” or “invade,” (A14,23) or deployed doomsday rhetoric using words such as “disastrous” or “catastrophic,” (A7,14) to describe impacts of AI. At the same time, these sensationalized ways of describing AI did not specify the scale and ways in which individuals can be impacted. Many of these extreme negative perspectives (1 article) were reported as quotes from executives of large tech companies. Despite the corpus of articles sensationalizing the use and impact of AI, several articles called for moderation of AI through policies, or other ways to think about AI critically.

4.1.3 Portrayal of AI in News (RQ1.3). In our review, most articles didn’t directly address the question “What is AI?,” likely assuming knowledge of AI among readerships. Articles that did elaborate on definitions took approaches like (2 articles) providing examples of AI systems (e.g. chatbots), or (2 articles) defining AI by its ability (“What can AI do?”) and (4 articles) defining AI by its mechanisms (“How does AI work?”). Only 1 article that discussed the regulation of AI attempted to answer “What should AI do?”. Discussions on this question are missing from other articles. Most descriptions of AI capabilities focused on well-known tools, such as ChatGPT, and their ability to generate written content or answer logical questions. 4 articles discussed multi-modal AI capabilities such as generating images, audio, or video. Articles seldom discussed the quality and accuracy of AI model outputs when showcasing their capabilities, and they tended to emphasize AI’s ability to generate audio or video “realistically” and “quickly” without including context on prompting, training, and computing power required to produce results. 5 articles personified the capabilities of AI systems, using words like “seeing” (A5), “feeling” (A23), and “critical thinking” (A9), and 2 articles provided descriptions of capabilities in vague terms, such as “to protect patients, we will use A.I. to develop cancer drugs that work better and cost less” (A17).

4.1.4 AI literacy Concerns for News Writing. Based on our analysis of articles, we produced a list of common AI literacy concerns seen in news writings about AI, which include: (1) Over-representation of a specific kind of AI, (2) Use of sensationalist language to describe benefits of AI, (3) Use of sensationalist language to describe harms

of AI, (4) Hypothetical scenarios about the future of AI ungrounded in reality, (5) Personifying the training and capabilities of AI, (6) Presenting capabilities of AI without mentioning its limitations, and (7) Describing AI capabilities with vague language. Appendix A shows detailed explanation of each concern, as they were presented during interviews.

4.2 Interviews with Journalists

We interviewed 10 participants, including 1 editor (P1), 2 senior reporters with more than 5 years of experience (P7, 9), 2 undergraduate students (P5, 6) and 5 early-to-mid career journalists with roles such as fact-checker (P4), news researcher (P8) and reporter (P2, 3, 10) (Appendix B). These interviews allowed us to develop themes describing existing practices for tech news production and factors influencing those practices in newsrooms (Table 1).

4.2.1 Newsroom Culture & Brand Identity. Newsrooms' cultures and brand identities dictate challenges that journalists face and influence production practices. Identities form both internally in how a newsroom recognizes its journalistic missions and targeted audiences, and externally in how other newsrooms and audiences perceive its coverage. There is an unspoken agreement among news practitioners about specific cultures and brand identities of different newsrooms – P10 pointed out: “there’s obviously people that are like, oh, Fox News is conservative and CNN is liberal. But yeah, I would say within tech too, it can get even more nuanced. And there’s also all these tech-specific publications, the information or tech crunch that each have their own kind of personality a little bit.” Participants explained how these characteristics dictate decisions when producing news, such as the type of news, tone, content covered, and available resources within the newsrooms. For example, P3 highlighted a critical newsroom culture surrounding AI: “the newsroom that I’ve worked in, we all have the same perspective of being cautious about artificial intelligence, but I think that also means that we often don’t write about places where AI is being used for good...” In contrast, P10’s newsroom adopted a more positive stance: “I think the reason why we don’t even write as many of those stories (critical of AI), is because we don’t attract those sorts of sources.”

4.2.2 Writing Practices. Participants recognized various writing practices as useful for presenting information on AI, while noting some challenges in producing “good” writing on AI, such as a lack of transparency from developers and researchers of AI (P1, 3, 5, 8). We observed a tension between wanting to avoid sensationalized language and wanting to accurately represent quotes from sources or press releases among participants. In some cases, sensationalized language was a preferred way to accurately represent the discourse on AI, and in other cases, quotes with sensationalized language were deemed unnecessary due to the effort needed for fact-checking and providing contextual information (P1, 2, 5, 6, 7, 8). All of our participants mentioned specificity as a goal to strive for when writing, which is achieved by providing details about facts, providing context, limitations, and methods used in reporting (for example, providing the code used for calling AI model APIs when used for reporting). Participants proposed a few strategies for presenting technological details of AI, including using simple

technical terms (like “recognize patterns in large amount of data”) (P2, 3, 7), referencing familiar technologies, like voice assistant or auto-complete (P8), and mentioning limitations of technologies, but these practices might differ based on newsroom branding.

4.2.3 Reporting Practices. Gathering, picking, and understanding information are important parts of the journalistic process for AI literacy interventions. Participants reported that they determine newsworthiness based on factors such as impact, timeliness, and geographical relevance. Thus, there are more incentives to report on well-known tech companies and individuals, though P8 points out that it’s important to consider stories outside of the popular discourse too. Participants indicated difficulties in reaching some sources, especially developers and researchers of AI, due to PR measures in tech companies, which constraint the inclusion of diverse voices in news stories: “I haven’t spoken with the people who develop these tools directly just because they’re hard to get in touch with. I think they’re very nervous about talking with reporters...” (–P8).

4.2.4 Balance in Representation. Reaching a balanced representation of opinions on AI, political leanings, or strengths and limitations of AI is a recurring theme. The need for balance exists in several tasks in news production cycles, like including quotes with different opinions (“So, basically what I would have my reporter to do is, leave this in because Sam Altman, he is a big figure, so I don’t think we should just delete what he’s saying...But then giving evidence from the other side basically saying this is overblown.” – P1); explaining factual information to contextualize sensationalized language (P3); and sourcing voices of actors from different sides of a story (P3).

4.2.5 Potential Interventions. We gathered participants’ ideas and reactions to designing AI literacy support for journalists. The most popular idea was to develop styleguides and educational resources for newsrooms, like the AP stylebook [25]. Most participants endorsed the common concerns about presentation in news about AI that we identified via our text analysis and believed a list of concerns could be a helpful resource to distribute in newsroom. We received mixed reactions to having an AI-based interactive tool to assist with writing articles. Some saw opportunities in using AI to identify AI literacy concerns (“I even wonder if you could almost take it a step further and build a sort of co-pilot editor that is trained on AI reporting best practices that could automatically add feedback the way like Grammarly does.” –P1) and others felt apprehension due to lack of trust in AI systems and emotional pressure of receiving feedback while writing (“something that went through and highlighted for me things that maybe needed rewriting would really freak me out.” –P4). Further, P8 suggested establishing a dedicated newsroom role for promoting AI literacy.

4.2.6 Constraints Specific to News. Lastly, we considered the impact of existing business models and human resource structures within newsrooms. Business incentives put pressure on journalists to generate click-through rates and revenue, while the diverse news roles, such as headline writer, reporter, editor, visual journalists, and etc, in the process of news production complicate designing for them. Time is also a unique constraint to consider when designing interventions that support journalists and foster public AI literacy.

Table 1: Values and Tensions for News Practitioners Writing about AI

Theme	Sub-themes
Writing Practices	(a) Tension in representing sensationalized vs. accurate information; (b) Value of being specific when presenting information; (c) Simple terms, familiar examples, and limitations are strategies for presenting technical aspects of technologies
Reporting Practices	(a) Newsworthiness determined by impact, timeliness, geography; (b) Gathering information and sources on AI can be challenging
Potential Interventions	(a) Mixed reactions to AI-assisted interactive tools; (b) Styleguides and educational resources popular interventions for AI literacy; (c) Structural changes needed in newsrooms to support AI literacy
Newsroom Culture & Brand Identity	-
Balance in Representation	-
Constraints Specific to News	-

As P6 noted, “Because the main problem with newsrooms is you need to get content out because content and views is how you get money, but then the more content you’re trying to get out, the less quality stories you can create...”

5 Discussion

Our results address our research questions by providing a synthesis of current representations of AI in the U.S.-based media (RQ1) and an analysis of existing practices for the production of news about AI (RQ2), which reveals opportunities for design (RQ3). In this section, we present a diagram illustrating how journalists seek information about AI at different stages of the news production cycle and suggest some possible design interventions to foster public AI literacy at each of these stages (see Figure. 1). Central to our organization are four news production stages: *sourcing*, *sensemaking*, *creating*, and *editing*. These stages are borrowed from a recent report on the use of generative AI in newsrooms [10], and we picked ones that are most prevalent in our interviews. News production cycles are iterative and multi-directional—each stage in the cycle could lead to any other. For **sourcing**, journalists often source their stories on AI from users, developers, experts, and breaking events. While prior computational journalism work looked at content discovery [22], a potential intervention that systematically audits source demographics within a newsroom could help establish standards of diverse sourcing. For **sensemaking**, prior HCI work on fact-checking [28] and explainable AI (XAI) [16, 27] may provide some clues to understanding practitioners’ considerations and mental models of AI. Moreover, we see possible interventions in interviewing (assist journalists with deciphering academic and technical sources), experimenting with AI (build infrastructure for easier and cheaper interactions with different models), and researching about AI online (aid with discovering, synthesizing, discerning information online). For **creating**, as originally intended, we mainly focus on writing as a task for design interventions. Though, we discover through the interviews that many stages of news production invite design interventions, and writing assistance is not particularly popular among some of our participants. We also consider interventions for multi-modal ways of presenting news stories (audio and video). Prior work in HCI has explored tools for ideation [24] and creating news reels [31]. We recommend AI literacy interventions with a focus

on practice, including developing and implementing industry-wide and newsroom-specific standards for writing about AI. Individual AI literacy trainings via journalism education institutions or newsroom workshops could also help improve journalists’ personal AI literacy. For **editing**, like *creating*, we believe establishing explicit newsroom standards and identities surrounding AI could be beneficial. We also think establishing ways to engage and measure audiences’ literacy level and being transparent to audiences about the reporting process for AI news could be an effective intervention.

Our recommendations extend prior work [6] by putting considerations for fostering AI literacy in news writing into context via interviews. We validate and expand upon the curricula proposed in prior work and identify specific points for intervention in the journalistic workflow. This research offers a basis for organizing future participatory design sessions with journalists as proposed in their work.

Our findings are limited by a sample of 10 U.S.-based journalists and 24 English articles. Future work can engage with a larger and more diverse set of practitioners and articles. Specifically, we are interested in perspectives from those holding more executive positions in newsrooms and journalists with more years of practice.

6 Conclusion and Future Work

We investigate common AI literacy concerns in current media on AI and common journalistic processes for producing them. We make recommendations for future interventions to foster public AI literacy through media.

We see several directions for future work. First, our interviews indicate that style-guides based on common AI literacy concerns in news writing may be an effective intervention. Many participants in the study expressed concerns about institutional structures for distributing and applying interventions in practice, which also invites further participatory research on incorporating AI literacy tools into a journalistic workflow. Additionally, the introduction of tools to help with fostering AI literacy might be more effective in some news production stages than others, for example, a framework for reflecting on the diversity of sources might be preferred over an in-line editing tool. Lastly, our work only focused on concerns and practices surrounding representation of AI in news. While other works have explored communicating other scientific concepts in




Figure 1: Phases where journalists seek information about AI through different strategies during the news production cycle

news [19], it is not yet clear whether the set of challenges and recommendations we discussed are specific to AI in news. Further work can be done to explore tools that support scientific literacy and media literacy at large.

Acknowledgments

We thank Nicholas Diakopoulos, Joe Germuska, Ágnes Horvát, and Daniel Romero for their expertise and assistance throughout all aspects of our study.

References

- [1] Regina Bernhaupt, Florian 'Floyd' Mueller, David Verweij, Josh Andres, Joanna McGrenere, Andy Cockburn, Ignacio Avellino, Alix Goguy, Pernille Bjørn, Shengdong (Shen) Zhao, Briane Paul Samson, Rafal Kocielnik, Han L Han, Miguel A Renom, Wendy E Mackay, and Michel Beaudouin-Lafon. 2020. Textlets: Supporting Constraints and Consistency in Text Documents. *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems* (2020), 1–13. doi:10.1145/3313831.3376804
- [2] Regina Bernhaupt, Florian 'Floyd' Mueller, David Verweij, Josh Andres, Joanna McGrenere, Andy Cockburn, Ignacio Avellino, Alix Goguy, Pernille Bjørn, Shengdong (Shen) Zhao, Briane Paul Samson, Rafal Kocielnik, Duri Long, and Brian Magerko. 2020. What is AI Literacy? Competencies and Design Considerations. *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems* (2020), 1–16. doi:10.1145/3313831.3376727
- [3] Maalvika Bhat and Duri Long. 2024. Designing Interactive Explainable AI Tools for Algorithmic Literacy and Transparency. In *Proceedings of the 2024 ACM Designing Interactive Systems Conference* (Copenhagen, Denmark) (DIS '24). Association for Computing Machinery, New York, NY, USA, 939–957. doi:10.1145/3643834.3660722
- [4] William Brannon, Doug Beeferman, Hang Jiang, Andrew Heyward, and Deb Roy. 2024. AudienceView: AI-Assisted Interpretation of Audience Feedback in Journalism. *arXiv* (2024). doi:10.48550/arxiv.2407.12613 arXiv:2407.12613
- [5] J. Scott Brennan, Philip N. Howard, and Rasmus Kleis Nielsen. 2018. *An Industry-Led Debate: How UK Media Cover Artificial Intelligence*. Technical Report. doi.org/10.60625/risj-v219-d676
- [6] Mandi Cai and Sachita Nishal. 2022. Motivations, Goals, and Pathways for AI Literacy for Journalism. In *CHI '23 Workshop on AI Literacy: Finding Common Threads between Education, Design, Policy, and Explainability*. https://api.semanticscholar.org/CorpusID:272332784
- [7] Vincent Conitzer, Gillian Hadfield, Shannon Vallor, Ching-Hua Chuan, Wan-Hsiu Sunny Tsai, and Su Yeon Cho. 2019. Framing Artificial Intelligence in American Newspapers. *Proceedings of the 2019 AAAI/ACM Conference on AI, Ethics, and Society* (2019), 339–344. doi:10.1145/3306618.3314285
- [8] Hannes Cools and Nicholas Diakopoulos. 2024. Uses of Generative AI in the Newsroom: Mapping Journalists' Perceptions of Perils and Possibilities. *Journalism Practice* ahead-of-print, ahead-of-print (2024), 1–19. doi:10.1080/17512786.2024.2394558
- [9] Hannes Cools, Baldwin Van Gorp, and Michael Opgenhaffen. 2024. Where exactly between utopia and dystopia? A framing analysis of AI and automation in US newspapers. *Journalism* 25, 1 (2024), 3–21. doi:10.1177/14648849221122647
- [10] Nicholas Diakopoulos, Hannes Cools, Charlotte Li, Natali Helberger, Ernest Kung, Aimee Rinehart, and Lisa Gibbs. 2024. *Generative AI in Journalism: The Evolution of Newswork and Ethics in a Generative Information Ecosystem*. Technical Report. The Associated Press. doi:10.13140/RG.2.2.31540.05765
- [11] Ethan Fast and Eric Horvitz. 2016. Long-Term Trends in the Public Perception of Artificial Intelligence. *arXiv* (2016). doi:10.48550/arxiv.1609.04904 arXiv:1609.04904
- [12] Katy Ilonka Gero, Vivian Liu, and Lydia Chilton. 2022. Sparks: Inspiration for Science Writing using Language Models. *Designing Interactive Systems Conference* (2022), 1002–1019. doi:10.1145/3532106.3533533
- [13] Yael de Haan, Eric van den Berg, Nele Goutier, Sanne Kruikemeier, and Sophie Lecheler. 2022. Invisible Friend or Foe? How Journalists Use and Perceive Algorithmic-Driven Tools in Their Research Process. *Digital Journalism* 10, 10 (2022), 1775–1793. doi:10.1080/21670811.2022.2027798
- [14] Bronwyn Jones, Rhianne Jones, and Ewa Luger. 2022. AI 'Everywhere and Nowhere': Addressing the AI Intelligibility Problem in Public Service Journalism. *Digital Journalism* 10, 10 (2022), 1731–1755. doi:10.1080/21670811.2022.2145328
- [15] Patrick Gage Kelley, Yongwei Yang, Courtney Heldreth, Christopher Moessner, Aaron Sedley, Andreas Kramm, David T. Newman, and Allison Woodruff. 2021.

- Exciting, Useful, Worrying, Futuristic: Public Perception of Artificial Intelligence in 8 Countries. In *Proceedings of the 2021 AAAI/ACM Conference on AI, Ethics, and Society* (Virtual Event, USA) (AIES '21). Association for Computing Machinery, New York, NY, USA, 627–637. doi:10.1145/3461702.3462605
- [16] Q. Vera Liao, Daniel Gruen, and Sarah Miller. 2020. Questioning the AI: Informing Design Practices for Explainable AI User Experiences. In *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems* (Honolulu, HI, USA) (CHI '20). Association for Computing Machinery, New York, NY, USA, 1–15. doi:10.1145/3313831.3376590
 - [17] Duri Long, Takeria Blunt, and Brian Magerko. 2021. Co-Designing AI Literacy Exhibits for Informal Learning Spaces. *Proc. ACM Hum.-Comput. Interact.* 5, CSCW2, Article 293 (Oct. 2021), 35 pages. doi:10.1145/3476034
 - [18] Duri Long, Jessica Roberts, Brian Magerko, Kenneth Holstein, Daniella DiPaola, and Fred Martin. 2023. AI Literacy: Finding Common Threads between Education, Design, Policy, and Explainability. In *Extended Abstracts of the 2023 CHI Conference on Human Factors in Computing Systems* (Hamburg, Germany) (CHI EA '23). Association for Computing Machinery, New York, NY, USA, Article 329, 6 pages. doi:10.1145/3544549.3573808
 - [19] Matthew C. Nisbet. 2009. Communicating Climate Change: Why Frames Matter for Public Engagement. *Environment: Science and Policy for Sustainable Development* 51, 2 (2009), 12–23. doi:10.3200/envt.51.2.12-23
 - [20] Sachita Nishal and Nicholas Diakopoulos. 2022. From Crowd Ratings to Predictive Models of Newsworthiness to Support Science Journalism. *Proceedings of the ACM on Human-Computer Interaction* 6, CSCW2 (2022), 1–28. doi:10.1145/3555542
 - [21] Sachita Nishal, Eric Lee, and Nicholas Diakopoulos. 2024. De-jargonizing Science for Journalists with GPT-4: A Pilot Study. *arXiv* (2024). doi:10.48550/arxiv.2410.12069
 - [22] Sachita Nishal, Jasmine Sinchai, and Nicholas Diakopoulos. 2024. Understanding Practices around Computational News Discovery Tools in the Domain of Science Journalism. *Proceedings of the ACM on Human-Computer Interaction* 8, CSCW1 (2024), 1–36. doi:10.1145/3637419
 - [23] Leila Ouchchy, Allen Coin, and Veljko Dubljević. 2020. AI in the headlines: the portrayal of the ethical issues of artificial intelligence in the media. *AI & SOCIETY* 35, 4 (2020), 927–936. doi:10.1007/s00146-020-00965-5
 - [24] Savvas Petridis, Nicholas Diakopoulos, Kevin Crowston, Mark Hansen, Keren Henderson, Stan Jastrzebski, Jeffrey V Nickerson, and Lydia B Chilton. 2023. AngleKindling: Supporting Journalistic Angle Ideation with Large Language Models. *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems* (2023), 1–16. doi:10.1145/3544548.3580907
 - [25] Associated Press. 2024. *The Associated Press Stylebook* (57th ed.). The Associated Press, New York, NY.
 - [26] publisher. ProQuest (Firm). 1996. US Newsstream.
 - [27] Jakob Schoeffer, Maria De-Arteaga, and Niklas Kühl. 2024. Explanations, Fairness, and Appropriate Reliance in Human-AI Decision-Making. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems* (Honolulu, HI, USA) (CHI '24). Association for Computing Machinery, New York, NY, USA, Article 836, 18 pages. doi:10.1145/3613904.3642621
 - [28] Connie Moon Sehat, Ryan Li, Peipei Nie, Tarunima Prabhakar, and Amy X. Zhang. 2024. Misinformation as a Harm: Structured Approaches for Fact-Checking Prioritization. *Proc. ACM Hum.-Comput. Interact.* 8, CSCW1, Article 171 (April 2024), 36 pages. doi:10.1145/3641010
 - [29] Joris Veerbeek and Nicholas Diakopoulos. 2024. Using Generative Agents to Create Tip Sheets for Investigative Data Reporting. *arXiv* (2024). doi:10.48550/arxiv.2409.07286
 - [30] Danding Wang, Qian Yang, Ashraf Abdul, and Brian Y. Lim. 2019. Designing Theory-Driven User-Centric Explainable AI. In *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems* (Glasgow, Scotland UK) (CHI '19). Association for Computing Machinery, New York, NY, USA, 1–15. doi:10.1145/3290605.3300831
 - [31] Sitong Wang, Samia Menon, Tao Long, Keren Henderson, Dingzeyu Li, Kevin Crowston, Mark Hansen, Jeffrey V Nickerson, and Lydia B Chilton. 2024. Reel-Framer: Human-AI Co-Creation for News-to-Video Translation. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems* (Honolulu, HI, USA) (CHI '24). Association for Computing Machinery, New York, NY, USA, Article 169, 20 pages. doi:10.1145/3613904.3642868
 - [32] Haitao Wen, Xinyu Lin, Rongkang Liu, and Chang Su. 2024. Enhancing College Students' AI Literacy through Human-AI Co-Creation: A Quantitative Study. In *Proceedings of the 2024 International Conference on Digital Society and Artificial Intelligence* (Qingdao, China) (DSAI '24). Association for Computing Machinery, New York, NY, USA, 123–129. doi:10.1145/3677892.3677913

A Concerns in News Writing on AI

- (1) **Over-representation of a specific kind of A.I.:** recent discussions of AI in news media mostly focus on generative AI systems (ChatGPT, DALL-E, Gemini, etc) produced by large

well-known tech companies (OpenAI, Google, Microsoft, etc) with few mentions of other prevalent AI technologies, such as facial recognition, computer vision, auto-complete, etc.

- (2) **Use of sensationalist language to describe benefits of A.I.:** A common rhetoric in this category is describing AI technologies as an unprecedented and significant innovation. Writings in this category appeals to ethos such as the inevitability of incorporation of AI, or FOMO (Fear-Of-Missing-Out) in a future of AI, as well as humans' inability to perform at the same level as AI systems, to urge adoptions of AI. This rhetoric is often sensationalist because AI research has been around for the past few decades and the incorporation of AI already exists in some aspects for a while, for example, auto-complete when texting.
- (3) **Use of sensationalist language to describe harms of A.I.:** Writings in this category appeal to exaggerated ways AI systems might bring harm to humans. They employ militant languages like “kill” or “invade.” or deploy doom’s day rhetoric using words such as “disastrous” or “catastrophic,” to describe impacts of AI without specifying how.
- (4) **Hypothetical scenarios about the future of A.I.:** Some writings propose hypothetical scenarios of the use of AI (e.g.,) without grounding it in its current capabilities.
- (5) **Personifying the training and capabilities of A.I.:** personification as a rhetorical device is often used in defining and describing the capabilities of AI. While some research shows that using metaphors that embody humans helps with understanding computational systems, some descriptions, such as describing A.I as having feelings and critical thinking skills, stray from realistic capabilities of AI and can lead to overtrust in AI systems.
- (6) **Presenting capabilities of A.I. without mentioning its limitations:** Some examples list various capabilities of A.I. systems (especially chatbots and generative AI) to show the versatility of these systems, but they do not discuss how accurately these systems might accomplish those tasks and how these systems might make mistakes.
- (7) **Describing A.I. capabilities with vague language:** Some writings describe capabilities of AI systems, such as “make better and cheap drugs,” without specifying how an AI system might do such things. This comes down to the lack of specificity in describing the use and limitations of AI systems.


B Interview Participant Overview

C Low-Fi Editing Prototype

Received 20 February 2007; revised 12 March 2009; accepted 5 June 2009

Table 2: Overview of Participants

Participant #	Experience Level	Newsroom Role
P1	over 5 years	Director & Editor
P2	1-5 years	Reporter
P3	1-5 years	Reporter
P4	1-5 years	Fact-checker
P5	less than 1 year	Reporter
P6	less than 1 year	Reporter
P7	over 5 years	Freelance Writer
P8	1-5 years	News Researcher
P9	over 5 years	Reporter
P10	1-5 years	Reporter

**Figure 2: Editing Exercise Prototype with Highlights and Comments**




Figure 3: Editing Exercise Prototype with Highlights




Figure 4: Editing Exercise Prototype with No Suggestions