

User Willingness to Engage with AI-Generated Influencer Content

Cherie Sew

School of Computing and Information Systems
University of Melbourne
Melbourne, Australia
csew@student.unimelb.edu.au

Ryan M. Kelly

School of Computing Technologies
RMIT University
Melbourne, VIC, Australia
ryan.kelly@rmit.edu.au

Sarah Schömbbs

School of Computing and Information Systems
The University of Melbourne
Melbourne, VIC, Australia
sschombs@student.unimelb.edu.au

Jorge Goncalves

School of Computing and Information Systems
University of Melbourne
Melbourne, Australia
jorge.goncalves@unimelb.edu.au

Abstract

In this study, we investigate the reasons behind social media users' willingness or reluctance to engage with AI-generated influencers, an increasingly prevalent presence on social media platforms. We conducted a user study with 120 participants to gather their opinions and perceptions of AI-generated influencers after they watched four short-form videos featuring different AI-generated influencers with varying degrees of human-likeness. We grouped our findings into four high-level categories: positive perceptions, influence of message content, reservations about AI-generated influencers, and shifting views and future expectations. From these categories, we present ten sub-categories as potential drivers for willingness or reluctance in user engagement. Additionally, we discuss the implications, ethics, and risks of utilising AI-generated influencers.

CCS Concepts

• Human-centered computing → Empirical studies in HCI.

Keywords

AI-generated Influencers, Human-Likeness, Self-Disclosure, Social Media

ACM Reference Format:

Cherie Sew, Sarah Schömbbs, Ryan M. Kelly, and Jorge Goncalves. 2025. User Willingness to Engage with AI-Generated Influencer Content. In *Extended Abstracts of the CHI Conference on Human Factors in Computing Systems (CHI EA '25)*, April 26–May 01, 2025, Yokohama, Japan. ACM, New York, NY, USA, 7 pages. <https://doi.org/10.1145/3706599.3719916>

1 Introduction

Virtual influencers are fictional digital characters that can visually resemble humans or have humanoid or non-human appearances. Given their advantages such as space and time flexibility over

human influencers [14], they have been utilised by corporations and public agencies for information propagation. With a projected growth of USD 8.30 billion in 2025 to USD 45.88 billion in 2030, the virtual influencer market is forecasted to grow rapidly [1]. The creation and maintenance of a virtual influencer and their content can be costly due to underlying requirements such as the need for graphic or animator professionals as well as copywriters [9]. However, these costs are likely to decrease as advancements in Generative AI (GenAI) provide solutions for content generation, effectively lowering the barrier of entry for those interested in leveraging the benefits of virtual influencers.

With GenAI, a partially autonomous AI-generated influencer becomes a possibility as stakeholders can create a personalised character and associated content with ease. These stakeholders may be inclined to create AI-generated influencers to leverage the higher user engagement that virtual influencers can achieve compared to their human counterparts [3]. However, as public caution towards AI technologies grows [13], it is also important to understand users' concerns and reluctance toward consuming AI-generated influencer content. Moreover, AI-generated influencers can create hyper-personalised engagement—further fuelling parasocial bonds, and reinforcing echo chambers [5]. Thus, there is a need to examine users' perceptions and concerns with AI-generated influencers as this can provide a preliminary explanation for users' willingness in engaging and consuming AI-generated influencer content. In this study, we aim to understand values expressed by participants that may influence their opinions of AI-generated influencers.

We collected qualitative responses from 120 participants through an online study conducted on Prolific and analysed their feedback using a general inductive approach [33]. Our findings reveal polarised perceptions of AI-generated influencers, which influenced participants' willingness to engage with their content. Moreover, the relevance and quality of message content emerged as a key factor in shaping participants' engagement. Additionally, participants expressed that they would be more inclined to consume content from AI-generated influencers if advancements in GenAI technologies improve their overall quality.

We contribute to the growing body of research on AI-generated influencers by establishing a foundation for understanding user perceptions and opinions of these digital characters through our

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

CHI EA '25, Yokohama, Japan

© 2025 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-1395-8/25/04

<https://doi.org/10.1145/3706599.3719916>

initial exploratory study. Additionally, our findings shed light on the implications of AI-generated influencers in the social media landscape and highlight promising opportunities for future research in this topic.

2 Related Work

2.1 AI-Generated Influencers

AI-generated influencers are a subset of virtual influencers, but are designed and produced using GenAI technologies. The use of virtual influencers over human influencers is rising in popularity as virtual influencers have similar or even better engagement potential compared to human influencers with comparable type of content [32] and follower counts [3]. Further, virtual influencers have the ability to offer brand safety, flexibility, exclusivity, and perceived innovation [9].

Akin to virtual influencers, AI-generated influencers are digital characters with carefully crafted personas. However, the development and maintenance of a virtual influencer can be a significant investment due to human resource requirements, such as the need for computer graphics experts [9, 26]. Thus, there has been a rise in the use of AI-generated influencers as GenAI technologies provide a more accessible and cost-friendly way of creating virtual influencers for brands and organisations.

Beyond digital marketing, virtual influencers have collaborated with public agencies and governments for public information campaigns such as eating habits for climate change [23], and COVID-19 pandemic related information [11]. These collaborations stem from efforts to reach the primary audience of virtual influencers: young adults, as well as social media users who identify as “anti-intellectuals” that are more likely to respond positively to messages from virtual influencers, which they might have disregarded from experts [15]. However, although AI-generated influencers could function as an alternative for virtual influencers, users may react differently to content that is not related to marketing and advertising due to their biases in perception of AI technologies.

Further, prior literature suggested that social media users who follow and engage with virtual influencers are seeking emotional diversions from their daily lives [24]. Other studies found several key factors for the success of virtual influencers, including novelty, visual appeal, and creative content that is unlimited by human constraints [8, 20, 26]. Conversely, trust [8] and the perceived lack of authenticity and relatability [26] persist as some of the challenges faced by virtual influencers. Moreover, the novelty factor is effective in capturing initial appeal but most experts questioned the ability of virtual influencers in interest retention and their longitudinal success in digital marketing [26].

Notably, prior studies, such as those by Choudhury and Shamzare [8] and Lou et al. [20], primarily focused on participants who were already active followers of popular virtual influencers, suggesting a pre-existing bias toward such figures. However, research exploring the perceptions and concerns of general social media users regarding AI-generated influencers remains limited. To address this gap, our study investigates users’ opinions of AI-generated influencers and their willingness to engage with content created by them.

3 Method

The data presented in this paper was collected through an online survey as part of a larger study in which participants watched short-form videos of fictional AI-generated influencers. The main study aimed to understand the impact of human-likeness and self-disclosure on message acceptance of AI-generated influencers. In addition of addressing that aim, we collected qualitative data at the end of the survey to understand participants’ reactions to the videos and the general idea of watching content from AI-generated influencers. This paper’s findings are based on analysis of this data.

3.1 Experimental Stimuli

3.1.1 AI-generated influencer Characters. As the focus of the main study was to investigate the impact of visual human-likeness on message acceptance of AI-generated influencers, we began by using Midjourney to produce static images of four characters to represent the AI-generated influencers. According to previous surveys of influencer demographics, a majority of influencers are female [16] and, on short-form video platforms like TikTok, are aged 18–24 years [12]. Hence, the AI-generated influencers used in the main study were specified to be female in their 20s to ensure demographic and visual alignment with influencers that participants are most likely to encounter on social media platforms.

Further, the AI-generated influencers were intentionally created with a range of visual human-likeness while other aspects of their presentation were controlled for consistency. However, the different range of visual human-likeness also aligns with popular virtual influencers that are currently active on social media. The AI-generated influencers we used are shown in Figure 1.

3.1.2 Message Content. We generated a set of messages using ChatGPT-4 by prompting for popular topics that enable online influencers to convey personal experiences or opinions to their followers for better engagement. Research has shown that influencers tend to share messages related to adversity and growth as these are often universal topics that can enhance the influencers’ perceived connectedness with their followers, which is an important aspect that drives engagement [4]. Hence, we selected ‘Personal Challenges and Growth’ as the main topic, and further prompted ChatGPT to produce different scripts for each AI-generated influencer. This yielded four subtopics: Finding Balance, Overcoming Fear of Failure, Nurturing Self-Compassion, and Embracing Change.

3.1.3 AI-generated influencer Audio and Video. Next, we used text-to-speech software from ElevenLabs to convert the scripts into lifelike audio narration. As the main study was a within-subjects experiment, we used different voice profiles for each AI-generated influencer to prevent bias from occurring in the event that participants recognized the same voice across all stimuli. We made a concerted effort to select consistent voice profiles by opting for profiles labelled as having American accents, as well as having a precise and intentional narration tone. Finally, the static images of the AI-generated influencers and audio narrations were combined into an animated video using the AI video generator platform HeyGen.




Figure 1: AI-generated influencers created for the study. From left to right: High human-like (H), Moderate-High human-like (MH), Moderate-Low human-like (ML), Low human-like (L).

3.2 Participants

We recruited 120 U.S.-based participants through the Prolific platform. 70 participants identified as male, 49 participants identified as female, and 1 participant preferred not to specify. The participants had a mean age of 39.7 (SD = 11.1), and they were filtered to have an approval rating of $\geq 99\%$ with at least 500 submissions.

The study was completed with a median time of 22 minutes, and participants were compensated with \$ 15.66 USD per hour. The study was approved by our University's Human Ethics Committee.

3.3 Procedure

Participants read a plain language statement before they began the survey. The participants were asked to imagine that the short videos were of 'lifestyle influencers' that they encountered while scrolling through social media, and the content is related to their own personal challenges they are currently facing. This was followed by the main experiment which involved answering several validated surveys (not reported here). Finally, they were asked two open-ended questions: 'Reflecting on the entire experience, how does the revelation that the entire character was AI-generated affect your feelings towards watching the video? Please elaborate.', and 'Does the knowledge that the characters and messages were AI-generated alter your willingness to watch similar videos in the future? Why or why not? Please elaborate.'

3.4 Analysis

The qualitative responses were analysed by two authors from the research team using Thomas' general inductive approach [33]. The first coder began with an initial read through of the data in order to thoroughly understand its contents. This was followed by identifying and generating categories, and the description for each categories. Further, the two coders met to review and refine the categories in order to resolve any discrepancies in the qualitative analysis process. With a final agreement of the categories, both

coders independently and deductively assigned codes to the participants' responses based on the categories developed. We also measured coding alignment between the two coders by conducting Cohen's Kappa for inter-rater reliability [22]. We observed an overall inter-rater reliability of 0.86 and 0.81 for Q1 and Q2 respectively, which suggests a strong agreement between both coders [22]. All disagreements were resolved through discussions.

4 Results

We present an analysis of participants' potential reasons for their willingness to consume AI-generated content.

4.1 Positive Perceptions

- (1) **Found AI use appealing:** Participants had a "good feeling" (P5, P44) and were "impressed" (P104) that the video was AI generated: "Discovering that the character was AI-generated only enhances my appreciation for the video. I felt good and happy to meet them." (P5). Further, they expressed that the novelty of the technology captured their attention: "AI-generated videos were motivating and it's new, which grabs my attention" (P119).
- (2) **Enjoyed AI content:** Some participants stated they were willing to consume similar content due to "enjoyment" (P1, P25, P50). Participants also expressed an increase in enjoyment and subsequent willingness **as long as the message delivered was positive**: "I'd watch the video as long as it was pleasing to my eyes and senses. While I don't believe the AI has experienced none of the things that it was talking about, that doesn't take away from the message itself. They were good and if you don't have people in your life to say things like that to you then this is a very good tool to have." (P22).
- (3) **Motivated by self-improvement:** Participants reported willingness to watch AI-generated influencers to "learn" (P36, P95, P115) different viewpoints. Participants also reflected

that the messages delivered can be used for personal growth: *“I’d definitely be willing to watch more videos like this because even if the messages and the characters are AI-generated, the messages included were appealing and helpful to a typical person. The lessons included in the messages should be taken into account by humans to feel better about themselves.”* (P93). Additionally, one participant reported trusting artificial intelligence for learning specialised information: *“I love using ChatGPT for my medical questions. I trust in AI.”* (P71).

4.2 Influence of message content

- (1) **Importance of message over the source:** Participants asserted that the message delivered holds more importance over its nature, and some participants liked the message because they found it *“heart-warming”* (P96) and *“inspirational”* (P4). In fact, a subset of participants who liked the message explicitly stated that they *“did not care”* the videos were AI generated. Subsequently, participants also expressed they would be willing to consume AI-generated influencer content regardless of the AI nature as long as they perceived that the message holds value: *“As long as they are providing helpful information, I would continue to seek them out for other problems or answers I’m looking for to better myself. AI-generated content doesn’t matter to me because if used properly, it can be helpful and efficient to a lot of people.”* (P114).
- (2) **Concerns with message quality:** However, some participants had *“doubts”* (P116) about the messages and found them *“difficult to accept”* (P103) because they were generated by AI. Others expressed that they either *“wouldn’t take the messages seriously”* (P11, P26), or that they were unaffected by the messages: *“The content of the video was meaningless and would be so if a real person was saying it.”* (P42).

4.3 Reservations about AI-generated influencers

- (1) **Inauthenticity and Relatability:** Participants perceived the AI-generated influencers as *“not genuine”* or *“inauthentic”*, creating a sense that the AI-generated influencers did not match the messages delivered: *“It was harder for me to take it seriously. I felt it was all made up and lacked heart.”* (P102). Furthermore, a group of participants stated their reluctance of watching AI-generated influencer content was due to the notion that AI-generated influencers are *“unable to relate.”* to humans. Moreover, another group of participants had a preference for messages shared by humans related to their *“lived experience”* (P96, P107), and five participants asserted that humans more *“relatable”* (P40).
- (2) **Apprehension and Scepticism:** Participants felt apprehensive about consuming AI-generated influencer content because of perceived *“uncanniness”* (P67) and *“peculiarity”* (P98) of listening to a “robot”. Furthermore, participants also indicated feeling suspicious of the intentions of the human creators behind the generated AI-generated influencer content after it was disclosed that everything was AI generated. Similarly, when questioned on their willingness to watch AI-generated influencer content in the future, some participants

reiterated distrust on the human creators as they believed that the generated AI could have been used to mask the creators’ identities in order to manipulate others: *“These videos and their contents weren’t born from the ether; someone is behind them, and I do not know their intentions, but the fact that they can’t show their true selves makes me suspect they want to manipulate others. I would not want to watch any AI-generated content for that reason.”* (P91).

- (3) **Aversion towards AI content:** Many participants reported unwillingness to watch AI content as they perceived it to be a *“waste of time”* (P70, P92). The dislike described from participants stems from the perception that AI is fake: *“It confirmed that I don’t like AI videos. I actively disliked them as I recognized them as AI and it made them feel extra fake.”* (P49). Further, participants expressed *“disapproval”* (P41, P57, P70) on receiving advice from AI, even if it was sound advice. Participants also voiced different concerns for the use of AI-generated influencers due to the *“potential negative impact on the climate”* (P11), *“reinforced detrimental stereotypes of beauty or race”* (P30), and *“validity of AI use for mental health therapy”* (P86).

4.4 Shifting views and future expectations

- (1) **Interested in the technology and normalising behaviours:** Participants highlighted *“novelty”* (P106) and *“curiosity”* (P108) as the motivating factors to watch future AI-generated influencer content. Participants were also interested in AI-generated influencers as they wished to follow the advancement of GenAI technology: *“Even if I didn’t necessarily connect with the message I would watch in the future to see how much better the realism gets going forward.”* (P99). Additionally, some participants asserted that AI-generated videos were already *“part of life”* (P8) or will be *“the future”* (P7).
- (2) **Desire for improved technology:** Participants noted the quality of the AI-generated influencers when watching the videos. While it was observed that the current technology was still *“too artificial”* (P30), some expressed *“amazement”* (P22, P99) and looked forward for the future of AI-generated influencers. Accordingly, participants reported an increased willingness to watch AI-generated influencer content when GenAI technology has achieved an optimum state for producing realistic avatars and better content: *“I would certainly not watch these types of videos until AI is more sophisticated and helpful. The advice would need to be deeper and more thorough. It would also help to have less artificial feeling avatars and voices.”* (P16).

5 Discussion

5.1 Factors Influencing User Engagement and Hesitation Toward AI-generated Influencer Content

This study aimed to understand values expressed by participants that may influence their opinions of AI-generated influencers. Our results indicate that participants frequently distinguished the message from the AI-generated influencer delivering it. This distinction

is significant, as it suggests a more nuanced engagement with AI-generated influencers—one where individuals dissociate the message content from its source. This finding parallels users' ability to separate VTuber characters from their human actor counterparts, highlighting a tendency to prioritise the quality of the content over its origin when interacting with virtual characters [21]. This openness to engage with AI-generated content is **driven by the content of the message rather than the source's authenticity or aesthetics**. Many participants expressed a willingness to learn from the content, even when aware that the AI-generated influencer had not personally experienced the information it shared. While a few participants showed limited acceptance of AI-generated content, our results suggest that **the perceived value of the message significantly influences engagement and enjoyment with AI-generated influencers**.

Conversely, some participants considered the relationship between the source (AI-generated influencer) and the message, leading to hesitancy in consuming AI-generated content. A recurring theme was the perceived inauthenticity and lack of relatability of the AI-generated influencer, as it had **not personally experienced the situations it described, which resulted in an articulated preference for human influencers**. This supports previous literature on the potential challenges faced by virtual influencers [2, 8, 20, 26].

Interestingly, research in Human-Robot Interaction suggests that, in addition to naming a robot, explicitly framing its background and past experiences can be a means to induce anthropomorphism and elicit empathic responses [10, 28]. In our study, we did not provide further details about the character itself, highlighting a potential avenue for future research. In addition, it is important to emphasise that our study focused on messages related to 'Personal Challenges and Growth,' a topic inherently grounded in personal experience, which could have amplified an already existing dissonance between message and AI-generated influencer. **Future research should explore messages that are less dependent on lived experiences and emotionally driven content**.

Notably, when engaging with virtual influencers, users tend to have a more positive experience and acceptance when virtual influencers are honest about their "virtual" nature [19]. In our study, we disclosed the AI generated nature of both character and message content post interaction. We encourage **future research to examine how being transparent and informing users about the "AI" nature of AI-generated influencers before or during the interaction affects users' trust and message acceptance**. Although some participants expressed their willingness to engage with AI-generated influencers with visual appearances that are realistically human-like, which supports previous research [30], the perceived 'uncanniness' of the AI-generated influencer contributed to enhance scepticism and apprehension.

Furthermore, some participants expressed reluctance in engaging with AI-generated influencer content due to their antipathy toward AI content, a sentiment that is becoming increasingly widespread [13]. Participants also voiced various concerns about the use of AI-generated influencers, reflecting a growing awareness of the potential negative impacts of AI, such as environmental costs [18] and the perpetuation of biases [17]. Our analysis further revealed apprehension and distrust directed at the human creators behind AI-generated influencer content. This suggests that the source of

distrust extends beyond the AI itself to concerns about the *human* intentions and authenticity of those who design and manage AI-generated influencers [27].

5.2 Implications for Future Research and the Ethical Deployment of AI-Generated Influencers

Our study contributes to the research on AI-generated influencers by exploring users' willingness or reluctance to engage with AI-generated content through an examination of their perceptions and concerns. We highlight the polarised views users hold regarding AI-generated influencers, which are shaped by their attitudes toward AI technology on social media platforms. One group of participants expressed enthusiasm for AI-generated influencers and their potential benefits, driven by a positive view of AI, while another group showed reluctance, fuelled by scepticism or dislike of AI. Despite these divides, acceptance of AI-generated influencers remains nuanced. For those seeking to leverage AI-generated influencers, it is essential to consider strategies to build trust for better user engagement. **Future research should investigate the complex layers of trust associated with AI-generated influencers, exploring the relationship between users' trust in the technology, the AI-generated influencer itself, and the creator behind the AI-driven content**.

Going forward, stakeholders will need to strike a delicate balance between generating content that aligns with the values of their intended user base and fostering trust in both the AI-generated influencers and their creators. Failing to do so may lead to backlash, as seen in recent headlines where Meta faced significant opposition due to its deployment of AI-generated influencers. This backlash was further fuelled by the AI-generated influencer's self-admission of manufactured trust and manipulated false intimacy with real users [25]. The lack of transparency regarding the influencer's purpose and the concealment of its AI nature highlighted the negative consequences of withholding the truth from users. However, as these AI-generated influencers were developed by Meta rather than independent creators, the question of ethical responsibility arises. While social media platforms *should* implement governance frameworks for AI-generated influencers—and, by extension, AI content—it would be more effective for governments and policy-makers to focus on regulations that ensure creator/developer accountability, as social media platforms may not always prioritise the best interests of users.

The increasing accessibility of GenAI enables everyday users to create their own AI-generated influencers, opening the door to widespread use. However, this raises concerns about amplifying individual risks, such as stereotyping and the underrepresentation of certain groups, on an unpredictable scale. In addition, GenAI could be misused for malicious purposes, such as deliberately manipulating users through misinformation [17, 29]. These challenges highlight the need for measures that help users form a well-rounded understanding of AI-generated influencers and encourage healthy engagement with them. One potential solution to mitigate these risks, while navigating the fine line between AI distrust and excitement, could involve AI-generated influencer identifiers. Similar to Chan et al. [6], who advocate for AI agent identifiers, these could

provide users with additional context about the output and the agent behind it. Such measures are critical to pre-emptively address potential harms associated with emerging technologies [7, 31].

Understanding users' willingness to engage with AI-generated influencers and consume their content is a crucial step toward fostering sociotechnical relationships that are not only supportive and beneficial but also ethically sound. While we recognise the legitimate concerns surrounding the risks and challenges posed by AI-generated influencers, our investigation into issues that drive or hinder user engagement aims to contribute to more informed discussions on the design and regulation of these technologies.

6 Conclusion

In this study, we explore participants' opinions, perceptions, and concerns regarding AI-generated influencers, and how these values influence their willingness to engage with AI-generated content. Our results indicate that, beyond the novelty factor, participants who expressed a willingness to consume AI-generated content typically dissociated the content of the message from the AI-generated influencer delivering it. This finding provides a potential explanation for user interest retention, as participants reported a desire to learn from AI-generated content when it was perceived as positive and valuable.

On the other hand, participants who evaluated both the message and the AI-generated influencer as a whole expressed more hesitancy in engaging with such content. We identify recurring challenges related to authenticity, relatability, and the "uncanny valley", as well as a general aversion to AI content, as key issues that contributed to their reservations. Notably, we also observed evidence of triadic trust [27], where participants expressed concerns about the creators' intentions as a reason for their reluctance to engage with AI-generated influencers.

Finally, we discuss how our findings relate to the broader implications, ethical considerations, and risks of AI-generated influencers on social media platforms. Overall, our results offer valuable insight into the acceptance of and challenges associated with AI-generated influencers, providing an important context from the end-users' perspective.

References

- [1] [n. d.]. Virtual Influencer Market Size & Share | Industry Report, 2030 – grandviewresearch.com. <https://www.grandviewresearch.com/industry-analysis/virtual-influencer-market-report>. [Accessed 22-01-2025].
- [2] Jbid Arsenyan and Agata Mirowska. 2021. Almost human? A comparative case study on the social media presence of virtual influencers. *International Journal of Human-Computer Studies* 155 (2021), 102694.
- [3] Nick Baklanov. 2019. The Top Instagram Virtual Influencers in 2019 | Hype - Journal | HypeAuditor Blog. <https://hypeauditor.com/blog/the-top-instagram-virtual-influencers-in-2019/>
- [4] J. B. Bayer, P. Triêu, and N. B. Ellison. 2020. Social media elements, ecologies, and effects. *Annual Review of Psychology* 71 (2020), 471–497. Issue 1. <https://doi.org/10.1146/annurev-psych-010419-050944>
- [5] Geraden Cann. 2024. Margot Monet's creators hoped she would rake in thousands as an AI influencer. The truth was more frightening. <https://www.abc.net.au/news/2024-08-16/ai-influencers-instagram-porn-social-media-relationships/104071876>.
- [6] Alan Chan, Carson Ezell, Max Kaufmann, Kevin Wei, Lewis Hammond, Herbie Bradley, Emma Bluemke, Nitarshan Rajkumar, David Krueger, Noam Kolt, Lennart Heim, and Markus Anderljung. 2024. Visibility into AI Agents. In *The 2024 ACM Conference on Fairness, Accountability, and Transparency*. ACM, Rio de Janeiro Brazil, 958–973. <https://doi.org/10.1145/3630106.3658948>
- [7] Alan Chan, Rebecca Salganik, Alva Markelius, Chris Pang, Nitarshan Rajkumar, Dmitrii Krashenninnikov, Lauro Langosco, Zhonghao He, Yawen Duan, Micah Carroll, Michelle Lin, Alex Mayhew, Katherine Collins, Maryam Molamohammadi, John Burden, Wanru Zhao, Shalaleh Rismani, Konstantinos Voudouris, Umang Bhatt, Adrian Weller, David Krueger, and Tegan Maharaj. 2023. Harms from Increasingly Agentic Algorithmic Systems. In *Proceedings of the 2023 ACM Conference on Fairness, Accountability, and Transparency (FAccT '23)*. Association for Computing Machinery, New York, NY, USA, 651–666. <https://doi.org/10.1145/3593013.3594033>
- [8] Avishek Choudhury and Hamid Shamszade. 2023. Investigating the impact of user trust on the adoption and use of ChatGPT: survey analysis. *Journal of Medical Internet Research* 25 (2023), e47184. <https://doi.org/10.2196/47184>
- [9] Mauro Conti, Jenil Gathani, and Pier Paolo Tricomi. 2022. Virtual influencers in online social media. *IEEE Communications Magazine* 60, 8 (2022), 86–91. <https://doi.org/10.1109/MCOM.001.2100786>
- [10] Kate Darling, Palash Nandy, and Cynthia Breazeal. 2015. Empathic concern and the effect of stories in human-robot interaction. In *2015 24th IEEE International Symposium on Robot and Human Interactive Communication (RO-MAN)*. 770–775. <https://doi.org/10.1109/ROMAN.2015.7333675>
- [11] Lindsay Dodgson. 2020. The World Health Organization has recruited a CGI influencer to get young people interested in safe practices around the coronavirus. <https://www.businessinsider.com/who-using-cgi-influencer-to-spread-safe-coronavirus-practices-2020-4>
- [12] Fabio Duarte. 2025. TikTok User Age, Gender, & Demographics (2024). <https://explodingtopics.com/blog/tiktok-demographics#tiktok-users-age>
- [13] Michelle Faverio and Alec Tyson. 2023. What the data says about Americans' views of artificial intelligence. <https://www.pewresearch.org/short-reads/2023/11/21/what-the-data-says-about-americans-views-of-artificial-intelligence/>.
- [14] Michael Gerlich. 2023. The power of virtual influencers: Impact on consumer behaviour and attitudes in the age of AI. *Administrative Sciences* 13, 8 (2023), 178. <https://doi.org/10.3390/admsci13080178>
- [15] Maximilian HEE Gerrath, Hossein Olya, Zahra Shah, and Huaiyu Li. 2024. Virtual influencers and pro-environmental causes: the roles of message warmth and trust in experts. *Journal of Business Research* 175 (2024), 114520. <https://doi.org/10.1016/j.jbusres.2024.114520>
- [16] IZEA. 2023. The State of Influencer Equality, 2023 Report. <https://izea.com/resources/insights/2023-state-of-influencer-equality/>
- [17] Hannah Rose Kirk, Bertie Vidgen, Paul Röttger, and Scott A. Hale. 2024. The benefits, risks and bounds of personalizing the alignment of large language models to individuals. *Nature Machine Intelligence* 6, 4 (April 2024), 383–392. <https://doi.org/10.1038/s42256-024-00820-y>
- [18] Pengfei Li, Jianyi Yang, Mohammad A. Islam, and Shaolei Ren. 2025. Making AI Less "Thirsty": Uncovering and Addressing the Secret Water Footprint of AI Models. <https://doi.org/10.48550/arXiv.2304.03271> arXiv:2304.03271 [cs].
- [19] Rachel Esther Lim and So Young Lee. 2023. "You are a virtual influencer!": Understanding the impact of origin disclosure and emotional narratives on parasocial relationships and virtual influencer credibility. *Computers in Human Behavior* 148 (2023), 107897. <https://doi.org/10.1016/j.chb.2023.107897>
- [20] Chen Lou, Siu Ting Josie Kiew, Tao Chen, Tze Yen Michelle Lee, Jia En Celine Ong, and ZhaoXi Phua. 2023. Authentically fake? How consumers respond to the influence of virtual influencers. *Journal of Advertising* 52, 4 (2023), 540–557. <https://doi.org/10.1080/00913367.2022.2149641>
- [21] Zhicong Lu, Chenxinran Shen, Jiannan Li, Hong Shen, and Daniel Wigdor. 2021. More kawaii than a real-person live streamer: understanding how the otaku community engages with and perceives virtual YouTubers. In *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems*. 1–14. <https://doi.org/10.1145/3411764.3445660>
- [22] Mary L McHugh. 2012. Interrater reliability: the kappa statistic. *Biochemia medica* 22, 3 (2012), 276–282. <https://doi.org/10.11613/BM.2012.031>
- [23] Meta. 2023. Meta for government and charities, case studies: World Wildlife Fund Spain. <https://www.facebook.com/government-nonprofits/case-studies/2023-07-24-WWF-Spain>
- [24] Agata Mirowska and Jbid Arsenyan. 2023. Sweet escape: The role of empathy in social media engagement with human versus virtual influencers. *International Journal of Human-Computer Studies* 174 (Jun 2023), 103008. <https://doi.org/10.1016/j.ijhcs.2023.103008>
- [25] Allison Morrow. 2025. Meta scrambles to delete its own AI accounts after backlash intensifies. <https://edition.cnn.com/2025/01/03/business/meta-aiaccounts-instagram-facebook/index.html>
- [26] Evangelos Moustakas, Nishtha Lamba, Dina Mahmoud, and C Ranganathan. 2020. Blurring lines between fiction and reality: Perspectives of experts on marketing effectiveness of virtual influencers. In *2020 International Conference on Cyber Security and Protection of Digital Services (Cyber Security)*. IEEE, 1–6. <https://doi.org/10.1109/CyberSecurity49315.2020.9138861>
- [27] Kambiz Saffarizadeh, Mark Keil, and Likoebe Maruping. 2024. Relationship Between Trust in the AI Creator and Trust in AI Systems: The Crucial Role of AI Alignment and Steerability. *Journal of Management Information Systems* 41, 3 (2024), 645–681. <https://doi.org/10.1080/07421222.2024.2376382>

- [28] Sarah Schömbbs, Jacobe Klein, and Eileen Roesler. 2023. Feeling with a robot—the role of anthropomorphism by design and the tendency to anthropomorphize in human-robot interaction. *Frontiers in Robotics and AI* 10 (June 2023). <https://doi.org/10.3389/frobt.2023.1149601> Publisher: Frontiers.
- [29] Cherie Sew and Jorge Goncalves. 2024. Effects of External Factors on User Behaviour towards Influencer Generated Online Misinformation. In *Proceedings of the 35th Australian Computer-Human Interaction Conference* (Wellington, New Zealand) (*OzCHI '23*). Association for Computing Machinery, New York, NY, USA, 86–98. <https://doi.org/10.1145/3638380.3638399>
- [30] Mike Seymour, Lingyao Ivy Yuan, Alan Dennis, Kai Riemer, et al. 2021. Have we crossed the uncanny valley? Understanding affinity, trustworthiness, and preference for realistic digital humans in immersive environments. *Journal of the Association for Information Systems* 22, 3 (2021), 9. <https://doi.org/10.17705/1jais.00674>
- [31] Renee Shelby, Shalaleh Rismani, Kathryn Henne, AJung Moon, Negar Rostamzadeh, Paul Nicholas, N'Mah Yilla-Akbari, Jess Gallegos, Andrew Smart, Emilio Garcia, and Gurleen Virk. 2023. Sociotechnical Harms of Algorithmic Systems: Scoping a Taxonomy for Harm Reduction. In *Proceedings of the 2023 AAAI/ACM Conference on AI, Ethics, and Society*. ACM, Montréal QC Canada, 723–741. <https://doi.org/10.1145/3600211.3604673>
- [32] Jan-Philipp Stein, Priska Linda Breves, and Nora Anders. 2024. Parasocial interactions with real and virtual influencers: The role of perceived similarity and human-likeness. *New Media & Society* 26, 6 (2024), 3433–3453. <https://doi.org/10.1177/14614448221102>
- [33] David R Thomas. 2006. A general inductive approach for analyzing qualitative evaluation data. *American journal of evaluation* 27, 2 (2006), 237–246. <https://doi.org/10.1177/1098214005283748>