

The You You Are: A Payphone Installation for Parallel-Self Dialogue Enabled by Voice Cloning

Bowen Xue*
The University of Tokyo
Tokyo, Japan
sammy.x@cyber.t.u-tokyo.ac.jp

Genta Ono*
The University of Tokyo
Tokyo, Japan
genta.ono@cyber.t.u-tokyo.ac.jp

Tsubasa Yoshida*
The University of Tokyo
Tokyo, Japan
tsubasa@cyber.t.u-tokyo.ac.jp

Yong-Hao Hu
The University of Tokyo
Tokyo, Japan
yh-haoareyou@cyber.t.u-tokyo.ac.jp

Takuji Narumi
The University of Tokyo
Tokyo, Japan
narumi@cyber.t.u-tokyo.ac.jp



Figure 1: *The You You Are* (UUR): a walk-up payphone installation enabling real-time parallel-self dialogue via voice cloning. (A) In situ booth setup. (B) Retrofitted push-button payphone unit. (C) A thermal receipt printed after the call as a tangible trace.

Abstract

We present *The You You Are* (UUR), a payphone installation enabling real-time dialogue with a voice-cloned "parallel-self." The experience follows a compact ritual: visitors answer three seed questions, converse via a handset, and receive a printed receipt as a memento. This design addresses common pitfalls in mental-health conversational AI (CAI), such as the uncanny discomfort of self-voice mismatch and the tendency of large language models to offer overly directive advice. By framing the CAI as a multiverse variant and using a familiar physical interface, UUR lowers the demands for a perfect clone required to sustain the interaction's believability. Developed through two deployments with 571 visitors, including the analysis of 232 conversation logs, the system demonstrates the potential of UUR's multiverse narrative to enhance user acceptance of voice and content mismatches, and that longer sessions encourage greater self-disclosure.

*These authors contributed equally to this work.



This work is licensed under a Creative Commons Attribution 4.0 International License. *CHI EA '26, Barcelona, Spain*

© 2026 Copyright held by the owner/author(s).
ACM ISBN 979-8-4007-2281-3/26/04
<https://doi.org/10.1145/3772363.3799159>

CCS Concepts

• **Human-centered computing** → **Human computer interaction (HCI)**; *Interactive systems and tools*; *Natural language interfaces*; • **Computing methodologies** → Artificial intelligence.

Keywords

voice cloning, self-dialogue, conversational agents, mental health support, interactive installation

ACM Reference Format:

Bowen Xue, Genta Ono, Tsubasa Yoshida, Yong-Hao Hu, and Takuji Narumi. 2026. The You You Are: A Payphone Installation for Parallel-Self Dialogue Enabled by Voice Cloning. In *Extended Abstracts of the 2026 CHI Conference on Human Factors in Computing Systems (CHI EA '26)*, April 13–17, 2026, Barcelona, Spain. ACM, New York, NY, USA, 5 pages. <https://doi.org/10.1145/3772363.3799159>

1 Introduction

While social and emotional barriers often prevent individuals from seeking psychological support [2, 21, 29] and exacerbate mental health risks [13, 17], conversational artificial intelligence (CAI) powered by large language models (LLMs) provides an accessible, non-judgemental medium for private support, helping users overcome these barriers [11, 19, 20, 35]. However, sustaining user attention and motivation during dialogues with such CAI remains

a challenge [8, 15, 32, 36]. Recent systems explore voice cloning, allowing users perceiving themselves having a dialogue with themselves, which potentially enhances users' engagement [4, 7], yet important limitations persist: (1) even minor mismatch between one's own voice and the cloned one can evoke discomfort [14, 16]; (2) previous attempts focused on framing a future-/ideal-self, but insufficient personal data for training can result in a persona that fails to meet user expectations, compromising interaction's believability [31]. Moreover, during conversation with CAI, (3) system output is often restricted to advice-giving instead of open discussion, narrowing the space for critical reflection and self-derived reframing [30, 34].

To address these challenges, we present *The You You Are* (UUR), an interactive payphone installation that invites users to place a call to a "parallel-self from the multiverse," a CAI playing a role as a divergent version of the user living in an alternate universe. During the experience, a visitor picks up a phone in a booth to converse with a "parallel self", who speaks in a clone voice of the visitor; after the session ends, a receipt is printed with a short summary of the conversation and a brief takeaway line. Framing the CAI as a "parallel" rather than identical self allows UUR to broaden user acceptance toward (1) voice discrepancies and (2) content misalignments between CAI's output and the user's expectation, mitigating discomfort and preserving authenticity of the experience. Additionally, UUR facilitates reflection by (3) prioritizing non-directive dialogue over advice-giving, supplemented by a printed receipt. Finally, the payphone metaphor contextualizes the experience of talking to a "parallel self".

UUR went through multiple rounds of iterative design and in situ testing, including two public deployments at a university media art exhibition. In the second deployment, with visitors' consent, we collected conversation logs from 232 visitors and conducted immediate post-experience interviews. These materials informed further refinements to the prompt design and interaction flow.

2 Related Work

2.1 CAI Usage in Mental Health Improvement

Research into CAI has demonstrated significant potential in mental health, ranging from reducing symptoms of depression and distress [20] to detecting mental health conditions and suicidal ideation [11, 35]. Furthermore, beyond direct therapy, CAIs can serve as effective clinical mediators; by fostering trust through reciprocity and self-disclosure, these tools can encourage deeper patient transparency that successfully transfers to human professionals, thereby enhancing the broader therapeutic alliance [19].

However, previous CAI attempts were mainly designed to produce directive advice or generic interventions rather than prompting users to engage in their own critical reflection. This raises concerns regarding the quality and safety of therapeutic communication [30, 34]. Accordingly, recent research emphasizes supportive, reflective, and non-directive strategies such as motivational interviewing (MI) to guide self-exploration while mitigating potential harms [1, 23].

2.2 Self-Dialogue and Voice Cloning

Reflecting on one's own experiences from a third-person perspective, known as self-distancing, enhances the objectivity required for wiser reasoning [10]. Studies have found that using names or non-first-person pronouns instead of "I" effectively regulates emotions and reduces affective reactivity [18, 24]. This concept aligns with Dialogical Self Theory, which views the self as multiple entities engaging in dialogue inside one's mind [9, 12, 33]. This internal process can be externalized by techniques like switching chairs when speaking from different perspectives, which is found to help moderate regretful emotions [22, 25].

Recent CAI systems implementing such "distanced-self" demonstrate increased user engagement if the CAI replicates the user with sufficient believability [31]. Furthermore, voice cloning enhances these interactions, as one's own voice serves as a highly personal perceptual cue that naturally captures attention [3, 28]. *InnerSelf* synthesizes a user's voice to deliver supportive conversations oriented around positive self-talk [4]. *Emotional Self-Voice* combines voice cloning with an ideal-self framing to improve affective and self-evaluative outcomes [7].

However, technical limitations continue to impact user experiences with these "distanced-self" CAIs. First, perceiving of one's own voice involves distinct neural and perceptual processing compared to that of others; consequently, deviations in a synthesized self-voice are more easily detected and can evoke user discomfort [14, 16]. Second, future-self or ideal-self framings set high expectations for believability that require extensive personal data to fulfill. Insufficient data can lead to a mismatch between the expected persona and the CAI's actual output, causing users to notice the underlying technology, thereby reducing engagement or leading them to question the interaction's authenticity [31].

2.3 Parallel-Self Framing

Counterfactual thinking is the cognitive process by which individuals simulate alternatives to past events [6, 26]. Similarly, narratives of multiverse fictions provide a culturally accessible framework for visualizing how divergent choices might have produced different lives [5]. Central to this genre is the "parallel-self", an alternate version of the individual residing in a different spacetime. A "parallel-self" keeps the original self's basic personality but shows how different life paths can create divergent traits, serving as a possible alternative to who they are now.

Drawing on this idea, we speculated that adopting a multiverse narrative within a conversation with CAI, framing the interaction as a dialogue with a "parallel-self," may help users reconcile discrepancies between their real-world self and the CAI-cloned voice or idealized persona. In this framework, the discrepancies are expected to be interpreted as narrative divergences inherent to the multiverse, rather than technical system failures.

3 The You You Are

UUR is a walk-up installation that embeds a voice-cloned CAI into the familiar metaphor of a public payphone. The core interaction reframes lifting the handset as placing a call to a parallel-self in the multiverse. After the call, UUR prints a thermal receipt containing




Figure 2: Overall interaction flow of UUR.




Figure 3: Hardware overview of UUR.

a short conversation summary and a single takeaway line designed as a memory cue rather than prescriptive advice.

At the system level, UUR is organized around four design decisions to support user engagement:

Parallel-Self Framing. UUR establishes a minimal narrative scaffold for the visitor based on their answers to three initial questions. This scaffold is referenced throughout the experience to generate CAI output. Any divergence between the output and the visitor’s expectations is framed as a plausible timeline divergence rather than a conversational breakdown.

Voice Cloning. UUR clones a visitor-specific voice from a short speech sample and renders the CAI’s replies in that voice via low-latency, streamed text-to-speech.

Non-Directive Dialogue. UUR encodes a motivational-interviewing (MI)-inspired stance in the prompt, prioritizing reflective listening and values elicitation over direct instructions.

The Payphone Metaphor. UUR leverages a familiar calling script and a recognizable boundary for private talk in public. Notably, because users expect low audio fidelity (e.g., band-limited speech) and lack visual cues on a phone call, they are less likely to notice synthesis artifacts or voice deviations. Additionally, the inherent nature of the phone as a long-distance communication tool [27] makes the concept

of talking to a “parallel-self” in another multiverse more intuitively acceptable to the user.

3.1 Interaction Flow

Figure 2 summarizes the overall interaction flow of UUR. A typical session proceeds as follows: **(1) Enter, pick up, and consent.** A visitor enters the booth and the handset starts ringing. After the visitor lifts the handset, a small display presents brief instructions and a consent notice. **(2) Dial and connect.** Holding the handset, the visitor dials on the push-button keypad to initiate the call. **(3) Introduce self and answer seed questions.** UUR first asks for the visitor’s name, then collects three spoken seed responses in this order: (i) what they ate yesterday and their thoughts about it, (ii) what they most want to achieve right now, and (iii) one regret from their life so far. These responses establish stable reference points for the ensuing dialogue and provide a short voice sample for cloning. **(4) Talk with the parallel-self.** The visitor speaks with the voice-cloned parallel-self in real time via the handset. **(5) Hang up and receive a receipt.** The session ends when the handset is returned; UUR prints a thermal receipt containing a brief summary and a single takeaway line as a tangible trace of the call.

3.2 Hardware Fabrication

Figure 3 gives an overview of the UUR hardware and signal flow between the host PC, the retrofitted payphone, and peripheral devices. UUR runs a real-time dialogue pipeline on a mini PC (GMKtec

NucBox 3 / KB3; Intel Celeron J4125, 8 GB DDR4 RAM, 128 GB SSD; Windows 11 Pro) mounted inside a retrofitted push-button payphone (NTT model 675PN-VB; manufactured in 1989). A Raspberry Pi Pico microcontroller (RP2040) provides low-latency sensing and actuation for the physical interface. The host PC drives audio playback, renders step-by-step guidance and the consent notice on a small 4.3-inch LCD display (mini-HDMI input), and controls a serial thermal printer for the post-call receipt.

Payphone hardware. The apparatus is organized into three modules. (1) *Handset audio I/O.* The handset integrates a microphone and a speaker to support full-duplex voice interaction. (2) *Retrofitted phone base.* The hook switch is instrumented with a microswitch to detect on-hook/off-hook transitions. The push-button keypad is scanned as a digital key matrix, and a passive buzzer provides button-click feedback to preserve the familiar dialing feel. (3) *I/O controller and peripherals.* The Pico reads the hook switch and keypad states and drives the buzzer via GPIO pins, reporting events to the host PC over USB. The host PC runs the dialogue pipeline and audio playback, presents instructions and the consent notice on the LCD display, and triggers the 58 mm thermal printer upon hang-up to produce a short conversation summary and a single takeaway line. Session metadata are logged only with participant consent.

Booth fabrication. For public deployment, the booth must be lightweight, stable, and quickly reconfigurable. UUR adopts a polyvinyl chloride (PVC) frame with clear vinyl sheeting to minimize transport overhead and enable rapid assembly. This design establishes a stable structure and a legible boundary for private talk within public space.

3.3 Software System

UUR organizes the experience as a single pipeline. The system first prompts the user with three short seed questions and records a brief voice sample. UUR then runs repeated conversation turns through a speech-to-text (ASR) → response generation (LLM) → text-to-speech (TTS) loop, and finally transitions to a hang-up routine that prints a thermal receipt.

Session initialization. At the beginning of a session, UUR asks three questions (see Sec. 3.1) covering recent, future, and past aspects of the user’s situation. UUR records the spoken responses as dialogue context and as a voice sample for cloning. UUR then proceeds to the main call phase, where each user utterance is processed as one conversational turn.

Turn-level processing pipeline. UUR processes each turn with a real-time pipeline comprising automatic speech recognition (ASR), LLM-based response generation, and neural text-to-speech (TTS) synthesis. Concretely, UUR transcribes the user’s speech into text, generates the parallel-self response with gpt-oss-120b¹ via Groq’s OpenAI-compatible interface, and renders the response in the cloned voice with streamed TTS so playback can begin as early as possible. Regarding the model selection of gpt-oss-120b, our iterative deployments identified it as the ideal trade-off between response quality and the low latency required to maintain a natural phone-like conversational rhythm.

Latency management. Even with streamed synthesis, UUR introduces a short latency on the order of 1–2 s in our deployment between response readiness and audible speech onset. To avoid silence and maintain turn-taking rhythm, UUR inserts brief filled pauses (e.g., “uh”, “well”, “I mean”) during this latency window. UUR stops these cues as soon as the synthesized reply begins, preventing overlap.

Session termination and receipt output. A session ends when the handset is returned; UUR additionally caps the number of turns to maintain exhibit throughput. UUR then plays a brief hang-up cue, displays a closing message, and prints a thermal receipt as a tangible trace of the call. The receipt includes a phone-bill-style header rendered in ASCII, a timestamp, a pseudonymous session ID, a call duration, a short LLM-generated summary, and a single takeaway line intended as a memory cue rather than prescriptive advice. To reinforce the payphone metaphor, the receipt also displays a fixed fee line of 15 Japanese yen; the experience itself is free to visitors.

4 Demonstration and User Feedback

UUR was deployed twice as a public, walk-up installation at a university media art exhibition, where 571 visitors experienced calls with their parallel selves. The installation was placed in an open gallery space with passing foot traffic. The installation consisted of a retrofitted payphone placed on a pedestal. The system operated continuously for the duration of the exhibition; staff intervened only for occasional maintenance, such as replacing receipt paper and performing resets. Participation required no sign-up. Visitors initiated a session by entering the booth and lifting the handset. A typical session lasted 8–10 minutes and was capped at 10 turns, ending when the handset was returned.

During the second deployment, with visitors’ consent, we recorded 232 conversation logs and conducted brief post-experience interviews. Key observations supported UUR’s design objectives: While around 20% of visitors noted the mismatches between cloned voice and their own, they still reported a sensation of “talking to themselves” without discomfort. This suggests that UUR’s parallel-self narrative potentially achieved to enhance acceptance to the vocal and content discrepancies. In addition, although the relatively brief average sessions (4–5 turns) limited direct observation of visitors’ critical reflection, we identified that visitors who engaged in longer conversations (7 turns or more) showed a tendency toward increased self-disclosure. This suggests UUR’s potential to build trust and engagement over time, fostering deeper conversations and, consequently, invoke more profound reflections.

5 Conclusion

This paper introduced *The You You Are* (UUR), a walk-up payphone installation that invites users to speak with a multiverse parallel-self through a voice-cloned CAI. UUR addresses engagement challenges in mental-health-oriented CAI by combining a parallel-self framing that tolerates divergence, voice cloning to support engagement, an MI-inspired non-directive dialogue stance to preserve user agency, and a payphone metaphor that leverages expectations of low-fidelity payphone audio. UUR also prints a receipt as a tangible trace of the call, including an LLM-generated summary and a brief takeaway line designed as a memory cue; the system design was

¹<https://huggingface.co/openai/gpt-oss-120b>

iteratively refined through two public deployments, where it was observed to potentially enhance user acceptance of voice and content mismatches and to demonstrate greater self-disclosure with increased sessions. Future work will examine UUR via longer-term deployments, focusing on return usage and longitudinal outcomes related to rumination and reframing.

Acknowledgments

This work was partially supported by JSPS KAKENHI (JP23K25758), JST Moonshot Research & Development Program (JPMJMS2013), JST ASPIRE (JPMJAP2404), and JST COI-NEXT (JPMJPF2201).

References

- [1] Andrew Brown, Ash Tanuj Kumar, Osnat Melamed, Imtihan Ahmed, Yu Hao Wang, Arnaud Deza, Marc Morcos, Leon Zhu, Marta Maslej, Nadia Minian, et al. 2023. A motivational interviewing chatbot with generative reflections for increasing readiness to quit smoking: iterative development study. *JMIR Mental Health* 10 (2023), e49132.
- [2] Emily M Buehler, Ashley M Peterson, and Andrew C High. 2018. Investigating the role of problem severity and channel publicness on the sophistication of support-seeking messages on social network sites: A replication study. *Communication Studies* 69, 3 (2018), 304–313.
- [3] Tatiana Conde, Oscar F Gonçalves, and Ana P Pinheiro. 2018. Stimulus complexity matters when you hear your own voice: Attention effects on self-generated voice processing. *International Journal of Psychophysiology* 133 (2018), 66–78.
- [4] Guang Dai, Pinhao Wang, Cheng Yao, and Fangtian Ying. 2025. InnerSelf: Designing Self-Deepfaked Voice for Emotional Well-being. *arXiv preprint arXiv:2503.14257* (2025).
- [5] Gabriele D'Amato and Luca Diani. 2024. Multiverse Fiction: A Narratological Approach to Infinite Worlds Narratives. *Between* 14, 27 (2024).
- [6] Kai Epstude and Neal J Roese. 2008. The functional theory of counterfactual thinking. *Personality and social psychology review* 12, 2 (2008), 168–192.
- [7] Cathy Mengying Fang, Phoebe Chua, Samantha WT Chan, Joanne Leong, Andria Bao, and Pattie Maes. 2025. Leveraging AI-Generated Emotional Self-Voice to Nudge People towards their Ideal Selves. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*. 1–20.
- [8] Daniel ZQ Gan, Lauren McGillivray, Jin Han, Helen Christensen, and Michelle Torok. 2021. Effect of engagement with digital interventions on mental health outcomes: a systematic review and meta-analysis. *Frontiers in digital health* 3 (2021), 764079.
- [9] Georgia Gkantona. 2025. The Construction of the Narrative Self: Applying the Internal Multi-Actor Performance Method for Children (IMAP-C) to Help Children Cope with Emotional Distress. *Journal of Constructivist Psychology* 38, 2 (2025), 280–296.
- [10] Igor Grossmann and Ethan Kross. 2014. Exploring Solomon's paradox: Self-distancing eliminates the self-other asymmetry in wise reasoning about close relationships in younger and older adults. *Psychological science* 25, 8 (2014), 1571–1580.
- [11] Zhijun Guo, Alvina Lai, Johan H Thygesen, Joseph Farrington, Thomas Keen, Kezhi Li, et al. 2024. Large language models for mental health applications: systematic review. *JMIR mental health* 11, 1 (2024), e57400.
- [12] Hubert JM Hermans. 2012. Dialogical Self Theory and the increasing multiplicity of I-positions in a globalizing society: An introduction. *New directions for child and adolescent development* 2012, 137 (2012), 1–21.
- [13] Ingrid Hogge and Paige Blankenship. 2020. Self-concealment and suicidality: Mediating roles of unmet interpersonal needs and attitudes toward help-seeking. *Journal of clinical psychology* 76, 10 (2020), 1893–1903.
- [14] Taishi Hosaka, Marino Kimura, and Yuko Yotsumoto. 2021. Neural representations of own-voice in the human auditory cortex. *Scientific Reports* 11, 1 (2021), 591.
- [15] Ahmad Ishqi Jabir, Xiaowen Lin, Laura Martinengo, Gemma Sharp, Yin-Leng Theng, and Lorainne Tudor Car. 2024. Attrition in conversational agent-delivered mental health interventions: systematic review and meta-analysis. *Journal of medical Internet research* 26 (2024), e48168.
- [16] Hye-jeong Jo, Chanmi Park, Eunyoung Lee, Jee Hang Lee, Jinwoo Kim, Sujin Han, Joohan Kim, Eun Joo Kim, Eosu Kim, and Jae-Jin Kim. 2024. Neural effects of one's own voice on self-talk for emotion regulation. *Brain Sciences* 14, 7 (2024), 637.
- [17] Brian TaeHyuk Keum, John L Oliffe, Simon M Rice, David Kealy, Zac E Seidler, Dan W Cox, Ronald F Levant, and John S Ogradniczuk. 2023. Distress disclosure and psychological distress among men: The role of feeling understood and loneliness. *Current Psychology* 42, 13 (2023), 10533–10542.
- [18] Ethan Kross, Emma Bruehlman-Senecal, Jiyoung Park, Aleah Burson, Adrienne Dougherty, Holly Shaback, Ryan Bremner, Jason Moser, and Ozlem Ayduk. 2014. Self-talk as a regulatory mechanism: how you do it matters. *Journal of personality and social psychology* 106, 2 (2014), 304.
- [19] Yi-Chieh Lee, Naomi Yamashita, and Yun Huang. 2020. Designing a chatbot as a mediator for promoting deep self-disclosure to a real mental health professional. *Proceedings of the ACM on Human-Computer Interaction* 4, CSCW1 (2020), 1–27.
- [20] Han Li, Renwen Zhang, Yi-Chieh Lee, Robert E Kraut, and David C Mohr. 2023. Systematic review and meta-analysis of AI-based conversational agents for promoting mental health and well-being. *NPJ Digital Medicine* 6, 1 (2023), 236.
- [21] Vivien KG Lim, Thompson SH Teo, and Xiuxi Zhao. 2013. Psychological costs of support seeking and choice of communication channel. *Behaviour & Information Technology* 32, 2 (2013), 132–146.
- [22] Małgorzata Lysiak. 2024. Dialogical temporal chair technique. In *Art that Tells the Truth*. Routledge, 189–197.
- [23] Selina Meyer and David Elswiler. 2025. LLM-based conversational agents for behaviour change support: A randomised controlled trial examining efficacy, safety, and the role of user behaviour. *International Journal of Human-Computer Studies* 200 (2025), 103514.
- [24] Jason S Moser, Adrienne Dougherty, Whitney I Mattson, Benjamin Katz, Tim P Moran, Darwin Guevarra, Holly Shaback, Ozlem Ayduk, John Jonides, Marc G Berman, et al. 2017. Third-person self-talk facilitates emotion regulation without engaging cognitive control: Converging evidence from ERP and fMRI. *Scientific reports* 7, 1 (2017), 4519.
- [25] Sandra C Paivio and Leslie S Greenberg. 1995. Resolving "unfinished business": efficacy of experiential therapy using empty-chair dialogue. *Journal of Consulting and Clinical psychology* 63, 3 (1995), 419.
- [26] Neal J Roese. 1997. Counterfactual thinking. *Psychological bulletin* 121, 1 (1997), 133.
- [27] Avital Ronell. 1989. *The telephone book: technology-schizophrenia-electric speech*. U of Nebraska Press.
- [28] Christine Rosa, Maryse Lassonde, Claudine Pinard, Julian Paul Keenan, and Pascal Belin. 2008. Investigations of hemispheric specialization of self-voice recognition. *Brain and cognition* 68, 2 (2008), 204–214.
- [29] Nina Schnyder, Radosław Panczak, Nicola Groth, and Frauke Schultze-Lutter. 2017. Association between mental health-related stigma and active help-seeking: systematic review and meta-analysis. *The British Journal of Psychiatry* 210, 4 (2017), 261–268.
- [30] Till Scholich, Maya Barr, Shannon Wiltsey Stirman, and Shriti Raj. 2025. A Comparison of Responses from Human Therapists and Large Language Model-Based Chatbots to Assess Therapeutic Communication: Mixed Methods Study. *JMIR Mental Health* 12, 1 (2025), e69709.
- [31] Mehrnoosh Sadat Shirvani, Jackie Liu, Thomas Chao, Suky Martinez, Laura Brandt, Ig-Jae Kim, and Dongwook Yoon. 2025. Talking to an AI Mirror: Designing Self-Clone Chatbots for Enhanced Engagement in Digital Mental Health Support. *arXiv preprint arXiv:2509.06393* (2025).
- [32] Camille E Short, Ann DeSmet, Catherine Woods, Susan L Williams, Carol Maher, Anouk Middelweerd, Andre Matthias Müller, Petra A Wark, Corneel Vandelandotte, Louise Poppe, et al. 2018. Measuring engagement in eHealth and mHealth behavior change interventions: viewpoint of methodologies. *Journal of medical Internet research* 20, 11 (2018), e292.
- [33] Maria Viou and Eugenie Georgaca. 2021. "Enriching Our Inner Dialogue": An Activity to Explore Compassionate Voices. *Journal of Constructivist Psychology* 34, 2 (2021), 207–217.
- [34] Yinzhou Wang, Yimeng Wang, Ye Xiao, Liabette Escamilla, Bianca Augustine, Kelly Crace, Gang Zhou, and Yixuan Zhang. 2025. Evaluating an LLM-Powered Chatbot for Cognitive Restructuring: Insights from Mental Health Professionals. *arXiv preprint arXiv:2501.15599* (2025).
- [35] Xuhai Xu, Bingsheng Yao, Yuanzhe Dong, Saadia Gabriel, Hong Yu, James Hendler, Marzyeh Ghassemi, Anind K Dey, and Dakuo Wang. 2024. Mental-llm: Leveraging large language models for mental health prediction via online text data. *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies* 8, 1 (2024), 1–32.
- [36] Lucy Yardley, Bonnie J Spring, Heleen Riper, Leanne G Morrison, David H Crane, Kristina Curtis, Gina C Merchant, Felix Naughton, and Ann Blandford. 2016. Understanding and promoting effective engagement with digital behavior change interventions. *American journal of preventive medicine* 51, 5 (2016), 833–842.