

Accurate but Not Confident or Confident but Not Accurate? Cognitive Offloading Impairs Confidence Calibration in Human-AI Teams

Charlotte Cheatle
School of Psychology
University of Surrey
Guildford, United Kingdom
c.cheatle@surrey.ac.uk

Adrian Banks
School of Psychology
University of Surrey
Guildford, United Kingdom
a.banks@surrey.ac.uk

Abstract

Appropriate AI reliance requires users to accurately evaluate their own performance in order to discern whether to retain or defer responsibility. While underconfidence is a known driver of automation bias, little is known about how collaboration with AI itself influences users' metacognition. We investigated how cognitive offloading different stages of the decision-making process affects confidence calibration. Participants completed diagnostic decision-making tasks with varying levels of memory and judgement support. Overall, cognitive offloading impaired confidence calibration, with unaided decision makers showing the greatest alignment between confidence and accuracy. Importantly, offloading different cognitive processes produced distinct metacognitive biases: judgement offloading led to overconfidence, whereas combined offloading of memory and judgement processes led to underconfidence. These findings demonstrate that AI support can disrupt user confidence calibration in systematic ways, depending on the type and extent of cognitive delegation. The results highlight metacognitive miscalibration as a critical and underexplored consequence of human-AI collaboration.

CCS Concepts

• **Human-centered computing** → Human computer interaction (HCI); Empirical studies in HC.

Keywords

AI-assisted decision making, confidence calibration, metacognition, human-AI collaboration

ACM Reference Format:

Charlotte Cheatle and Adrian Banks. 2026. Accurate but Not Confident or Confident but Not Accurate? Cognitive Offloading Impairs Confidence Calibration in Human-AI Teams. In *Extended Abstracts of the 2026 CHI Conference on Human Factors in Computing Systems (CHI EA '26)*, April 13–17, 2026, Barcelona, Spain. ACM, New York, NY, USA, 5 pages. <https://doi.org/10.1145/3772363.3798384>



This work is licensed under a Creative Commons Attribution 4.0 International License. *CHI EA '26, Barcelona, Spain*

© 2026 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-2281-3/2026/04

<https://doi.org/10.1145/3772363.3798384>

1 Introduction

The rapid integration of artificial intelligence (AI) into everyday life is reshaping how people approach cognitive tasks. Nevertheless, inappropriate reliance remains a barrier to synergistic team performance, with human-AI teams often failing to outperform the best of either agent alone [31]. Appropriate reliance requires an accurate mental model of the distributed cognitive system [1, 28], that is, an individual must have an awareness of their own and the AI's capabilities to discern whether to retain or defer responsibility. Accordingly, much empirical research has been dedicated to understanding and improving user awareness of AI behaviour, for example, through explanations [32] or AI confidence levels [12, 17]. Yet, these interventions have yielded limited benefits [19, 25]. Less attention has been given to the humans' ability to monitor and evaluate their own capabilities in augmented decision making, known as metacognition. Metacognitive evaluations play a crucial role in cognitive offloading strategies, with underconfidence driving overreliance on AI systems [20]. Critically, however, how offloading and reliance on cognitive aids in turn shape metacognition is not well understood, despite the profound implications this has for fostering appropriate reliance and optimising human-AI collaboration. Therefore, the present study addresses this gap. We focus on confidence calibration—the degree of alignment between subjective confidence and objective accuracy—in an AI-assisted decision-making task.

1.1 Cognitive Offloading

Cognitive offloading refers to the process of shifting cognitive demands onto external tools [24], enhancing both immediate recall [5] and secondary task performance [11, 29]. The advent of generative AI extends traditional offloading from simple information storage to the dynamic delegation of cognitive tasks. AI systems can now support judgement and reasoning, reducing mental load and improving decision accuracy [4, 25], benefits which are amplified under higher task difficulty [22, 32]. Research to date has examined this phenomenon with memory or judgement, in isolation. However, based on a normative model of optimal multi-attribute choice, decision making involves sequential stages: identifying, recalling, weighting, integrating, and comparing cues [21]. Both memory and judgement therefore represent core cognitive components: accurate recall of cues and the effective weighting of their importance jointly determine decision quality. Thus, the significance of offloading in decision-making contexts is twofold: external aids can support with memory and/or judgement. Accordingly, we

investigate increasing levels of cognitive delegation, comparing memory offloading, judgement offloading, and combined memory and judgement offloading, to identify their distinct contributions in a decision-making task.

1.2 Metacognition

Metacognition, often termed ‘thinking about thinking’, refers to how individuals evaluate their own cognition. Metacognitive beliefs have primarily been considered a determinant of reliance on external aids [10, 16, 20], however, the inverse relationship, how external support reciprocally influences metacognition, is not well understood. Cognitive offloading can distort general, long-term assumptions about one’s cognitive abilities [27, 33], with individuals overestimating their abilities even when the AI is harmful to performance [6]. Still, very little is known about the influence of cognitive support on concurrent, task-specific metacognitive evaluations. Taudien and colleagues (p. 251) [30] claim that AI recommendations provide “false comfort”, reporting elevated confidence for both correct and incorrect AI-supported judgements. Nevertheless, it is unclear whether these increases in confidence are justified due to coinciding increases in accuracy provided by AI. This lack of evidence surrounding the influence of offloading on confidence *calibration* presents a critical gap in our understanding of human-AI teaming and informs the basis of the present experiment. Here, we offer two competing hypotheses for the potential effects of cognitive offloading on confidence calibration.

On the one hand, evidence for offloading’s role in inflating metacognitive beliefs would lead us to believe that offloading would similarly inflate metacognitive experiences. Literature on group decision making and confidence calibration would support this account [18, 23]. Therefore, we could predict that participants collaborating with AI will exhibit overconfidence. An alternative perspective, grounded in research on automation bias, could argue that individuals uncritically accept AI outputs rather than engaging analytically [2, 7]. This reflects a broader pattern of cognitive disengagement observed when external systems augment cognition [8, 14, 27], which can reduce users’ perceived ownership over the final outcome [15]. Therefore, in delegating the decision-making process, individuals may lose insight into the underlying reasoning and, consequently, lose confidence that their decision is correct. Therefore, we could instead predict that while AI advice may improve objective accuracy, this may not bring with it increased confidence, translating into underconfidence.

1.3 The Present Study

Together, these competing perspectives capture a critical question in human-AI teaming about the effects that augmentation is having on the way we make decisions and evaluate those decisions. The present study directly evaluates these opposing hypotheses to determine whether cognitive offloading of memory and judgement processes produces overconfidence or underconfidence—a question with significant implications for how humans calibrate reliance on AI systems in future decision-making contexts. To test these predictions, we conducted a study using a diagnostic decision-making paradigm. Participants were randomly assigned to one of four conditions manipulating offloading (memory, judgement, combined

memory and judgement, no offloading) and completed tasks at varying levels of difficulty. Confidence calibration was then compared between groups.

2 Method

2.1 Participants and Design

We recruited 153 participants through Prolific. Data from 47 participants were removed due to pre-defined exclusion criteria. The final sample comprised 106 participants (67% female; $M_{\text{age}} = 39.1$, $SD_{\text{age}} = 13.7$). The experiment used a 4 (cognitive offloading) $\times$ 2 (task difficulty) mixed design within a decision-making paradigm based on medical diagnostic vignettes developed for the study. We investigated differences in decision accuracy, confidence, and confidence calibration. Confidence calibration was computed as the mean difference between subjective confidence and objective accuracy across trials, providing an over/underconfidence statistic (O/U). Scores of zero indicate perfect calibration, with positive and negative values reflecting overconfidence and underconfidence, respectively [34].

2.2 Stimuli

Sixteen medical diagnostic vignettes were developed for this experiment (8 easy, 8 hard). Each vignette presented approximately 25–30 medical cues across six categories (patient information, medical history, symptoms, physical exam, vitals, lab results), each containing between two and four salient cues indicative of the correct diagnosis. We selected common medical conditions with symptoms that could be described without specialised medical terminology. Vignettes in the hard condition included ten numerical lab results (vs. five in the easy condition) to manipulate memory difficulty. They also included red herrings to suggest at least three plausible diagnoses (vs. a single likely diagnosis in the easy condition) to manipulate judgement difficulty. Pilot testing of vignettes confirmed the difficulty manipulation ($N = 22$).

2.3 Procedure and Materials

The experiment was conducted via Qualtrics. Participants were instructed about the task, provided consent, and were randomly assigned to an offloading condition. An example vignette was shown to familiarise participants with the task before they completed eight experimental trials (four easy, four hard) presented in random order. Each of the six sections of the vignettes was displayed on a separate page, and participants could not return to previous pages, creating a deliberate memory burden. The subsequent procedure varied based on the offloading condition. In the no offloading condition ($n = 28$), participants read all the information and made a diagnosis without any support. In the memory-only offloading condition ($n = 28$), participants could select key pieces of information to be presented again at the point of diagnosis to aid memory. In the judgement-only offloading condition ($n = 23$), participants read the vignettes, made an initial diagnosis, and recalled from memory key information. This was submitted to GPT-4o (temperature 0.2), which returned a percentage likelihood for the participant’s diagnosis, along with two alternative diagnoses and their respective likelihoods. Participants then submitted a final diagnosis. Finally, in the combined memory and judgement offloading condition ($n =$




Figure 1: Accuracy (A) and confidence (B) across offloading conditions and difficulty. Error bars represent standard error.

27), participants selected relevant information while reading, which was then passed to GPT-4o alongside their initial diagnosis. Again, GPT-4o then returned likelihood estimates for the initial and two alternative diagnoses, before participants submitted a final decision. GPT-4o has been found to exceed the accuracy required to pass standardised medical licensing exams [13] and so was deemed to be a useful partner in this decision task.

For each trial, final diagnoses were selected from 10 multiple-choice options. Participants rated their confidence in their final decision on a 7-point scale (1=not at all confident, 7=extremely confident) and rated the difficulty of the diagnostic task (0=extremely easy, 100=extremely difficult).

3 Results

Performance and confidence were compared between offloading groups using 2×4 mixed-model ANOVAs. Decision accuracy was significantly higher for easy tasks compared to hard tasks ($F(1, 102) = 250.49, p < .001, \eta^2 = .711$), but comparable between offloading conditions ($F(3, 102) = 2.22, p = .091, \eta^2 = .061$). However, a significant interaction effect revealed that offloading improved performance under high task difficulty ($F(3, 102) = 3.86, p = .012, \eta^2 = .102$). Specifically, those in the combined memory and judgement offloading condition ($M = .56, SD = .26$) outperformed unaided decision makers ($M = .32, SD = .21$) on hard tasks ($p = .009$; Bonferroni corrected) (Fig. 1).

Analyses of confidence revealed a significant main effect of difficulty ($F(1, 102) = 156.67, p < .001, \eta^2 = .606$) and of offloading condition ($F(3, 102) = 3.44, p = .020, \eta^2 = .092$), driven by a greater confidence in the judgement offloading condition relative to the memory offloading condition ($p = .019$). A significant interaction effect was also observed ($F(3, 102) = 3.01, p = .034, \eta^2 = .081$). Under low task difficulty, confidence levels of those in the judgement offloading condition ($M = 5.4, SD = 1.1$) exceeded the memory offloading condition ($M = 4.5, SD = 1.3$) ($p = .044$). Under high task difficulty, confidence was significantly higher for judgement offloading ($M = 4.4, SD = 1.1$) compared to unaided decision makers ($M = 3.3, SD = 1.1$) ($p = .018$), and this difference approached significance compared to those who offloaded memory ($M = 4.5, SD = 1.3$) ($p = .051$) (Fig. 1).

Analyses of confidence calibration revealed that the unaided group were the most accurately calibrated ($M = -.05, SD = .19$),




Figure 2: Confidence calibration across offloading conditions. Error bars represent standard error.

with a significant correlation between accuracy and confidence for easy tasks, $r(26) = .40, p = .037$. No such relationship was observed for any of the offloading conditions, suggesting miscalibration. The direction of miscalibration differed between offloading conditions. Participants in the judgement offloading condition exhibited overconfidence, whereas all other groups reported underconfidence, with the combined-offloading group showing the greatest underconfidence (Fig. 2). Furthermore, participants tended to be overconfident for hard tasks ($M = .05, SD = .35$) and underconfident for easy tasks ($M = -.16, SD = .25$), reflecting the well-established hard-easy effect [9].

Analysis of AI reliance revealed participants' final decisions matched the AI recommendation in 90.4% of cases. However, this often reflected agreement between participants and the AI prior to feedback: 47.3% involved alignment between the initial diagnosis, AI recommendation, and final decision that resulted in a correct outcome, while 28.6% involved aligned but incorrect outcomes. Such convergence occurred more frequently in the judgement offloading condition (89.0%) than in the combined offloading condition (63.7%).

4 Discussion

Our results offer the first empirical evidence that AI decision support, while enhancing performance, impairs confidence calibration. Specifically, unaided decision makers showed the only alignment between confidence and accuracy, while all offloading groups demonstrated miscalibration. Interestingly, the direction of miscalibration differed by AI condition: offloading of the evaluative process (judgement-only) produced overconfidence, whereas combined offloading of both cue storage and evaluation (memory and judgement) produced underconfidence. Together, these findings reveal that offloading affects not only how people decide, but also how well they evaluate their decision making—a dynamic that carries significant implications for augmented human cognition.

One explanation for the underconfidence observed in the combined offloading condition lies in cognitive engagement. As presented earlier, extensive offloading may lead individuals to disengage in the decision-making process [7, 15], attenuating their sense of ownership and understanding. Anticipating AI support throughout the process may have reduced their attention to relevant cues and subsequent internal evaluative processes [14, 27]. Consequently, while participants in the combined offloading condition achieved higher accuracy through greater delegation, they may have lacked a sense of why their decision was correct, and so, relative to their improved performance, they appear underconfident. By contrast, judgement-only offloading preserved a greater degree of active involvement, as participants still needed to prioritise and encode relevant cues, supporting more informed self-evaluations.

A second mechanism concerns how participants shaped AI input. Working memory constraints in the judgement offloading group likely led to selective cue sampling aligned with participants' initial hypotheses, biasing AI recommendations toward agreement in almost 90% of cases. This high agreement rate likely inflated confidence without improving accuracy, producing overconfidence. In contrast, combined offloading enabled more comprehensive cue input, improving AI accuracy but increasing disagreement with initial judgements. Although beneficial for performance, this disagreement may have undermined confidence. Consistent with advice-taking research [3, 26], AI agreement inflated confidence whereas disagreement reduced it, suggesting that metacognitive evaluations were driven more by AI agreement than by actual team performance.

4.1 Conclusions

This study provides the first empirical evidence that AI support, while enhancing performance, can impair users' confidence calibration, with the direction of miscalibration depending on the type and extent of cognitive delegation. Offloading judgement alone inflated confidence, while combined offloading of memory and judgement diminished it. Both patterns reflect poor metacognitive calibration—one group accurate but not confident, the other confident but not accurate. This dissociation underscores a critical metacognitive cost of AI support. We propose two mechanisms which may underlie these effects, warranting further investigation and verification. Findings carry important implications for designing human-AI systems. Interventions to foster appropriate

reliance must account for the dynamics of the distributed cognitive system, targeting overconfidence under certain conditions and underconfidence for others.

References

- [1] Gagan Bansal, Besmira Nushi, Ece Kamar, Walter S. Lasecki, Daniel S. Weld, and Eric Horvitz. 2019. Beyond accuracy: The role of mental models in human-AI team performance. *Proceedings of the AAAI Conference on Human Computation and Crowdsourcing* 7: 2–11. <https://doi.org/10.1609/hcomp.v7i1.5285>
- [2] Zana Bućinca, Maja Barbara Malaya, and Krzysztof Z. Gajos. 2021. To trust or to think: cognitive forcing functions can reduce overreliance on AI in AI-assisted decision-making. *Proceedings of the ACM on Human-Computer Interaction* 5, CSCW1: 1–21. <https://doi.org/10.1145/3449287>
- [3] David V. Budescu, Adrian K. Rantilla, Hsiu-Ting Yu, and Tzur M. Karelitz. 2003. The effects of asymmetry among advisors on the aggregation of their opinions. *Organizational Behavior and Human Decision Processes* 90, 1: 178–194. [https://doi.org/10.1016/S0749-5978\(02\)00516-2](https://doi.org/10.1016/S0749-5978(02)00516-2)
- [4] Ricardo Buettner. 2013. Cognitive workload of humans using artificial intelligence systems: Towards objective measurement applying eye-tracking technology. In *KI 2013: Advances in Artificial Intelligence*, Ingo J. Timm and Matthias Thimm (eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 37–48. https://doi.org/10.1007/978-3-642-40942-4_4
- [5] Lois K. Burnett and Lauren L. Richmond. 2025. Meta-analytic investigations of the effect of cognitive offloading on memory-based task performance and interindividual variability. *Memory & Cognition*. <https://doi.org/10.3758/s13421-025-01743-8>
- [6] Matthew Fisher and Daniel M. Oppenheimer. 2021. Who knows what? Knowledge misattribution in the division of cognitive labor. *Journal of Experimental Psychology: Applied* 27, 2: 292–306. <https://doi.org/10.1037/xap0000310>
- [7] Krzysztof Z. Gajos and Lena Mamykina. 2022. Do people engage cognitively with AI? Impact of AI Assistance on incidental learning. In *Proceedings of the 27th International Conference on Intelligent User Interfaces*, 794–806. <https://doi.org/10.1145/3490099.3511138>
- [8] Michael Gerlich. 2025. AI tools in society: Impacts on cognitive offloading and the future of critical thinking. *Societies* 15, 1: 6. <https://doi.org/10.3390/soc15010006>
- [9] Gerd Gigerenzer, Ulrich Hoffrage, and Heinz Kleinbölting. 1991. Probabilistic mental models: A Brunswikian theory of confidence. *Psychological Review* 98, 4: 506–528. <https://doi.org/10.1037/0033-295X.98.4.506>
- [10] Sam J. Gilbert, Arabella Bird, Jason M. Carpenter, Stephen M. Fleming, Chhavi Sachdeva, and Pei-Chun Tsai. 2020. Optimal use of reminders: Metacognition, effort, and cognitive offloading. *Journal of Experimental Psychology: General* 149, 3: 501–517. <https://doi.org/10.1037/xge0000652>
- [11] Sandra Grinschgl, Frank Papenmeier, and Hauke S. Meyerhoff. 2023. Mutual interplay between cognitive offloading and secondary task performance. *Psychonomic bulletin & review* 30, 6: 2250–2261. <https://doi.org/10.3758/s13423-023-02312-3>
- [12] Nanami Ishizu, Wen Liang Yeoh, Hiroshi Okumura, and Osamu Fukuda. 2024. The effect of communicating AI confidence on human decision making when performing a binary decision task. *Applied Sciences* 14, 16: 7192. <https://doi.org/10.3390/app14167192>
- [13] Alousious Kasagga, Aayam Sapkota, Gichin Changaramkumarath, Jane M. Abucha, Mekdes M. Wollel, Nethra Somannagari, Malik Y. Husami, Kirubel T. Hailu, and Ecem Kasagga. 2025. Performance of ChatGPT and Large Language Models on Medical Licensing Exams Worldwide: A Systematic Review and Network Meta-Analysis With Meta-Regression. *Cureus*. <https://doi.org/10.7759/cureus.94300>
- [14] Mona S. Kleinberg and Richard R. Lau. 2021. Googling politics: How offloading affects voting and political knowledge. *Political Psychology* 42, 1: 93–110. <https://doi.org/10.1111/pops.12689>
- [15] Nataliya Kosmyna, Eugene Hauptmann, Ye Tong Yuan, Jessica Situ, Xian-Hao Liao, Ashly Vivian Beresnitzky, Iris Braunstein, and Pattie Maes. 2025. Your Brain on ChatGPT: Accumulation of Cognitive Debt when Using an AI Assistant for Essay Writing Task. <https://doi.org/10.48550/arXiv.2506.08872>
- [16] Alisa Küper, Georg Christian Lodde, Elisabeth Livingstone, Dirk Schandendorf, and Nicole Krämer. 2025. Psychological factors influencing appropriate reliance on AI-enabled clinical decision support systems: experimental web-based study among dermatologists. *Journal of Medical Internet Research* 27. <https://doi.org/10.2196/58660>
- [17] Doyeon Lee, Joseph Pruitt, Tianyu Zhou, Jing Du, and Brian Odegaard. 2025. Metacognitive sensitivity: The key to calibrating trust and optimal decision making with AI. *PNAS Nexus* 4, 5. <https://doi.org/10.1093/pnasnexus/pgaf133>
- [18] Michael D. Lee and Matthew J. Dry. 2006. Decision making and confidence given uncertain advice. *Cognitive Science* 30, 6: 1081–1095. https://doi.org/10.1207/s15516709cog0000_71
- [19] Jingshu Li, Yitian Yang, Q. Vera Liao, Junti Zhang, and Yi-Chieh Lee. 2025. As confidence aligns: Understanding the effect of AI confidence on human self-confidence in human-AI decision making. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*, 1–16. <https://doi.org/10.1145/3706598.3713336>

- [20] Shuai Ma, Xinru Wang, Ying Lei, Chuhan Shi, Ming Yin, and Xiaojuan Ma. 2024. “Are you really sure?” Understanding the effects of human self-confidence calibration in AI-assisted decision making. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems*, 1–20. <https://doi.org/10.1145/3613904.3642671>
- [21] John W. Payne, Eric J. Johnson, and James R. Bettman. 1993. *The adaptive decision maker*. Cambridge University Press, Cambridge. <https://doi.org/10.1017/CBO9781139173933>
- [22] Phil Peper, Durna Alakbarova, and B. Hunter Ball. 2023. Benefits from prospective memory offloading depend on memory load and reminder type. *Journal of Experimental Psychology: Learning, Memory, and Cognition* 49, 4: 590–606. <https://doi.org/10.1037/xlm0001191>
- [23] Judith M. Puncochar and Paul W. Fox. 2004. Confidence in individual and group decision making: When “two heads” are worse than one. *Journal of Educational Psychology* 96, 3: 582–591. <https://doi.org/10.1037/0022-0663.96.3.582>
- [24] Evan F. Risko and Sam J. Gilbert. 2016. Cognitive Offloading. *Trends in Cognitive Sciences* 20, 9: 676–688. <https://doi.org/10.1016/j.tics.2016.07.002>
- [25] Max Schemmer, Patrick Hemmer, Maximilian Nitsche, Niklas Kühl, and Michael Vössing. 2022. A Meta-Analysis of the Utility of Explainable Artificial Intelligence in Human-AI Decision-Making. In *Proceedings of the 2022 AAAI/ACM Conference on AI, Ethics, and Society*, 617–626. <https://doi.org/10.1145/3514094.3534128>
- [26] Jack B. Soll, Asa B. Palley, and Christina A. Rader. 2022. The bad thing about good advice: Understanding when and how advice exacerbates overconfidence. *Management Science* 68, 4: 2949–2969. <https://doi.org/10.1287/mnsc.2021.3987>
- [27] Betsy Sparrow, Jenny Liu, and Daniel M. Wegner. 2011. Google effects on memory: Cognitive consequences of having information at our fingertips. *Science* 333, 6043: 776–778. <https://doi.org/10.1126/science.1207745>
- [28] Divya Srivastava, J. Mason Lilly, and Karen M. Feigh. 2025. Exploring the role of judgement and shared situation awareness when working with AI recommender systems. *Cognition, Technology & Work* 27, 1–2: 1–18. <https://doi.org/10.1007/s10111-024-00771-9>
- [29] Benjamin C. Storm and Sean M. Stone. 2015. Saving-enhanced memory: The benefits of saving on the learning and remembering of new information. *Psychological Science* 26, 2: 182–188. <https://doi.org/10.1177/0956797614559285>
- [30] Anna Taudien, Andreas Fügener, Alok Gupta, and Wolfgang Ketter. 2022. The effect of AI advice on human confidence in decision-making. <https://doi.org/10.24251/HICSS.2022.029>
- [31] Michelle Vaccaro, Abdullah Almaatouq, and Thomas Malone. 2024. When combinations of humans and AI are useful: A systematic review and meta-analysis. *Nature Human Behaviour* 8, 12: 2293–2303. <https://doi.org/10.1038/s41562-024-02024-1>
- [32] Helena Vasconcelos, Matthew Jörke, Madeleine Grunde-McLaughlin, Tobias Gerstenberg, Michael S. Bernstein, and Ranjay Krishna. 2023. Explanations can reduce overreliance on AI systems during decision-making. *Proceedings of the ACM on Human-Computer Interaction* 7, CSCW1: 1–38. <https://doi.org/10.1145/3579605>
- [33] Adrian F. Ward. 2021. People mistake the internet’s knowledge for their own. *Proceedings of the National Academy of Sciences* 118, 43. <https://doi.org/10.1073/pnas.2105061118>
- [34] Nathan Weber and Neil Brewer. 2004. Confidence-accuracy calibration in absolute and relative face recognition judgments. *Journal of Experimental Psychology: Applied* 10, 3: 156–172. <https://doi.org/10.1037/1076-898X.10.3.156>