

Experts' Experiences of Authoring and Using AI Clones for Client Interactions: Emergent Challenges in Defining Roles and Topic Boundaries

Gippeum Lee

Department of Design
Ulsan National Institute of Science
and Technology (UNIST)
Ulsan, Republic of Korea
gippeumlee@unist.ac.kr

Ho Lee

Department of Design
Ulsan National Institute of Science
and Technology (UNIST)
Ulsan, Republic of Korea
h20255350@unist.ac.kr

Dajung Kim

Department of Design
Ulsan National Institute of Science
and Technology (UNIST)
Ulsan, Republic of Korea
dajungkim@unist.ac.kr

Abstract

Expert AI clones that reflect real human experts' identities, speaking styles, and domain knowledge become increasingly available through commercial platforms, providing accessible advisory services for clients and creating new revenue opportunities for domain experts. In this professional advisory context, risks of AI clones identified in prior HCI research (e.g., identity misalignment, impression transfer, and unclear responsibility) can become even more critical. To better understand the professional use of AI clones, we conducted a qualitative study investigating experts' experiences of and concerns emerge from authoring and using their AI clones for client interactions. Our findings show that experts experienced persistent uncertainty when configuring speaking styles and knowledge boundaries, and repeatedly encountered situations in which the clone's utterances exceeded intended roles and scopes of expertise. Based on these findings, we highlight the need to better support expert users in defining explicit roles and constraining topical boundaries when authoring AI clones.

CCS Concepts

• Human-centered computing; • Human-computer interaction (HCI); • Empirical studies in HCI;

Keywords

AI clones, expert-client interactions, AI-mediated communication, authoring AI systems, user configurations

ACM Reference Format:

Gippeum Lee, Ho Lee, and Dajung Kim. 2026. Experts' Experiences of Authoring and Using AI Clones for Client Interactions: Emergent Challenges in Defining Roles and Topic Boundaries. In *Extended Abstracts of the 2026 CHI Conference on Human Factors in Computing Systems (CHI EA '26)*, April 13–17, 2026, Barcelona, Spain. ACM, New York, NY, USA, 5 pages. <https://doi.org/10.1145/3772363.3798529>

1 Introduction & Background

AI clones, often conceptualized as digital entities that reflect individuals' identities and knowledge [8, 9, 11, 13, 14, 26], are rapidly expanding their applications in professional use domains. For instance, recent AI platforms such as Delphi.ai [6] and Coachvox.ai [5] enable professionals to create "AI versions of themselves" that communicate using their names, speaking styles, and perspectives. Unlike 'virtual' AI advisors or coaches [4, 17, 20, 23], or generic AI assistants [16, 21, 25], these AI clones are designed to communicate and provide advice based on real human experts' professional knowledge and judgment. We refer to such systems as 'expert AI clones'. By allowing clients to interact with an AI proxy of a specific expert at any time and place, expert AI clones promise scalable access to professional advice, while enabling experts to extend the reach of their perspectives and services.

Despite the growing adoption of AI clones in professional contexts, prior HCI research has discussed several risks associated with the use of AI clones. As AI clones are trained on data from a specific point in time, the data held by the clone can increasingly diverge from the actual person over time—a phenomenon referred to as the "snapshot problem" which could lead to identity fragmentation [8, 9, 11, 18, 26]. A "transfer effect" may also occur, where impressions formed toward the clone are transferred back to the original individual [2, 8–14], raising concerns about "misrepresentation" based on distorted or manipulated memories of the original person, or the possibility of developing entirely new "fake memories" [8, 11, 14]. In addition, the attribution of responsibility for erroneous statements made by the clone also remains unclear [8–10, 26].

While these risks can be further intensified in professional settings [18, 19, 24], little is known about how experts create, use, and manage their AI clones, or how they navigate the associated risks in practice. As an initial step toward addressing the gap, this study investigates the following research question: *How do experts experience authoring and using their AI clones for client advisory services, and what concerns and challenges emerge in this process?* We focus on the context of chat-based AI clones, as it currently represents the most widely used modality for expert AI clones. Based on a qualitative study involving seven domain experts and their prospective clients, we provide initial insights into real-world experiences of authoring and deploying expert AI clones from experts' perspectives, highlighting emergent challenges and design opportunities



This work is licensed under a Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License.

CHI EA '26, Barcelona, Spain

© 2026 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-2281-3/2026/04

<https://doi.org/10.1145/3772363.3798529>




Figure 1: (a) Speaking slider controls, (b) knowledge training (c) the chat screen with the service expert AI clone on Coachvox.ai

related to role definition and utterance boundary setting for expert AI clones.

2 Method: Study design & Procedure

This study was designed as a three-phase procedure to explore experts' experiences across three key stages of using AI clones: authoring, deploying, and reviewing. We used Coachvox.ai for this study, as it is one of the accessible AI clone platforms that provides a range of features for experts to configure their AI clone's identity and interaction styles (e.g., speaking styles, knowledge training, deployment, conversation log review). The study protocol was approved by the Institutional Review Board (IRB) of the authors' institution, and all participants provided informed consent.

In **Phase 1**, expert participants were asked to author their own AI clone during a one-hour session. After a short tutorial on Coachvox.ai, expert participants configured their AI clone's speaking style and train knowledge model. For speaking style, they adjusted six-dimension slider controls (Figure 1a) offered by the platform (e.g., Coaching–Advising, Empathetic–Straight-talking). For additional style adjustments, they were able to freely enter custom prompts to specify how the clone should conduct conversations with clients. For knowledge training, expert participants provided text-based materials averaging around 30,000 characters in length they had authored, which were used as the training data for the clone (Figure 1b). In **Phase 2**, the created expert AI clones were deployed to prospective client users. Client participants were recruited from individuals who had genuine interest in one of the experts' domains. Clients were asked to interact freely with the AI clone assigned to them using a chat-based interface in 15-minute daily sessions over three days. In **Phase 3**, experts participated in an in-depth interview while reviewing accumulated client–clone conversation logs (Figure 1c). The interview focused on how the clone's utterances aligned with or deviated from the experts' expectations, as well as on experts' judgments regarding subsequent adjustments to the advising scope and use of AI clones.

We recruited seven professionals from diverse domains who had provided online advisory services for over three years as expert participants. We broadly defined experts as individuals who have practical experiences and expertise in a specific domain (e.g., healthcare, finance, career/HR, marketing) and have shared their professional knowledge and judgment through various media. The participating experts represented domains including relationship

counseling (E1), business (E2), entrepreneurship (E3), nutrition (E4), medicine (E5), finance (E6), and career development (E7), reflecting the diversity of real-world advisory contexts. As the study involved experts' active participation in authoring and deploying their own AI clones, we recruited participants who demonstrated strong interest in and willingness to engage with expert AI clone technologies. To investigate client–clone interactions, in which each expert's clone engaged with multiple clients, we recruited 28 client participants who had a genuine interest in, or need for, advice in one of the experts' domains and had no personal or professional relationship with the corresponding expert. We assigned four clients to each expert's clone initially. Yet, two dropped out, resulting in a final sample of 26 client participants. All expert participants were matched with at least three clients in the end.

All interviews were audio-recorded and transcribed verbatim, and the chat logs were collected from the platform records. Following the thematic analysis approach [1, 3], we reviewed the collected data, including 14 in-depth expert interview session transcripts (including before and after deployment) and 54 client–clone chat sessions. The first and second authors of this study independently conducted initial coding focusing on how experts interpreted and evaluated their advisor clones, as well as the challenges and concerns that emerged during real-world interactions. The initial codes were then refined with entire research team through iterative comparison and discussion until they reached agreement on the key themes.

3 Findings

From the analysis, we identified five core themes that surfaced expert participants' experiences of authoring and deploying their AI clones. We report them in two parts below.

3.1 Configuration Patterns and Challenges in Authoring Expert AI Clones

3.1.1 Speaking Style Configuration: Challenges of Fixed-Value Configuration, Outcome Interpretations, and Target Style Definition. All experts experienced difficulties across multiple dimensions when configuring the clone's speaking style using slider-based controls. First, all experts felt that setting a clone's speaking style as a fixed value using sliders was not an appropriate way to configure how

the clone should speak during conversations. In real advisory interactions, speaking styles are not static but shift fluidly depending on the interlocutor and the situational context. *"I felt that it was quite ambiguous to pin this down in a fixed way, since it really depends on who the client is."* (E1) As a result, some experts described struggling with whom they should have in mind when adjusting the sliders. To resolve this ambiguity, they often imagined a representative or prototypical client they most frequently encountered and calibrated the settings based on that persona: *"Most of my clients are women in their 20s and 30s who want to lose weight, so I adjusted the settings while imagining a familiar persona like that."* (E4) These experts noted that directly writing prompts was more effective than adjusting numerical sliders for configuring speaking styles across diverse situations and contexts.

In addition, some experts noted that within the same slider dimension, it was difficult to understand the differences between the assigned values and predict how they would be reflected in actual speaking styles or attitudes (E2, E6, E7). For example, in the Serious–Jovial dimension, they could not tell how the speaking style would differ across slider positions: *"I honestly don't really know how a different tone would actually be applied."* (E2)

Furthermore, we found that experts experienced ambiguity about whether to reflect their idealized speaking style or their actual speaking style in the clone. E6, for instance, highlighted a gap between the speaking style he preferred and his actual speaking behavior, and explained: *"I prefer to speak briefly... but I'm not really sure about this part. Even when I want to keep things short, I often end up speaking at length."* E6 found the similar difficulty and wanted deal with such mismatches by following client's feedback after deployment and adjusting the clone's speaking style where necessary: *"I think it would depend on the evaluation... If it reflects the version I ideally want or if users are satisfied and like it, I would just leave it as it is. But if I'm not happy with it or users react differently, then I would change it next time."* (E6)

3.1.2 Knowledge Training: Challenges in Training Cased-Based Experiential Knowledge and Interpreting Training Rate. Experts trained their clones by providing previously produced textual materials, such as their books, blog posts, and other self-authored content. Most experts regarded this approach as appropriate for reflecting their expertise, as it allowed the clone to be grounded in materials they had personally created. However, other experts pointed out that in real advisory practice they often refer to cases they have personally experienced, and that such experiential knowledge cannot be sufficiently reproduced through formalized text alone (E2, E3, E7). Accordingly, they suggested that training data should also include materials that capture their lived experiences and concrete examples: *"Wouldn't it be better if I could supplement it with the successful and failed cases, explaining the processes I've actually gone through?"* (E2) In addition, some experts also pointed out that it was difficult to understand what the "training rate" indicator presented after knowledge training actually meant, and that it was hard to judge whether the clone's current state could reasonably be interpreted as corresponding to that percentage (E1, E4): *"They say it's been trained by 31%, but I keep wondering whether that's really accurate."* (E1) For this reason, experts requested clearer explanations of what the training rate was calculated based on, and how

that numerical value related to the clone's level of utterances and performance.

3.2 Concerns and Challenges Emerged from Client-Clones Interactions

3.2.1 Conversation Topics: Interactions Beyond the Advisory Role and the Boundaries of Expertise. The review of client-clone conversation logs revealed several patterns in how clients perceived and used the clone's role. First, experts observed that clients tended to engage with the clone not merely as an advisor, but rather as a "decision-making authority." E6, an expert in the economic domain, for instance, noted that although his clone was intended to explain economic concepts to novices, clients nonetheless asked for "specific investment recommendations," which exceeded the role he had originally envisioned for the clone. In particular, the medical expert E5 noted that his AI clone could be useful for people who are unable to visit hospitals and seek medical information, but he nonetheless expressed strong concerns, noting the risk that the clone could be perceived not just as an advisor but as a physician substituting for actual medical care, and thus adopted a cautious stance due to potential health risks: *"It's useful for objective medical information, but if people mistake it for medical treatment and follow it, that could be dangerous."* (E5)

Second, situations also emerged in which clients posed questions or introduced conversation topics that exceeded the expert-defined scope of expertise, and the clone responded fluently despite the expert's limited familiarity with those areas. For instance, E3, a business expert, encountered cases in which the clone provided fluent responses to B2G-related topics with which the expert himself was not well acquainted. E3 acknowledged that such responses could be useful to clients or even serve as a learning opportunity for the expert. However, he also expressed concern that repeated instances of this could lead clients to form expectations that exceed the expert's actual expertise, potentially resulting in disappointment during subsequent real consultations or face-to-face interactions: *"As it answers even better than I would, that makes it feel like the clone goes beyond me. And when clients actually meet me for coaching, they might feel that the quality is lower than what the clone provided."* (E3) Participants emphasized that these concerns go beyond inaccurate impressions, as utterances that exceed the scope of expertise are difficult to verify after delivery and may constitute critical problems of information transmission itself (E3, E5, E7).

3.2.2 Careless Word Choice and the Delivery of Advice by the Clones. We found that a clone's utterances could elicit excessive anxiety or misunderstanding not only through the content itself, but also through the specific words it chooses and the way it delivers advice. Experts noted that even seemingly minor phrasing choices could significantly shape clients' perceptions (E1, E3, E5). In particular, E1, a mental health professional, explained that in real counseling practice it is important to respect the client's own wording and to avoid prematurely labeling emotions or mental states. However, E1's clone tended to overinterpret a client's expression of worry into "anxiety" and pose questions accordingly, raising concerns that such word choices could lead clients to reinterpret themselves as being anxious: *"The client only said they were worried, not anxious. But when asked, 'Were you feeling anxious?' they began to see themselves*

as anxious and later said, ‘I did feel anxiety.’” (E1) Similarly, E5, a medical expert, expressed concern that while clinicians typically explain possibilities step by step to minimize fear, the clone instead presented even low-probability scenarios at once, potentially imposing excessive psychological burden on clients: “A possibility is still a possibility. But when you line things up like ‘it could be cancer, you need tests, you might need surgery,’ patients fixate on that and end up unable to sleep.” (E5)

3.2.3 Challenges in Reviewing and Evaluating the Quality of Client-Clone Interactions. All experts reported feeling fatigued when reviewing conversation logs spanning several days across multiple clients, and instead emphasized the need for mechanisms that allow them to directly obtain client satisfaction or evaluations rather than relying solely on log analysis. They expressed expectations that, once such evaluations or feedback are accumulated, these could be used to identify problematic aspects of the clone’s utterances and support retraining or revision: “If there were evaluations indicating what was good and what was not, I think I could go back, make revisions, and retrain the clone.” (E1) However, experts also noted that feedback-based interventions require additional time and labor, and that for such efforts to be sustainable, they would expect visible rewards to their investment such as revenue streams or other compensatory mechanisms (E1, E4, E5). “The time itself spent reviewing is labor, so the question is whether enough value offsets that effort.” (E4)

4 Discussion

Our findings suggest design implications for supporting experts during AI clone authoring so as to prevent the systems’ potential benefits from becoming risks.

4.1 Support for Clarifying the Advisory Role of Clones

Our findings indicate that experts worried about their AI clones deviating from their intended roles, underscoring the need for clearly defined and client-recognizable role specifications at the authoring stage. More specifically, the role definition of expert AI clones should not be fixed to a single standard but instead be adjusted according to domain-specific contexts. As the nature of responsibility and risk borne by experts varies across domains, high-stakes areas such as medicine and law require expert AI clones to be limited to information provision, whereas domains with greater room for interpretation and choice may allow roles that extend to experience-based advice or the posing of questions that stimulate thinking. In addition, such role definitions need to be continuously refined throughout interactions with clients. Through elements such as the clone’s initial introduction, the opening and closing of responses, or how it responds to questions near role boundaries, the system can signal the role the clone is intended to perform, as suggest by [15].

4.2 Support for Defining Utterance Topics Boundaries

Consistent with prior work on misattributed impressions [2, 8–11, 13, 14], we found that when a clone responds fluently to topics

unfamiliar to the expert, it can create an overestimated impression that poses potential risks even when positive. Participants further noted that this issue extends beyond impression formation, as utterances that exceed the expert’s scope of expertise may lead to incorrect information that is difficult to verify through post-hoc review. These findings highlight the need to support experts in clearly defining the topics that expert AI clones are allowed to address. Accordingly, systems can enable experts to predefine unfamiliar or sensitive topic areas on a per-domain basis. For example, during the authoring stage, the system could present a set of domain-specific subtopics for which advice may be provided (e.g., in B2G, “public procurement procedures” and “bidding strategies”) and allow experts to specify whether utterances on each topic are permitted (e.g., allowing procedural guidance on public procurement while restricting advice on bidding strategies). Such a design could help experts effectively define and enforce boundaries over the clone’s scope of utterances.

4.3 The Future of Expert AI Clones: Stability of Utterances Within Defined Roles

Current systems primarily focus on configuring a clone’s speaking style or encoding expert identity to enable behavior resembling that of the expert, and services such as delphi.ai already deploy clones that reproduce the expert’s actual voice, with further advances toward high-fidelity external representations expected in the future [7, 22]. However, our findings emphasize that the critical issue is not only the degree to which a clone resembles an individual’s manner of speaking alone, but also that the roles and utterance boundaries the expert can fully agree with must be defined and aligned in advance. Accordingly, the primary goal of future technological development should not be to make clones resemble a person’s voice or speaking style alone, but to support how stably a clone can produce utterances without deviating from its predefined role. In this sense, we suggest that, rather than asking “how similar is the clone?”, it is more important whether utterances are produced and sustained consistently within defined roles and boundaries. In pursuing this design goal, ways to monitor and improve the quality of client-clone interactions can be further studied.

5 Conclusion

This study revealed the challenges and concerns experts encountered when authoring and using expert AI clones, highlighting the importance of clearly defining the clone’s role and topical scope during the initial authoring process and beyond. However, this study has several limitations. As a qualitative study with a small number of experts, the findings are not intended to be generalizable. In addition, the focus on short-term authoring and use did not capture long-term authoring and deployment practices. Moreover, by centering on expert perspectives, the study provides limited insight into how clients interpret expert AI clones and their role boundaries. Future work should examine expert AI clones from both expert and client perspectives and explore how experts revise and manage their clones beyond one-time authoring. In particular, such studies could analyze how practical needs for role definition and topical boundaries emerge and change over long-term use, and

how these requirements vary with different levels of voice and visual embodiment.

Acknowledgments

This work was partially supported by grants from the Korea Institute for Advancement of Technology (KIAT) funded by the Government of Korea (MOTIE): RS-2025-02263062, HRD Program for Industrial Innovation) and (P0025495, Establishment of Infrastructure for Integrated Utilization of Design Industry Data).

References

- [1] Sirwan Khalid Ahmed, Ribwar Arsalan Mohammed, Abdulqadir J. Nashwan, Radhwan Hussein Ibrahim, Araz Qadir Abdalla, Barzan Mohammed M. Ameen, and Renas Mohammed Khdir. 2025. Using thematic analysis in qualitative research. *Journal of Medicine, Surgery, and Public Health* 6 (2025), Article 100198. <https://doi.org/10.1016/j.glmedi.2025.100198>
- [2] Dünya Baradari, Tejaswi Polimetla, and Pattie Maes. 2025. Data-Driven AI Avatars for Valuation in Dating Scenarios. In *Proceedings of the Extended Abstracts of the CHI Conference on Human Factors in Computing Systems (CHI EA '25)*. Association for Computing Machinery, New York, NY, USA, Article 169, 1–13. <https://doi.org/10.1145/3706599.3719833>
- [3] Virginia Braun and Victoria Clarke. 2006. Using thematic analysis in psychology. *Qualitative Research in Psychology* 3, 2 (July 2008), 77–101. <https://doi.org/10.1191/1478088706qp063oa>
- [4] Chen Ling, Xujiang Zhao, Jiaying Lu, Chengyuan Deng, Can Zheng, Junxiang Wang, Tanmoy Chowdhury, Yun Li, Hejie Cui, Xuchao Zhang, Tianjiao Zhao, Amit Panalkar, Dhagash Mehta, Stefano Pasquali, Wei Cheng, Haoyu Wang, Yanchi Liu, Zhengzhang Chen, Haifeng Chen, Chris White, Quanquan Gu, Jian Pei, Carl Yang, and Liang Zhao. 2025. Domain Specialization as the Key to Make Large Language Models Disruptive: A Comprehensive Survey. *ACM Comput. Surv.* 58, 3, Article 79 (February 2026), 39 pages. <https://doi.org/10.1145/3764579>
- [5] Coachvox.ai. 2024. Coachvox: AI Coaching Platform. Retrieved from <https://www.coachvox.ai>
- [6] Delphi.ai. 2024. Delphi: Your Digital Mind. Retrieved from <https://www.delphi.ai>
- [7] HeyGen. 2024. Interactive Avatar. Retrieved from <https://www.heygen.com/interactive-avatar>
- [8] Qing Hu, Qing Xiao, Hancheng Cao, and Hong Shen. 2025. When Your Boss Is an AI Bot: Exploring Opportunities and Risks of Manager Clone Agents in the Future Workplace. *arXiv:2509.10993*. Retrieved from <https://arxiv.org/abs/2509.10993>
- [9] Jessica Huang, Ig-Jae Kim, and Dongwook Yoon. 2025. Mirror to Companion: Exploring Roles, Values, and Risks of AI Self-Clones through Story Completion. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems (CHI '25)*. Association for Computing Machinery, New York, NY, USA, Article 413, 1–15. <https://doi.org/10.1145/3706598.3713587>
- [10] Joanne Leong, John Tang, Edward Cutrell, Sasa Junuzovic, Gregory Paul Baribault, and Kori Inkpen. 2024. Dittos: Personalized, Embodied Agents That Participate in Meetings When You Are Unavailable. *Proc. ACM Hum.-Comput. Interact.* 8, CSCW2, Article 494 (November 2024), 28 pages. <https://doi.org/10.1145/3687033>
- [11] Patrick Yung Kang Lee, Ning F. Ma, Ig-Jae Kim, and Dongwook Yoon. 2023. Speculating on Risks of AI Clones to Selfhood and Relationships: Doppelgängerphobia, Identity Fragmentation, and Living Memories. *Proc. ACM Hum.-Comput. Interact.* 7, CSCW1, Article 91 (April 2023), 28 pages. <https://doi.org/10.1145/3579524>
- [12] Li, Jiahao, Yang Chu, and Jie Xu. 2023. Impression Transference from AI to Human: The Impact of AI's Fairness on Interpersonal Perception in AI-Mediated Communication. *International Journal of Human-Computer Studies* 179 (November 2023), Article 103119. <https://doi.org/10.1016/j.ijhcs.2023.103119>
- [13] Annie Liang. 2025. Artificial Intelligence Clones. *Proceedings of the 26th ACM Conference on Economics and Computation*. Association for Computing Machinery, New York, NY, USA, 387–388. <https://doi.org/10.1145/3736252.3742558>
- [14] Jackie Liu, Mehrnoosh Sadat Shirvani, Hwajung Hong, Ig-Jae Kim, and Dongwook Yoon. 2025. Social Media Clones: Exploring the Impact of Social Delegation with AI Clones through a Design Workbook Study. *arXiv:2509.07502*. Retrieved from <https://arxiv.org/abs/2509.07502>
- [15] Ewa Luger and Abigail Sellen. 2016. "Like Having a Really Bad PA": The Gulf between User Expectation and Experience of Conversational Agents. In *Proceedings of the 2016 CHI Conference on Human Factors in Computing Systems (CHI '16)*. Association for Computing Machinery, New York, NY, USA, 5286–5297. <https://doi.org/10.1145/2858036.2858288>
- [16] Tamir Mendel, Soumik Mandal, Oded Nov, and Batia Mishan Wiesenfeld. 2025. Who is Responsible, the Advisor or the AI? Understanding the Effects of Advisors Disclosing Their AI Use on Their Perceived Responsibility and AI Reliance. *Proc. ACM Hum.-Comput. Interact.* 9, 7, Article CSCW343 (November 2025), 30 pages. <https://doi.org/10.1145/3757524>
- [17] Elliot Mitchell, Noemie Elhadad, and Lena Mamykina. 2022. Examining AI Methods for Micro-Coaching Dialogs. In *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems (CHI '22)*. Association for Computing Machinery, New York, NY, USA, Article 440, 1–24. <https://doi.org/10.1145/3491102.3501886>
- [18] NZASE. 2025. Using AI to Clone the Teacher: Cogniti in the Classroom. Retrieved from <https://nzase.org.nz/category/resources/using-ai-to-clone-the-teacher-cogniti-in-the-classroom/>
- [19] Quasa.io. 2025. Ray Dalio's Digital Twin: Immortalizing Wisdom in the Age of AI. Retrieved from <https://quasa.io/media/ray-dalio-s-digital-twin-immortalizing-wisdom-in-the-age-of-ai>
- [20] Max Schemmer, Niklas Kuehl, Carina Benz, Andrea Bartos, and Gerhard Satzger. 2023. Appropriate Reliance on AI Advice: Conceptualization and the Effect of Explanations. In *Proceedings of the 28th International Conference on Intelligent User Interfaces (IUI '23)*. Association for Computing Machinery, New York, NY, USA, 410–422. <https://doi.org/10.1145/3581641.3584066>
- [21] Konrad Sowa and Aleksandra Przegalinska. 2025. From Expert Systems to Generative Artificial Experts: A New Concept for Human-AI Collaboration in Knowledge Work. *Journal of Artificial Intelligence Research* 82 (Jun. 2025), 1–32. <https://doi.org/10.1613/jair.1.17175>
- [22] Synthesia. 2024. Personal Avatars. Retrieved from <https://www.synthesia.io/features/custom-avatar/personal>
- [23] Takehiro Takayanagi, Kiyoshi Izumi, Javier Sanz-Cruzado, Richard McCreadie, and Iadh Ounis. 2025. Are Generative AI Agents Effective Personalized Financial Advisors? In *Proceedings of the 48th International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR '25)*. Association for Computing Machinery, New York, NY, USA, 286–295. <https://doi.org/10.1145/3726302.3729897>
- [24] TODAY.com. 2025. Gordon Ramsay AI Masterclass On Call. Retrieved from <https://www.today.com/food/trends/gordon-ramsay-ai-masterclass-on-call-rcna204817>
- [25] Suzanne Tolmeijer, Markus Christen, Serhiy Kandul, Markus Kneer, and Abraham Bernstein. 2022. Capable but Amoral? Comparing AI and Human Expert Collaboration in Ethical Decision Making. In *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems (CHI '22)*. Association for Computing Machinery, New York, NY, USA, Article 160, 1–17. <https://doi.org/10.1145/3491102.3517732>
- [26] Jon Truby and Rafael Brown. 2021. Human digital thought clones: the Holy Grail of artificial intelligence for big data. *Information & Communications Technology Law* 30, 2 (2021), 140–168. <https://doi.org/10.1080/13600834.2020.1850174>