

Developmentally Safe Generative AI Environment for Youth

Jake Chanenson
Amyoli Internet Research Lab
University of Chicago
Chicago, Illinois, USA
jchanen1@uchicago.edu

Yaman Yu*
School of Information Sciences
University of Illinois at Urbana
Champaign
Champaign, Illinois, USA
yamanyu2@illinois.edu

Jessica Vitak
College of Information
University of Maryland
College Park, Maryland, USA
jvitak@umd.edu

Sheena Erete
College of Information Studies
University of Maryland College Park
College Park, Maryland, USA
serete@umd.edu

Tamara Clegg
College of Information Studies /
College of Education
University of Maryland
College Park, Maryland, USA
tclegg@umd.edu

Diana Freed
Department of Computer Science
Brown University
Providence, Rhode Island, USA
diana_freed@brown.edu

Marshini Chetty
Department of Computer Science
University of Chicago
Chicago, Illinois, USA
marshini@uchicago.edu

Abstract

Generative AI (GenAI) systems such as ChatGPT, Gemini, Character.AI, and LLaMA are becoming tightly woven into youths' everyday lives. Young people now turn to these tools not only for homework help and creative projects, but also for companionship, identity exploration, and emotionally charged conversations. These always-available, highly responsive systems can support learning and self-expression, yet they also introduce new and poorly understood risks for cognitive, emotional, social, moral, and identity development. At the same time, caregivers, educators, designers, and policymakers often have limited visibility into how youth actually use GenAI and few shared frameworks for supporting developmentally appropriate, safe, and empowering engagement. This workshop aims to critically explore the holistic impact of GenAI on youth and to build a shared, developmentally grounded agenda for research and design. Our goals are to: (1) identify promising research approaches and ethical, participatory methods for studying youth-GenAI interactions in situ; (2) surface open challenges, hidden risks, and unintended impacts of GenAI use across diverse cultures and contexts; (3) explore community-based and longitudinal research models that connect youth, caregivers, educators, designers, and policymakers; (4) co-develop shared infrastructures and evaluation metrics to guide design, education, and policy, such as datasets, toolkits, forums, and indicators of developmental safety and well-being. We will bring together researchers and practitioners

from HCI, education, child development, security and privacy, and policy to collaboratively sketch actionable strategies for developmentally safe GenAI environments that balance risk mitigation with youth agency, autonomy, and space for exploration and growth.

CCS Concepts

• **Human-centered computing**; • **Security and privacy** → **Human and societal aspects of security and privacy**;

Keywords

Youth, Generative AI, AI Safety, Privacy

ACM Reference Format:

Jake Chanenson, Yaman Yu, Jessica Vitak, Sheena Erete, Tamara Clegg, Diana Freed, and Marshini Chetty. 2026. Developmentally Safe Generative AI Environment for Youth. In *Extended Abstracts of the 2026 CHI Conference on Human Factors in Computing Systems (CHI EA '26)*, April 13–17, 2026, Barcelona, Spain. ACM, New York, NY, USA, 5 pages. <https://doi.org/10.1145/3772363.3778682>

Background

Generative AI (GenAI) systems such as ChatGPT, Gemini, Character AI, and LLaMA are rapidly changing how youth interact with technology. Beyond productivity and learning support [22], these systems are increasingly integrated into youths' lives as creative partners, conversational companions, and sources of emotional connection [33]. Many youth now engage with GenAI through conversations that feel personal, responsive, and meaningful [6], introducing new opportunities and challenges for well-being and development [32].

Prior work has mapped online risks through the “4Cs” and related youth-safety frameworks [11, 18–20, 31], and broader AI risk taxonomies highlight bias, toxicity, privacy, and misinformation concerns [12, 13, 24, 27, 34]. Unlike earlier digital technologies,

*Corresponding author



This work is licensed under a Creative Commons Attribution 4.0 International License. *CHI EA '26, Barcelona, Spain*

© 2026 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-2281-3/26/04

<https://doi.org/10.1145/3772363.3778682>

GenAI systems generate real-time, personalized responses that evolve based on user input, creating experiences that are dynamic and adaptive [8, 9, 18, 31]. The fluid, adaptive nature of GenAI can blur boundaries between reality and simulation, influence values and behavior, and shape how young people see themselves and others [15, 16, 33]. Additionally, recent research shows that youth are using GenAI in increasingly nuanced and emotionally charged ways. The shift to always-available conversational agents reshapes these risks in form and intensity, with early evidence of emotionally charged uses such as seeking social support and identity exploration, alongside over-trust and “empathy-gap” concerns [5, 15, 16, 21, 25, 32]. Incident repositories further surface emergent harms in youth contexts [1–4]. Youth-centered risk taxonomy is beginning to shed light on these dynamics [32], but they also reveal growing concerns about how GenAI may shape cognitive, emotional, social, moral, and identity development [14–17, 23, 25]. These findings highlight an urgent need to guide and safeguard youth in an era of always-available and seemingly empathetic AI systems.

At the same time, caregivers and educators often lack visibility into these experiences and face challenges in offering appropriate support [9, 11, 33]. Despite growing interest in the space, many existing strategies for supporting youth remain fragmented [26]. Families often rely on trial-and-error approaches such as limiting screen time or checking chat histories, which rarely address the nuanced risks of GenAI use or uphold youths’ autonomy and agency [28–31]. Educational institutions, meanwhile, are still adapting to the broader digital shift and often lack the tools or frameworks needed to respond to this new mode of interaction. Platform-level safeguards tend to focus on content moderation or usage restrictions, without fully addressing the developmental contexts in which youth are engaging.

These gaps highlight that no single stakeholder, whether family, school, or platform, can fully account for the risks and opportunities of GenAI in youths’ lives. Youth move fluidly across home, school, peer, and online spaces, and their experiences are shaped by the interplay of these environments. Prior work demonstrates that community-based research can address complex socio-technical challenges by centering those most affected and fostering collaboration across diverse stakeholders [7, 10]. To meet the moment, we must move beyond fragmented responses and foster a collaborative approach that brings together researchers, caregivers, educators, designers, and other youth-focused professionals to co-create safe, developmentally appropriate, inclusive and empowering GenAI experiences for the next generation. Advancing such an approach calls for grounded, interdisciplinary research that not only examines the risks and challenges of youth engagement with GenAI, but also supports the creation of shared frameworks and actionable strategies. Key priorities include developing research methods that are ethical, participatory, and adaptable; identifying risks and benefits through real-world use; and building shared infrastructures to support design, education, and policy development. This workshop aims to offer a space to collectively advance these goals and shape a research agenda rooted in care, collaboration, and youth agency.

Workshop Goals

This workshop aims to critically explore the holistic impact of GenAI on youth, examining the ways in which these technologies influence social, cognitive, and emotional development, as well as the overall well-being of young people. The overarching objective is to assemble a diverse group of researchers and practitioners who are invested in youth-centered inquiry, thereby fostering new dialogue and collaborations.

Through this workshop, participants will have cross-disciplinary discussions that seek to advance research agendas grounded in the needs, rights, and agency of youth in relation to GenAI. Working collectively, the group will identify, adapt, and promote best practices for both the study and design of GenAI systems. Together, we will engage in critical discussions about ethical research methodologies, responsible approaches, and the key priorities that should guide the examination and implementation of GenAI technologies with youth in real-world environments.

The workshop is guided by the following research questions:

- (1) **Promising research approaches:** What research methods and frameworks are most effective for studying youth interactions with GenAI in-situ? How can we facilitate ethical, meaningful, and safe data contributions from youth participants?
- (2) **Open challenges and future directions:** Which critical research gaps remain in understanding youth engagement with GenAI systems? How can we collectively define, prioritize, and address these open challenges?
- (3) **Community-based research models:** In what ways can community workshops, longitudinal partnerships, or participatory networks accelerate ethical, equitable, and scalable GenAI research with youth?
- (4) **Surfacing risks and unintended impacts:** How can we collaboratively identify hidden risks or unintended consequences of GenAI use among youth, especially across diverse cultures and contexts?
- (5) **Sustaining research infrastructures:** What types of infrastructures (e.g., shared datasets, toolkits, forums) are needed to support and connect the community working at the intersection of GenAI and youth well-being?
- (6) **Safe and effective evaluation metrics:** What are appropriate metrics and indicators for evaluating the safety and positive impact of GenAI usage for youth?

Workshop

Workshop Activities

Workshop Length. This workshop will run as a single in-person event consisting of two sessions (180 minutes plus a short break), structured to maximize participation, knowledge exchange, and collaborative problem-solving. See Table 1 for the schedule overview.

Session 1. The first session sets the stage by establishing shared context and connections among participants. We begin with a short welcome and framing overview that links back to the goals of the workshop: addressing developmental safety, surfacing open challenges, and building collective approaches to youth and generative

Table 1: The workshop schedule is presented in relative time intervals, with 0:00 marking the start of the workshop. All subsequent activities are shown as durations measured from this starting point rather than by absolute clock times.

Time	Activity
0:00–0:20	Overview, Opening Keynote and Logistics
0:20–1:00	Lightning Talks
1:00–1:30	Synergy Circles
1:30–1:40	Break
1:40–2:45	Breakout Discussions
2:45–3:00	Report Out & Future Directions

AI. A keynote from Jessica Vitak will then highlight current empirical and policy perspectives, offering provocation for further discussion.

Participants will then engage in Lightning Talks, with each person given three minutes to share their research, practical experience, or perspective. This format introduces the diversity of expertise in the room, surfaces common themes, and sparks ideas for deeper exploration. Participants will then engage in Synergy Circles, a structured space that enables participants to connect quickly with others who share similar interests or complementary approaches based on the Lightning Talks. We close Session 1 with a short break.

Session 2. The second session forms the heart of the workshop and is designed to maximize active engagement. Rather than remaining at the level of abstract discussion, participants will work together in small groups through a series of carefully structured activities. Based on a brief pre-workshop survey, participants will be grouped by shared interests and complementary perspectives to ensure inclusive and engaging discussions. Each group will be assigned a facilitator and a list of discussion prompts to guide the conversation. Each breakout group will be organized around one of the workshop’s core themes, ensuring that diverse expertise is distributed and that every participant has an opportunity to contribute.

We will begin by posing a concrete design or problem-solving prompt to each group. For example, groups examining research approaches will be asked to sketch a sample study design for in-situ youth–AI research and annotate how ethical safeguards could be embedded at each stage. Groups considering future challenges will be given case cards that illustrate real-world dilemmas (e.g., a teenager seeking mental health support from an AI chatbot) and will work to identify unanswered research questions these scenarios raise.

To make the session interactive and generative, each group will be provided with visual mapping tools (sticky notes, templates, markers) to capture ideas in real time. Groups focused on community-based models will map stakeholders such as schools, libraries, or youth clubs, and then draw network diagrams to show how partnerships could be structured. Groups tackling risks and unintended impacts will construct scenario timelines, tracing how small AI interactions might escalate into larger developmental or cultural harms. Others focusing on research infrastructures will draft mock-up outlines of shared datasets, benchmarks, or toolkits,

while the evaluation metrics group will brainstorm candidate indicators, sorting them into short-term and long-term measures of youth safety and well-being.

In the final stage, each group will present back to the full room with a short set of visuals and recommendations. A facilitator will then guide a collective synthesis, inviting cross-group reflections and discussion of connections. This closing activity will not only consolidate insights into a coherent whole but also culminate in the drafting of a shared vision statement for Youth×AI research, alongside a list of actionable next steps and commitments for collaboration beyond the workshop. These next steps will foster community and connection subsequent to the workshop.

Accessibility

The organizing team is committed to making the workshop fully accessible for all participants. All presentations and discussion prompts will include live captioning via PowerPoint, and written materials will be distributed in advance in both Markdown and PDF formats. We will also:

- (1) Ensure slides and handouts use accessible fonts, high-contrast color schemes, and alt text for all images.
- (2) Provide clear verbal descriptions of visuals during presentations.
- (3) Arrange room layouts to support wheelchair accessibility and inclusive group discussions.
- (4) Offer flexible participation modes within group activities (e.g., written, verbal, or visual contributions) to accommodate different communication preferences.
- (5) Designate an accessibility contact person on the organizing team to respond to requests and provide support before and during the workshop.

Plans to Publish Proceedings

We will designate note takers for all breakout discussions to ensure detailed documentation of insights and recommendations. After the workshop, the organizing team will synthesize these notes into a summary that highlights best practices, emerging challenges, and future research directions. The synthesized workshop summary will be shared on the workshop website for broader visibility, and we will consider submitting an extended version of the outcomes to a conference or journal to further advance community discussion. In addition, all accepted submissions (position papers, design proposals, policy case studies, and statements of interest) will be published on our workshop website with participant slide decks and materials.

Offline Materials

PDFs of the lightning talk slides, submitted by participants in advance, will be made available. In addition, PDFs of the breakout session notes will be provided to allow attendees to review the material afterward. The session will not be recorded, in order to foster an environment where participants can engage fully and comfortably.

Organizers

The organizers of this workshop bring together deep expertise in children's privacy, online safety, and emerging work on AI and AI-related risks. Their backgrounds make them uniquely positioned to guide the discussion and ensure the workshop addresses the complex intersections of youth well-being, safety, and Generative AI.

Jake Chanenson

Jake is a Ph.D. student in Computer Science at the University of Chicago, specializing in law-informed research at the intersection of human-computer interaction (HCI), privacy, and technology policy. With a Masters in Law from the University of Chicago Law School, Jake combines technical and legal expertise to address complex questions related to youth privacy, AI, and policy. Jake is an affiliate at the University of North Carolina's Center for Information, Technology, and Public Life (CITAP) and their research has been published in top academic venues, including ACM's CHI, PETS, and has received Best Paper Honorable Mention. Their research interests center around children's online safety and privacy.

Yaman Yu

Yaman is a Ph.D. candidate in Information Sciences at the University of Illinois Urbana-Champaign. She is also a 2025 Google PhD Fellow and MIT EECS Rising Star. Her research explores the intersection of responsible AI and governance, security and privacy, and HCI, with a particular focus on safeguarding youth in generative AI environments. She investigates risks, designs safeguards, and builds frameworks that promote youth-centered safety and agency in AI systems. Her work has been published in top venues including the IEEE S&P, CCS, and ACM CHI, and has received a CHI Best Paper Honorable Mention.

Jessica Vitak

Jessica is a Professor in the College of Information at the University of Maryland, College Park and an affiliate Professor in UMD's Communication Department and the Joint Program in Survey Methodology (JPSM). She is the faculty director of the iSchool's PhD program and Director of the Human-Computer Interaction Lab (HCIL). She is also affiliated with the Social Data Science Center (SoDA), Center for Trustworthy AI in Law and Society (TRAILS), and the Values-Centered AI (VCAI) group at UMD. Her research focuses on networked privacy and pervasive data ethics. She has a long history of working to understand the social impacts of new technologies on youth and how that impacts their privacy and security.

Sheena Erete

Sheena is an Associate Professor in the College of Information at the University of Maryland, College Park, the Associate Director of Research for the Artificial Intelligence Interdisciplinary Institute at Maryland (AIM), and the founding director of the Community Research, Equity, and Design (CREED) Collective. Her research focuses on co-designing socio-cultural technologies, practices, and policies with community residents to amplify their local efforts related to topics such as community safety, education, civic engagement, and health.

Tamara (Tammy) Clegg

Tammy is a Professor in the College of Information at the University of Maryland, College Park, the Director of the Bachelor of Arts degree in technology and information design, and the Director of the University of Maryland's Computing Catalyst. Her work creates innovative learning experiences and technologies to help youth and underrepresented learners engage in STEM through everyday life experiences. She also studies ways such learning environments help learners develop STEM identities and dispositions.

Diana Freed

Diana is an Assistant Professor of Computer and Data Science at Brown University. She is also a CRA and Microsoft Trustworthy AI Research Fellow, a Visiting Scholar at Harvard Law School's Petrie-Flom Center, and a Faculty Associate at the Berkman Klein Center for Internet & Society. Her research is at the intersection of HCI, computer security and privacy, and technology policy. A central thread of her work investigates how generative AI impacts youth digital safety, raising new challenges for security, privacy, and policy.

Marshini Chetty

Marshini is an Associate Professor in the Department of Computer Science at the University of Chicago. Her research focuses on making the Internet a more trustworthy place including helping to educate youth about online privacy and security.

Call for Participation

Generative AI is rapidly redefining the digital environment of youth, presenting both exciting opportunities for creativity and learning, as well as new challenges for privacy, safety, and agency. As platforms like ChatGPT and Character.AI become embedded in everyday life, questions multiply: What distinguishes actual risk from ordinary adolescent experimentation? In what ways do generative AI tools amplify youth voice and autonomy, and where do they present unforeseen vulnerabilities? What frameworks can help educators, researchers, designers, and guardians navigate this evolving landscape responsibly?

This workshop invites contributions from researchers and designers who are interested in exploring the complex interplay between generative AI, risk, and youth experience. We seek to bring together diverse perspectives, including empirical research, design prototypes, and policy innovations. Specific questions the workshop will address include: how youth engage with generative AI platforms; the real-world risks, and opportunities these tools present; and methods for conducting responsible, privacy-preserving research and intervention. As such, we welcome all interested attendees and especially encourage participation from researchers working in child-centered computing, safety/security/privacy, or AI. If you are exploring new directions in your research and are interested in Youth × AI Safety, we also invite you to apply.

We invite submissions in the form of short position papers (one to two pages), design proposals, statement of interest, or policy case studies. All participants are encouraged to actively contribute and help build a shared roadmap for responsible, ethical, and empowering youth-AI engagement. The submissions will be reviewed

by the workshop organizers and program committee. At least one author of each accepted submission must attend the workshop, which consists of two 90-minute sessions. All accepted submissions will be published through our workshop website and presented at the workshop. We expect around 40 in-person attendees for the workshop, all of whom will participate onsite. Join us to co-create the future of youth and AI, distinguishing risk from agency, and carving out space for youth exploration and growth.

References

- [1] [n. d.]. AIAAIC. <https://www.aiaaic.org/home>.
- [2] [n. d.]. AVID Incident Database. <https://avidml.org/database/>.
- [3] AI Algorithmic and Automation Incidents Repository. 2023. Student Violates 50 Lancaster County School Students with Nude Deepfakes. <https://tinyurl.com/ykcad37r>. Accessed: 2025-02-12.
- [4] AI Algorithmic and Automation Incidents Repository. 2024. Character AI Hosts Paedophile and Suicide Chatbots. <https://tinyurl.com/bdyayc4z>. Accessed: 2025-02-12.
- [5] Carlos Caetano, Gabriel O dos Santos, Caio Petrucci, Artur Barros, Camila Laranjeira, Leo Sampaio Ferraz Ribeiro, Júlia Fernandes de Mendonça, Jeferson A dos Santos, and Sandra Avila. 2025. Neglected Risks: The Disturbing Reality of Children's Images in Datasets and the Urgent Call for Accountability. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency*. 2542–2553.
- [6] Common Sense Media. 2025. Nearly 3 in 4 Teens Have Used AI Companions, New National Survey Finds. <https://www.commonsensemedia.org/press-releases/nearly-3-in-4-teens-have-used-ai-companions-new-national-survey-finds>. Press release.
- [7] Sasha Costanza-Chock. 2020. *Design Justice: Community-Led Practices to Build the Worlds We Need*. The MIT Press, Cambridge, MA.
- [8] Lorrie Faith Cranor, Adam L Durity, Abigail Marsh, and Blase Ur. 2014. {Parents'} and {Teens'} Perspectives on Privacy In a {Technology-Filled} World. In *10th Symposium On Usable Privacy and Security (SOUPS 2014)*. 19–35.
- [9] Maria Eira, Amir Kaveh Rasouli, and Vicky Charisi. 2025. Parents' Perceptions About the Use of Generative AI Systems by Adolescents. In *Proceedings of the 24th Interaction Design and Children*. 927–931.
- [10] Sheena Erete, Aarti Israni, and Tawanna Dillahunt. 2018. An Intersectional Approach to Designing in the Margins. *Interactions* 25, 3 (April 2018), 66–69. <https://doi.org/10.1145/3194349> <https://dl.acm.org/doi/10.1145/3194349>.
- [11] Diana Freed, Natalie N Bazarova, Sunny Consolvo, Eunice J Han, Patrick Gage Kelley, Kurt Thomas, and Dan Cosley. 2023. Understanding digital-safety experiences of youth in the US. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*. 1–15.
- [12] Isabel O Gallegos, Ryan A Rossi, Joe Barrow, Md Mehrab Tanjim, Sunghul Kim, Franck Dernoncourt, Tong Yu, Ruiyi Zhang, and Nesreen K Ahmed. 2024. Bias and fairness in large language models: A survey. *Computational Linguistics* (2024), 1–79.
- [13] Samuel Gehman, Suchin Gururangan, Maarten Sap, Yejin Choi, and Noah A. Smith. 2020. RealToxicityPrompts: Evaluating Neural Toxic Degeneration in Language Models. In *Findings*. <https://api.semanticscholar.org/CorpusID:221878771>
- [14] Ariel Han, Xiaofei Zhou, Zhenyao Cai, Shenshen Han, Richard Ko, Seth Corrigan, and Kylie A Pepler. 2024. Teachers, Parents, and Students' perspectives on Integrating Generative AI into Elementary Literacy Education. In *Proceedings of the CHI Conference on Human Factors in Computing Systems*. ACM, 1–17. <https://doi.org/10.1145/3613904.3642438>
- [15] Nomisha Kurian. 2024. 'No, Alexa, no!': designing child-safe AI and protecting children from the risks of the 'empathy gap' in large language models. *Learning, Media and Technology* (2024). <https://api.semanticscholar.org/CorpusID:271158326>
- [16] Nomisha Kurian. 2025. AI's empathy gap: The risks of conversational Artificial Intelligence for young children's well-being and key ethical considerations for early childhood education and care. *Contemporary Issues in Early Childhood* 26, 1 (2025), 132–139.
- [17] Hannah R Lawrence, Renee A Schneider, Susan B Rubin, Maja J Matarić, Daniel J McDuff, and Megan Jones Bell. 2024. The opportunities and risks of large language models in mental health. *JMIR Mental Health* 11, 1 (2024), e59479.
- [18] Sonia Livingstone. 2008. Taking risky opportunities in youthful content creation: teenagers' use of social networking sites for intimacy, privacy and self-expression. *New media & society* 10, 3 (2008), 393–411.
- [19] Sonia Livingstone and Leslie Haddon. 2008. Risky experiences for children online: Charting European research on children and the internet. *Children & society* 22, 4 (2008), 314–323.
- [20] Sonia Livingstone, Leslie Haddon, Anke Görzig, and Kjartan Ólafsson. 2011. Risks and safety on the internet: the perspective of European children: full findings and policy implications from the EU Kids Online survey of 9–16 year olds and their parents in 25 countries. (2011).
- [21] Masab Mansoor, Ali Hamide, and Tyler Tran. 2025. Conversational AI in Pediatric Mental Health: A Narrative Review. *Children* 12, 3 (2025), 359.
- [22] Ofcom. 2023. Gen Z Driving Early Adoption of Gen AI, Our Latest Research Shows. <https://www.ofcom.org.uk/news-centre/2023/gen-z-driving-early-adoption-of-gen-ai> Accessed: 2024-06-03.
- [23] Jinkyung Park, Vivek Singh, and Pamela Wisniewski. 2024. Toward Safe Evolution of Artificial Intelligence (AI) based Conversational Agents to Support Adolescent Mental and Sexual Health Knowledge Discovery. *arXiv.org* (2024). <https://doi.org/10.48550/ARXIV.2404.03023>
- [24] Peter Slattery, Alexander K Saeri, Emily AC Grundy, Jess Graham, Michael Noetel, Risto Uuk, James Dao, Soroush Pour, Stephen Casper, and Neil Thompson. 2024. The ai risk repository: A comprehensive meta-review, database, and taxonomy of risks from artificial intelligence. *arXiv preprint arXiv:2408.12622* (2024).
- [25] Jaemarie Solyst, Ellia Yang, Shixian Xie, Jessica Hammer, Amy Ogan, and Motahare Eslami. 2024. Children's Overtrust and Shifting Perspectives of Generative AI. *arXiv* (2024). <https://doi.org/10.48550/ARXIV.2404.14511>
- [26] Elizabeth A Sweigart, Jinkyung Katie Park, and Pamela J Wisniewski. 2025. It takes a village: Youth online safety research highlights need for interdisciplinary, multistakeholder solutions. *Journal of Advances in Information Technology* 16, 1 (2025).
- [27] Laura Weidinger, John Mellor, Maribeth Rauh, Conor Griffin, Jonathan Uesato, Po-Sen Huang, Myra Cheng, Mia Glaese, Borja Balle, Atoosa Kasirzadeh, et al. 2021. Ethical and social risks of harm from language models. *arXiv preprint arXiv:2112.04359* (2021).
- [28] Pamela Wisniewski, Arup Kumar Ghosh, Heng Xu, Mary Beth Rosson, and John M Carroll. 2017. Parental control vs. teen self-regulation: Is there a middle ground for mobile online safety?. In *Proceedings of the 2017 ACM conference on computer supported cooperative work and social computing*. 51–69.
- [29] Pamela Wisniewski, Haiyan Jia, Na Wang, Saijing Zheng, Heng Xu, Mary Beth Rosson, and John M Carroll. 2015. Resilience mitigates the negative effects of adolescent internet addiction and online risk exposure. In *Proceedings of the 33rd Annual ACM Conference on Human Factors in Computing Systems*. 4029–4038.
- [30] Pamela Wisniewski, Heng Xu, Mary Beth Rosson, Daniel F Perkins, and John M Carroll. 2016. Dear diary: Teens reflect on their weekly online risk experiences. In *Proceedings of the 2016 CHI Conference on Human Factors in Computing Systems*. 3919–3930.
- [31] Pamela J Wisniewski, Heng Xu, Mary Beth Rosson, and John M Carroll. 2014. Adolescent online safety: the "moral" of the story. In *Proceedings of the 17th ACM conference on Computer supported cooperative work & social computing*. 1258–1271.
- [32] Yaman Yu, Yiren Liu, Jacky Zhang, Yun Huang, and Yang Wang. 2025. Youth-Centered GAI Risks (YAIR): A Taxonomy of Generative AI Risks from Empirical Data. (2025), 150–165.
- [33] Yaman Yu, Tanusree Sharma, Melinda Hu, Justin Wang, and Yang Wang. 2025. Exploring parent-child perceptions on safety in generative AI: concerns, mitigation strategies, and design implications. In *2025 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2735–2752.
- [34] Yi Zeng, Kevin Klyman, Andy Zhou, Yu Yang, Minzhou Pan, Ruoxi Jia, Dawn Song, Percy Liang, and Bo Li. 2024. Ai risk categorization decoded (air 2024): From government regulations to corporate policies. *arXiv preprint arXiv:2406.17864* (2024).