

Human Expertise for AI Red-Teaming and Scalable Evaluation

Alice Qian
Human-Computer Interaction
Institute
Carnegie Mellon University
Pittsburgh, Pennsylvania, USA
aqzhang@andrew.cmu.edu

Srravya Chandhramowuli
Institute for Design Informatics
University of Edinburgh
Edinburgh, United Kingdom
srravya.c@ed.ac.uk

Laura Dabbish
Human-Computer Interaction
Institute
Carnegie Mellon University
Pittsburgh, Pennsylvania, USA
dabbish@cs.cmu.edu

Hong Shen
Human-Computer Interaction
Institute
Carnegie Mellon University
Pittsburgh, Pennsylvania, USA
hongs@cs.cmu.edu

Alex S Taylor
Design Informatics
University of Edinburgh
Edinburgh, Scotland, United Kingdom
alex.taylor@ed.ac.uk

Ding Wang
Google
Atlanta, Georgia, USA
ding.w9911@gmail.com

Theodora Skeadas
Humane Intelligence
Cambridge, Massachusetts, USA
theodora@humane-intelligence.org

Bolor-Erdene Jagdagdorj
Microsoft
seattle, Washington, USA
bjagdagdorj@microsoft.com

Abstract

Rapid adoption of generative AI has outpaced the infrastructure needed to red team systems responsibly. This workshop tackles a core tension: scaling AI red teaming while centering human expertise and well-being. We convene academic, industry, and nonprofit practitioners for two threads. (A) Vision: surface high-level goals and principles for effective, humane red teaming. (B) Build: identify opportunities to support human-AI red teaming, such as scenario libraries, role prompts for red teamers, and calibration methods that align automated efforts with human expertise. Through this workshop, we will develop a vision for the future of effective AI red teaming that leverages and protects human expertise while meeting the needs of evaluation at scale.

CCS Concepts

• Security and privacy; • Human-centered computing → Collaborative and social computing; Collaborative and social computing; • Social and professional topics → Computing / technology policy; Computing industry; Computing / technology policy;

Keywords

artificial intelligence, red teaming, AI red teaming, labor, fairness, well-being, security, AI safety, AI ethics

ACM Reference Format:

Alice Qian, Srravya Chandhramowuli, Laura Dabbish, Hong Shen, Alex S Taylor, Ding Wang, Theodora Skeadas, and Bolor-Erdene Jagdagdorj. 2026.



This work is licensed under a Creative Commons Attribution 4.0 International License.
CHI EA '26, Barcelona, Spain
© 2026 Copyright held by the owner/author(s).
ACM ISBN 979-8-4007-2281-3/26/04
<https://doi.org/10.1145/3772363.3778702>

Human Expertise for AI Red-Teaming and Scalable Evaluation. In *Extended Abstracts of the 2026 CHI Conference on Human Factors in Computing Systems (CHI EA '26)*, April 13–17, 2026, Barcelona, Spain. ACM, New York, NY, USA, 5 pages. <https://doi.org/10.1145/3772363.3778702>

1 Motivation

Rapid adoption of generative AI has brought not only new opportunities but also new risks and harms. As these systems enter sensitive domains from education and health to companionship and work, they are implicated in psychological harm, misinformation, and exploitation. Recent lawsuits and investigations have highlighted cases of people harmed by unsafe or manipulative outputs [1, 4], as well as the toll on those tasked with moderating or evaluating harmful content generated by these systems [15, 16]. These incidents underscore that evaluation is not an academic afterthought: it is central to responsible AI development and deployment.

AI red teaming has emerged as one of the most important practices for identifying risks in generative systems before they escalate. Red teaming refers to adversarial testing, where evaluators intentionally probe models for vulnerabilities ranging from unsafe outputs to misuse potential [7]. It draws inspiration from military and cybersecurity contexts, where simulated attacks are used to reveal weaknesses before they can be exploited. In the AI context, red teaming is now becoming standard practice in large technology companies, where in-house teams, outsourced contractors, and crowdsourced workers are employed to test models [25]. National governments have also promoted red teaming as a cornerstone of AI safety and security [2], and public initiatives such as the DEFCON 2023 red teaming challenge and NIST's ARIA competition have further demonstrated its importance [14, 23]. For the CHI community, this positions AI red teaming as a site where interaction design, user research, and evaluation methods can materially shape safety outcomes.

Despite its promise, AI red teaming today faces a fundamental tension: the need to scale while retaining quality and expertise. Generative models operate across countless contexts, languages, and domains, making it difficult for a limited group of human evaluators to anticipate every possible failure mode. Scaling requires broader participation, yet effective red teaming often benefits from deep expertise such as domain knowledge, contextual awareness, creativity, and sensitivity to harms [25]. Balancing throughput and expertise is therefore a socio-technical challenge that invites HCI approaches to methods, tooling, recruitment, and workflow design [20].

At the same time, the well-being of red teamers themselves has often been overlooked. Because red teaming requires evaluators to engage directly with harmful, biased, or manipulative model outputs, they face risks of psychological stress, burnout, or trauma [15, 16]. These dynamics echo prior findings in content moderation and data annotation research, where workers bore hidden costs in maintaining platform safety [3, 9, 10, 21, 22, 27]. Without deliberate design for risk disclosure, rotation, training, recovery, and participatory safeguards, efforts to scale red teaming risk reproducing these harms at a larger scale.

At a broader level, scaling AI red teaming presents a design and methods challenge for CHI. Policy mentions have outpaced agreement on what counts as sufficient practice, how results should be captured, and how evaluator well-being is protected. Current practices range from short “bug bounty” contests to extensive in-house evaluations, making it difficult for policymakers or the public to know whether a system has been responsibly tested [13, 26]. Engagement from the HCI and CHI community can motivate future collaborative efforts across policymakers, academics, and industry practitioners to define usable protocols, reporting artifacts, and participatory safeguards so that AI red teaming operates as an accountable socio-technical practice.

This workshop responds to these gaps by convening a cross-sector community from academia, industry, and nonprofits to define the future of scalable, human-centered red teaming. We aim to:

- (1) Surface shared challenges and articulate high-level goals for red teaming infrastructures;
- (2) Identify the tooling and well-being needs of red teamers, including supports for sustainable labor and mechanisms to integrate automation responsibly.

2 Workshop Themes

Building on the challenges outlined above, this workshop will structure discussion and activities around several key themes. Each theme highlights a different aspect of scaling AI red teaming while keeping human expertise and well-being at the center.

- **Automation levels and trade-offs in red teaming:** What is the right balance between human and automated evaluation? We will explore different levels of automation in red teaming pipelines, from model-in-the-loop probing to fully manual expert review, and debate the trade-offs between scalability, coverage, validity, and trustworthiness. A central question is when automation can augment expertise and when human oversight is indispensable.
- **Evaluator well-being and labor support:** Red teamers are regularly exposed to harmful or sensitive model outputs,

which can lead to psychological stress and burnout. This theme will address how we can better support the people doing this work, including risk disclosure, rotation and recovery practices, compensation structures, and training. We will consider lessons from related domains such as content moderation and annotation labor.

- **Unmet tooling and design needs:** Today’s red teaming practices are often improvised, lacking standardized workflows, interfaces, or collaborative supports. This theme will focus on identifying gaps in the current ecosystem—such as scenario libraries, annotation tools, triage and adjudication pipelines, and dashboards—that would enable more effective and humane red-teaming at scale. We will also discuss how to design these tools to amplify expertise rather than deskill or overburden evaluators.

3 Organizers

This workshop brings together a diverse group committed to effective evaluation and auditing of AI systems. The team’s academic work examines safety practices in organizations and crowd settings [16, 19], labor and infrastructures in data production and their entanglements with responsibility and power [28], collective intelligence and user driven auditing for AI systems [6, 17, 24], and the social and organizational foundations of emerging technologies including open source sustainability, collaborative work, and the role of expertise in responsible AI [5, 8, 12, 29]. Our industry organizers contribute hands-on AI red teaming and trust and safety program leadership across major platforms and nonprofits, connecting practice and research. Collectively, we have organized workshops and special interest groups at HCI venues such as CSCW and HCOMP. Bringing these perspectives together, we aim to connect HCI and AI communities, center labor and data governance alongside safety and auditing, and catalyze concrete collaboration toward human-centered AI red teaming practices.

Alice Qian is a PhD student at Carnegie Mellon University’s Human-Computer Interaction Institute, where her research focuses on examining methods of leveraging and amplifying the human infrastructure of responsible AI development. She has specific experience studying AI safety practices within organizational settings [16] as well as through crowdsourcing efforts [19]. She has co-organized two previous CSCW workshops focusing on AI red teaming. She also has experience red teaming in practice as well as in the analysis of results.

Srravya Chandhiramowuli examines the role of human values in data annotation and AI development. Her current research examines the work of data annotation for AI, paying particular attention to systemic challenges and frictions, to envision and inform equitable futures of AI. She is a PhD candidate in the Institute for Design Informatics at the University of Edinburgh.

Alex Taylor has been contributing to Science & Technology Studies and HCI for over twenty years. His interests are in how digital technologies are co-constitutive of forms of knowing and doing, and, as a consequence, provide a basis for fundamental transformations in society. He is a Reader in Design Informatics at the University of Edinburgh.

Bolor-Erdene Jagdagdorj is a member of the AI Red Team at Microsoft, where her work focuses on ensuring AI safety and security and ethical technologies. Prior to joining Microsoft, she earned a Master's Degree in Information Security from Carnegie Mellon University with her thesis in measuring and analyzing abortion data privacy through third party tracking. She has presented about AI Red Teaming at multiple conferences including USENIX and Grace Hopper Conference.

Ding Wang is a HCI researcher at Google Research, focusing on the intersection of HCI and AI. Her work examines the labor, infrastructures, and human practices in data production for AI systems, exploring issues of responsibility, representation, and power [28]. Previously, she was a postdoctoral researcher at Microsoft Research India, working on projects related to the future of work and digital health. Ding earned her PhD from Lancaster University, where her research critically examined smart city development and alternative approaches to system design and evaluation.

Theodora Skeades is Chief of Staff at Humane Intelligence. Previously at Twitter, she managed the Trust and Safety Council and a program for human rights defenders. She also worked in national security at Booz Allen Hamilton, focusing on disinformation, and with nonprofits in Morocco, Turkey, Greece, and Costa Rica.

Hong Shen is an Assistant Professor at Carnegie Mellon University, where she directs the CARE Lab and conducts research on collective intelligence and responsible AI practices. Her past work has explored how to better support users and the public in auditing, red teaming, and aligning frontier AI systems [6, 17, 24]. She has organized workshops and Special Interest Groups on Red Teaming at CSCW 2025, on User-Driven AI Auditing at HCOMP 2024, and on Participatory Generative AI Design at AAAI 2024.

Laura Dabbish is a Professor at Carnegie Mellon University's Human-Computer Interaction Institute and Director of the Connected Experience Lab, which investigates the social, organizational, and policy dimensions of emerging technologies with a focus on the invisible human labor behind AI systems. She has conducted research on open source software sustainability [8], collaborative work practices [5, 12], and the role of human expertise in responsible AI development [29].

4 Plans to Publish Workshop Proceedings

We anticipate approximately 40–60 attendees. To ensure the selection of workshop participants with relevant experiences, we ask participants to submit a position paper (1–4 pages) on their motivation for attending the workshop, relevant prior experience in either AI red teaming or research, and expected contribution to the workshop. We designed this format to elicit participation so as to make the workshop accessible to both academic participants and those from industry, such as AI red teamers. Workshop organizers will put out calls to solicit submissions to communities of academics, AI red teamers, and other relevant spaces in industry and non-profit sectors. Submissions will be reviewed by the organizer and selected based on quality and for diversity of participant experiences and backgrounds. Authors of selected submissions will be invited to participate in our workshop. Each participant must submit their

own statement of interest. Through this method, we hope to cultivate a participant group with varying perspectives that will better enrich discussions.

The aim of this workshop is to convene a community concerned with effective scaling of human-centered AI red teaming. As such, throughout the workshop we will note ideas and discussion topics in a shared document. All participants will have continued access to this document after the workshop. We will also create an online space for participants to continue discussions and collaborations (e.g., a Discord server). To facilitate open discussion beyond the workshop, accepted papers from participants will be published on *ceur-ws*¹. Additionally, we plan to facilitate and encourage attendees to extend their work for future publication in HCI journals.

5 Workshop Length

We propose a long workshop format totaling 180 minutes of interactive sessions. A longer format is essential for balancing high-level visioning with concrete design work. The first session will introduce the concept of AI red teaming and facilitate broader reflection on the challenges and goals of scaling it responsibly. The second session will focus more narrowly on specific opportunities for developing tools and human–AI interaction methods to support red teaming in practice. This structure allows participants to first align on shared goals and then transition into actionable design discussions. The workshop will also include a keynote, time for participant introductions, and informal networking (e.g., over lunch), ensuring a mix of framing, collaboration, and community-building.

6 Workshop Activities

We design the workshop to balance framing, interactive discussion, and concrete agenda-setting. The tentative schedule is as follows:

- **Introduction (30 min):** The organizers will welcome participants, introduce the workshop goals, and facilitate a round of participant introductions. We will also provide an overview of the current state of AI red teaming, highlighting both the promise and the challenges of scaling while centering human expertise.
- **Session 1 — Envisioning AI Red Teaming at Scale (60 min):** In this session, participants will be divided into small breakout groups assigned based on prior experience to discuss different levels of automation and human-in-the-loop design for red teaming. Groups will identify trade-offs in coverage, validity, and expertise, and propose guiding principles for when automation is appropriate and when expert oversight is required. Each group will share back to the large group, and facilitators will synthesize themes into a draft set of key future directions for scaling AI red teaming responsibly.
- **Break (30 min):** During the break between the two 90-minute workshop sessions, workshop participants will be encouraged to engage in informal networking as a way to continue conversations.
- **Keynote (30 min):** Organizers with AI red teaming experience will present insights on the future of AI evaluation and human–AI collaboration, drawing on experience from

¹<https://ceur-ws.org/>

industry red-teaming practices and/or HCI research. A short Q&A will follow to allow participants to connect the keynote to the workshop themes.

- **Session 2 – Building Tools and Support Infrastructure (60 min):** The second session will focus on concrete opportunities to develop red-teaming tools and infrastructures that amplify expertise while protecting well-being. Participants will work in breakout groups on specific prompts such as: What new tools are needed to support human-AI collaboration in red teaming? How might we design workflows to protect evaluator well-being? What governance or policy interventions could help? Groups will produce short sketches (e.g., protocol outlines, tool wishlists) and share 1-3 takeaways with the large group.

The day will close with a brief reflection, where organizers will summarize key takeaways and invite participants to continue conversations via the workshop website and shared communication channels.

7 Workshop Outcomes

We expect this workshop to generate both tangible and long-term outcomes. In the short term, we will synthesize the discussions into a shared report that summarizes key challenges, guiding principles for scaling AI red teaming, and a set of proposed directions for tool development and evaluator support. Information about workshop participants will be made available on the workshop website. In the longer term, we aim to foster a cross-sector community of practice around human-centered red teaming. Participants will be invited to collaborate on future outputs, such as an ACM Interactions article or white paper, a repository of red-teaming protocols and tool wishlists, and potential collaborative grant proposals. By convening stakeholders from academia, industry, and nonprofits, the workshop will help seed sustained collaborations and provide a roadmap for scaling red teaming responsibly, without sacrificing expertise, rigor, or care.

8 Accessibility

We will adhere to ACM SIGCHI accessibility guidelines. We will survey registered attendees in advance to identify any specific needs (e.g., live captioning, sign language interpreters, alternative text). We plan to provide live captioning for spoken content (e.g., keynote and discussions) and to make any notes available asynchronously for review. These measures will help ensure inclusive participation and accommodate diverse attendees.

9 Call for Participation

AI red teaming stands as a key method for identifying risks in AI systems, but it remains unclear how AI red teaming can be scaled to meet organizational demands while centering on human expertise and retaining the support of red teamers. To explore this problem, we invite participants to join a workshop at CHI on effective human-involved AI red teaming. The workshop will combine keynote presentations, breakout discussions, and collaborative design activities to address the challenges of expanding red teaming capacity while protecting evaluator expertise and well-being. Our goals are (1) to articulate shared challenges and high-level visions for scalable AI

red teaming, and (2) to identify concrete opportunities for tooling, governance, and labor support. This workshop welcomes participants from academia, industry, government, and nonprofits who are engaged in AI evaluation, safety, and governance. We encourage submissions from AI developers, AI red teamers, HCI researchers, policymakers, and community advocates.

Submission Requirements: To participate, please submit a position paper (1-4 pages) in the ACM Master Article Format² describing motivation for participation, relevant background, and expected contribution to the workshop discussion on our workshop website³. Expected contributions could range from ideas based on lived experiences, red teaming, to proposed research methods to apply to unanswered questions in the problem of effectively scaling human-centered red teaming. Submissions will be reviewed by organizing members to be selected based on quality and to ensure participation from participants with a wide range of backgrounds. We expect to accept 40-60 participants based on physical space constraints. At least one author of each submission must be able to attend the workshop in person. Accepted papers will be published via the workshop website.

10 Note About Past Workshops

This workshop builds on previous successfully facilitated workshops at HCI and CSCW venues focused on human-centered AI red teaming, including “The Human Factor in AI Red Teaming” [18], “The Work of AI: Mapping Human Labour in the AI Pipeline” [11], and “The Work of AI Red Teaming,” which explored human involvement in AI red teaming and the broader labor infrastructures behind AI development including AI red teaming. Our workshop extends these efforts by focusing specifically on the challenges of scaling red-teaming efforts and ensuring human-centered practices.

Acknowledgments

Efforts in organizing this workshop have been supported by the National Science Foundation (NSF) Graduate Research Fellowship Program (GRFP) under Award No. DGE2140739 and the National Institute of Standards and Technology (NIST) under Federal Award ID Number 60NANB24D231 and Carnegie Mellon University⁴ AI Measurement Science and Engineering Center (AIMSEC).

References

- [1] BBC News. 2025. *Parents of teenager who took his own life sue OpenAI*. <https://www.bbc.com/news/articles/cgerwp7rdlvo>
- [2] Joseph R. JR Biden. 2023. Executive Order on the Safe, Secure, and Trustworthy Development and Use of Artificial Intelligence. <https://bidenwhitehouse.archives.gov/briefing-room/presidential-actions/2023/10/30/executive-order-on-the-safe-secure-and-trustworthy-development-and-use-of-artificial-intelligence/>
- [3] Sakine Mohamadi Bozorg. 2024. *In the Colony of Blooded Eyes: A Trilogy on Working Conditions as a Content Moderator in Germany*. Data Workers' Inquiry. <https://data-workers.org/sakine/>
- [4] CBS News. 2025. *Parents of teens who died by suicide after AI chatbot interactions testify in Congress*. <https://www.cbsnews.com/news/ai-chatbots-teens-suicide-parents-testify-congress/>
- [5] L. Dabbish, C. Stuart, J. Tsay, and J. Herbsleb. 2012. Social coding in GitHub: transparency and collaboration in an open software repository. In *Proceedings of the ACM 2012 conference on computer supported cooperative work*. 1277–1286.
- [6] Xianzhe Fan, Qing Xiao, Xuhui Zhou, Jiaxin Pei, Maarten Sap, Zhicong Lu, and Hong Shen. 2025. User-Driven Value Alignment: Understanding Users'

²<https://chi2025.acm.org/submission-guides/chi-publication-formats/>

³<https://hearts-workshop.github.io>

⁴<https://ror.org/05x2bcf33>

- Perceptions and Strategies for Addressing Biased and Discriminatory Statements in AI Companions. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*. 1–19.
- [7] Michael Feffer, Anusha Sinha, Wesley H Deng, Zachary C Lipton, and Hoda Heidari. 2024. Red-Teaming for generative AI: Silver bullet or security theater?. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, Vol. 7. 421–437.
 - [8] H. Frluckaj, L. Dabbish, D. G. Widder, H. S. Qiu, and J. D. Herbsleb. 2022. Gender and participation in open source software development. *Proceedings of the ACM on Human-Computer Interaction* 6, CSCW2 (2022), 1–31.
 - [9] Fasica Berhane Gebrekidan. 2024. *Content Moderation: The harrowing, traumatizing job that left many African data workers with mental health issues and drug dependency*. Data Workers' Inquiry. <https://data-workers.org/fasica/>
 - [10] Shagun Jhaver, Amy Bruckman, and Eric Gilbert. 2019. Does transparency in moderation really matter? User behavior after content removal explanations on reddit. *Proceedings of the ACM on Human-Computer Interaction* 3, CSCW (2019), 1–27.
 - [11] Airi Lampinen, Rob Comber, Sravya Chandhiramowuli, Naja Holten Møller, and Alex S Taylor. 2024. The Work of AI: Mapping Human Labour in the AI Pipeline. In *Companion Publication of the 2024 Conference on Computer-Supported Cooperative Work and Social Computing*. 728–731.
 - [12] R. Li, P. Pandurangan, H. Frluckaj, and L. Dabbish. 2021. Code of conduct conversations in open source software projects on GitHub. *Proceedings of the ACM on Human-computer Interaction* 5, CSCW1 (2021), 1–31.
 - [13] Jacob Metcalf, Elizabeth Anne Watkins, Emanuel Moss, and Madeleine Clare Elish. 2023. Recommendations for Using Red-Teaming for AI Accountability. <https://datasociety.net/wp-content/uploads/2023/10/Recommendations-for-Using-Red-Teaming-for-AI-Accountability-PolicyBrief.pdf> Policy Brief.
 - [14] Will Oremus. 2023. *Meet the hackers who are trying to make AI go rogue*. <https://www.washingtonpost.com/technology/2023/08/08/ai-red-team-defcon/>
 - [15] Sachin R Pendse, Darren Gergle, Rachel Kornfield, Jonah Meyerhoff, David Mohr, Jina Suh, Annie Wescott, Casey Williams, and Jessica Schleider. 2025. When Testing AI Tests Us: Safeguarding Mental Health on the Digital Frontlines. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency*. 1793–1804.
 - [16] Alice Qian, Judith Amores, Hong Shen, Mary Czerwinski, Mary L Gray, and Jina Suh. 2025. AURA: Amplifying Understanding, Resilience, and Awareness for Responsible AI Content Work. *Proceedings of the ACM on Human-Computer Interaction* 9, CSCW2 (2025), 1–45.
 - [17] Alice Qian, Judith Amores, Hong Shen, Mary Czerwinski, Mary L Gray, and Jina Suh. 2025. AURA: Amplifying understanding, resilience, and awareness for responsible AI content work. *Proceedings of the ACM on Human-Computer Interaction* 9, 2 (2025), 1–45.
 - [18] Alice Qian, Ryland Shaw, Jacy Reese Anthis, Ashlee Milton, Emily Tseng, Jina Suh, Lama Ahmad, Ram Shankar Siva Kumar, Julian Posada, Benjamin Shestakofsky, et al. 2024. The human factor in ai red teaming: Perspectives from social and collaborative computing. In *Companion Publication of the 2024 Conference on Computer-Supported Cooperative Work and Social Computing*. 712–715.
 - [19] Alice Qian, Ryland Shaw, Laura Dabbish, Jina Suh, and Hong Shen. 2025. Locating Risk: Task Designers and the Challenge of Risk Disclosure in RAI Content Work. *arXiv preprint arXiv:2505.24246* (2025).
 - [20] Alice Qian, Jina Suh, Mary L Gray, and Hong Shen. 2025. Effective Automation to Support the Human Infrastructure in AI Red Teaming. *Interactions* 32, 4 (2025), 58–61.
 - [21] Bothokwa Ranta. 2024. *The Unknown Women of Content Moderation*. Data Workers' Inquiry. <https://data-workers.org/ranta/>
 - [22] Sarah T Roberts. 2019. *Behind the screen*. Yale University Press.
 - [23] Reva Schwartz, Gabriella Waters, Razvan Amironesei, Craig Greenberg, Jon Fiscus, Patrick Hall, Anya Jones, Shomik Jain, Afzal Godil, Kristen Greene, et al. 2024. The Assessing Risks and Impacts of AI (ARIA) Program Evaluation Design Document. (2024).
 - [24] Hong Shen, Alicia DeVos, Motahhare Eslami, and Kenneth Holstein. 2021. Everyday algorithm auditing: Understanding the power of everyday users in surfacing harmful algorithmic behaviors. *Proceedings of the ACM on Human-Computer Interaction* 5, CSCW2 (2021), 1–29.
 - [25] Ranjit Singh, Borhane Bili-Hamelin, Carol Anderson, Emnet Tafesse, Briana Vecchione, Beth Duckles, and Jacob Metcalf. 2025. Red-Teaming in the Public Interest. *New York: Data & Society Research Institute* (2025).
 - [26] Araz Taeihagh. 2025. Governance of Generative AI. *puaf001* pages.
 - [27] Ding Wang, Shantanu Prabhat, and Nithya Sambasivan. 2022. Whose AI Dream? In search of the aspiration in data annotation.. In *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems* (New Orleans, LA, USA) (CHI '22). Association for Computing Machinery, New York, NY, USA, Article 582, 16 pages. doi:10.1145/3491102.3502121
 - [28] Ding Wang, Shantanu Prabhat, and Nithya Sambasivan. 2022. Whose AI Dream? In search of the aspiration in data annotation.. In *Proceedings of the 2022 CHI conference on human factors in computing systems*. 1–16.
 - [29] D. G. Widder, D. Nafus, L. Dabbish, and J. Herbsleb. 2022. Limits and possibilities for “ethical ai” in open source: A study of deepfakes. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency*. 2035–2046.