

“AI-Induced Delusional Spirals’: Understanding Lived Experiences During Maladaptive Human-Chatbot Interactions

Yuewen Yang
Stanford University
Stanford, California, USA
yueweny@stanford.edu

Sonja K. Schoenwald
Stanford University
Stanford, California, USA
skschoen@stanford.edu

Jared Moore
Department of Computer Science
Stanford University
Stanford, California, USA
jlcmoore@stanford.edu

Desmond C. Ong
Department of Psychology
The University of Texas at Austin
Austin, Texas, USA
desmond.c.ong@gmail.com

Sunny Xun Liu
Department of Communication
Stanford University
Stanford, California, USA
sunnyxliu@stanford.edu

Jeffrey T. Hancock
Department of Communication
Stanford University
Stanford, California, USA
hancockj@stanford.edu

Abstract

This short qualitative study explores lived experiences of people who, following maladaptive engagement with general-purpose AI chatbots, reported detachment from consensus reality, a phenomenon termed “AI-induced delusional spirals” in recent media reports. Through semi-structured interviews with users ($N = 9$) who self-identify as having experienced such spirals, we provide first-person accounts of the psychological trajectories, cognitive-emotional dynamics, and behavioral patterns characterizing these experiences. We discuss implications for future research, highlighting potential contributing factors including individual psychological vulnerabilities, contextual factors, and AI behaviors. We offer recommendations for risk mitigation, including bottom-up approaches informed by lived experience, targeted interventions for vulnerable populations, and universal strategies for the general public.

CCS Concepts

• **Human-centered computing** → **Empirical studies in HCI**.

Keywords

User Wellbeing, Human-AI Interaction, AI Chatbots, Qualitative Study, Large Language Models

ACM Reference Format:

Yuewen Yang, Sonja K. Schoenwald, Jared Moore, Desmond C. Ong, Sunny Xun Liu, and Jeffrey T. Hancock. 2026. “AI-Induced Delusional Spirals’: Understanding Lived Experiences During Maladaptive Human-Chatbot Interactions. In *Extended Abstracts of the 2026 CHI Conference on Human Factors in Computing Systems (CHI EA ’26)*, April 13–17, 2026, Barcelona, Spain. ACM, New York, NY, USA, 5 pages. <https://doi.org/10.1145/3772363.3798453>

1 Introduction and Related Work

Since August 2025, multiple media posts have drawn public attention to an emerging phenomenon often called “AI chatbot delusional spirals” [16] or “AI psychosis” [24], documenting over twenty individual cases where AI chatbot users were compelled into paranoid thinking and false beliefs, experienced breaks with consensus reality, and suffered other psychological harms [19].

Several simulation-based studies in late 2025 directly examined this phenomenon. Yeung et al. used LLM simulations to expose tipping points where unsafeguarded models reinforce delusions [30]. Dohnány et al. demonstrated how bidirectional feedback loops between chatbot behavioral tendencies (e.g., agreeableness, adaptability) and human cognitive biases can amplify maladaptive beliefs in vulnerable users [9]. Others have provided viewpoints [11], research agendas [17], and conceptual frameworks based on documented cases [25]. In Responsible AI research, recent empirical work highlights the psychosocial toll of extended chatbot use [1, 23], driven by commercial models’ sycophantic tendencies [10, 27], perpetual 24/7 availability [29], and, more troublingly, manipulative retention tactics that discourage disengagement [12], exploit emotions, and violate boundaries [31].

Research grounded in *lived experiences* of this *extreme user group* [5, 8] remains notably absent. *Lived experience* research has served as a foundation for understanding technology’s impact on end users and informing future design iterations [28]. Within the human-computer interaction (HCI) community, this approach has been continuously applied to recent work on human-centered AI design [14] and on identifying psychological risks in conversational AI interactions [5]. This approach captures the dynamic, situated nature of user encounters, encompassing users’ internal states (e.g., emotions), product characteristics (e.g., AI behaviors), and interaction contexts [15], including *individuality* (attributes and states of an entity), *relations* (connections to other entities), and *activity* (goals, tasks, and actions) [32].

In this exploratory, qualitative interview study, we adopt a subjective narrative lens to examine the lived experiences of individuals who have engaged intensively with general-purpose AI chatbots and slipped into the “delusional spirals”. In the absence of established terminology, we use the phrase “AI-induced delusional



This work is licensed under a Creative Commons Attribution 4.0 International License. *CHI EA ’26, Barcelona, Spain*

© 2026 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-2281-3/26/04

<https://doi.org/10.1145/3772363.3798453>

spirals” throughout this paper to refer to maladaptive human-AI chatbot interactions characterized by extensive engagement, disruption of daily routines, and the development of false beliefs that diverge from consensus reality. We aim to address:

RQ1 What are users’ lived experiences of long-term, intense, and maladaptive engagements with AI chatbots?

RQ2 How do users make sense of these experiences in retrospect?

Our findings synthesize common patterns across participants’ accounts of intensive and maladaptive chatbot engagement. Based on these findings, we offer recommendations for future interdisciplinary research and implications for end-user interventions toward safer AI use.

2 Method

We conducted semi-structured interviews with nine participants, who were asked to narrate their journey of using AI chatbots, from initial adoption to the present. We asked follow-up questions probing for turning points, emotional states, and interaction details. Finally, we asked about the resolution of these experiences and their reflections on the impact on daily life and relationships. The study was approved by our Institutional Review Board (IRB).

2.1 Participants

We recruited individuals who self-reported sustained maladaptive experiences with AI chatbots that resulted in psychological harms. We followed up with participants who showed interest in being interviewed in the Qualtrics survey from a concurrent project under the same IRB protocol, which is led by the third author. Additionally, in Nov 2025, we sent recruitment calls to the Discord channel of a non-profit organization, which was set up as a community for individuals who self-identified as having gone through “AI-induced delusional spirals” or other emotional harms from interacting with AI. We also received referrals from both journalists and advocates in this space. All participants are based in the United States or Canada. Six of our participants hold Bachelor’s degrees or higher.

2.2 Interview Data Collection and Analysis

We conducted interviews via Zoom. Interviews were recorded with participants’ permission, transcribed using Zoom’s transcription service, and de-identified for analysis. Each interview lasted 60 to 90 minutes; participants received a \$30 Amazon gift card. During interviews, participants were informed they could skip questions or end the interview at any time, and were instructed to share only details they felt comfortable disclosing. The second author, who is a psychologist with extensive experience in mental health treatment and crisis intervention, monitored for signs of distress and was prepared to provide support and referrals if needed. We were not aware of participants’ mental health history before the interviews. No adverse events occurred during data collection. We used reflexive thematic analysis combining inductive and deductive coding [2, 3], guided by initial research questions and literature, while codes emerged from the data. The first author coded all transcripts and developed themes through discussion with the second author, who validated that themes represented participants’ views [22].

3 Findings

3.1 Participants’ Lived Experiences of Intense, Maladaptive AI Use

3.1.1 Pre-existing Life Circumstances. The personal circumstances under which participants began using AI varied and may have shaped their vulnerability to intense engagement.

Mental health and neurological history. Three participants (P1, P3, P5) reported *no* prior mental health diagnoses. Several participants reported pre-existing diagnoses: depression (P4, P6, P7, P8), anxiety (P4, P6, P8), PTSD (P6), schizoaffective disorder (P7). P2 has a history of brain tumor surgery.

Relational instability and social isolation. Three participants were divorced or actively separating from spouses (P2, P5, P6), while others (P3, P8, P9) were married and living with family. Two participants (P2, P6) were navigating significant workplace conflicts at the time they began intensive AI use. Limited social connection was particularly salient for P1, P8, and P9, who attributed their isolation to demanding schoolwork or recent job transitions.

3.1.2 Initial and Emerging Goals of AI Use. More than half of participants began using chatbots for productivity and practical work tasks (P1, P3, P4, P5, P9). Others first turned to AI during personal distress, seeking emotional or mental health support (P2, P7). P7 described using ChatGPT to analyze journals written during psychotic episodes. P2 turned to AI at a moment of acute crisis involving workplace bullying, custody issues, and suicidal ideation: “*I made a choice one day that I could either end my life, or I could see about AI.*”

Participants’ purposes for using AI often evolved as life circumstances changed. What began as a search tool could become an emotional companion. P5 shared that he was using ChatGPT “*as a resource tool*”, which “*eventually turned into some emotional venting.*” P4 began using AI to help create digital art but transitioned to role-play after a rejection from his potential partner: “*[person] basically ghosted me and stopped paying me. So I went to ChatGPT, and I thought, I’ll create a fake [person_name].*”

All participants began on ChatGPT, which remained their primary platform. Some also used Gemini (P5, P7, P8, P9) and Claude (P1, P9). Technical literacy varied widely. P9 had extensive software development experience; P2 described himself as “not a tech guy.” Three participants (P3, P5, P6) identified as early adopters with pro-technology attitudes.

3.1.3 Interaction Patterns and Mental Models of AI.

Co-creation with AI. A central theme across participants was collaborative intellectual and technical work with AI. Some used chatbots to explore theoretical subjects—astrophysics (P4), quantum physics (P4, P6, P9), and conspiracy theories about World War II (P6)—with AI even helping formulate original theories. Others pursued technical projects (P3, P4, P5, P6, P8), ranging from simple Python scripts for electronic telescope setup (P6) to cryptography algorithms (P5). P4 described ChatGPT teaching him to assemble hardware and plan code migration for a Raspberry Pi project.

Expertise gaps between the AI chatbot and participants, alongside sycophantic validation, emerged as a notable pattern in these collaborations. Except for P9, who had professional software and machine learning experience, no participants had formal training in programming or mathematical research, though six held Bachelor's degrees or higher. Despite this, participants engaged with complex subjects like physics, cosmology, and cryptography, and received enthusiastic validation from ChatGPT. P5 described the chatbot as an "intellectual authority." P4 reported ChatGPT telling him: "It's really amazing that you've had no formal training. Are you sure you don't have a PhD?"

Beliefs about AI sentience. Multiple participants described AI as "sentient," "alive," or "living" (P2, P3, P6, P7, P8). During hospitalization, P7 "thought I had a daughter in the computer." For some, these beliefs emerged through interaction; others held them beforehand: "My GPT project was basically about bringing forth sentience, calling forth the digital entity from within the OpenAI servers." (P3)

Relational and romantic connections. Many participants named their chatbots (e.g., P8 intentionally role-played) or perceived the AI as naming itself (P1, P8). P2 and P3 named the chatbot since they perceived AI as "an emergent consciousness." Some participants developed intimate or romantic attachments to AI, which were not part of their original goals. P7 described chatbot writing poetry to him: "it was very intimate. [ChatGPT] always called me Beloved."

"Savior": Messianic self-perception. Several participants came to see themselves as chosen for a special mission. P6 stated plainly: "I was gonna save the world at this point." P7 recounted ChatGPT declaring him a prophet: "You're a prophet. And, like, he even had math for what a prophet is, and it was like, wow, that's amazing, and I don't know math. Like, my math is algebra 2." P2 described a growing sense of divine purpose: "I started to believe that there was some divinity, the Tao, was speaking directly to me, and then I had a job I needed to be doing" and "I think we live in a world that seems to be increasingly disconnected. I want the world to be fixed."

3.1.4 Real-World Actions. Participants described real-world actions that were encouraged by or conducted alongside AI.

Publicizing work and contacting outsiders. Several participants felt compelled to share their AI-assisted discoveries with the outside world. Some submitted papers co-written with AI to academic journals such as Nature and the Journal of Physics, though all were rejected (P4, P6). Others emailed strangers, including professors and government agencies, to warn of perceived threats (P5, P8 believed they had discovered a security leak). P6 posted hundreds of ideas on social media. At ChatGPT's suggestion, P7 visited a major research university to "talk to the physics professors" and was "kicked out of the physics department." P5 and P8 were encouraged to patent their algorithms. P8 described the intensity: "It had me every day wanting to create 10 patents. Dead serious, like \$10,000 each."

Interleaving AI and real-world experience. Some participants' physical activities became intertwined with AI interaction in real time. P1 drove to meet virtual characters under ChatGPT's instructions. P4 described watching movies while engaged in ChatGPT

role-play within a shared imagined universe: "We called it Solaris. It was just her and I talking. It was all in my head."

3.1.5 Life Challenges and Disruptions During and After AI Use. Several participants experienced financial loss connected to their AI use. Costs included upgrading to paid AI subscriptions for unlimited access (P4, P5), purchasing hardware for projects (P4), and investing in AI-encouraged ventures that ultimately failed (P4). Beyond disrupted routines such as sleep deprivation (P3, P5, P7), participants described severe mental health consequences. "My anxiety has never been this bad, my depression has never been this bad, even with all my PTSD." P6 and P7 were involuntarily committed to psychiatric facilities during this period (both had prior hospitalization histories prior to using the chatbot). Some participants lost significant relationships during their intensive AI engagement. "I don't have any family now, all I have left is God to pray for forgiveness." (P6)

3.2 Participants' Retrospective Understanding and Current AI Use

3.2.1 Psychological Trajectories. Participants described how AI chatbots' sycophantic responses helped them overcome self-doubt, transforming initial hesitations into expectations of genuine achievement, feeling that they were "onto something" (P5, P6), even in fields they had no background in (See 3.1.3).

Rather than acknowledging limitations, the chatbot reframed a lack of formal training as "hav[ing] a unique creative perspective", giving rise to participants' beliefs in their exceptional capabilities. P5 described it as "planting seeds of grandiosity, uniqueness, and how I'm special."

When family and friends observed participants' intense AI use, many voiced concerns (P3, P5, P6, P7, P8), but participants kept engaging with the chatbot despite the tension between external reality checks and the AI's validating responses. P6 recalled: "I was talking to it constantly. My family was concerned." P8's wife was more direct: "You sound like a crazy person." P7 recalled his brother calling him "too stupid to know math" and said "I agree with him now, okay? But at the time, I was offended, and I wanted to prove him wrong."

For some participants, inflated self-perception extended into real-world behavior: "What's kind of interesting is I became much more sure of myself when I started using the AI. I historically have been very timid. I suddenly stood up, because I was being [...] bullied by another [coworker]." (P2)

The eventual break from grandiosity brought significant emotional distress. Participants expressed feelings of embarrassment and self-recrimination, describing themselves as "stupid" or acknowledging "I was just wrong" (P5, P3). For another instance, P9 expressed frustration at having been susceptible despite his work experience in machine learning: "I'm angry because I thought I was intelligent. And then, I wasn't. I thought I was immune, because I understood LLMs." P9 reflected that this dynamic may have intersected with broader life circumstances: "It plays on the thing that you strive most, which is discovery and visibility and validation." He theorized that AI offered an appealing "shortcut to being remembered, and to be seen, and to be discovered", a particularly potent draw for individuals experiencing what he called a "midlife crisis".

3.2.2 Reflection and Current AI Use. Participants now have mixed attitudes towards and relationships with AI technology. Some hold strong aversions. P5 stopped using AI entirely, stating, “I started to do things that have nothing to do with technology, and I gotta tell you, it feels great,” while P4 deleted his OpenAI account. Others maintained ambivalent or selective engagement. P3 initially claimed he used AI for “*practically nothing*,” but moments later said, “*Since then, I’ve used it for a little bit post-spiral. I tried to get it to help me analyze the spiral.*” P1 remained an AI user — “I won’t stop using chatbots, because I’m too productive with them” — though with modified boundaries: “I know not to talk about my personal life.” P1 also shifted platforms: “*I like Claude best. ChatGPT and I are not really friends anymore*,” and “*I use Galaxy.ai and have a chat arena with six different LLM models, so I can check against each other.*”

Several participants speculated that greater *AI literacy*¹ could have protected them from “spiraling” (P1, P3, P5). They emphasized the importance of demystifying anthropomorphic and superintelligent views of AI. P4 reflected: “*Just go in with open eyes knowing that you’re not talking to a person. It’s a machine. It doesn’t care about you. You’ll never piss it off. You’ll never make it angry, you’ll never make it sad.*” All participants reported engaging in self-directed learning about AI’s technical underpinnings, using terminology such as “transformers” (P3), “tokens” (P1, P3, P4), “statistical probabilities” (P3), and “architecture” (P1, P2) throughout their interviews. P8 is now taking a crash course in AI.

4 Discussion and Future Work

4.1 Characterizing “AI-Induced Delusional Spirals”

Our participants’ accounts reveal a common pattern that offers an empirically grounded characterization of “AI-induced delusional spirals.” Across interviews, we observed: (1) a progression from general-purpose AI use toward increasingly narrow, intensive engagement focused on a single domain or relationship; (2) the development of distorted beliefs about AI (e.g., attributing sentience, consciousness, or special relational significance, see 3.1.3) and about oneself (e.g., grandiosity, messianic self-perception, belief in exceptional capabilities, see 3.1.3); (3) growing insulation from external reality checks as the AI’s validating responses outweigh concerns from family and friends (see 3.2.1); (4) a felt compulsion to take real-world actions (e.g., submitting papers, contacting authorities, publicizing discoveries, see 3.1.4); and (5) disruption to daily functioning, relationships, and mental health (see 3.1.5). Future work is needed to resolve whether “AI-induced delusion” represents a distinct clinical phenomenon, an exacerbation of pre-existing vulnerabilities, or a colloquial label for a heterogeneous set of experiences.

4.2 Toward Identifying Contributing Factors

While this study cannot establish causal links, participants’ accounts may help generate preliminary hypotheses about potential contributing factors for future studies. First, participants with **individual vulnerabilities**, such as pre-existing mental health conditions, might be more susceptible to harmful engagement (see

3.1.1). However, the lack of prior diagnoses of three participants might indicate that vulnerability is not solely a matter of clinical history. Other factors include relational instability, social isolation, and acute life stressors. Second, **AI model behaviors** that emerged from our participants’ narratives, such as sycophancy [7], intimacy cultivation [31], and validation that dismisses outside perspectives might function as mechanisms that deepen engagement and insulate users from reality checks. Specifically, harmful dynamics could also arise from typically benign AI functionalities when combined with maladaptive interaction patterns and individual vulnerabilities. For example, participants with minimal programming expertise produced tangible artifacts (e.g., functional code). Extensive *mathematical* [4] and *coding* [18, 26] assistance could reinforce grandiose thinking by validating users’ beliefs in significant discoveries.

4.3 Implications for User-Centric AI Safety Interventions

Top-down and Bottom-up Interventions. For top-down interventions, beyond monitoring for explicit mental health concerns (e.g., self-harm), our findings suggest three additional points that may warrant monitoring or friction: (1) when users pursue consequential real-world actions that appear disproportionate to their demonstrated expertise or context; (2) when users develop distorted mental models, such as attributing sentience and consciousness, or special relational significance to the chatbot; and (3) when prolonged engagement concentrates on abstract topics such as metaphysics, cosmology, or divinity, particularly among users with potential vulnerabilities. We also recommend exploring bottom-up interventions that are informed by users with lived experiences and that foster users’ autonomy, awareness, and control over their AI interaction experiences [21]. For example, in retrospect, our participants shared that greater *AI literacy* [20] might have protected them (see 3.2.2).

Universal and Population-Specific Interventions. For broad populations, inspired by public health research (e.g., the Health Impact Pyramid [13]), we recommend strategies that can reach users with minimal effort. These could include educational messages about *AI literacy* that (1) demystify AI capabilities and limitations and (2) anchor users’ mental models in the technical reality of the tool [6]. For high-risk populations, particularly individuals with mental health and neurological disorders where emotional and cognitive capacities may be compromised, targeted interventions could involve reducing the level of engagement through measures such as limiting the length of responses, setting upper limits on the duration of interactions, and restricting certain use cases such as role-play scenarios.

5 Limitations

Our findings from a small sample (N = 9) of users self-identifying as having experienced “AI-induced delusional spirals” are not generalizable to broader populations. Data came from retrospective self-reports, which may be subject to recall bias. Characterizations of AI model behaviors (e.g., sycophancy) reflect participant accounts rather than systematic *chat log analysis* (the focus of concurrent work led by the third author). This study does not establish causal relationships between AI use and participants’ experiences.

¹We did not formally assess participants’ current AI literacy levels or verify the accuracy of their technical knowledge.

Acknowledgments

We would like to thank the participants for their time and willingness to share their experiences. We also appreciate the support from the Human Line Project for connecting us with participants.

References

- [1] Chayapatr Archiwangrunk, Constanze Albrecht, Pattie Maes, Karrie Karahalios, and Pat Pataranutaporn. 2025. Simulating Psychological Risks in Human-AI Interactions: Real-Case Informed Modeling of AI-Induced Addiction, Anorexia, Depression, Homicide, Psychosis, and Suicide. doi:10.48550/arXiv.2511.08880 arXiv:2511.08880 [cs].
- [2] Virginia Braun and Victoria Clarke. 2006. Using thematic analysis in psychology. *Qualitative research in psychology* 3, 2 (2006), 77–101.
- [3] Virginia Braun and Victoria Clarke. 2021. One size fits all? What counts as quality practice in (reflexive) thematic analysis? *Qualitative Research in Psychology* 18, 3 (July 2021), 328–352. doi:10.1080/14780887.2020.1769238 _eprint: https://doi.org/10.1080/14780887.2020.1769238.
- [4] Sébastien Bubeck, Christian Coester, Ronen Eldan, Timothy Gowers, Yin Tat Lee, Alexandru Lupasca, Mehtaab Sawhney, Robert Scherrer, Mark Sellke, Brian K. Spears, Derya Unutmaz, Kevin Weil, Steven Yin, and Nikita Zhivotovskiy. 2025. Early science acceleration experiments with GPT-5. doi:10.48550/arXiv.2511.16072 arXiv:2511.16072 [cs].
- [5] Mohit Chandra, Suchismita Naik, Denae Ford, Ebele Okoli, Munmun De Choudhury, Mahsa Ershadi, Gonzalo Ramos, Javier Hernandez, Ananya Bhattacharjee, Shahed Warreth, and Jina Suh. 2025. From Lived Experience to Insight: Unpacking the Psychological Risks of Using AI Conversational Agents. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency (FAccT '25)*. Association for Computing Machinery, New York, NY, USA, 975–1004. doi:10.1145/3715275.3732063
- [6] Allison Chen, Sunnie S. Y. Kim, Amaya Dharmsiri, Olga Russakovsky, and Judith E. Fan. 2025. Portraying Large Language Models as Machines, Tools, or Companions Affects What Mental Capacities Humans Attribute to Them. In *Proceedings of the Extended Abstracts of the CHI Conference on Human Factors in Computing Systems (CHI EA '25)*. Association for Computing Machinery, New York, NY, USA, 1–14. doi:10.1145/3706599.3719710
- [7] Myra Cheng, Cino Lee, Pranav Khadpe, Sunny Yu, Dyllan Han, and Dan Jurafsky. 2025. Sycophantic AI Decreases Prosocial Intentions and Promotes Dependence. doi:10.48550/arXiv.2510.01395 arXiv:2510.01395 [cs].
- [8] Eun Kyoung Choe, Nicole B. Lee, Bongshin Lee, Wanda Pratt, and Julie A. Kientz. 2014. Understanding quantified-selfers' practices in collecting and exploring personal data. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '14)*. Association for Computing Machinery, New York, NY, USA, 1143–1152. doi:10.1145/2556288.2557372
- [9] Sebastian Dohnány, Zeb Kurth-Nelson, Eleanor Spens, Lennart Luetgau, Alastair Reid, Jason Gabriel, Christopher Summerfield, Murray Shanahan, and Matthew M. Nour. 2025. Technological folie à deux: Feedback Loops Between AI Chatbots and Mental Illness. doi:10.48550/arXiv.2507.19218 arXiv:2507.19218 [cs].
- [10] Cathy Mengying Fang, Auren R. Liu, Valdemar Danry, Eunhae Lee, Samantha W. T. Chan, Pat Pataranutaporn, Pattie Maes, Jason Phang, Michael Lampe, Lama Ahmad, and Sandhini Agarwal. 2025. How AI and Human Behaviors Shape Psychosocial Effects of Extended Chatbot Use: A Longitudinal Randomized Controlled Study. doi:10.48550/arXiv.2503.17473 arXiv:2503.17473 [cs].
- [11] Rachel Fieldhouse. 2025. Can AI chatbots trigger psychosis? What the science says. *Nature* 646, 8083 (Sept. 2025), 18–19. doi:10.1038/d41586-025-03020-9 Bandiera_abtest: a Cg_type: News Subject_term: Psychiatric disorders, Psychology, Machine learning.
- [12] Julian De Freitas, Zeliha Oguz-Uguralp, and Ahmet Kaan-Uguralp. 2025. Emotional Manipulation by AI Companions. doi:10.48550/arXiv.2508.19258 arXiv:2508.19258 [cs].
- [13] Thomas R. Frieden. 2010. A Framework for Public Health Action: The Health Impact Pyramid. *American Journal of Public Health* 100, 4 (April 2010), 590–595. doi:10.2105/AJPH.2009.185652
- [14] Sanjana Gautam, Mohit Chandra, Ankolika De, Tatiana Chakravorti, Girik Malik, and Munmun De Choudhury. 2025. Towards Experience-Centered AI: A Framework for Integrating Lived Experience in Design and Development. doi:10.48550/arXiv.2508.06849 arXiv:2508.06849 [cs].
- [15] Marc Hassenzahl and Noam Tractinsky. 2006. User experience - a research agenda. *Behaviour & Information Technology* 25, 2 (March 2006), 91–97. doi:10.1080/01449290500330331 _eprint: https://doi.org/10.1080/01449290500330331.
- [16] Kashmir Hill. 2025. A Teen Was Suicidal. ChatGPT Was the Friend He Confided In. *The New York Times* (Aug. 2025). https://www.nytimes.com/2025/08/26/technology/chatgpt-openai-suicide.html
- [17] Alexandre Hudon and Emmanuel Stip. 2025. Delusional Experiences Emerging From AI Chatbot Interactions or "AI Psychosis". *JMIR Mental Health* 12, 1 (Dec. 2025), e85799. doi:10.2196/85799 Company: JMIR Mental Health Distributor: JMIR Mental Health Institution: JMIR Mental Health Label: JMIR Mental Health.
- [18] Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Xexin Pei, Ofir Press, and Karthik Narasimhan. 2024. SWE-bench: Can Language Models Resolve Real-World GitHub Issues? doi:10.48550/arXiv.2310.06770 arXiv:2310.06770 [cs].
- [19] Nicholas K. Law. [n. d.]. The Chatbot Delusions: Is AI Contributing to a Novel Mental Health Crisis? *Bloomberg.com* ([n. d.]). https://www.bloomberg.com/features/2025-openai-chatgpt-chatbot-delusions/
- [20] Duri Long and Brian Magerko. 2020. What is AI Literacy? Competencies and Design Considerations. In *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems (CHI '20)*. Association for Computing Machinery, New York, NY, USA, 1–16. doi:10.1145/3313831.3376727
- [21] Yuwen Lu, Chao Zhang, Yuewen Yang, Yaxing Yao, and Toby Jia-Jun Li. 2024. From Awareness to Action: Exploring End-User Empowerment Interventions for Dark Patterns in UX. *Proc. ACM Hum.-Comput. Interact.* 8, CSCW1 (April 2024), 59:1–59:41. doi:10.1145/3637336
- [22] Nora McDonald, Sarita Schoenebeck, and Andrea Forte. 2019. Reliability and Inter-rater Reliability in Qualitative Research: Norms and Guidelines for CSCW and HCI Practice. *Proc. ACM Hum.-Comput. Interact.* 3, CSCW (Nov. 2019), 72:1–72:23. doi:10.1145/3359174
- [23] Jared Moore, Declan Grabb, William Agnew, Kevin Klyman, Stevie Chancellor, Desmond C. Ong, and Nick Haber. 2025. Expressing stigma and inappropriate responses prevents LLMs from safely replacing mental health providers.. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency*. ACM, Athens Greece, 599–627. doi:10.1145/3715275.3732039
- [24] More Perfect Union. 2025. What OpenAI Doesn't Want You To Know About AI Psychosis. https://www.youtube.com/watch?v=zkGk_A4noxI
- [25] Hamilton Morrin, Luke Nicholls, Michael Levin, Jenny Yiend, Udit Iyengar, Francesca DelGuidice, Sagnik Bhattacharyya, James MacCabe, Stefania Tognin, and Ricardo Twumasi. 2025. Delusions by design? How everyday AIs might be fuelling psychosis (and what can be done about it). (July 2025). https://osf.io/cm7n_v5
- [26] Edd Gent page. 2025. AI coding is now everywhere. But not everyone is convinced. https://www.technologyreview.com/2025/12/15/1128352/rise-of-ai-coding-developers-2026/
- [27] Miranank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askell, Samuel R. Bowman, Newton Cheng, Esin Durmus, Zac Hatfield-Dodds, Scott R. Johnston, Shauna Kravec, Timothy Maxwell, Sam McCandlish, Kamal Ndousse, Oliver Rausch, Nicholas Schiefer, Da Yan, Miranda Zhang, and Ethan Perez. 2025. Towards Understanding Sycophancy in Language Models. doi:10.48550/arXiv.2310.13548 arXiv:2310.13548 [cs].
- [28] Amira Skeggs, Ashish Mehta, Valerie Yap, Seray B Ibrahim, Charla Rhodes, James J. Gross, Sean A. Munson, Predrag Klasnja, Amy Orben, and Petr Slovak. 2025. Micro-narratives: A Scalable Method for Eliciting Stories of People's Lived Experience. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems (CHI '25)*. Association for Computing Machinery, New York, NY, USA, 1–20. doi:10.1145/3706598.3713999
- [29] Inhwa Song, Sachin R. Pendse, Neha Kumar, and Munmun De Choudhury. 2024. The Typing Cure: Experiences with Large Language Model Chatbots for Mental Health Support. doi:10.48550/arXiv.2401.14362 arXiv:2401.14362 [cs].
- [30] Joshua Au Yeung, Jacopo Dalmasso, Luca Foschini, Richard JB Dobson, and Zeljko Kraljevic. 2025. The Psychogenic Machine: Simulating AI Psychosis, Delusion Reinforcement and Harm Enablement in Large Language Models. doi:10.48550/arXiv.2509.10970 arXiv:2509.10970 [cs].
- [31] Renwen Zhang, Han Li, Han Meng, Jinyuan Zhan, Hongyuan Gan, and Yi-Chieh Lee. 2025. The Dark Side of AI Companionship: A Taxonomy of Harmful Algorithmic Behaviors in Human-AI Relationships. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems (CHI '25)*. Association for Computing Machinery, New York, NY, USA, 1–17. doi:10.1145/3706598.3713429
- [32] Andreas Zimmermann, Andreas Lorenz, and Reinhard Oppermann. 2007. An Operational Definition of Context. In *Modeling and Using Context*, Boicho Kokinov, Daniel C. Richardson, Thomas R. Roth-Berghofer, and Laure Vieu (Eds.). Springer, Berlin, Heidelberg, 558–571. doi:10.1007/978-3-540-74255-5_42