

Inside the Newsroom: How Media Company AI Strategies Fall Short in Practice

Tine Rønning Pedersen
Computer Science
Aarhus University
Aarhus, Denmark
tine@cs.au.dk

Klara Øvlisen
Computer Science
Aarhus University
Aarhus, Denmark
klara@cs.au.dk

Marianne Graves Petersen
Computer Science
Aarhus University
Aarhus, Denmark
mgraves@cs.au.dk

Abstract

LLMs have entered the newsrooms, introducing a threat to the democratic role of journalism. Therefore, it is important that journalists use AI tools responsibly and safely. Through contextual interviews with ten journalists across four media companies, we study how journalists engage with the strategies media companies employ to ensure safe and responsible use and how these unfold in practice. We find three strategies used across all companies in this study: 1) develop in-house tools, 2) establish guidelines for AI use, and 3) provide journalists with knowledge on AI. We find that these strategies fall short in practice: the journalists prefer ChatGPT to in-house tools, are unaware of the guidelines, or do not consider them useful, and display a high-level, use-oriented understanding. We argue that HCI may play a role in shaping future tools, best practices, and reflective dialogues about LLMs in the newsroom.

CCS Concepts

• Human-centered computing → Empirical studies in HCI.

Keywords

Journalism, Large Language Models, Contextual Interviews, Qualitative Methods

ACM Reference Format:

Tine Rønning Pedersen, Klara Øvlisen, and Marianne Graves Petersen. 2026. Inside the Newsroom: How Media Company AI Strategies Fall Short in Practice. In *Extended Abstracts of the 2026 CHI Conference on Human Factors in Computing Systems (CHI EA '26)*, April 13–17, 2026, Barcelona, Spain. ACM, New York, NY, USA, 5 pages. <https://doi.org/10.1145/3772363.3799296>

1 Introduction

Journalism and the production of high-quality information are foundational for a democracy [28]. Therefore, it is of utmost importance that journalism remains independent and autonomous [14, 32]. However, generative AI may challenge this independence and autonomy [32, 33], and news organizations therefore typically employ several strategies to ensure safe and responsible use.

In this work, we examine how journalists in practice experience the strategies news organizations employ to ensure safe and responsible use of generative AI in journalism. We do so through ten contextual interviews with journalists from four different media

companies in Scandinavia. We find that media companies typically employ three strategies: 1) the development of in-house tools for journalists to use, 2) the establishment of ethical guidelines for the use of generative AI, and 3) the education of journalists on generative AI. We find that all three strategies did not have the intended effect in practice, as experienced by the journalists in this study. We argue that HCI can provide interface design that fosters standardized practices and critical reflection in journalism, help embody journalistic values in system design, and empower journalists to question and engage with existing and future practices.

2 Related Work

In this section, we review previous research on AI in journalism, including the motivations and challenges, as well as the prerequisites identified in the literature for ensuring safe and responsible use.

2.1 AI's Role, Benefits, and Challenges in Journalism

In the last couple of years, a growing body of research has emerged focusing on how AI affects journalism, which is a profession defined and formed by values such as trustworthiness, autonomy, objectivity, and transparency [14, 16, 17, 23, 27, 28]. The focus of this new body of research has primarily been on how the new tools are used and the possibilities and challenges they bring [3, 12, 16, 27, 28, 32, 33].

One common benefit AI offers is the efficiency and scalability of routine tasks, which reduce costs and allow journalists more time for creative work [3, 12, 27, 28]. AI often assumes roles such as collaborator or editor, and using generative AI for text content production is found to be the most common use case [16]. Beckett and Yaseen [3] highlight that accessibility, low requirements for technical skills, and generative AI's ability to understand context distinguish it from other AI technologies, making it easy to adopt [3]. While AI may offer opportunities to free up journalists' time, the tools also pose risks [3, 12, 16, 27]. Often mentioned perils of using AI in journalism include hallucinations and biases inherent in AI models [12, 16, 23, 27, 28], risk of losing control over journalistic autonomy [12], and fear of power being handed over to tech companies that develop these tools [3, 32, 33].

2.2 Approaches for Responsible AI Use in the News Industry

One common strategy to ensure safe and responsible use is for news organizations to develop comprehensive and ethical guidelines [1, 2, 12, 16]. Common themes in the guidelines include transparency, human oversight, banned vs. allowed uses, hallucinations, bias, and



This work is licensed under a Creative Commons Attribution 4.0 International License. *CHI EA '26, Barcelona, Spain*

© 2026 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-2281-3/26/04

<https://doi.org/10.1145/3772363.3799296>

Name (Gender)	Years of Experience	Media Company
Louis (M)	20	Media Company 1, small local media
Kenneth (M)	20	
Maya (W)	5	
Ellie (W)	5.5	Media Company 2, large national media
Mason (M)	12	
Max (M)	9	Media Company 3, large national media
Rachel (W)	6	
Miles (M)	6	Media Company 4, medium sized local media
Tyler (M)	15	
Kelly (W)	4	

Table 1: Overview of participants and media companies

copyright and intellectual property [1, 2, 11, 16]. Becker [1] finds that guidelines differ widely among news organizations and that no standardized model exist, while Diakopoulos et al. [16] find that guidelines are too high-level and difficult to apply in practice.

One often missing guideline concerns external collaborators and technical dependencies [2]. Simon [33] describes how it is difficult for media companies to develop their own AI tools with the same quality as the big tech companies, as this requires high costs and enormous resources which often make media companies dependent on external collaborators [33]. Dependency on external collaborations is described as ‘media capture’ by Simon [32], where a news organization is being ‘‘captured’ by another agent (e.g., a government or a business) and subjected to their interests, thus losing some or all of its autonomy’’ [32, p. 1839]. As a result, researchers have taken up the challenge and are developing a transparent, open-source LLM for and with journalists [35].

Another condition for responsible use of AI tools is AI literacy [15, 16, 22, 25]. Deuze and Beckett [15] describe three key components of AI literacy specifically for journalists: Knowledge about AI, the ability to recognize instances where AI may be usefully and creatively applied, and skills to help, coach, or teach others [15]. However, Jones et al. [22] find a widespread lack of AI literacy among journalists [22].

We expand on this research by investigating which strategies media companies in a Scandinavian context employ to ensure safe and responsible use and how these strategies unfold in practice. We do this through contextual interviews to gain insights into how the strategies affect the concrete everyday practice of journalists collaborating with LLMs.

3 Methodology

In this study, we use contextual interviews [4, 19] to observe concrete practices. Two authors conducted ten contextual interviews with journalists from four diverse media companies in Scandinavia, which were then analyzed using thematic analysis [6, 10]. An overview of the participants and media companies is presented in table 1, and all participants in this work are pseudonomized.

The interviews were conducted at the participants’ workplace and lasted between 30 and 90 minutes. The interviews at the shorter end of this range were limited by journalists’ significant time constraints. Before the contextual interview began, the journalists completed a brief survey consisting primarily of demographic questions. In three of the contextual interviews, we directly observed the journalists performing the tasks they were engaged in at that moment. The remaining journalists provided retrospective accounts of tasks in which they had recently used AI. These retrospective accounts included accessing the concrete chat history on ChatGPT and in-house tools, in which they discussed their tasks, intentions, and frustrations. Finally, we conducted a brief semi-structured interview to address topics that could not be revealed in the contextual part, including knowledge of guidelines and AI. The interview questions were motivated by themes seen in the literature [1, 2, 11, 16, 22, 32, 33]. The data gathered included surveys, contextual and semi-structured interview transcripts, and field notes.

To analyze the data, two authors wrote down their impressions immediately after the interview. The interviews were then transcribed and analyzed using thematic analysis [6, 10] and an emergent coding process [24]. Three authors, including one who had not participated in conducting the interviews, reviewed a selected interview together line by line, highlighting interesting parts, and developed a codebook consisting of 31 codes. Immediately thereafter, two authors independently coded the next interview. After finishing, they discussed whether any codes should be added to the codebook and calculated Cohen’s kappa value [24]. The two authors determined that the codebook was sufficient and obtained a Cohen’s kappa of 0.797, indicating satisfactory agreement [24]. All interviews were then coded. Before going to step three in the thematic analysis process [6], two authors read through all coded interviews and original field notes. After doing this, they started searching for themes by using affinity diagramming [4]. The themes were iteratively refined and when satisfied with the grouping of the codes, they began searching for representative names and identified three main themes and nine subthemes.

4 Findings

Generative AI has entered the newsroom, and all journalists participating in this study used it for tasks such as proofreading, transcription, headline generation, idea generation, and article drafting, where eight out of ten journalists used it daily and two out of ten used it weekly. In this section, we will describe how the journalists in this study engage with the strategies employed by media companies to ensure safe and responsible use of generative AI. We identify three common strategies that journalists experience: developing in-house tools for journalists, generating ethical guidelines for AI use, and equipping journalists with knowledge of these tools. We will describe each strategy and how it takes place in practice.

4.1 Negative Perception of In-House Generative AI Tools

All four media companies had at some point developed their own LLM-infused tool. The motivation behind the in-house tools is to alleviate journalists’ burden of prompting and to provide standardized methods for routine tasks, such as proofreading, by offering

prewritten prompts within the tools. All in-house tools seen in this study are built on OpenAI's API, and at the time of the interview, three out of four were in use. The development of the in-house tool at Media Company 1 was discontinued after management concluded that ChatGPT outperformed it. Instead, the company provided all journalists with their own paid-for ChatGPT login. As expressed by Louis: "Now we have this one, and it works really well. There is no reason to use the other one." (Louis), the journalists are happy with this solution.

Media Company 2's GDPR-compliant in-house tool includes different prewritten prompts and allows journalists to engage in dialogue with the LLM. Ellie is dissatisfied with Media Company 2's tool because it is not up to date with recent information and can only use information from before 2021, and because the interface does not allow the journalists to upload any files to it. She stated that she will use the in-house tool when handling sensitive data, and otherwise, she prefers to use ChatGPT and NotebookLM. Mason believes ChatGPT is superior and prefers to use it, as he views the in-house tool primarily as a tool for catching errors: "[In-house tool] is more for errors" (Mason).

Media Company 3 had developed a GDPR-compliant in-house tool and also provided closed access to ChatGPT. However, neither of the two journalists participating in this study had used the in-house tool, nor had they investigated its usefulness. A consequence of this is that we did not actually see the tool's interface or features. Instead, both journalists preferred to use ChatGPT, as Rachel expressed: "I just think that ChatGPT works pretty well. So, I have never actually looked into it" (Rachel). Rachel used the company-provided login, and Max used his personal login to the free version of ChatGPT.

The in-house tool developed by Media Company 4 greatly resembles the tool from Media Company 2. However, this tool does not allow journalists to chat with the model, and the tool is not GDPR-compliant. Here, one journalist, Kelly, preferred the in-house tool since it had been approved by management. However, the two other journalists preferred ChatGPT, because they can go into dialogue with it, as expressed by Miles: "[In-house tool] is locked, such that you can't continue talking with it. It gives one reply, and then it has forgotten what it has done afterwards" (Miles).

4.2 Guidelines for Generative AI Use Are Not Useful

In addition to developing in-house tools, all media companies represented in this study have developed guidelines for the use of generative AI. This aligns with the literature, which highlights guidelines as a prerequisite for using generative AI responsibly in journalism [11, 16]. In this study, we find that the guidelines are either unknown to the journalists or of little practical usefulness.

Two of the journalists in this study, Tyler and Louis, are aware of where they can find the guidelines and what they include. Five of the journalists, Ellie, Kenneth, Kelly, Max, and Maya, do not know the content of the guidelines. Kelly expressed a wish to get some information about guidelines: "I miss information about our ethical guidelines" (Kelly). Kelly, Kenneth, and Ellie do not believe their company has any guidelines, whereas Max and Maya are aware that guidelines likely exist but do not know what they include.

Max, for example, expresses: "Oh dear, I should probably have read them" (Max). The last three journalists, Rachel, Miles, and Mason, are aware of the guidelines but do not consider them to be useful. Rachel believes that they are self-evident, high-level, and obvious. Miles and Mason consciously work around them. Mason describes how the journalists at Media Company 2 are not allowed to have a chat history in ChatGPT. He has uploaded some of his previously well-received articles and other material to ChatGPT to adapt it to his writing style, even though he knows this is not permitted under the guidelines. Similarly, Miles describes how his media company has a guideline stating that he may only upload publish-ready material to generative AI. He describes how his and his leader's interpretations of this rule differ, and he uses his own interpretation. These examples illustrate a current and deep misalignment between guidelines and journalists practice.

4.3 Lack of Deep Knowledge about Generative AI

A third strategy adopted by the media companies in this study is to provide journalists with knowledge of AI and its use. In this regard, we observe differences among the media companies. While Media Company 1 and 4 have required all journalists to attend a mandatory course on generative AI, Media Company 2 and 3 have offered voluntary courses to the journalists. At both Media Company 2 and 3, we observed that one in two journalists had sought out and attended these courses, whereas the other two had not. Max, for example, argued: "I haven't really thought that it was relevant" (Max). At Media Company 2, an introduction to their own in-house tool had been given, whereas the journalists at Media Company 3 have not received a mandatory introduction to their in-house tool. Rather than building knowledge through courses, the journalists describe how they often acquire it by talking with their colleagues and how useful advice spreads from one colleague to another. Some journalists also believe that the best way to learn it is to use it, and that, in doing so, they become better over time.

Two of our questions in the semi-structured interview were to examine their AI literacy, including how an LLM works and the origin of hallucinations in the output. Regarding the first questions, all journalists provided high-level answers, such as: "It is built on statistics and experience" (Rachel), "It is something about how it turns it into numbers, and then something comes out of it" (Mason), and "It is some kind of algorithm, a machine, a robot, which has been coded to react to what I am telling it" (Maya). Kenneth, in contrast, focused on how he can use it: "It works in a way that can help me with things... It is everything from headlines and subheadings to fact-boxes" (Kenneth). This illustrates that even the journalists who have attended courses on ChatGPT still do not have a clear understanding of how the model works [29]. Regarding the second question on hallucinations, Ellie, Kelly, Miles, and Maya had no idea why they occur, and the other six provided relatively accurate answers [21, 38], such as that AI will always give an answer, even if it does not know: "It will rather provide an answer than admit that it cannot" (Rachel), and that there may be mistakes in the training data: "Because in one way or another, it is trained on some data in which there are some errors." (Max). This indicates that the majority of the journalists provide relatively accurate explanations

of some of the reasons hallucinations may occur [21, 38], while still using guesswork to fill in the blanks.

5 Discussion

In this section, we comment on the success of the media companies' strategies observed in this study and discuss how HCI can engage in the challenges observed.

5.1 Unsuccessful Company Strategies in Practice

As described in section 4, we find that there are three common strategies employed by media companies to ensure safe and responsible use: in-house tools, guidelines, and building AI knowledge. We find that these strategies are unsuccessful in practice to varying degrees.

As described in section 4.2, all media companies in this study have developed a set of guidelines, however, they do not seem to have the intended influence in practice across all four media companies. Some journalists are unaware of their existence, some deem them to be unnecessary, and some even deliberately work against them. This is critical if the guidelines are intended to ensure safe and ethical use of AI.

Diakopoulos et al. [16] furthermore mention AI literacy as a condition for responsible use of generative AI tools. We observe that all media companies in this study have adopted this strategy by offering AI courses to their journalists, however, two of the media companies did not make attendance mandatory. Placing the responsibility for acquiring AI literacy on individual journalists may lead to some journalists not prioritizing it, since it is an extra task on top of their already busy workloads. Furthermore, we observed that even journalists who had attended AI courses did not appear to have gained an understanding of the tools, as illustrated by the high-level, use-oriented responses about how a language model works described in section 4.3.

A third strategy adopted by the news companies in this study is to develop in-house LLM-infused tools. While the development of these tools represents a significant investment, we observed how the journalists ultimately preferred to use ChatGPT, undercutting the media companies' strategies. Compounding this, the examples of in-house LLM-infused tools we saw in this study were all systems built on top of OpenAI's API. This indicates media capture [32] by OpenAI, who gain influence despite media companies' attempts to counteract this. OpenAI themselves represent a big tech company driven by commercial interests [31] (i.e., not aligned with journalistic or democratic goals), and its models are typically described as 'black boxes' [31]. As such, workers in the news industry have little knowledge of what occurs within these models and little influence over their development. This risks a homogenization of news products and threatens journalistic autonomy, recognized as an important value of journalism by Deuze [14]. Overall, this introduces a threat to quality journalism, which is recognized as a key pillar of a democratic society [28].

5.2 Call for HCI to Engage in Challenges

Overall, the unsuccessful strategies leave generative AI use highly individualized, and there appears to be a lack of agreed-upon best

practices. Here, HCI may help by developing new strategies for interface design for this new group of tools [18, 20] and tools in which best practices are embedded in the system interaction design, thereby helping to ensure that they are followed. This could include concretizing guidelines through interface design [30], making outputs appear less authoritative by showing the uncertainty in the LLM's output [5, 30, 39], or providing different alternative outputs. Furthermore, interaction design could foster AI Literacy and critical reflection [13] and conceptualize some of the underlying mechanisms through the interface. A concrete design provocation could be highlighting that the sender of the output in ChatGPT is a large proprietary tech company. This could help demonstrate how OpenAI may become a co-creator of journalistic products if outputs are not critically reflected on. Another design idea could be to explore how to visualize the attention mechanisms in the models in a way that makes sense for journalists to improve AI literacy [30, 36, 37] or to visualize how human involvement influences the training of the models [29, 34]. Additionally, we believe that embodying journalistic values in system design [23, 26] and a Participatory Design approach [8] could empower journalists to question their existing practices and help them influence the development of alternatives [7, 9]. One such example is described by Tseng et al. [35] who adopt a participatory co-design approach to explore the potential of a journalist-led and -controlled LLM. We call upon the HCI community to take up this timely and important challenge to support journalists in maintaining control of their news products.

6 Conclusion

In conclusion, we find that journalists engage with three strategies to ensure safe and responsible AI use employed by media companies in Scandinavia: 1) develop in-house tools, 2) establish guidelines for AI use, and 3) equip journalists with knowledge on generative AI. Through contextual interviews, we find that, in practice, these strategies are unsuccessful to varying degrees. Almost all journalists in this study preferred ChatGPT to the company's in-house tools, the guidelines were either unknown or not deemed useful in practice, and the journalists displayed a high-level, use-oriented understanding of the tools. This results in an individualized approach to generative AI in journalism with no standardized, agreed-upon best practice. Future work includes conducting more contextual interviews across different contexts to see whether this is generalizable. We argue that HCI may play a role in establishing a best practice and fostering critical reflection through interface design, embodying journalistic values into systems, and empowering journalists through Participatory Design.

Acknowledgments

This work was funded by the Danish Independent Research Foundation, grant ID: 10.46540/4264-00168B. We thank all participants in this study for generously devoting their valuable time and insights.

References

- [1] Kim Becker. 2023. New game, new rules. An investigation into editorial guidelines for dealing with artificial intelligence in the newsroom. *Journalism Research* 6(2) (07 2023), 133–152. doi:10.1453/2569-152X-22023-13404-en

- [2] Kim Björn Becker, Felix M. Simon, and Christopher Crum. 2025. Policies in Parallel? A Comparative Study of Journalistic AI Policies in 52 Global News Organisations. *Digital Journalism* (2025), 1–21. doi:10.1080/21670811.2024.2431519
- [3] Charlie Beckett and Mira Yaseen. 2023. Generating change: A global survey of what news organisations are doing with AI. (2023). <https://researchonline.lse.ac.uk/id/eprint/129643>
- [4] Hugh Beyer. 2010. *User-centered agile methods*. Morgan & Claypool Publishers.
- [5] Umang Bhatt, Javier Antorán, Yunfeng Zhang, Q Vera Liao, Prasanna Sattigeri, Riccardo Fogliato, Gabrielle Melançon, Ranganath Krishnan, Jason Stanley, Omesh Tickoo, et al. 2021. Uncertainty as a form of transparency: Measuring, communicating, and using uncertainty. In *Proceedings of the 2021 AAAI/ACM Conference on AI, Ethics, and Society*. 401–413. doi:10.1145/3461702.3462571
- [6] Virginia Braun and Victoria Clarke. 2006. Using thematic analysis in psychology. *Qualitative Research in Psychology* 3, 2 (2006), 77–101. doi:10.1191/1478088706qp063oa
- [7] Susanne Bødker. 2003. A for alternatives. *Scand. J. Inf. Syst.* 15, 1 (Jan. 2003), 87–89.
- [8] Susanne Bødker, Ole Sejer Iversen, Christian Dindler, and Rachel Charlotte Smith. 2021. *Participatory Design*. Synthesis Lectures on Human-Centered Informatics, Vol. 14/5. Morgan & Claypool Publishers. doi:10.2200/S01136ED1V01Y202110HCI052
- [9] Susanne Bødker and Morten Kyng. 2018. Participatory design that matters—Facing the big issues. *ACM Transactions on Computer-Human Interaction (TOCHI)* 25, 1 (2018), 1–31. doi:10.1145/3152421
- [10] Victoria Clarke and Virginia Braun. 2016. Thematic analysis. *The Journal of Positive Psychology* 12, 3 (2016), 297–298. doi:10.1080/17439760.2016.1262613
- [11] Hannes Cools and Nicholas Diakopoulos. 2023. Towards Guidelines for Guidelines on the use of Generative AI in Newsrooms. <https://generative-ai-newsroom.com/towards-guidelines-for-guidelines-on-the-use-of-generative-ai-in-newsrooms-55b0c2c1d960> Accessed: 2025-06-02.
- [12] Hannes Cools and Nicholas Diakopoulos. 2024. Uses of Generative AI in the Newsroom: Mapping Journalists' Perceptions of Perils and Possibilities. *Journalism Practice* (2024), 1–19. doi:10.1080/17512786.2024.2394558
- [13] Peter Dalsgaard. 2025. Creative Ambiguity and Cognitive Tension in Generative AI Tools. In *Proceedings of the 36th Annual Conference of the European Association of Cognitive Ergonomics (EACE)*. 1–10. doi:10.1145/3746175.3746183
- [14] Mark Deuze. 2005. What is journalism? Professional identity and ideology of journalists reconsidered. *Journalism* 6, 4 (2005), 442–464. doi:10.1177/1464884905056815
- [15] Mark Deuze and Charlie Beckett. 2022. Imagination, Algorithms and News: Developing AI Literacy for Journalism. *Digital Journalism* 10, 10 (2022), 1913–1918. doi:10.1080/21670811.2022.2119152
- [16] Nicholas Diakopoulos, Hannes Cools, N Helberger, C Li, Ernest Kung, and Aimee Rinehart. 2024. Generative AI in Journalism: The Evolution of Newswork and Ethics in a Generative Information Ecosystem. (2024). doi:10.13140/RG.2.2.31540.05765
- [17] Konstantin Nicholas Dörr and Katharina Hollnbuchner. 2016. Ethical Challenges of Algorithmic Journalism. *Digital Journalism* 5, 4 (2016), 404–419. doi:10.1080/21670811.2016.1167612
- [18] Jonathan Grudin. 2009. AI and HCI: Two fields divided by a common focus. *AI magazine* 30, 4 (2009), 48–48. doi:10.1609/aimag.v30i4.2271
- [19] Karen Holtzblatt and Hugh Beyer. 2014. *Contextual design: evolved*. Morgan & Claypool Publishers.
- [20] Kasper Hornbæk, Per Ola Kristensson, and Antti Oulasvirta. 2025. Rethinking HCI Education for the Era of AI. *Interactions* 33, 1 (2025), 40–43. doi:10.1145/3777517
- [21] Lei Huang, Weijiang Yu, Weitao Ma, Weihong Zhong, Zhangyin Feng, Haotian Wang, Qianglong Chen, Weihua Peng, Xiaocheng Feng, Bing Qin, et al. 2025. A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions. *ACM Transactions on Information Systems* 43, 2 (2025), 1–55. doi:10.1145/3703155
- [22] Bronwyn Jones, Rhianne Jones, and Ewa Luger. 2022. AI 'Everywhere and Nowhere': Addressing the AI Intelligibility Problem in Public Service Journalism. *Digital Journalism* 10 (12 2022), 1731–1755. doi:10.1080/21670811.2022.2145328
- [23] Tomoko Komatsu, Marisela Gutierrez Lopez, Stephann Makri, Colin Porlezza, Glenda Cooper, Andrew MacFarlane, and Sondess Missaoui. 2020. AI should embody our values: Investigating journalistic values to inform AI technology design. In *Proceedings of the 11th Nordic Conference on Human-Computer Interaction: Shaping Experiences, Shaping Society*. 1–13. doi:10.1145/3419249.3420105
- [24] J. Lazar, J. H. Feng, and H. Hochheiser. 2017. *Research Methods in Human-Computer Interaction*. Morgan Kaufmann Publishers.
- [25] Duri Long and Brian Magerko. 2020. What is AI literacy? Competencies and design considerations. In *Proceedings of the 2020 CHI conference on human factors in computing systems*. 1–16. doi:10.1145/3313831.3376727
- [26] Marisela Gutierrez Lopez, Colin Porlezza, Glenda Cooper, Stephann Makri, Andrew MacFarlane, and Sondess Missaoui. 2023. A question of design: Strategies for embedding AI-driven tools into journalistic work routines. *Digital Journalism* 11, 3 (2023), 484–503. doi:10.1080/21670811.2022.2043759
- [27] Amaya Noain-Sánchez. 2022. Addressing the Impact of Artificial Intelligence on Journalism: the perception of experts, journalists and academics. *Communication & Society* 35, 3 (2022), 105–121. doi:10.15581/003.35.3.105-121
- [28] Andreas L Opdahl, Bjørnar Tessem, Duc-Tien Dang-Nguyen, Enrico Motta, Vinay Setty, Eivind Throndsen, Are Tverberg, and Christoph Trattner. 2023. Trustworthy journalism through AI. *Data & Knowledge Engineering* 146 (2023). doi:10.1016/j.datak.2023.102182
- [29] OpenAI. 2022. Introducing ChatGPT. <https://openai.com/index/chatgpt> Accessed: 2025-06-02.
- [30] Tine Rønning Pedersen, Klara Øvlsen, Luke Connelly, Ira Assent, and Marianne Graves Petersen. 2026. Breaking News or Breaking Trust? Exploring Challenges and a Design Space for Trustworthy LLM Integration in Journalism. In *Proceedings of the 2026 CHI Conference on Human Factors in Computing Systems* (Barcelona, Spain) (CHI '26). Association for Computing Machinery, New York, NY, USA. doi:10.1145/3772318.3791457
- [31] Johannes Schneider. 2024. Explainable generative AI (GenXAI): A survey, conceptualization, and research agenda. *Artificial Intelligence Review* 57, 11 (2024), 289. doi:10.1007/s10462-024-10916-x
- [32] Felix M Simon. 2022. Uneasy bedfellows: AI in the news, platform companies and the issue of journalistic autonomy. *Digital Journalism* 10, 10 (2022), 1832–1854.
- [33] Felix M Simon. 2024. Escape me if you can: How AI reshapes news organisations' dependency on platform companies. *Digital Journalism* 12, 2 (2024), 149–170. doi:10.1080/21670811.2023.2287464
- [34] Nisan Stiennon, Long Ouyang, Jeffrey Wu, Daniel Ziegler, Ryan Lowe, Chelsea Voss, Alec Radford, Dario Amodei, and Paul F Christiano. 2020. Learning to summarize with human feedback. *Advances in neural information processing systems* 33 (2020), 3008–3021. doi:10.48550/arXiv.2009.01325
- [35] Emily Tseng, Meg Young, Marianne Aubin Le Quéré, Aimee Rinehart, and Harini Suresh. 2025. "Ownership, Not Just Happy Talk": Co-Designing a Participatory Large Language Model for Journalism. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency*. 3119–3130. doi:10.1145/3715275.3732198
- [36] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. *Advances in neural information processing systems* 30 (2017). doi:10.48550/arXiv.1706.03762
- [37] Jesse Vig. 2019. A Multiscale Visualization of Attention in the Transformer Model. *CoRR abs/1906.05714* (2019). arXiv:1906.05714 <http://arxiv.org/abs/1906.05714>
- [38] Yue Zhang, Yafu Li, Leyang Cui, Deng Cai, Lemao Liu, Tingchen Fu, Xinting Huang, Enbo Zhao, Yu Zhang, Yulong Chen, et al. 2025. Siren's Song in the AI Ocean: A Survey on Hallucination in Large Language Models. *Computational Linguistics* (2025), 1–46. doi:10.1162/COLLa.16
- [39] Yunfeng Zhang, Q. Vera Liao, and Rachel K. E. Bellamy. 2020. Effect of Confidence and Explanation on Accuracy and Trust Calibration in AI-Assisted Decision Making. *CoRR abs/2001.02114* (2020). arXiv:2001.02114 <http://arxiv.org/abs/2001.02114>