

One Kiss: Emojis as Agents of Genre Flux in Generative Comics

Xiruo Wang*
University College London
London, London, United Kingdom
xiruo.wang.22@ucl.ac.uk

Xinyi Jiang
Tsinghua University
Beijing, China
xinyijiang2002@163.com

Ziqi Lyu
Beijing Forestry University
Beijing, China
lzqleisure@gmail.com




Figure 1: The One Kiss Interface. Users define structure via Spatial Pacing (sketching the frame for Panel 4) and modulate tone via Affective Injection (pairing keywords like “Gunshot” with a “Worried” emoji). The system’s Genre Flux algorithm updates the affective radar chart (top right), shifting the narrative style toward “Mystery” and “Chaos” for the next generation.

Abstract

Generative AI has made visual storytelling widely accessible, yet current prompt-based interactions often force users into a trade-off between precise control and creative flow. We present **One Kiss**, a co-creative comic generation system that introduces “Affective Steering”. Instead of writing text prompts, users guide the tone of their story through emoji inputs, whose semantic ambiguity becomes a resource rather than a limitation. Unlike traditional text-to-image tools that rely on explicit descriptions, One Kiss uses a dual-stream input in which users define structural pacing by sketching panel frames and set atmospheric tone by pairing keywords with emojis. This mechanism enables “Genre Flux,” where emotional inputs accumulate across panels and gradually shift the genre of a story. A preliminary study ($N = 6$) suggests that this soft steering approach may reframe the user’s role from prompt

engineer to narrative director, with ambiguity serving as a source of creative surprise rather than a loss of control.

CCS Concepts

• **Human-centered computing** → HCI design and evaluation methods; **Interactive systems and tools**; *Interaction design*.

Keywords

generative AI, co-creation, emoji, affective computing, comic generation, human-AI interaction

ACM Reference Format:

Xiruo Wang, Xinyi Jiang, and Ziqi Lyu. 2026. One Kiss: Emojis as Agents of Genre Flux in Generative Comics. In *Extended Abstracts of the 2026 CHI Conference on Human Factors in Computing Systems (CHI EA '26)*, April 13–17, 2026, Barcelona, Spain. ACM, New York, NY, USA, 4 pages. <https://doi.org/10.1145/3772363.3798758>

*First author and corresponding author.



This work is licensed under a Creative Commons Attribution 4.0 International License. *CHI EA '26, Barcelona, Spain*

© 2026 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-2281-3/26/04

<https://doi.org/10.1145/3772363.3798758>

1 Introduction

Text-to-image models now allow novices to produce complex imagery without traditional artistic skills [3]. However, the prevailing interaction paradigm of text-based prompting imposes a substantial cognitive load on users. To maintain narrative consistency or achieve a specific atmospheric tone, users are often forced to engage

in laborious prompt engineering, an iterative process that resembles debugging code rather than creative expression [11]. This gap widens when users attempt to articulate implicit affective qualities, such as “a melancholic silence,” rather than explicit physical descriptions. While text is precise for specifying objects, it is often too rigid for the non-verbal qualities of emotional storytelling.

We take a different approach: rather than increasing prompt precision, we treat ambiguity as a design resource. Drawing on Gaver et al.’s framework of ambiguity as a resource for design [6], we suggest that semantic ambiguity can open spaces for interpretation and surprise. In digital communication, emojis carry rich emotional meaning in a compact, non-verbal form [1]. Unlike text prompts that specify what the AI should draw, emojis offer a way to steer how the narrative feels, providing a form of control that feels intuitive rather than technical.

In this work, we introduce **One Kiss**, an interactive comic generation system built around what we call **Affective Steering**, designed for creative exploration rather than production. The system decouples narrative control into two streams. First, users define structural rhythm through spatial sketching. Second, they shape narrative tone through emoji inputs. This mechanism enables **Genre Flux**, where affective inputs accumulate and gradually shift the genre of the story. Our preliminary evaluation ($N = 6$) suggests that this approach reduces the anxiety of blank-page generation and may reframe the user’s role from prompt engineer to narrative director. This paper contributes the One Kiss prototype and empirical insights into how soft control maintains creative flow during AI co-creation.

2 Related Work

2.1 Steering Interaction in Co-Creative Narrative Systems

Text prompting remains the dominant interface for generative AI, yet non-expert users often struggle to express affective intentions through text alone [11]. Co-creative systems have responded by introducing alternative forms of steering, but these have focused on two areas: narrative structure and visual composition.

On the structural side, *TaleBrush* lets users sketch fortune curves to control how a character’s arc rises and falls across a story [3], while dialogue-based systems like *ID.8* [2] and “*1001 Nights*” [5] use planning conversations and role-play agents to support plot development. On the visual side, *Opal* [8] and *AI Instruments* [10] offer multimodal cues such as layouts and reference images to guide image generation without relying on text alone.

What these approaches share is that they steer *what happens* in a story or *what a scene looks like*, but not *how the story feels over time*. Even in sequential visual story generation, where models like StoryDALL-E [9] maintain cross-frame visual consistency, each panel is generated in a tonally independent way. There is no mechanism for emotional signals to accumulate across a sequence and gradually reshape the genre. One Kiss addresses this gap.

2.2 Ambiguity as a Resource for Affective Steering

To steer tone rather than structure, we need an input modality that is expressive but deliberately imprecise. Gaver et al.’s framing of ambiguity as a design resource [6] suggests that interpretive openness can invite engagement rather than frustrate it. Emojis fit this description: large-scale sentiment analyses have shown that emojis carry quantifiable emotional valence [7], and experimental evidence confirms that they effectively convey emotionality in digital contexts [4], yet their meaning remains sufficiently open to support interpretive flexibility [1].

One Kiss applies this through Affective Steering: users pair keywords with emojis, and these pairings accumulate over time, shifting the visual genre rather than specifying plot events.

3 System Implementation

We developed One Kiss as an interactive research prototype to implement Affective Steering. The system architecture consists of a dual-stream frontend for multimodal input and a state-aware backend for narrative accumulation.

3.1 The Dual-Stream Interface

Spatial Pacing Stream: Users begin a session by sketching a bounding box and selecting an initial narrative keyword to generate the first panel. Subsequent panels are added iteratively, with each sketch defining the visual rhythm for that beat. The dimensions of the box serve as a structural prompt, instructing the image generator to adjust composition complexity. For example, a wide, short box biases the model towards panoramic, cinematic shots, whereas a tall, narrow box triggers close-up character portraits. This allows users to control the camera language without technical vocabulary.

Affective Injection Stream: Rather than writing text prompts for every detail, users pair a keyword from the interface vocabulary (e.g., “Gunshot”, “Broken Glass”) with an emoji (e.g., “Wilted Flower”). This pairing drives the generation. The interface provides visual feedback through a particle system that changes color based on the injected emoji, so users can see how their input affects the system’s state.

3.2 Algorithmic Steering and Genre Flux

Unlike stateless generation models that treat each prompt in isolation, One Kiss maintains a dynamic **Narrative State Vector** (V_t). Central to this is the Genre Flux Algorithm, which tracks the cumulative emotional trajectory of the story.

Let S_t be the emotional score vector at panel t , comprising four dimensions: {*Romance*, *Tragedy*, *Chaos*, *Mystery*}. These four dimensions were chosen for the current prototype to cover a broad range of genre directions; they are not intended as a fixed taxonomy and could be extended in future iterations. The state update rule is defined as:

$$S_t = S_{t-1} \cdot \alpha + (W_{kw} + W_{emoji}) \cdot \beta \quad (1)$$

Where α is a decay factor empirically tuned during pilot testing to balance narrative continuity and responsiveness to new inputs;




Figure 2: Technical Workflow. The system processes inputs through two parallel streams: (1) A Spatial Stream that converts sketches into ControlNet guidance, and (2) an Affective Stream where the Genre Flux Algorithm converts emojis into dynamic style modifiers. These streams converge in the Prompt Synthesizer to guide the Stable Diffusion model.

a value of 0.8 was found to allow gradual genre drift without jarring tonal shifts. β is a rarity multiplier ranging from 1.0 to 3.0, derived from the semantic rarity of the user’s keyword choice, on the assumption that unusual inputs signal intentional narrative pivots.

The system continuously monitors S_t . When a specific dimension dominates (e.g., $S_{tragedy} > \theta$, where θ is the flux threshold), the system activates a **Style Modifier**. This modifier does not act merely as a color filter. It injects genre-specific constraints into both the prompt and negative prompt, influencing compositional density and lighting contrast. For instance, crossing the Tragedy threshold might inject “monochrome blue palette, high contrast shadows, film noir grain” while simultaneously suppressing “bright colors, cheerful expressions,” making the genre shift visible.

To address the common challenge of character hallucination in sequential generation, the system implements a **Global Character Anchor**. A fixed visual description prompt is prepended to every generation request, ensuring that while the genre style shifts dynamically, the protagonist’s identity features remain consistent across panels. The current prototype focuses on single-protagonist narratives, allowing the Genre Flux mechanism to maintain a coherent visual identity throughout the comic. The protagonist’s visual description is specified by the user at the start of each session.

4 Preliminary Evaluation

We conducted an exploratory user study to test Affective Steering in practice and understand how users perceive the control-ambiguity trade-off.

4.1 Participants and Procedure

We recruited 6 participants (3 female, 3 male, aged 22–26) with varying levels of experience in generative AI tools. The study was conducted remotely using a think-aloud protocol. The procedure consisted of three phases. First, a 10-minute tutorial introduced users to the spatial sketching and emoji interface. Second, users were given a creative task: create a 6-panel comic that starts with a neutral tone and “drifts” into a specific genre (e.g., Horror, Romance,

or Tragedy) using only the provided tools. Finally, a post-task semi-structured interview explored their sense of agency and creative flow.

4.2 Results: From Engineering to Directing

Participants used the emoji inputs to steer narrative trajectories. Two authors independently reviewed the 36 generated panels, assessing whether the dominant visual style, as indicated by color palette, lighting contrast, and compositional density, had shifted to match each participant’s stated target genre by the fourth panel. This analysis confirmed that all six participants successfully triggered a Genre Flux, with disagreements resolved through discussion.

Qualitative feedback pointed to a shift in how users thought about the interaction. One participant (P2) noted that they did not have to explicitly describe the sadness or instruct the AI to make the scene dark. Instead, dropping a wilted flower emoji caused the lighting to shift naturally. They described this interaction as akin to “giving an actor a mood cue rather than writing a technical script.” This suggests that semantic ambiguity lets users work at the level of narrative mood rather than struggling with word choice and prompt syntax. Four of six participants described the interaction as feeling more like directing a story than prompting a system, and several noted that compared to their prior experience with text-to-image tools, the emoji-based input felt less intimidating when starting from a blank canvas.

4.3 Serendipity and Perceived Agency

Participants did not see the ambiguity of emojis as a drawback. Instead, they reported that the lack of granular control encouraged a sense of discovery. For example, P4 paired a “Clown” emoji with a “Rain” keyword, expecting a literal interpretation. The system instead generated a surreal and chaotic background with distorted colors, which the user embraced as a creative plot twist. At the sequence level, P3 used emoji inputs to steer a six-panel comic through a thriller arc that concluded with the protagonists kissing, an outcome they had not planned but described as surprising yet

coherent. This suggests that Genre Flux can produce surprises that span the whole story, not just individual panels.

These observations suggest a shift in how agency is distributed. Users give up control over pixel-level rendering while keeping influence over pacing and genre. Several participants also described a sense of agency over the narrative direction. This may relate to the fact that the primary structural decisions, such as panel layout and emotional pivots through emoji choice, originated from the user, even though the pixel-level rendering was automated.

Participants also identified tensions in the spatial pacing stream. One user noted that early keyword choices constrained later scenes. Selecting “Coffee Shop” as an initial context caused subsequent panels to remain anchored to that setting, even when the intended narrative had moved on. Another observed that consecutive panels often defaulted to similar close-up compositions, and wanted more control over how camera framing varied across panels. These observations point to a limitation of the current prototype. While the affective stream supports gradual drift, the spatial stream does not yet support the same kind of gradual evolution. Given the small sample size of six participants, these findings should be interpreted as exploratory.

5 Discussion

Beyond the specific prototype, One Kiss raises a broader question: how can semantic ambiguity serve as a resource in generative systems?

5.1 The Tension Between Precision and Ambiguity

This work highlights a trade-off between lexical precision and affective expression. Traditional prompt-based systems prioritize object-level accuracy (e.g., “a red chair”), which often fails to capture atmospheric nuance without extensive engineering. In contrast, Affective Steering sacrifices this granular control for high-level narrative resonance. Our study suggests that for creative exploration, users gravitated toward the “soft control” of emojis, which act as semantic scaffolds. This suggests a shift from an “architectural” view of creativity toward something closer to gardening, planting seeds and seeing what grows.

5.2 Balancing Consistency and Drift

A critical tension in our design is the balance between narrative drift and character consistency. While the Genre Flux algorithm intentionally destabilizes the visual style, the Global Character Anchor is essential to prevent the narrative from collapsing into incoherence. Future iterations will need to refine this balance, perhaps by allowing users to lock specific visual elements that should remain immune to the flux.

6 Conclusion and Future Work

We introduced One Kiss, a system that uses emoji ambiguity to steer the genre of generative comics. Our preliminary results suggest that this soft steering approach can support creative flow and unexpected discovery. Future work will explore fine-grained steering mechanisms, such as continuous intensity sliders for emoji weights, and investigate therapeutic applications where non-verbal emoji

inputs could help users express difficult emotions through digital storytelling.

References

- [1] Hamza Alshenqeeti. 2016. Are Emojis Creating a New or Old Visual Language for New Generations? A Socio-semiotic Study. *Advances in Language and Literary Studies* 7, 6 (2016), 56–69.
- [2] Victor Nikhil Antony and Chien-Ming Huang. 2024. ID.8: Co-creating visual stories with generative AI. *ACM Transactions on Interactive Intelligent Systems* 14, 3 (2024), 1–29. doi:10.1145/3672277
- [3] John Joon Young Chung, Woosook Kim, Kang Min Kang, Ece Kamar, and Walter S Lasecki. 2022. TaleBrush: Sketching Stories with Generative Pre-trained Language Models. In *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems*. ACM, 1–19. doi:10.1145/3491102.3502118
- [4] Thorsten M. Erle, Karoline Schmid, Simon H. Goslar, and Jared D. Martin. 2022. Emojis as Social Information in Digital Communication. *Emotion* 22, 7 (2022), 1529–1543. doi:10.1037/emo0000992
- [5] Jiaying Fu, Xiruo Wang, Kate Vi, Zhouyi Li, Chuyan Xu, and Yuqian Sun. 2025. “I Like Your Story!”: A Co-Creative Story-Crafting Game with a Persona-Driven Character Based on Generative AI. In *Extended Abstracts of the 2025 CHI Conference on Human Factors in Computing Systems (CHI EA '25)*. ACM. doi:10.1145/3706599.3721163 Interactivity.
- [6] William W Gaver, John Bowers, and Andrew Boucher. 2003. Ambiguity as a resource for design. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. ACM, 233–240. doi:10.1145/642611.642653
- [7] Petra Kralj Novak, Jasmina Smilović, Borut Sluban, and Igor Mozetič. 2015. Sentiment of Emojis. *PLOS ONE* 10 (12 2015), 1–22. doi:10.1371/journal.pone.0144296
- [8] Vivian Liu, Hanrui Zhang, and Lydia B Chilton. 2022. Opal: Multimodal Image Generation for News Illustration. In *Proceedings of the 35th Annual ACM Symposium on User Interface Software and Technology*. ACM, 1–14. doi:10.1145/3526113.3545621
- [9] Adyasha Maharana, Darryl Hannan, and Mohit Bansal. 2022. StoryDALL-E: Adapting Pretrained Text-to-Image Transformers for Story Continuation. In *Computer Vision – ECCV 2022*. Springer, 70–87. doi:10.1007/978-3-031-19836-6_5
- [10] Nathalie Henry Riche, Anna Offenwanger, Frederic Gmeiner, David Brown, Hugo Romat, Michel Pahud, Nicolai Marquardt, Kori Inkpen, and Ken Hinckley. 2025. AI-Instruments: Embodying Prompts as Instruments to Abstract and Reflect Graphical Interface Commands as General-Purpose Tools. In *Proceedings of the CHI Conference on Human Factors in Computing Systems (CHI '25)*. ACM, Article 1104, 18 pages. doi:10.1145/3706598.3714259
- [11] JD Zamfirescu-Pereira, Richmond Y Wong, Bjoern Hartmann, and Qian Yang. 2023. Why Johnny Can’t Prompt: How Non-AI Experts Try (and Fail) to Design LLM Prompts. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*. ACM, 1–21. doi:10.1145/3544548.3581388