

The Agentic Ideology: The Rise of the AI Agent

Zhaoxiang Tong
Institute of Network Society
China Academy of Art
Hangzhou, Zhejiang, China
tongzhaoxiang23@gmail.com

Yisi Liu
Mask Network
Shanghai, China
Institute of Network Society
China Academy of Art
Hangzhou, Zhejiang, China
yisiliu@mask.io

Abstract

This paper examines the techno-political transformation behind the rise of autonomous AI agents. Through three recent cases: the “Doubao Phone” interface conflict, Meta’s acquisition of Manus AI, and Alibaba’s task-oriented Qwen agents, we analyze an emerging governance discourse, called the Agentic Ideology. This discourse combines neoliberal empowerment rhetoric with techno-feudal rent extraction, reorganizing human action through agent mediation.

We conduct a comparative case analysis structured around three analytical layers: protocol, interface, and governance. At protocol level, we examine Anthropic’s Model Context Protocol (MCP) and its primitives for structuring delegation and decision visibility. At interface level, we compare structured parsing approaches with “pure vision” GUI agents such as UI-TARS. At governance level, we analyze interface conflicts between agents and major platforms.

We argue that the shift from Human-Computer Interaction to Agent-Computer Interaction redistributes control within digital infrastructure, and HCI research should address how public interfaces can support accountable human-agent collaboration.

CCS Concepts

• **Human-centered computing** → **HCI theory, concepts and models.**

Keywords

AI agents, Model Context Protocol, GUI agent, ghost work, agentic ideology

ACM Reference Format:

Zhaoxiang Tong and Yisi Liu. 2026. The Agentic Ideology: The Rise of the AI Agent. In *Extended Abstracts of the 2026 CHI Conference on Human Factors in Computing Systems (CHI EA '26)*, April 13–17, 2026, Barcelona, Spain. ACM, New York, NY, USA, 5 pages. <https://doi.org/10.1145/3772363.3798864>

1 Introduction

The contemporary internet was built for human eyes. Graphical user interfaces (GUI) assume a user who can see, click, and navigate, someone who reads screens rather than parses structures. This creates a central tension for Agent-Computer Interaction (ACI) [20]: increasingly capable and intellectual agents are expected to act

in digital environments whose interaction primitives were never designed for machine perception or machine reasoning. Industry responses have been swift and wide-ranging, from agent startups that promise end-to-end task execution, to emerging protocol standards meant to bridge models and systems, to vision-based GUI agents that operate directly on pixels. The rise of Manus AI, the release of the Model Context Protocol (MCP) [9], and the appearance of models such as UI-TARS [12] can be regarded as milestones in this shift.

However, most discussions of AI agents remain narrowly framed, either as questions of capability (how well agents complete tasks) or as generalized ethics (what agents should or should not do) [1, 15]. What is less developed is an account of interfaces and protocols as instruments of political-economic governance [16]. Existing work on platform capitalism has clarified how platforms commodify data and consolidate gate-keeping power, but it does not fully explain what changes when agents, rather than humans, become the primary “consumers” of interfaces [10]. In that setting, power no longer concentrates only around attention, but around preference: who can observe, infer, and shape the function that turns a vague intent into an executable plan. Similarly, the debate about standards such as MCP often stays at the level of interoperability, overlooking how the primitives of the protocol distribute authority and responsibility across models, applications, and users. Finally, celebratory narratives of agent automation routinely obscure the labor that makes these systems possible, including the large-scale data annotation and operational work often discussed as ghost work.

This paper adopts a three-layer analytical perspective, protocol, interface, and governance, to structure our analysis of how the current wave of AI agents is reconfiguring technical architectures and, with them, the distribution of control in digital infrastructure. Section 2 examines the rise of agent technologies and the “rebellion myth” that accompanies them. It then provides a close reading of MCP as a protocol that formalizes new divisions of agency through its primitives, and traces how GUI agents connect to labor regimes that remain largely invisible in mainstreams. The paper also analyzes emerging “interface wars”, where platforms respond to agentic automation with new forms of technical containment. We argue that if the public character of digital infrastructure is not revisited, agents are unlikely to function as straightforward liberators. They may instead become gentle overseers in a new kind of digital plantation—polite in interaction, but deeply consequential in how they reorganize power.



This work is licensed under a Creative Commons Attribution 4.0 International License. *CHI EA '26, Barcelona, Spain*

© 2026 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-2281-3/26/04

<https://doi.org/10.1145/3772363.3798864>

2 The Rise of the Agent: A New Myth of Rebellion

On December 29, 2025, Meta acquired Manus AI [11], a company specializing in agents capable of executing complex tasks. The deal was followed by geopolitical scrutiny, an early signal that agent technologies are likely to be treated as strategic assets rather than ordinary consumer software. Manus was founded in China, later relocated to Singapore, and reportedly achieved rapid growth in annual recurring revenue (ARR), reflecting strong demand for agents that can handle complex situations. This case points to a broader shift in the competitive landscape: for customer-facing AI, the advantage increasingly depends on their grounding capability, being able to connect language with concrete actions, tools, and workflows in real settings.

The name "Manus", derived from "mens et manus" ("mind and hand") [11], encapsulates a familiar creed: knowledge matters insofar as it can be made to act in the world. Contemporary AI agents extend this creed into a seductive story of rebellion. The promise is immediate intent satisfaction. Users are told they no longer need to handle the friction and enclosure of applications; they can state an intention, and an agent, cast as a loyal servant, will traverse walled gardens on their behalf. In this narrative, the agent becomes an invisible "hand", a digital extension of the user's will that supposedly frees them from the labor of graphical interfaces. This rhetoric echoes older counter-cultural imaginations of computing, including Jeffersonian fantasies of digital democracy in which powerful tools restore autonomy to individuals against institutional control [3]. Today, users are encouraged to fully unlock the power of agents, as a means to challenge the predominant and monopolistic platform governance.

2.1 Protocol: The Architecture of New Sovereignty

The rise of AI agents has been enabled by rapid advances in large language models (LLMs). Yet prior to 2025, the ecosystem exhibited a persistent disconnect. On one side, increasingly capable models remained confined to text boxes. On the other, personal and organizational data, along with everyday workflows, were locked behind the fences of specialized software and platform ecosystems. In late 2024, Anthropic introduced the Model Context Protocol (MCP) as an attempt to bridge this gap.

At a technical level, MCP uses a JSON-RPC 2.0 compatible message framing to support cross-language and cross-environment interoperability between agent hosts and external services [9]. The reference architecture follows a Host-Client-Server pattern, introducing a configurable boundary between model execution and user systems. MCP supports local transport (stdio) as well as remote transport (SSE/HTTP), enabling modular connectivity across deployment contexts. Interaction is organized around three primitives, tools, resources, and prompts, which structure what agents can do, what they can see, and how their behavior is controlled [8].

- **Tools (model-mediated):** Tools allow an LLM to discover and invoke functions based on context. In practice, this shifts parts of operational decision-making from the user's direct manipulation toward delegated action selection and sequencing. It also introduces new security risks, as execution shifts




Figure 1: Delegation of User Control through MCP's three primitives: tools, resources, and prompts.

from deterministic code to probabilistic model outputs. This in turn makes it easier for agent providers to favor their own affiliated tool chains.

- **Resources (application-governed):** Resources expose user data to the agent, including files, messages, personal information, histories, etc. MCP also supports subscription-like mechanisms for updates, moving from static retrieval toward dynamic context awareness. This creates a new channel through which providers can observe not only outcomes but also traces of intent and workflow, which might be described as the extraction of "cognitive surplus" from the process of doing, not just the result. Users work on their own devices, yet continuously yield behavioral data as a form of digital rent.
- **Prompts (nominally user-configurable):** Prompts provide reusable templates that guide model interaction. While often treated as a locus of user control, their effects are bounded by model training, instruction hierarchies, and system-level constraints. As agents take on more complex workflows, the user can drift from operator to manager, overseeing delegated labor whose internal reasoning and execution are only partially legible, producing a form of separation between the person and the process.

Despite the promise of MCP, early measurements suggest a fragmented and brittle ecosystem. One study [6] of 8,401 servers across multiple MCP registries reports that over 50.9% of public projects are invalid or low-value, and only 6.9% are indexed across major marketplaces. The study also identifies structural security risks, including dependency monocultures and description-code inconsistencies in approximately 13% of servers, where the tool's natural-language description diverges from its actual execution logic. These findings imply that agents remain confined to a small, manually configured tool set—an adoption and trust bottleneck that constrains generalization.

2.2 Interface: The Final Imitation Game

In 1950, Alan Turing's "Computing Machinery and Intelligence" reframed questions of intelligence through the "Imitation Game", asking whether machines could perform convincingly in human

terms [18]. Today, GUI agents revive this logic in a distinctly infrastructural register. Rather than persuading a human interlocutor, they must pass the practical tests that define contemporary access: authentication checks, anti-bot systems, and boring tasks that separate permitted users from automated machines.

Early web automation typically relied on structured representations such as the Document Object Model (DOM) or operating-system accessibility trees. Agents could read HTML or XML-like hierarchies, identify specific elements, and trigger associated events. This pathway, however, is fragile: it depends on stable structure, compatible APIs, and continued platform tolerance. To escape these constraints, a different route has gained traction. GUI agents such as ByteDance's UI-TARS pursue "pure vision perception", taking screenshots as input and discarding reliance on underlying HTML or view hierarchies [12].

Furthermore, on December 1, 2025, ByteDance released the "Doubao Phone", a system-level agentic approach that sought to automate cross-application workflows through low-level permissions and simulated user operations (clicking, screen reading, and navigation). The attempt quickly met technical countermeasures from major platforms, including execution environment checks and account restrictions when the agent tried to invoke key workflows within super-app ecosystems such as WeChat and Alipay. Whatever the specific implementation details, this approach is signaling a broader shift: when agents operate through the same surface channels as humans, platform boundaries become contested at the level of interface behavior and enforcement.

In this conflict, platforms such as Tencent and Alibaba publicly frame their resistance in terms of user privacy and regulatory compliance, arguing that invasive agents pose systemic risks by accessing sensitive payment data and personal information without adequate authorization. Yet beneath this legitimate security rhetoric lies a more structural motive: defending the rent-extracting logic of the traditional app economy. In that model, platforms command the user's full attention, that users must open applications, navigate conversion funnels, view in-feed advertisements, and submit to recommendation algorithms, generating vast advertising and transaction revenues, all by themselves. The Doubao Phone's agentic approach threatened to let users complete these workflows silently in the background, bypassing splash-screen ads and private-domain marketing channels entirely. In effect, agent automation erodes the attention-based foundations of platform monetization.

This "interface war" matters because pixel-level agents can, in principle, route around formal platform APIs. Pure-vision approaches effectively demote applications into back-end stores of capability while positioning the agent as the primary front-end mediator. Training such agents requires large-scale human interaction data. Some of this data can be gathered through accessibility APIs and user feedback programs. More broadly, however, agent ecosystems depend on the labor of data production and annotation, work that is rarely visible in product narratives [5, 13]. The cliché industry saying that "every bit of AI intelligence hides a bit of human labor" captures a real contradiction: frontier capabilities are frequently built on low-paid, repetitive work at scale.

Data annotation sits at the base of the AI supply chain, and its organization often resembles labor-intensive industrial models. Dataset quality shapes both the upper bound of model performance

and the character of downstream bias; as a result, firms commonly enforce strict quality-control processes to keep error within acceptable limits. In China, the sector has also developed a distinct spatial pattern of value extraction. Large-scale, standardized annotation work tends to move toward lower-cost regions in central and western provinces, while higher-value roles closer to clients—project management, model tuning, and system integration—concentrate in eastern hubs and capture unbalanced, high premiums. To illustrate this spatial pattern, we retrieved approximately 1,500 firms registered under data-annotation related keywords from the public business registry platform "Qichacha". While not exhaustive, the resulting distribution offers a rough picture of the sector's geographic concentration across China (see Figure 2).



Figure 2: Geographic distribution of data-annotation firms in China.

2.3 Governance: The Political Economy of the Agentic Era

On January 15, 2026, Alibaba announced an updated version of Qwen centered on an "AI Task Assistant" [2]. The model was deeply integrated across Alibaba's ecosystem, Taobao, Alipay, Fliggy, Amap, and other applications that anchor everyday life in China. With user authorization, the assistant could complete end-to-end workflows from decision to payment within a single conversational surface, reducing the need to traverse separate applications. Yet this smoothness is achieved through enclosure. Qwen does not seamlessly operate within WeChat, nor does it freely draw on JD.com data. In practice, it resembles a supercharged concierge for Alibaba's own ecosystem rather than an independent agent acting across platforms. This carries symbolic weight: competition increasingly shifts from traffic capture to control over the infrastructure through which intent is translated into action. Where platforms historically built walled gardens through closed APIs, agents now attempt to climb those walls through simulated interaction.

The rise of ACI presses on the most sensitive nerve of platform governance: data monopoly. When agents cross application boundaries to complete tasks, they collide with established governance regimes, producing what we describe as interface wars. Alibaba's ability to shape underlying cloud infrastructure, combined with the constrained mobility of agents outside its ecosystem, illustrates a form of enclosure that exceeds ordinary market rivalry. In the language of techno-feudalism, cloud capital becomes the "land" of the digital economy, and accumulation shifts from profit-making toward rent extraction [4, 19]. Closed ecosystems and proprietary infrastructure create dependencies through which firms must pay tribute to "cloud lords" in order to operate. Under the banner of security and user experience, open-internet narratives give way to authoritarian rule: protection in exchange for compliance, and access in exchange for surplus labor. While agents might have been expected to erode such fragmentation by enabling freer data flows, the more immediate tendency is consolidation: powerful actors use agents to reinforce, rather than tear down, the walls.

The attention economy of Web 2.0 is increasingly supplemented by an agentic economy [14]. Agents do not watch ads; they execute preferences. Scarcity shifts from human attention toward user preference that determines how an agent interprets a command such as "book the best ticket back home for me tomorrow." Whoever controls that interpretive layer, its rules, constraints, and monetization, holds unbalanced power. The liberation promised by agent narratives is therefore paradoxical. An agent might bypass a platform's interface and reduce friction for the user, but in doing so it can also turn that platform into a mere pipe, just like how SSL and TLS turned internet service providers into a pure infrastructure [17]. More importantly, the agent itself becomes a new, potentially more centralized intermediary.

3 Conclusion: The Agentic Ideology

As the title suggests, technical architectures are never only technical. The three cases examined in this paper trace a progression: from an independent startup absorbed into a platform giant, to an agentic system met with cross-platform resistance, to an enclosed ecosystem where agents reinforce monopoly. Together, they pull the discussion of AI agents from narrow technical ethics toward political-economic critique. The current push toward ACI is animated by a hybrid ideology that borrows the language of empowerment while reproducing enclosure through new infrastructural bottlenecks. It promises liberation by delegating subjectivity outward to agents, yet it may instead deepen dependency on platform power. This paper treats that ideology as an analytic lens rather than a settled claim, opening a line of inquiry that we intend to develop with fuller empirical and theoretical grounding.

At present, we are standing at a crossroads. The current configuration, platforms blocking agents and agents mimicking humans, resembles a Hobbesian war of all against all [7]. One possible resolution is the emergence of "interface agreements" such as MCP: a compromise in which platforms open controlled APIs and agents pay tolls or share an "intent tax." The result would be an internet increasingly organized around machine-to-machine transactions, with humans reduced to initiators of intent and final underwriters of value. Two trajectories appear plausible. The first is state-mandated

interoperability: as in prior regulatory interventions that required dominant platforms to relax link-blocking, the state could enforce standards that force apps to provide certified agents with standardized access, treating agents as regulated infrastructure. The second is oligopolistic peace: major firms could reach non-aggression pacts or cross-licensing agreements that partition the intent economy and squeeze independent entrants out of the market.

Absent public intervention, AI productivity may advance while social structure regresses toward more extreme monopoly. A small number of firms controlling "super agents" could become a new clergy, monopolizing the principal channel through which people interact with the digital world. The alternative is not to choose between platform capitalism and agentic enclosure, but to reassert the public character of digital infrastructure. If the agent is becoming the key to digital life, that key should not be held solely by ByteDance or Tencent. It should be treated as a democratically governed public utility, with enforceable algorithmic transparency and meaningful user data sovereignty. As Barbrook warned decades ago, the choice is between building a new Athens in cyberspace and inheriting a more efficient plantation, and this time supervised by an agent that asks, politely, what you would like it to do.

References

- [1] Deepak Bhaskar Acharya, Karthigeyan Kuppan, and B Divya. 2025. Agentic ai: Autonomous intelligence for complex goals—a comprehensive survey. *IEEE Access* (2025).
- [2] Alibaba Group. 2026. *Alibaba's Qwen App Advances Agentic AI Strategy by Turning Core Ecosystem Services into Executable AI Capabilities-Alibaba Group*. <https://www.alibaba.com/en-US/document-1948497434959151104>
- [3] Richard Barbrook and Andy Cameron. 1996. The californian ideology. *Science as culture* 6, 1 (1996), 44–72.
- [4] Cédric Durand. 2020. *Technofeudalisme: Critique de l'économie numérique*. Zones.
- [5] Mary L Gray and Siddharth Suri. 2019. *Ghost work: How to stop Silicon Valley from building a new global underclass*. Harper Business.
- [6] Hechuan Guo, Yongle Hao, Yue Zhang, Minghui Xu, Peizhuo Lv, Jiezhai Chen, and Xiuzhen Cheng. 2025. A Measurement Study of Model Context Protocol Ecosystem. *arXiv preprint arXiv:2509.25292* (2025).
- [7] Thomas Hobbes. 2017. *Leviathan*. Penguin Classics.
- [8] Xinyi Hou, Yanjie Zhao, Shengao Wang, and Haoyu Wang. 2025. Model context protocol (mcp): Landscape, security threats, and future research directions. *arXiv preprint arXiv:2503.23278* (2025).
- [9] LLC LF Projects. 2025. *Model Context Protocol*. <https://modelcontextprotocol.io/>
- [10] Junting Lu, Zhiyang Zhang, Fangkai Yang, Jue Zhang, Lu Wang, Chao Du, Qingwei Lin, Saravan Rajmohan, Dongmei Zhang, and Qi Zhang. 2025. Axis: Efficient human-agent-computer interaction with api-first llm-based agents. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 7711–7743.
- [11] Manus. 2025. *Manus Joins Meta for Next Era of Innovation*. <https://manus.im/blog/manus-joins-meta-for-next-era-of-innovation>
- [12] Yujia Qin, Yining Ye, Junjie Fang, Haoming Wang, Shihao Liang, Shizuo Tian, Junda Zhang, Jiahao Li, Yunxin Li, Shijue Huang, et al. 2025. Ui-tars: Pioneering automated gui interaction with native agents. *arXiv preprint arXiv:2501.12326* (2025).
- [13] ATM Rahman and Sharifa Sultana. 2025. Ghostcrafting AI: Under the Rug of Platform Labor. *arXiv preprint arXiv:2512.21649* (2025).
- [14] David M Rothschild, Markus Mobius, Jake M Hoffman, Eleanor W Dillon, Daniel G Goldstein, Nicole Immorlica, Sonia Jaffe, Brendan Lucier, Aleksandr Slivkins, and Matthew Vogel. 2025. The Agentic Economy. *arXiv preprint arXiv:2505.15799* (2025).
- [15] Ranjan Sapkota, Konstantinos I Roumeliotis, and Manoj Karkee. 2025. Ai agents vs. agentic ai: A conceptual taxonomy, applications and challenges. *arXiv preprint arXiv:2505.10468* (2025).
- [16] Yonadav Shavit, Sandhini Agarwal, Miles Brundage, Steven Adler, Cullen O'Keefe, Rosie Campbell, Teddy Lee, Pamela Mishkin, Tyna Eloundou, Alan Hickey, et al. 2023. Practices for governing agentic AI systems. *Research Paper, OpenAI* (2023).
- [17] Wazen M Shbair, Thibault Chole, Jérôme François, and Isabelle Chrisment. 2020. A survey of https traffic and services identification approaches. *arXiv preprint arXiv:2008.08339* (2020).

- [18] Alan M Turing. 2007. Computing machinery and intelligence. In *Parsing the Turing test: Philosophical and methodological issues in the quest for the thinking computer*. Springer, 23–65.
- [19] Yanis Varoufakis. 2024. *Technofeudalism: What Killed Capitalism*. Melville House.
- [20] John Yang, Carlos E Jimenez, Alexander Wettig, Kilian Lieret, Shunyu Yao, Karthik Narasimhan, and Ofir Press. 2024. Swe-agent: Agent-computer interfaces enable automated software engineering. *Advances in Neural Information Processing Systems* 37 (2024), 50528–50652.