

Towards a Structured Framework for Large-Scale AI-Driven Thematic Analysis on Domain-Specific Customer Feedback

Jade Logan
Swansea University
Swansea, United Kingdom
850852@swansea.ac.uk

Nicholas Micallef
Swansea University
Swansea, United Kingdom
nicholas.micallef@swansea.ac.uk

Daniele Doneddu
Swansea University
Swansea, United Kingdom
d.doneddu@swansea.ac.uk

Kevin McLafferty
Swansea University
Swansea, United Kingdom
kevin.mclafferty@swansea.ac.uk

Muneeb Ahmad
Swansea University
Swansea, United Kingdom
m.i.ahmad@swansea.ac.uk

Abstract

The recent surge in Large Language Models (LLMs) has led to their use in large-scale qualitative analysis, yet many approaches remain opaque, methodologically weak, and low in reflexivity, limiting their validity. This paper presents a structured framework for fully automated, LLM-supported thematic analysis grounded in Braun and Clarke's method, evaluating two prompting strategies: iterative analysis and progressive batch expansion. A pilot study with financial-services customer experience (CX) feedback compares LLM-generated themes against independently coded human themes across Net Promoter Score comments. Results show that iterative analysis yields Jaccard indices of 0.71–0.86 with human themes, compared to 0.44–0.71 for batch processing, with higher alignment for dominant CX issues but weaker coverage of lower-frequency concerns. The framework offers a structured, scalable approach to CX analysis and clarifies where LLM-led thematic analysis converges with—and diverges from—human judgement.

CCS Concepts

• Human-centered computing → Empirical studies in HCI.

Keywords

Large Language Models, LLMs, Qualitative Analysis, Thematic Analysis, Customer Experience

ACM Reference Format:

Jade Logan, Nicholas Micallef, Daniele Doneddu, Kevin McLafferty, and Muneeb Ahmad. 2026. Towards a Structured Framework for Large-Scale AI-Driven Thematic Analysis on Domain-Specific Customer Feedback. In *Extended Abstracts of the 2026 CHI Conference on Human Factors in Computing Systems (CHI EA '26)*, April 13–17, 2026, Barcelona, Spain. ACM, New York, NY, USA, 5 pages. <https://doi.org/10.1145/3772363.3798789>

1 Introduction

The digitalisation of many services has led to a substantial increase in unstructured textual customer experience (CX) data, creating significant scalability challenges for qualitative analysis. Manual thematic analysis remains valuable but slow, resource-intensive, and difficult to apply consistently at the scale of modern CX datasets [2, 5, 6]. These pressures have fuelled growing academic interest in large language models (LLMs) to support or automate qualitative analysis tasks, including thematic analysis [15].

Emerging work suggests that LLM-assisted workflows can accelerate coding and surface themes that overlap with human analysis, particularly when prompts are aligned with established qualitative methods [7–9, 15, 18]. At the same time, scholars highlight serious concerns about opacity, bias, loss of reflexivity, and the risk of treating model outputs as authoritative in the absence of transparent procedures and systematic comparison to human analysis [1, 11, 12, 18]. Many studies focus on relatively small, well-structured datasets, use LLMs primarily as assistants, rely on single-pass analyses, or provide limited detail about how interactions with the model are structured, leaving open questions about their behaviour on short, unstructured CX comments collected at scale [2, 7–9, 15].

This study works with a large financial institution to examine fully automated, LLM-driven thematic analysis of Net Promoter Score (NPS) comments, consisting of promoters, passives, and detractors. It embeds Gemini 2.5 Pro within Braun and Clarke's six-phase framework [3], combines probability-based stopword removal to handle domain-specific terms [20], and compares two prompting strategies, iterative analysis (repeated passes on fixed samples) and progressive batch expansion (incrementally larger batches across full corpus), against human-led thematic analyses. We therefore address the following research questions: *To what extent do fully automated, LLM-powered thematic analyses align with human-led thematic analysis on large-scale CX data, and how do different prompting strategies (iterative analysis vs. progressive batch expansion) affect this alignment?*

Across three NPS subsets, our findings show that the structured, iterative LLM workflow produces moderate to high overlap with human themes on dominant CX issues, but under-represents several lower-frequency, interpretative concerns and focuses on general themes as batch size increases. On this basis, the paper contributes: (1) a structured, fully automated LLM workflow for thematic analysis of large-scale CX data; (2) an empirical comparison



This work is licensed under a Creative Commons Attribution 4.0 International License. *CHI EA '26, Barcelona, Spain*

© 2026 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-2281-3/26/04

<https://doi.org/10.1145/3772363.3798789>

of iterative analysis and progressive batch expansion; and (3) reflections on where LLM-supported thematic analysis approximates human-derived themes and where it remains limited.

2 Related Work

Traditional manual thematic analysis remains valued for producing rich and nuanced insights [16], but is labour-intensive and slow for large datasets [5, 6], which is increasingly impractical in industrial settings that demand rapid decision-making. Most critically, in high-volume domains such as financial services, the scale of routinely collected qualitative data now exceeds what can realistically be analysed through purely manual approaches [2]. Metrics such as Net Promoter Score (NPS) are widely used to measure customer satisfaction; however, research suggests these metrics often fail to capture the complexity of customer sentiment and may be misinterpreted by users [13]. As a result, analysing the qualitative comments accompanying these metrics remains essential for understanding the underlying drivers of customer experience.

LLMs have been proposed as one way to address these scalability and efficiency challenges [14]. Empirical studies report that LLM-assisted workflows can substantially reduce analysis time and support broader thematic coverage, with LLMs generating first-pass codes and candidate themes that often overlap with human coding and assist with data exploration, codebook development, and triage of large qualitative corpora [19]. Further work demonstrates that LLMs can reliably identify both semantic and latent meanings, with strong agreement between LLM-generated codes and expert human coders, including in the analysis of sensitive datasets such as patient complaints [17]. Beyond automation, recent work in HCI and qualitative methods increasingly positions LLMs as “research assistants” that, under structured prompting, can surface alternative framings and produce internally consistent coding rationales [7, 8, 15, 18].

Meanwhile, the use of LLMs in qualitative analysis remains strongly contested. Braun and Clarke argue that LLMs raise environmental concerns due to their computational demands, lack of lived experience for meaningful interpretation, and risk displacing qualitative analysis as a distinctly human practice [11]. Further critiques warn that general-purpose LLMs can introduce serious bias and may not reflect the epistemological commitments of qualitative research [1]. These concerns have called for transparent, systematic workflows in which LLMs are constrained by established qualitative frameworks and evaluated against human analysis rather than assumed to be valid by default [9, 12, 18]. Existing studies typically use LLMs as assistants or in single-pass analyses, giving little attention to how prompting strategies affect thematic stability and alignment, especially for short, unstructured customer comments. Our study addresses this gap.

3 Methodology

This study adopts a structured qualitative research design to explore whether fully automated, LLM-powered thematic analysis can produce outcomes that meaningfully align with human-led thematic analysis on large-scale financial services CX data. LLM analysis is embedded within Braun and Clarke’s six-phase thematic analysis framework [3], incorporating iterative analysis, progressive batch expansion, and probability-based stopword removal to

support analytical transparency and coherence. Automated outputs are compared against independently coded human thematic analyses to assess alignment and to characterise process limitations rather than to claim equivalence.

3.1 Dataset

The initial dataset comprised of 10,533 free-text customer comments accompanied by corresponding NPS scores. Following manual cleaning, the dataset was reduced to 6,561 anonymised comments. These were stratified into three subsets based on their NPS category: promoters, passives, and detractors. To enable human analysis within a reasonable timeframe, a random sample of 50 comments was selected from each, yielding 150 comments for manual thematic analysis. The LLM analysed the full 6,561-comment corpus, whereas human coders worked only on this stratified sample; all quantitative comparisons between human and LLM themes are therefore restricted to the sampled comments.

3.2 Data Preprocessing – Stopword Removal

To reduce distortion from context specific high-frequency terms, we conducted a targeted evaluation of stopword removal inspired by Xue et al’s., probability-based approach to identify non-informative terms in topic models [20]. This was due to our corpus containing domain-specific vocabulary (e.g., the financial institution’s name) that risked dominating theme generation, we adapted this probability-based logic to our CX data. Term-level probability distributions were computed over the corpus and inspected to identify highly frequent but semantically uninformative tokens. These terms were added to a custom stopword list, which was applied prior to analysis. The aim was to suppress context-specific noise without removing meaningful customer experience language, so that resulting themes were driven primarily by customer sentiment rather than artefacts of organisational terminology or dataset structure.

3.3 LLM Analytical Process

3.3.1 LLM selection. Model selection was constrained by the financial institution’s internal security and compliance policies. Only pre-approved cloud-based models could be used, so Gemini 2.5 Pro was selected for this study. The goal was not to benchmark models but to investigate how a contemporary, general-purpose LLM performs when embedded within a structured, systematic thematic analysis workflow under realistic organisational constraints.

3.3.2 Methodological grounded prompting. Prompt development evolved iteratively throughout the study. Initial experiments used persona-based prompting, a strategy employed in early-stage LLM-assisted qualitative analysis [15, 21]. Three role-playing prompts were trialled: a student learning qualitative analysis, a CX practitioner working within a country specific financial institution, and a qualitative researcher applying Braun and Clarke’s method. These personas were intended to encourage the model to adopt perspectives aligned with the analytical aims of the study. While this provided some structure, outputs didn’t reliably follow established qualitative procedures and interpretative depth varied considerably between runs. To address these inconsistencies, persona-based prompts were replaced with methodologically grounded prompts that explicitly followed Braun and Clarke’s method. The

final prompt specified each phase (familiarisation, coding, theme construction, theme review, defining and naming themes, and reporting) and required the model to outline its decisions at each stage. This design aimed to enforce a transparent, procedurally aligned workflow and to constrain the model’s internal heuristics within a recognised qualitative framework [3, 4].

3.3.3 Qualitative analysis method selection. Braun and Clarke’s six-phase thematic analysis method was selected as it provides a clear, systematic framework that can be operationalised in prompt form and is suitable for short, high-volume qualitative datasets [4], such as NPS comments. The method’s focus on iterative engagement with the data and explicit theme refinement aligns with recent calls for structured, reflexive use of LLMs in qualitative research [9, 15, 18].

3.3.4 Iterative and progressive analysis. To respond to critiques that LLMs lack reflexivity and may generate superficially coherent but analytically shallow themes [1, 11], two complementary experiments were conducted within this prompt framework, both of which are provided in the supplementary materials:

Iterative analysis. For each NPS subset, the model re-analysed the same full corpus across multiple iterations using a structured prompt to determine whether repeated passes produced more stable or conceptually refined themes. This procedure was intended to approximate the repeated revisiting and refinement of themes characteristic of human-led reflexive thematic analysis, and to examine how iteration affected theme stability and granularity.

Progressive batch expansion. Separately, the full corpus was presented to the LLM in progressively larger batches (e.g., 50, 100, and 200 comments). With each batch expansion, the model was instructed to integrate the newly analysed comments into its existing theme structure, revising or consolidating themes as necessary. This design mirrors how human researchers often build analytic understanding over time, as additional data are incorporated, and allowed us to examine how increasing batch size influences the number, specificity, and stability of themes.

These experiments tested whether structured iteration and corpus expansion enhance theme stability and interpretability, rather than resolving concerns about LLM reflexivity.

3.4 Analysis Validation

To reduce bias, the lead author, who had prior experience working with Gemini on this dataset, did not participate in manual coding. Two other authors independently conducted thematic analysis on the 150-comment sample using Braun and Clarke’s framework and a third domain-expert author (with over ten years of experience in business management) reviewed and reconciled the human themes to produce a consensus theme set for each NPS subset. Cohen’s kappa was used for inter-coder agreement between the two pairs of human annotators providing a baseline for human–human reliability [4]. Agreement between human and LLM themes was then assessed at the theme-set level using the Jaccard index, which quantifies the proportion of shared themes relative to the total number of distinct themes identified across both human and LLM analyses [10]. This dual-metric approach distinguishes (a) the extent of human–human agreement from (b) the overlap between LLM-derived theme and human-derived themes.

3.5 Ethical Consideration

As the analysis employed the cloud-based Gemini model, strict anonymisation procedures were implemented to minimise any risk of identifying individuals or organisations. Prior to any upload, the first author manually anonymised the dataset, replacing all personal names and organisational identifiers, and other potentially identifying information with the mask ‘*****’. No raw customer identifiers or account details were transmitted to the model.

4 Findings and Discussion

Across the three NPS-stratified subsets, the LLM analysed 6,561 customer comments (2,369 promoters, 1,304 passives, and 2,888 detractors). For validation, human analysts independently coded a stratified sample of 150 comments (50 per subset). Overall, LLM-generated themes from the iterative method showed stronger alignment with human-identified themes than the batch-based approach, particularly for high-frequency experiential issues.

4.1 Human–LLM Thematic Alignment

Inter-coder agreement between the two pairs of human annotators on the sample was $\kappa = 0.76, 0.87$, indicating moderate/substantial reliability for human thematic analysis. Agreement between human consensus and LLM-generated themes, measured using Jaccard index, was consistently higher for the iterative method than for the batch method. For promoters, passives, and detractors respectively, Jaccard indices were 0.44, 0.71, and 0.64 for batch processing, compared to 0.78, 0.86, and 0.71 under iterative analysis, indicating stronger overlap with human themes across all datasets.

Qualitatively, human and LLM analyses converged on several core themes. For promoters, both identified digital banking experience and transparency with the LLM expressing the latter as “product fees and transparency”. For detractors, both captured poor digital banking experiences and the importance of human relationships, framed by the LLM as “decline of relationship and in-person banking”. Passives showed the highest alignment, with both approaches identifying branch accessibility, digital banking usability, and relationship management.

Misalignment arose for lower-frequency or more interpretative themes. In detractor data, human additionally identified fraud safeguarding, perceived risk aversion, know-your-customer concerns, lack of a unique selling point, and barriers to switching financial institutions. The LLM instead introduced themes such as positive sentiment within negative feedback (detractors) and problem resolution (promoters), suggesting that while LLMs reliably surface dominant patterns, they may abstract away from infrequent, contextually nuanced concerns.

Promoter, passive, and detractor datasets also differed in heterogeneity and affective intensity. Promoter comments were largely positive and homogeneous, with occasional tensions (e.g., valuing digital convenience yet missing human contact) embedded in overall favourable assessments. Detractor comments were more heterogeneous and analytically dense, often combining multiple operational, relational, and strategic issues with strong affective language, increasing interpretative complexity and reducing convergence. Passive comments were more neutral and narrower in scope mapping

onto service- and process-oriented themes and producing the highest human–LLM alignment. These patterns indicate that dataset heterogeneity and emotional intensity seem to shape human–LLM alignment, with lower-salience feedback more amenable to automated interpretation than highly positive or critical responses.

4.2 Effects of Iterative Analysis

Iterative analysis improved thematic specificity and stability across all NPS categories, but with different convergence dynamics. For promoters, the LLM stabilised quickly: after two iterations on the same 50-comment subset, it indicated that no further refinement was required, and later iterations produced only minor wording changes. Early passes distinguished staff interaction, digital experience, and long-term loyalty; subsequent iterations clarified boundaries (e.g., separating general satisfaction and loyalty from staff interaction and service quality). For passives, theme counts increased from seven to nine across iterations, with labels remaining stable once introduced; here, the model focused on accommodating previously misfitting comments by adding themes when necessary and explicitly treated contextual descriptors as non-evaluative.

Detractor comments followed a more extended refinement trajectory. Across four iterations, the model repeatedly revisited boundaries between relationship-based and channel-based problems, restructuring initial themes (e.g., branch decline, difficulty contacting humans, poor remote channels) into broader categories such as “decline of relationship and in-person banking” and “poor remote and digital channels”. Outlying comments combining digital frustration with feelings of being undervalued prompted revisions to “general negative sentiment and service failures” and “poor customer service and feeling undervalued” until all cases could be accommodated. Across all datasets, theme counts showed initial expansion followed by consolidation, with the LLM articulating rationales for merges and renamings (e.g., combining “staff quality and helpfulness” with “in-branch and human interaction” into “staff interaction and service quality”). This suggests that structured iterative prompting can enforce a procedural form of re-examination and refinement, even though interpretive depth remains constrained by the model’s training and prompt design.

4.3 Effects of Progressive Batch Expansion

Progressive batch expansion produced increasingly stable but more generic thematic structures. For promoters, a batch size of 50 led the LLM to identify seven themes, several overlapping with human coding, including relationship- and process-focused themes. As additional batches were added, themes increased to nine and introduced distinctions such as separating general satisfaction and loyalty from staff interaction and service quality. However, at batch sizes of 100 and 200 comments, these nuances were folded into broader categories (e.g., a single general satisfaction theme), reducing granularity. Detractors had a similar pattern: at 50-comment batches, the LLM differentiated specific dissatisfaction themes (e.g., difficulty contacting humans, branch decline, feeling undervalued), but larger batches subsumed these under broader categories like general negative sentiment and service failures, while interpretative human-only themes (perceived risk aversion, lack of unique selling point, barriers to switching) remained absent.

Passive comments displayed early thematic convergence and little subsequent differentiation under progressive batch expansion. At 50-comment batches, the LLM identified six themes, with no new themes emerging as further batches were added. At 100 and 200 comments, seven themes were identified in the initial batch and remained unchanged, with the model reporting analysis of fewer comments than the passive dataset contained and providing only high-level descriptions of its process. Across all datasets, progressive batch expansion led to rapid stabilisation but at the cost of reduced transparency and sensitivity to emerging distinctions: the LLM did not consistently surface every coded comment or expose intermediate reasoning, and the number of supporting quotations was lower than the number of analysed comments. This constrains interpretability and limits the extent to which progressive batch expansion can be considered transparent or methodologically robust, despite the apparent stability of resulting theme structures.

4.4 Observed Limitations and Failure Modes

Several limitations and failure modes were observed, more pronounced for analytical transparency than thematic coverage. The batch approach was most problematic, as it did not surface all analysed comments or expose intermediate reasoning, and prompt refinement did not lead LLM to consistently provide complete rationales, undermining reproducibility. In contrast to concerns about under-representation of rare issues, the iterative method handled edge cases appropriately. For example, feedback about a deceased customer was explicitly flagged as a singular instance not fitting existing themes, a nuance absent from both the batch condition and human coding. Moreover, human analysts identified several themes the LLM did not surface (e.g., fraud safeguarding, barriers to switching), indicating that LLMs may not fully capture contextually embedded or latent concerns without explicit prompting. A further limitation is the small number of manually coded comments used to benchmark LLM results. Thus, the findings should be considered indicative rather than definitive. Future research should include larger manually coded samples for stronger validation. Additionally, only one model was tested in a single context; other models may introduce distinct biases.

5 Conclusion

This study examined how far fully automated, LLM-powered thematic analysis of CX data can approximate human-led thematic analysis. Findings show moderate to high alignment on dominant, high-frequency themes, but weaker coverage of lower-frequency and more interpretative concerns, and increasing genericity as batch size grows. Iterative analysis yielded more granular and human-aligned themes than progressive batch expansion, underscoring the importance of prompting strategy when deploying LLMs for thematic analysis. Overall, a systematic, fully automated LLM workflow can produce partially human-comparable theme structures for short, unstructured CX comments, while also revealing clear limitations around nuance, transparency, and reflexivity. LLM-supported thematic analysis appears promising as a component of scalable qualitative CX workflows, but not yet sufficient for standalone use without human oversight.

References

- [1] Julian Ashwin, Aditya Chhabra, and Vijayendra Rao. 2023. Using Large Language Models for Qualitative Analysis can Introduce Serious Bias. *Sociological Methods & Research* 0, 0 (2023), 00491241251338246. doi:10.1177/00491241251338246
- [2] Cauã Ferreira Barros, Bruna Borges Azevedo, Valdemar Vicente Graciano Neto, Mohamad Kassab, Marcos Kalinowski, Hugo Alexandre D. Do Nascimento, and Michelle C.G.S.P. Bandeira. 2025. Large Language Model for Qualitative Research: A Systematic Mapping Study. In *2025 IEEE/ACM International Workshop on Methodological Issues with Empirical Studies in Software Engineering (WSESE)*, 48–55. doi:10.1109/WSESE66602.2025.00015
- [3] Virginia Braun and Victoria Clarke. 2006. Using thematic analysis in psychology. *Qualitative Research in Psychology* 3, 2 (2006), 77–101. doi:10.1191/1478088706qp063oa
- [4] Virginia Braun and Victoria Clarke. 2021. One size fits all? What counts as quality practice in (reflexive) thematic analysis? *Qualitative Research in Psychology* 18, 3 (2021), 328–352. doi:10.1080/14780887.2020.1769238
- [5] Ashley Castleberry and Amanda Nolen. 2018. Thematic analysis of qualitative research data: Is it as easy as it sounds? *Currents in Pharmacy Teaching and Learning* 10, 6 (2018), 807–815. doi:10.1016/j.cptl.2018.03.019
- [6] Arif A Cevik and Fikri M Abu-Zidan. 2025. Utilizing AI-Powered Thematic Analysis: Methodology, Implementation, and Lessons Learned. *Cureus* 17, 6 (2025), e85338.
- [7] Zackary Okun Dunivin. 2024. Scalable Qualitative Coding with LLMs: Chain-of-Thought Reasoning Matches Human Performance in Some Hermeneutic Tasks. arXiv:2401.15170 [cs.CL] <https://arxiv.org/abs/2401.15170>
- [8] Tim Fischer and Chris Biemann. 2024. Exploring Large Language Models for Qualitative Data Analysis. In *Proceedings of the 4th International Conference on Natural Language Processing for Digital Humanities*, Mika Härmäläinen, Emily Öhman, So Miyagawa, Khalid Alnajjar, and Yuri Bizzoni (Eds.). Association for Computational Linguistics, Miami, USA, 423–437. doi:10.18653/v1/2024.nlp4dh-1.41
- [9] Adam S. Hayes. 2025. “Conversing” With Qualitative Data: Enhancing Qualitative Research Through Large Language Models (LLMs). *International Journal of Qualitative Methods* 24 (2025), 16094069251322346. doi:10.1177/16094069251322346
- [10] Nilesh Jain, Seyi Adeyinka, Leor Roseman, and Aza Allsop. 2025. Multi-LLM Thematic Analysis with Dual Reliability Metrics: Combining Cohen’s Kappa and Semantic Similarity for Qualitative Research Validation. arXiv:2512.20352 [cs.CL] <https://arxiv.org/abs/2512.20352>
- [11] Tanisha Jowsey, Virginia Braun, Victoria Clarke, Deborah Lupton, and Michelle Fine. 2025. We reject the use of generative artificial intelligence for reflexive qualitative research. *Qualitative Inquiry* (2025), 10778004251401851. doi:10.2139/ssrn.5676462
- [12] Chaewon Kim, Fengfeng Ke, Alex Barrett, and Nuodi Zhang. 2025. From Code to Insight: How LLMs Help and Hinder Qualitative Research. In *Artificial Intelligence in Education. Posters and Late Breaking Results, Workshops and Tutorials, Industry and Innovation Tracks, Practitioners, Doctoral Consortium, Blue Sky, and WideAIEd*, Alexandra I. Cristea, Erin Walker, Yu Lu, Olga C. Santos, and Seiji Isotani (Eds.). Springer Nature Switzerland, Cham, 28–35.
- [13] Jade Logan, Daniele Doneddu, Kevin McLafferty, Muneeb I. Ahmad, and Nicholas Micallef. 2025. Behind Customer Satisfaction Metrics: Exploring User Perceptions of Net Promoter Score (NPS) as a Measure of Satisfaction. In *Human-Computer Interaction – INTERACT 2025: 20th IFIP TC 13 International Conference, Belo Horizonte, Brazil, September 8–12, 2025, Proceedings, Part II* (Belo Horizonte, Brazil). Springer-Verlag, Berlin, Heidelberg, 135–156. doi:10.1007/978-3-032-05002-1_7
- [14] Walter S Mathis, Sophia Zhao, Nicholas Pratt, Jeremy Weleff, and Stefano De Paoli. 2024. Inductive thematic analysis of healthcare qualitative interviews using open-source large language models: How does it compare to traditional methods? *Computer Methods and Programs in Biomedicine* 255 (2024), 108356. doi:10.1016/j.cmpb.2024.108356
- [15] Kien Nguyen-Trung. 2025. ChatGPT in Thematic Analysis: Can AI become a research assistant in qualitative research? *Quality & Quantity* 59 (2025), 4945–4978.
- [16] Lorelli S. Nowell, Jill M. Norris, Deborah E. White, and Nancy J. Moules. 2017. Thematic Analysis: Striving to Meet the Trustworthiness Criteria. *International Journal of Qualitative Methods* 16, 1 (2017), 1609406917733847. doi:10.1177/1609406917733847
- [17] Maximo R Prescott, Samantha Yeager, Lillian Ham, Carlos D Rivera Saldana, Vanessa Serrano, Joey Narez, Dafna Paltin, Jorge Delgado, David J Moore, and Jessica Montoya. 2024. Comparing the Efficacy and Efficiency of Human and Generative AI: Qualitative Thematic Analyses. *JMIR AI* 3 (2 Aug 2024), e54482. doi:10.2196/54482
- [18] Hope Schroeder, Marianne Aubin Le Quéré, Casey Randazzo, David Mimmo, and Sarita Schoenebeck. 2025. Large Language Models in Qualitative Research: Uses, Tensions, and Intentions. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems (CHI '25)*. Association for Computing Machinery, New York, NY, USA, Article 481, 17 pages. doi:10.1145/3706598.3713120
- [19] Ansh Sharma and James R Wallace. 2025. DeTAILS: Deep Thematic Analysis with Iterative LLM Support. In *Proceedings of the 7th ACM Conference on Conversational User Interfaces (CUI '25)*. Association for Computing Machinery, New York, NY, USA, Article 28, 7 pages. doi:10.1145/3719160.3735657
- [20] Shuangjiang Xue, Pierre Le Bras, David A. Robb, Mike J. Chantler, and Stefano Padilla. 2025. Visual Exploration of Stopword Probabilities in Topic Models. arXiv:2501.10137 [cs.HC] <https://arxiv.org/abs/2501.10137>
- [21] He Zhang, Chuhao Wu, Jingyi Xie, Yao Lyu, Jie Cai, and John M. Carroll. 2025. Harnessing the power of AI in qualitative research: Exploring, using and re-designing ChatGPT. *Computers in Human Behavior: Artificial Humans* 4 (May 2025), 100144. doi:10.1016/j.chbah.2025.100144