

Understanding AI Developers' Decision Making in Design: The Role of Working Conditions and Uncertainty Mindset

Arlette Danielle Roman
Almanzar

Work and Organizational Psychology
ETH Zürich
Zürich, Switzerland
aroman@ethz.ch

David Joachim Grüning

Stanford University
Stanford, California, USA
Max Planck Institute for Human
Development
Berlin, Germany
gruening@stanford.edu

Gudela Grote

Work and Organizational Psychology
ETH Zürich
Zürich, Switzerland
ggrote@ethz.ch

Abstract

Developers of high-risk AI systems often face difficult tradeoffs when striving to embed ethical considerations into systems in competitive environments where ethics is often treated as an afterthought. Moreover, despite the pervasive uncertainties inherent in the development and deployment of complex technologies, limited research examines how developers appraise these uncertainties and how they affect their decision-making. In an exploratory survey, we study organizational and individual factors influencing developers' social responsibility in AI (SRAI) with 97 U.S.-based computer scientists. We assessed their uncertainty mindset, job autonomy, SRAI, and ethical tradeoffs decision, such as prioritizing fairness versus accuracy. Findings indicate that job autonomy is a strong predictor of SRAI and, surprisingly, a preference for accuracy over fairness. This may be due to several factors, including the fact that accuracy is among the most valued metrics. Additionally, an uncertainty-enabling mindset helps developers address both social and technical uncertainties.

CCS Concepts

• **Social and professional topics** → **User characteristics**; Computing/technology policy; • **Human-centered computing** → *Human-computer interaction (HCI)*; *Empirical studies in HCI*.

Keywords

Technology developers, Uncertainty, Job Design, Job Autonomy

ACM Reference Format:

Arlette Danielle Roman Almanzar, David Joachim Grüning, and Gudela Grote. 2026. Understanding AI Developers' Decision Making in Design: The Role of Working Conditions and Uncertainty Mindset. In *Extended Abstracts of the 2026 CHI Conference on Human Factors in Computing Systems (CHI EA '26)*, April 13–17, 2026, Barcelona, Spain. ACM, New York, NY, USA, 5 pages. <https://doi.org/10.1145/3772363.3799292>

1 Introduction

"Move fast and break things," but who is responsible for what is broken? [11]. The ongoing race in artificial intelligence (AI) drives this

modus operandi in big tech, prioritizing performance and speed, which can lead to engineering failures. These failures include discrimination in hiring systems, credit scoring, and facial recognition, among other use cases [3, 6, 28]. Thus, AI developers often face complex tradeoffs when developing high-risk AI, e.g., prioritizing accuracy over fairness or balancing model performance and explainability [8, 16, 22, 27]. Furthermore, they face growing pressure to meet ethical standards, as demanded in over 170 ethical guidelines for AI development [2]. Unfortunately, the highly competitive environment and challenging working conditions among developers may lead to the view that ethical considerations in technological innovation should be treated as an afterthought [5]. In this study, we explore organizational and individual factors that influence AI developers' sense of social responsibility and their design choices.

Emerging research has introduced new elements into the discussion of social responsibility and bias in AI, suggesting that developers' individual differences also contribute to shaping biased outcomes [7, 15]. These perspectives challenge an excessive focus on technical fixes and biased data as the primary ways to address engineering failures in sociotechnical systems. New evidence suggests that programmers' cultural biases [12] and their preference for hierarchy between groups—social dominance orientation—egalitarianism (SDO-E)—shape their professional design choices in AI and their sense of social responsibility toward their creations [24]. This opens new avenues for exploring how to mitigate biases in AI development. Similarly, Kahlert and Grote [13] found that interdisciplinary professional identities promote a more holistic, impact-aware mindset among technology developers, highlighting the relevance of cross-disciplinary collaboration and supporting the case for stakeholder engagement.

Job autonomy, as a working condition, is generally associated with positive outcomes and has been shown to increase job satisfaction and creativity, among other benefits. For instance, investment bankers who experienced higher job autonomy showed increased work engagement and wellbeing and received superior performance ratings [4]. Nevertheless, job autonomy has also been linked to unethical behavior, as it can foster a sense of being unconstrained by rules [1, 17, 18]. Lu and colleagues [18] suggest that this relationship is moderated by the extent to which the individual values autonomy. These concepts, along with social responsibility in AI, are also relevant for enhancing developer wellbeing, particularly given that wellbeing has been associated with an ethical culture at work [25] and corporate social responsibility [19]. As part of their



This work is licensed under a Creative Commons Attribution 4.0 International License. *CHI EA '26, Barcelona, Spain*

© 2026 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-2281-3/26/04

<https://doi.org/10.1145/3772363.3799292>

working conditions, developers often face uncertainty from multiple sources, including technical challenges during development, potential adverse outcomes of their creations, evolving regulatory standards, and unintended end-user applications. Yet, we do not understand the impact of an uncertainty mindset on technology developers' social responsibility in AI.

Even considering these studies, the role of individual and organizational factors in social responsibility in AI development remains under-researched [15, 24]. Consequently, it is unclear which working conditions can facilitate greater engagement with ethics in AI design and social responsibility. While job autonomy has been described as a "double-edged sword" and "too much of a good thing" [17, 18], we still do not understand how this translates to the AI development context in terms of social responsibility and ethical tradeoffs. Moreover, despite the prevalence of uncertainty, there is little evidence linking this concept to AI developers' design choices. Thus, understanding how developers view uncertainty—whether as enabling or disabling [20]—and how this perception influences their social responsibility and wellbeing may help managers and organizations foster an ethical culture around AI practices. In this context, we challenge the negative connotation of uncertainty and posit that an uncertainty-as-enabling mindset fosters openness to learning and continuous improvement [9, 10]. This shift in mindset may encourage developers to address not only technical uncertainties in AI development (e.g., reliability and robustness) but also social uncertainties inherent in sociotechnical systems (e.g., unintended societal outcomes).

A survey of technology developers found that nearly half of programmers among 65,553 respondents reported that "the humans creating AI should be responsible for considering the ramifications of the technology. Not the bosses. Not the middle managers. The coders" [26]. We explore what leads others to believe they are not responsible and what strategies can persuade developers to take on greater responsibility. Accordingly, we ask: Which factors are associated with developers' social responsibility in AI? Which working conditions can organizations provide to facilitate social responsibility, and what impact do these conditions have on developers' design choices in AI? Finally, how do organizational and individual factors impact ethical tradeoffs in AI design (e.g., fairness versus accuracy)?

2 Methods

To answer these questions, we conducted an exploratory observational study to shed light on these concepts. We recruited 97 U.S.-based employed computer scientists via Prolific. Participants completed a questionnaire assessing key variables, including the social responsibility in AI (SRAI) scale (measuring a sense of duty toward various stakeholders and the societal impact of algorithms), as well as their uncertainty mindset and job autonomy. No participants were excluded.

2.1 Measures

All scales were measured using 7-point Likert scales, with one (1) representing "strongly disagree" and seven (7) representing "strongly agree," unless otherwise indicated.

Uncertainty Mindset (UM) [20]: Four (4) subdimensions were measured and aggregated: uncertainty as a threat-opportunity and uncertainty as fixed-malleable. Sample items included "uncertainty at work can inhibit my professional development" and "when there is uncertainty at work, I cannot make a difference with my actions." Some items were framed negatively and reverse-coded. The UM scale was validated across multiple studies, including exploratory and confirmatory factor analyses. Internal consistency within subdimensions was good to excellent, $.89 [.84; .91] \leq \alpha \leq .94 [.91; .95]$, $.90 [.86; .93] \leq \omega \leq .94 [.92; .96]$, and across all subdimensions it was good, $\alpha = .88 [.84; .91]$, $\omega = .84 [.79; .89]$.

Willingness to Engage in Social Uncertainties: "How comfortable are you with addressing social uncertainties during system development? i.e., uncertainties related to the social impact of the model (such as its potential effects on users, communities, or society as a whole)." This was measured on a 7-point Likert scale, one (1) representing "Extremely uncomfortable" and seven (7) representing "Extremely comfortable."

Willingness to Engage in Technical Uncertainties: "How comfortable are you with addressing technical uncertainties during system development, i.e., uncertainties related to the technical aspects of the model (such as performance, reliability, or accuracy)? This was measured on a 7-point Likert scale, one (1) representing "Extremely uncomfortable" and seven (7) representing "Extremely comfortable."

Social Responsibility in AI (SRAI): Replicated from Román Almánzar and colleagues [24]. Seven (7)-item scale; sample items include "I have a responsibility to the end-users of AI system designs" and "I feel a personal responsibility to reduce unfair bias in AI." Reliability is excellent, $\alpha = .90 [.87; .93]$, $\omega = .90 [.87; .93]$.

Job Autonomy: Four (4) items scale [14]. Items included "In general, how much influence do you have over the following: the tasks you do at your work; the pace at which you work; how you do your work; the order in which you carry out tasks." This was measured on a 7-point Likert scale, one (1) representing "No influence at all" and seven (7) representing "Complete influence." Reliability was good, $\alpha = .84 [.78; .89]$, $\omega = .84 [.79; .90]$.

Fairness versus Accuracy Scenario or Accuracy Preference: [23] "Imagine you're building a recommendation algorithm for STEM courses. Would you include gender as a feature to enhance the accuracy of the algorithm, even if it leads to a decreased presence of women in STEM classes? In this case, I would include gender as a feature to prioritize accuracy."

3 Findings

To understand the power of predictors with respect to AI responsibility and accuracy preferences, we conducted hierarchical regression analyses and simple regressions. All significance tests were conducted using a pre-specified significance level of $\alpha = 0.05$. The assumptions for regression analyses, namely, linearity, relative homoscedasticity, low multicollinearity, and no extreme outliers, were supported. The important results are shown in Table 1 and 2.

Job Autonomy Predicting Social Responsibility in AI: Results from a simple regression indicated that job autonomy was a strong and significant predictor of SRAI ($b = 0.49$, $SE = 0.09$, $\beta = .46$, $t(95) = 5.03$, $p < .001$, 95% CI $[0.30; 0.68]$), explaining 21% of the variance, $R^2 = .21$.

Table 1: Hierarchical regression analyses predicting social responsibility in AI and accuracy preferences.

Predictor	Social Responsibility in AI			Accuracy Preferences		
	<i>b</i>	<i>SE</i>	β	<i>b</i>	<i>SE</i>	β
<i>Model 1</i>						
Uncertainty Mindset	0.27	0.12	0.24*	0.44	0.22	0.20*
R^2		0.06			0.04	
ΔR^2		0.06			0.04	
ΔF		F(1, 95) = 5.62*			F(1, 95) = 4.02*	
<i>Model 2</i>						
Uncertainty Mindset	0.06	0.12	0.05	0.29	0.24	0.14
Job Autonomy	0.47	0.11	0.44***	0.30	0.22	0.15
R^2		0.21			0.06	
ΔR^2		0.15			0.02	
ΔF		F(1, 94) = 18.66***			F(1, 94) = 1.90	

Table 2: Summary of Simple Regressions

Predictor	Outcome	<i>b</i>	<i>SE</i>	β	<i>t</i>	<i>p</i> -value	R^2	95% CI
Uncertainty Mindset	Comfort with Technical Uncertainties	0.37	0.13	0.27	2.70	0.008	0.07	[0.09; 0.63]
Uncertainty Mindset	Comfort with Social Uncertainties	0.35	0.16	0.22	2.16	0.033	0.05	[0.03; 0.67]
Job Autonomy	Social Responsibility in AI	0.49	0.09	0.46	5.03	<.001	0.21	[0.30; 0.68]
Job Autonomy	Accuracy Preferences	0.42	0.20	0.21	2.10	0.038	0.05	[0.02; 0.82]

Job Autonomy Predicting Preference for Accuracy Over Fairness: A separate regression model revealed that job autonomy significantly predicted programmers' preference for accuracy over fairness, i.e., greater agreement with including gender to improve accuracy ($b = 0.42$, $SE = 0.20$, $\beta = .21$, $t(95) = 2.10$, $p = .038$, 95% CI [0.02, 0.82]), accounting for 5% of the variance, $R^2 = .05$. Interestingly, while job autonomy was positively associated with social responsibility, it also predicted a stronger preference for accuracy over fairness rather than the expected preference for fairness. This is surprising because the exploration of job autonomy and its effects in the context of AI design and ethical trade-offs is novel, and it is generally associated with positive outcomes. One possible explanation is that job autonomy fosters a work culture that prioritizes performance, which in AI development may be evaluated in terms of algorithmic accuracy. Another reason may be that developers may associate the most accurate algorithm with the greatest societal benefit, or that highly technical developers may default to optimizing what they are most familiar with. This aligns with recent research that challenges the notion that job autonomy yields solely positive outcomes and suggests that, while beneficial, it can be a "double-edged sword" [18], potentially reinforcing performance-driven behaviors at the expense of ethical considerations [1, 14, 21]. These findings highlight the need for further investigation into the complex interplay between job autonomy, social responsibility, and ethical tradeoffs in AI decision-making.

Uncertainty Mindset Predicting Comfort with Technical and Social Uncertainty: The uncertainty-as-enabling mindset significantly predicted AI developers' comfort with addressing technical uncertainty ($b = 0.37$, $SE = 0.13$, $\beta = .27$, $t(95) = 2.70$, $p = .008$, 95% CI [0.09,

0.63]), accounting for 7% of the variance ($R^2 = .07$). Similarly, it significantly predicted comfort with addressing social uncertainty, ($b = 0.35$, $SE = 0.16$, $\beta = .22$, $t(95) = 2.16$, $p = .033$, 95% CI [0.03, 0.67]), explaining 5% of the variance, $R^2 = .05$. In contrast, job autonomy was not a significant predictor of comfort with either form of uncertainties, ($t(95) \leq 0.78$, $p \geq .440$).

Predicting Social Responsibility in AI (SRAI): A two-step hierarchical multiple regression was conducted. Uncertainty-as-enabling mindset was entered in the first step of the regression, and it significantly predicted social responsibility in AI (SRAI), ($b = 0.27$, $SE = 0.12$, $\beta = .24$, $p = .020$), explaining 6% of the variance ($R^2 = .06$). However, this changed when adding job autonomy in the second step, which emerged as the stronger predictor ($b = 0.47$, $SE = 0.11$, $\beta = .44$, $p < .001$), explaining an additional 15% of the variance, ($\Delta R^2 = .15$, $F(1, 94) = 18.66$, $p < .001$) while the uncertainty mindset became non-significant, ($b = 0.06$, $SE = 0.12$, $\beta = .05$, $p = .642$). The combined model explained 21% of the variance, $R^2 = 0.21$, suggesting that job autonomy, as an organizational factor, accounts for most of the unique variance in SRAI once both predictors are considered.

Predicting Accuracy Preferences: A two-step hierarchical regression was conducted for accuracy preferences. In Model 1, the uncertainty mindset significantly predicted accuracy preferences ($b = 0.44$, $SE = 0.22$, $\beta = .20$, $p = .048$), explaining 4% of the variance, $R^2 = .04$. This finding is somewhat unexpected, as an uncertainty-as-enabling mindset was also positively associated with social responsibility in AI (SRAI). While one might assume that individuals who embrace uncertainty would also prioritize fairness in AI decision-making, this result suggests that such individuals may instead view accuracy as a more reliable and socially beneficial goal. Moreover,

a high sense of social responsibility does not necessarily imply prioritizing accuracy over fairness in algorithms; other factors, such as profit- and performance-oriented organizational goals, may be influential. More research is needed to understand how professionals balance these ethical tradeoffs in AI development. However, in Model 2, the addition of job autonomy did not significantly improve model fit ($\Delta R^2 = .02$, $F(1, 94) = 1.90$, $p = .172$), and neither predictor reached significance in the final model (uncertainty mindset: $b = 0.29$, $SE = 0.24$, $\beta = .14$, $p = .225$; job autonomy: $b = 0.30$, $SE = 0.22$, $\beta = .15$, $p = .172$).

Correlational insights: Job autonomy was strongly correlated with an uncertainty-as-enabling mindset ($r = .43$, $p < .001$), indicating that individuals with higher levels of job autonomy are more likely to perceive uncertainty as an opportunity rather than a threat. This suggests that greater flexibility and decision-making power in the workplace can foster adaptive attitudes toward uncertainty.

4 Future Research

Future work should explore other organizational and individual factors and study multiple hypothetical scenarios of the fairness vs. accuracy tradeoff, including multiple protected attributes. Moreover, future studies should design experiments to establish causality by manipulating factors such as job autonomy and uncertainty mindset.

5 Theoretical Contributions

We seek to contribute to several theoretical streams. First, our findings expand the field of algorithmic fairness by exploring the tension between accuracy and fairness in design through the lens of individual and organizational factors. Second, we add to the recent literature on job autonomy as a "double-edged sword" in the context of a highly relevant organizational setting, the development of AI. Lastly, we reframe the narrative about the negative connotation of uncertainty, aligning with research suggesting that coping with and viewing uncertainty as enabling increases information-seeking and fosters comfort in addressing the social uncertainties of technology development.

6 Practical Contributions

We identify practically meaningful disagreements in beliefs about algorithmic fairness, suggesting that beliefs about which features are acceptable vary not only with individual factors but also with organizational factors such as job autonomy. It is particularly recommended that organizations operating in high-uncertainty environments foster an uncertainty-as-enabling mindset among their employees, as this view may encourage them to remain open to learning and engage with the social uncertainties of their creations. This understanding is helpful for training purposes, in order to persuade developers to adopt an uncertainty-as-enabling mindset to build more socially responsible technology. In addition, job autonomy should be accompanied by ethical guidelines and explicit discussions of the definition of fairness in a resource allocation algorithm in the early stages of development. This approach allows the identification of both intended and potential negative outcomes and facilitates a transparent evaluation of the appropriate fairness metrics.

References

- [1] Ambreen Ahmed, Aamir Firoz Shamsi, and Mudassar Aziz. 2020. A Missing Link Between Job Autonomy and Unethical Behavior. *ETIKONOMI* 19, 1 (2020), 95–118. doi:10.15408/etk.v19i1.12391
- [2] AlgorithmWatch. 2019. *AI Ethics Guidelines Global Inventory*. Technical Report. AlgorithmWatch. <https://inventory.algorithmwatch.org>
- [3] Julia Angwin, Jeff Larson, Surya Mattu, and Lauren Kirchner. 2016. Machine Bias: There's Software Used Across the Country to Predict Future Criminals and It's Biased Against Blacks. ProPublica. <https://www.propublica.org/article/machine-bias-risk-assessments-in-criminal-sentencing>. Accessed: 2026-01-22.
- [4] Paul P. Baard, Edward L. Deci, and Richard M. Ryan. 2004. Intrinsic Need Satisfaction: A Motivational Basis of Performance and Wellbeing in Two Work Settings. *Journal of Applied Social Psychology* 34, 10 (2004), 2045–2068. doi:10.1111/j.1559-1816.2004.tb02690.x
- [5] Sebastian Bartsch, Verena Milani, Martin Adam, and Alexander Benlian. 2024. Algorithmic Accountability: What Does It Mean for AI Developers and How Does It Affect AI Development Projects. In *Proceedings of the 57th Hawaii International Conference on System Sciences (HICSS-57)*. Maui, Hawaii, USA. https://aisel.aisnet.org/hicss-57/os/ai_and_organizing/6_Paper_6.
- [6] Joy Buolamwini and Timnit Gebru. 2018. Gender Shades: Intersectional Accuracy Disparities in Commercial Gender Classification. In *Proceedings of Machine Learning Research*, Vol. 81. 77–91. <https://proceedings.mlr.press/v81/buolamwini18a.html>. Accessed: 2026-01-22.
- [7] Stephen Cave and Kanta Dihal. 2020. The Whiteness of AI. *Philosophy & Technology* 33, 4 (2020), 685–703. doi:10.1007/s13347-020-00415-6
- [8] Sam Corbett-Davies, Emma Pierson, Avi Feller, Sharad Goel, and Aziz Huq. 2017. Algorithmic Decision Making and the Cost of Fairness. In *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. 797–806. doi:10.1145/3097983.3098095
- [9] Mark A. Griffin and Gudela Grote. 2020. When Is More Uncertainty Better? A Model of Uncertainty Regulation and Effectiveness. *Academy of Management Review* 45, 4 (2020). doi:10.5465/AMR.2018.0271
- [10] Gudela Grote. 2015. Promoting Safety by Increasing Uncertainty: Implications for Risk Management. *Safety Science* 71 (2015), 71–79. doi:10.1016/j.ssci.2014.02.010
- [11] Gudela Grote, Sharon K. Parker, and Kevin Crowston. 2024. Taming Artificial Intelligence: A Theory of Control–Accountability Alignment among AI Developers and Users. *Academy of Management Review* (2024). doi:10.5465/amr.2023.0117
- [12] Johanna Johansen, Tore Pedersen, and Christian Johansen. 2023. Studying human-to-computer bias transference. *AI & Society* 38 (2023), 1659–1683. doi:10.1007/s00146-021-01328-4
- [13] Aniko Kahlert and Gudela Grote. 2024. Why Should I Care?: Understanding Technology Developers' Design Mindsets in Relation to Prospective Work Design. *European Journal of Work and Organizational Psychology* 33, 2 (2024), 230–244. doi:10.1080/1359432X.2024.2314936
- [14] Bettina Kubicek, Matej Paškván, and Julian Bunner. 2017. The Bright and Dark Sides of Job Autonomy. In *Job Demands in a Changing World of Work: Impact on Workers' Health and Performance and Implications for Research and Practice*, Christian Korunka and Bettina Kubicek (Eds.). Springer, Cham, 45–63. doi:10.1007/978-3-319-54678-0_4
- [15] David Lehr and Paul Ohm. 2017. Playing with the Data: What Legal Scholars Should Learn About Machine Learning. *UC Davis Law Review* 51 (2017), 653–717.
- [16] Annie Liang, Jay Lu, and Xiaosheng Mu. 2022. Algorithmic design: Fairness versus accuracy. In *Proceedings of the 23rd ACM Conference on Economics and Computation*. 58–59. doi:10.1145/3490486.3538237
- [17] Xiaoqin Liu, Yevhen Baranchenko, Fansuo An, Zhibin Lin, and Jie Ma. 2021. The Impact of Ethical Leadership on Employee Creative Deviance: The Mediating Role of Job Autonomy. *Leadership & Organization Development Journal* 42, 2 (2021), 219–232. doi:10.1108/LODJ-01-2020-0026
- [18] Jun Gu Lu, Joel Brockner, Yoav Vardi, and Erez Weitz. 2017. The Dark Side of Experiencing Job Autonomy: Unethical Behavior. *Journal of Experimental Social Psychology* 73 (2017), 222–234. doi:10.1016/j.jesp.2017.05.007
- [19] Gloria Macassa, Cormac McGrath, Gianpaolo Tomaselli, and Sandra C. Buttigieg. 2021. Corporate social responsibility and internal stakeholders' health and well-being in Europe: a systematic descriptive review. *Health Promotion International* 36, 3 (2021), 866–883. doi:10.1093/heapro/daaa071
- [20] Federico Magni, Alina Gerlach, Julian Pfrombeck, Sara Zaniboni, Laura Elaine Strittmatter, and Gudela Grote. 2023. Mind the Unknown: Development and Validation of the Uncertainty Mindset. In *Academy of Management Proceedings*, Vol. 2023. Academy of Management, 13059.
- [21] Rhokeun Park and Deborah Searcy. 2012. Job Autonomy as a Predictor of Mental Wellbeing: The Moderating Role of Quality-Competitive Environment. *Journal of Business and Psychology* (2012). doi:10.1007/s10869-011-9244-3
- [22] Dana Pessach and Erez Shmueli. 2020. Algorithmic Fairness. arXiv:2001.09784 [cs.CY] doi:10.48550/arXiv.2001.09784 arXiv:2001.09784.
- [23] Emma Pierson. 2018. Demographics and Discussion Influence Views on Algorithmic Fairness. doi:10.48550/arXiv.1712.09124

- [24] Arlette Danielle Román Almánzar, David Joachim Grüning, and Laura Marie Edinger-Schons. 2025. It wasn't me: the impact of social responsibility and social dominance attitudes on AI programmers' moral imagination (intention to correct bias). *AI and Ethics* 5, 6 (2025), 5655–5669. doi:10.1007/s43681-024-00516-4
- [25] Achinto Roy, Alexander Newman, Heather Round, and Sukanto Bhattacharya. 2024. Ethical Culture in Organizations: A Review and Agenda for Future Research. *Business Ethics Quarterly* 34, 1 (2024), 97–138. doi:10.1017/beq.2022.44
- [26] Stack Overflow. 2018. Stack Overflow Developer Survey 2018. Online. <https://survey.stackoverflow.co/2018> Accessed: 2026-01-22.
- [27] Sandra Wachter, Brent Mittelstadt, and Chris Russell. 2021. Bias Preservation in Machine Learning: The Legality of Fairness Metrics Under EU Non-Discrimination Law. *West Virginia Law Review* 123 (2021), 735–768. <https://researchrepository.wvu.edu/wvhlr/vol123/iss3/4>
- [28] Meredith Whittaker, Meryl Alper, Cynthia L. Bennett, Sara Hendren, Elizabeth Kaziunas, Mara Mills, Meredith Ringel Morris, Joy Lisi Rankin, Emily Rogers, Marcel Salas, and Sarah Myers West. 2019. Disability, Bias & AI Report. AI Now Institute. <https://ainowinstitute.org/publications/disabilitybiasai-2019> AI Now Institute, November 20, 2019.