

Herding CATs: Making Sense of Creative Activity Traces

Max Kreminski
Storytelling Lab
Midjourney
San Francisco, California, USA
mkreminski@midjourney.com

Kihoon Son
School of Computing
KAIST
Daejeon, Republic of Korea
kihoon.son@kaist.ac.kr

Noor Hammad
Human-Computer Interaction
Institute
Carnegie Mellon University
Pittsburgh, Pennsylvania, USA
nhammad@cs.cmu.edu

Amy Smith
Queen Mary University of London
London, United Kingdom
amyelizabethsmith01@gmail.com

Qian Yang
Computing and Information Science
Cornell University
Ithaca, New York, USA
qy242@cornell.edu

Eric Rawn
Electrical Engineering and Computer
Sciences
University of California, Berkeley
Berkeley, California, USA
erawn@berkeley.edu

John Joon Young Chung
Storytelling Lab
Midjourney
San Francisco, California, USA
jchung@midjourney.com

Sang Won Lee
Virginia Tech
Blacksburg, Virginia, USA
sangwonlee@vt.edu

Shm Garanganao Almeda
Electrical Engineering and Computer
Sciences
University of California, Berkeley
Berkeley, California, USA
shm.almeda@berkeley.edu

J.D. Zamfirescu-Pereira
Computer Science
University of California, Los Angeles
Los Angeles, California, USA
zamfi@cs.ucla.edu




Figure 1: Prior analyses of creative activity traces, including a creative sensemaking curve (A) [8]; a visualization of trajectories through a design space (B) [1]; a fuzzy linkograph (C) [30]; and an analysis of version control commit histories (D) [31].

Abstract

This workshop aims to advance the analysis of *creative activity traces*, particularly those captured through user interaction with software creativity support tools (CSTs). Traces of creative activity

constitute a rich resource for identifying the impacts of CSTs—especially AI-based CSTs—on the creative process, and may also inform general-purpose process theories of creativity. Several new approaches to making sense of these traces have been introduced in the past few years, but many of these approaches have emerged from largely disjoint research communities, hindering the development of a shared analytical toolkit. We propose to gather HCI and creativity researchers, including proponents of several different trace analysis techniques, to sketch out a *technique design space* to guide future empirically grounded research on creativity support.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

CHI EA '26, Barcelona, Spain

© 2026 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-2281-3/26/04

<https://doi.org/10.1145/3772363.3778743>

CCS Concepts

• **Human-centered computing** → **HCI design and evaluation methods**; • **Applied computing** → **Arts and humanities**.

Keywords

creativity support tools, human-AI interaction

ACM Reference Format:

Max Kreminski, Amy Smith, John Joon Young Chung, Kihoon Son, Qian Yang, Sang Won Lee, Noor Hammad, Eric Rawn, Shm Garanganao Almeda, and J.D. Zamfirescu-Pereira. 2026. Herding CATs: Making Sense of Creative Activity Traces. In *Extended Abstracts of the 2026 CHI Conference on Human Factors in Computing Systems (CHI EA '26)*, April 13–17, 2026, Barcelona, Spain. ACM, New York, NY, USA, 5 pages. <https://doi.org/10.1145/3772363.3778743>

1 Motivation

Software creativity support tools (CSTs) [7, 11, 28] play an increasingly central role in a wide variety of creative domains, and evidence is mounting that these tools can exert substantial *influence* on the creative process [30]—whether by generating material that guides human ideation [2, 22, 26], establishing a *normative ground* that shapes creative activity [4, 18], or selectively adding and removing friction around certain kinds of *expressive communication* [19]. Despite the growing influence of CSTs, however, techniques for understanding how CST design shapes creative activity—and, indeed, for evaluating CSTs at all [25]—remain relatively limited.

Because software CSTs are software, they permit the automatic capture of *activity traces* reflecting the creative activity of large numbers of CST users at scale—including usage logs [29, 33], sequences of artifact [1, 15, 30] or interactor [8, 9] state snapshots, and commit histories [14, 23, 24, 31]. These traces may then be analyzed—either alone or alongside observational and self-report data from traditional user studies—to develop a sophisticated understanding of tool influences on the creative process. Similar approaches have already seen successful application to the understanding of co-creative *interaction dynamics* [8, 9, 30]; the characterization of designer *trajectories* within large spaces of possible artifacts [1, 14, 15]; and the assessment of *idea development* [10, 17, 32]. These results are exciting, not least because they suggest a compelling way forward against the long-established difficulty [13] of evaluating CSTs.

Best practices for analyzing creative activity traces, however, are still emerging. Trace analyses published in venues as diverse as AAAI (general AI), CHI (general HCI), IUI (human-AI interaction), C&C (creative HCI), EMNLP (NLP), FDG (game AI), and IEEE ToG (game AI) have revealed numerous intriguing parallels between activity traces across diverse creative domains—including the detectability of “curious users” who actively probe CST boundaries [15, 21, 27, 30]; recognizable patterns of goal formation by CST users [10, 15, 30]; and the applicability of clustering to identify designer archetypes [1, 30]. Nevertheless, discrepancies between trace analysis approaches also abound: some analyses focus on interactors themselves, while others focus on created artifacts; the nature and granularity of “design moves” that make up a trace may vary widely from one analysis to the next. Furthermore, although trace analysis is typically conducted retrospectively, there is now growing interest in “feeding back” the results of computational

trace analysis into the creative process while interaction is still unfolding, for instance to facilitate reflection [5, 30] or improve the quality of proactive machine interventions [3, 30].

We therefore believe the time is ripe for a workshop intended to advance the science and art of interpreting creative activity traces. We propose to gather HCI and creativity researchers to develop a shared set of *research questions* that can be addressed via creative activity trace analysis and the beginnings of a *technique design space* for trace analysis techniques, trace-aware CST design, and related topics. Researchers who have worked on human-AI interaction or specific “creativity-adjacent” topics (such as computer-supported cooperative work or software engineering) may also play a valuable role in extending the conversation to new research questions—e.g., *How might activity trace analysis inform general-purpose process models of creativity?* [12]—and new creative domains.

2 Length of the Workshop

Two 90-minute sessions (long).

3 Organizers

Collectively, the organizing team has worked with creative activity traces in a wide variety of different forms (including linkographs, CST interaction logs, live improvisations, version control histories, and personalization data) across a wide variety of different creative domains (including visual art, writing, worldbuilding, interactive narrative, programming, music, and video creation). Roughly half of the organizers are established research group leaders in creativity support; the other half are up-and-coming PhD students with a specific focus on creative activity traces as a design material, subject of analysis, or resource for making sense of the creative process. Organizers represent universities and industry research institutions from three different continents and plan to distribute the workshop CfP widely across their diverse collaborative networks, reaching many core HCI researchers but also researchers in relevant subfields of AI (e.g., computational creativity; game AI) and scholar-practitioners in a number of distinct creative fields.

3.1 Organizer Biographies

Max Kreminski is the director of the Storytelling Lab at Midjourney and an incoming assistant professor of Design Tech at Cornell Tech. As an HCI and AI researcher with a particular focus on creativity support tools, Max has worked to characterize and visualize patterns of turntaking and influence in human-AI co-creativity; to introduce new methods for evaluating CSTs; and to improve understanding of how tool design shapes the temporal unfolding of co-creative interaction. Max’s research has been featured in outlets such as *The New Yorker*, *New Scientist*, and *The Verge*; published and exhibited at top HCI and AI conferences, including CHI, UIST, NeurIPS, and EMNLP; and honored with a variety of awards, including the Best Paper award at the ACM Conference on Creativity & Cognition.

Amy Smith is a PhD candidate in Computer Science at Queen Mary University of London. Her research combines artificial intelligence, HCI and computational creativity to investigate how generative technologies transform creative practice, authorship, and cultural production. A former professional tattoo artist, she

develops artist-centered tools that fine-tune text-to-image models and studies collaborative sense-making with AI. Her work has been presented at AAAI, CHI, and the International Conference on Computational Social Science.

John Joon Young Chung is a research scientist at Midjourney. He completed his Computer Science and Engineering Ph.D. at the University of Michigan, Ann Arbor, and his Electrical Engineering BS at Seoul National University. He interned at Microsoft Research, Naver AI Lab, and Adobe Research during his Ph.D. years, and was a full-time researcher at Kixlab KAIST before he started his Ph.D. John's main research interests are AI-powered creativity support tools for art-making and storytelling, while his interests also span educational technologies, crowdsourcing, and natural language processing. He is a recipient of paper awards from ACM CHI and CSCW.

Kihoon Son is a fourth year Ph.D. student in the School of Computing at KAIST and a member of KIXLAB (the KAIST Interaction Lab). His Ph.D. research explores how designers' tacit knowledge can be captured, surfaced, and reused within creative workflows to enhance AI-supported design practice. To this end, Kihoon combines qualitative inquiry, the development of generative AI-based interactive systems, and the design of AI agents to explore representations of tacit knowledge and their reuse.

Qian Yang is an assistant professor of information science at Cornell University. She is a human-AI interaction designer and researcher. With her students, Yang designs novel AI applications and services that demonstrate new ways of bridging together frontier AI technologies, user needs, and societal betterment goals. For this work, she has received many academic and industry awards, including Schmidt Futures' Inaugural AI2050 Fellowship. She co-directs the Cornell Digital and AI Literacy Initiative and is a senior fellow at the Cornell Brooks School Tech Policy Institute.

Sang Won Lee is an Associate Professor of Computer Science at Virginia Tech and a Visiting Researcher at NAVER AI Lab. His research lies at the intersection of HCI, computer-mediated empathy, and creative computing. He investigates how technologies, such as conversational agents, VR, and personal informatics, can foster empathy, support self-reflection, and promote inclusive collaboration. His work often blends empirical studies with the design and evaluation of interactive systems, and has been recognized with multiple best paper awards at top venues like ACM CHI and Creativity & Cognition. In addition to academic research publications, he has presented his research in the form of musical performances at peer-reviewed venues, including NIME, Art-CHI, and ICMC. He is a winner of the International Computer Music Association–Music Award 2016 for his composition Live Writing: Gloomy Streets.

Noor Hammad is a PhD student at Carnegie Mellon University's Human-Computer Interaction Institute. Her research focuses on designing and developing tools that empower digital creators. Noor is particularly interested in how systems can leverage creator data to support more authentic, personalized, and expressive creative workflows, and has explored this interest through projects in live streaming, game development, and augmented reality. She has also interned at Adobe, where she developed AI creative assistants for video content creators. Noor is a recipient of the NSERC Postgraduate Fellowship and the Generation Google Scholarship.

Eric Rawn is a PhD student at UC Berkeley. They study programming practice: how people use programs to learn about the world and think through their work. Using long-term qualitative methods, like participant observation, codesign, and in-situ deployments, Eric develops programming language interventions and programming systems for scientists, data analysts, and artists. They have also worked specifically on the role of version control systems in creative practice.

Shm Garanganao Almeda is a Ph.D. student and artist in the Berkeley Institute of Design Lab at UC Berkeley. They study ecosystems of artistic support, with a focus on the way technologies change *how* a community supports artists. Their research includes empirical study of creative communities (NFT artists, voice actors, AI and data artists) and the development of support tools as probes to better elicit and study creatives' interactions with new mediums — towards an understanding of creative activity, and new media technology, led by artistic practices. Shm's work has been published at ACM DIS, CHI, and FAccT, and recognized with a Best Paper Honorable Mention at CHI 2025.

J.D. Zamfirescu-Pereira is a PhD candidate in EECS at UC Berkeley and incoming assistant professor of computer science at UCLA, researching how to improve the ways designers — visual artists, chatbot designers, programmers — work with, and through AI and LLMs. J.D.'s work has been published in CHI, UIST, FAccT, and HRI, and recognized with Best Paper Honorable Mention awards. He holds bachelor's and master's degrees in Computer Science from MIT, and has previous experience working at Google, founding a YC-backed startup, and advising new startups on tech direction and strategy.

4 Plans to Publish Workshop Proceedings

We will encourage authors to submit accepted position papers to arXiv, using a workshop-wide report number scheme to identify these arXiv uploads as part of the workshop proceedings. However, we cannot guarantee arXiv will accept all position papers for upload. We will also upload accepted position papers to the workshop's website and link to these papers from the workshop's online schedule, similarly to the GenAICHI workshop at CHI 2025 [20].

5 Accessibility

We do not anticipate any special accessibility needs for this workshop: the format is typical for a CHI workshop. Before the workshop, we will invite attendees to share accessibility accommodation requests, and prepare accordingly. During presentation-style workshop activities (e.g., organizers addressing the room as a whole; presentations of individual participants' submissions) we will attempt to provide captioning via live captioning software.

6 Offline Materials

Accepted position papers from workshop participants will be distributed after the session via the workshop website: see Section 4 above. The organizers will also upload any slides they present to the workshop attendees, alongside a retroactive writeup of the workshop's outcomes and next steps. Publication of these materials will be announced on a Discord server created to host workshop and post-workshop discussions.

7 Workshop Activities

Opening (15 min). To get participants on the same page, workshop organizers will briefly present the planned schedule and provide an overview of the workshop’s background and goals. Exemplars of past approaches to creative activity trace analysis will be highlighted here. We will also provide participants with access to a Discord server that can be used for Q&A, discussion group note-taking, sharing of discussion group outputs, and post-workshop coordination on next steps.

Lightning Presentations (30 min). Next, participants will rapidly present their accepted submissions, introducing their research interests and past relevant work.

Group Work: Creative Domains (45 min). The first of our group work periods will see participants split into several smaller groups to ideate on the forms that activity traces might take in different creative domains; the potential applicability of activity trace analysis to understanding creativity in these domains; and the limits of current knowledge and practice around the domains in question. In the final 5 minutes, groups will present their insights. We will then take a break before the second session of workshop activities.

Group Work: Research Questions (45 min). After the break, we will reconvene for a second group work period on *open research questions* that might be especially amenable to study via creative activity traces. This discussion will likely build on the set of relevant creative domains introduced in the first group work period, but participants will organize themselves in different groups to promote cross-pollination of ideas. As before, each breakout group will briefly present its insights in the last 5 minutes.

Group Discussion: Technique Design Space (30 min). A more synthetic period of facilitated group work will then see participants attempting to bring together insights from earlier discussions into a shared *technique design space* for creative activity trace analysis techniques. We will identify relevant dimensions of variation between different prior approaches to trace analysis; reexamine the variety of trace types that might emerge most naturally from different creative domains; consider how existing or novel CSTs might be instrumented to capture useful activity traces; and summarize our view of the space.

Conclusion and Next Steps (15 min). Workshop organizers will wrap things up and begin to coordinate next steps for further developing and distributing the group’s insights.

8 Post-Workshop Plans

We hope to continue using the workshop Discord server to begin coordinating collaboration on a full-length design space paper developing insights from the workshop into a more substantial form, similar to the design space paper [16] that emerged from the (similarly structured) second In2Writing workshop [6].

9 250-word Call for Participation

Creativity support tools (CSTs) now play a central role in a wide variety of creative domains — but techniques for evaluating them, and understanding how CST design shapes creative activity, are

relatively limited. This workshop aims to advance the analysis of Creativity Activity Traces (CATs), particularly those captured through user interaction with software CSTs. In the past few years, we have seen several approaches to making sense of CATs emerge, but from largely disjoint research communities. This workshop aims to gather HCI and creativity researchers and develop a shared analytical toolkit. Participation in the workshop is open to both experienced and emerging creativity researchers, including core HCI researchers, those in relevant AI subfields (e.g., computational creativity and game AI) and creative scholar-practitioners. Through the presentation of position papers and the collaborative drafting of a research agenda, this workshop will sketch out a *technique design space* for CAT analysis that can guide the future of empirically-grounded creativity research.

To attend, participants must submit a position paper of up to 4 pages in single-column ACM format (with additional pages available for references). Papers should describe a research question, direction, argument, or proposal related to creative activity traces. Participants will be encouraged to submit their accepted position papers to arXiv. Accepted submissions will also be made available from the online schedule on the workshop website. At least one author per accepted submission must attend the workshop in person; we expect 20-50 attendees total.

For more information, see the workshop website at <https://herding-cats-ws.github.io>.

10 Expected Size of Attendance

20-50 in-person participants (medium).

References

- [1] Alberto Alvarez, Jose Font, and Julian Togelius. 2022. Toward Designer Modeling Through Design Style Clustering. *IEEE Transactions on Games* 14, 4 (2022), 676–686.
- [2] Barrett R Anderson, Jash Hemant Shah, and Max Kreminski. 2024. Homogenization effects of large language models on human creative ideation. In *Proceedings of the 16th Conference on Creativity & Cognition*. 413–425.
- [3] Alexandra Bremers and Wendy Ju. 2024. Can Machines Tell What People Want? Bringing Situated Intelligence to Generative AI. In *Proceedings of the Halfway to the Future Symposium*.
- [4] Alex Calderwood, John Joon Young Chung, Yuqian Sun, Melissa Roemmele, and Max Kreminski. 2025. Phraselette: A Poet’s Procedural Palette. In *Proceedings of the 2025 ACM Designing Interactive Systems Conference*.
- [5] Dashiell Carrera and Sang Won Lee. 2022. Watch me write: exploring the effects of revealing creative writing process through writing replay. In *Proceedings of the 14th Conference on Creativity and Cognition*. 146–160.
- [6] Minsuk Chang, John Joon Young Chung, Katy Ilonka Gero, Ting-Hao Kenneth Huang, Dongyeop Kang, Mina Lee, Vipul Raheja, and Thiemo Wambsganss. 2023. The Second Workshop on Intelligent and Interactive Writing Assistants. In *Extended Abstracts of the 2023 CHI Conference on Human Factors in Computing Systems*.
- [7] John Joon Young Chung, Shiqing He, and Eytan Adar. 2021. The intersection of users, roles, interactions, and technologies in creativity support tools. In *Designing Interactive Systems Conference 2021*. 1817–1833.
- [8] Nicholas Davis, Chih-Pin Hsiao, Kunwar Yashraj Singh, Brenda Lin, and Brian Magerko. 2017. Creative sense-making: Quantifying interaction dynamics in co-creation. In *Proceedings of the 2017 ACM SIGCHI Conference on Creativity and Cognition*. 356–366.
- [9] Manoj Deshpande, Jisu Park, Supratim Pait, and Brian Magerko. 2024. Perceptions of Interaction Dynamics in Co-Creative AI: A Comparative Study of Interaction Modalities in Drawcto. In *Proceedings of the 16th Conference on Creativity & Cognition*. 102–116.
- [10] Shachar Don-Yehiya, Leshem Choshen, and Omri Abend. 2023. Human Learning by Model Feedback: The Dynamics of Iterative Prompting with Midjourney. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, Houma Bouamor, Juan Pino, and Kalika Bali (Eds.). Association for

- Computational Linguistics, Singapore, 4146–4161. doi:10.18653/v1/2023.emnlp-main.253
- [11] Jonas Frich, Lindsay MacDonald Vermeulen, Christian Remy, Michael Mose Biskjaer, and Peter Dalsgaard. 2019. Mapping the landscape of creativity support tools in HCI. In *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems*.
- [12] Adam E Green, Roger E Beaty, Yoed N Kenett, and James C Kaufman. 2024. The process definition of creativity. *Creativity Research Journal* 36, 3 (2024), 544–572.
- [13] Tom Hewett, Mary Czerwinski, Michael Terry, Jay Nunamaker, Linda Candy, Bill Kules, and Elisabeth Sylvan. 2005. Creativity support tool evaluation methods and metrics. In *Creativity Support Tools: A workshop sponsored by the National Science Foundation*. 10–24.
- [14] Rilla Khaled, Jonathan Lessard, and Pippin Barr. 2018. Documenting trajectories in design space: a methodology for applied game design research. In *Proceedings of the 13th International Conference on the Foundations of Digital Games*.
- [15] Max Kreminski, Isaac Karth, Michael Mateas, and Noah Wardrip-Fruin. 2022. Evaluating mixed-initiative creative interfaces via expressive range coverage analysis. In *IUI Workshops*. 34–45.
- [16] Mina Lee, Katy Ilonka Gero, John Joon Young Chung, Simon Buckingham Shum, Vipul Raheja, Hua Shen, Subhashini Venugopalan, Thiemo Wambsgans, David Zhou, Emad A. Alghamdi, Tal August, Avinash Bhat, Madiha Zahrah Choksi, Senjuti Dutta, Jin L. C. Guo, Md Naimul Hoque, Yewon Kim, Simon Knight, Seyed Parsa Neshaei, Agnia Sergeyuk, Antonette Shibani, Disha Shrivastava, Lila Shroff, Jessi Stark, Sarah Stermann, Sitong Wang, Antoine Bosselut, Daniel Buschek, Joseph Chee Chang, Sherol Chen, Max Kreminski, Joonsuk Park, Roy Pea, Eugenia H. Rho, Shannon Zejiang Shen, and Pao Siangliulue. 2024. A Design Space for Intelligent and Interactive Writing Assistants. In *Proceedings of the CHI Conference on Human Factors in Computing Systems*. doi:10.1145/3613904.3642697
- [17] Seung Won Lee, Tae Hee Jo, Semin Jin, Jiin Choi, Kyungwon Yun, Sergio Bromberg, Seonghoon Ban, and Kyung Hoon Hyun. 2024. The Impact of Sketch-guided vs. Prompt-guided 3D Generative AIs on the Design Exploration Process. In *Proceedings of the CHI Conference on Human Factors in Computing Systems*.
- [18] Jingyi Li, Eric Rawn, Jacob Ritchie, Jasper Tran O'Leary, and Sean Follmer. 2023. Beyond the artifact: power as a lens for creativity support tools. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*.
- [19] Ryan Louie, Jesse Engel, and Cheng-Zhi Anna Huang. 2022. Expressive communication: Evaluating developments in generative models and steering interfaces for music creation. In *Proceedings of the 27th International Conference on Intelligent User Interfaces*. 405–417.
- [20] Michael Muller, Lydia B Chilton, Mary Lou Maher, Charles Patrick Martin, Minsik Choi, Greg Walsh, and Anna Kantosalo. 2025. GenAICHI 2025: Generative AI and HCI at CHI 2025. In *Proceedings of the Extended Abstracts of the CHI Conference on Human Factors in Computing Systems*.
- [21] Mark J Nelson, Swen E Gaudl, Simon Colton, and Sebastian Deterding. 2018. Curious users of casual creators. In *Proceedings of the 13th International Conference on the Foundations of Digital Games*.
- [22] Vishakh Padmakumar and He He. 2023. Does writing with language models reduce content diversity? *arXiv preprint arXiv:2309.05196* (2023).
- [23] Ei Pa Pa Pe-Than, Laura Dabbish, and James Herbsleb. 2021. Open collaborative writing: Investigation of the fork-and-pull model. *Proceedings of the ACM on Human-Computer Interaction* 5, CSCW1 (2021).
- [24] Eric Rawn, Jingyi Li, Eric Paulos, and Sarah E Chasins. 2023. Understanding Version Control as Material Interaction with Quickpose. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*.
- [25] Christian Remy, Lindsay MacDonald Vermeulen, Jonas Frich, Michael Mose Biskjaer, and Peter Dalsgaard. 2020. Evaluating creativity support tools in HCI research. In *Proceedings of the 2020 ACM Designing Interactive Systems Conference*. 457–476.
- [26] Melissa Roemmele. 2021. Inspiration through observation: Demonstrating the influence of automatically generated text on creative writing. In *Proceedings of the Twelfth International Conference on Computational Creativity (ICCC 2021)*.
- [27] Hope Schroeder, Amy Smith, Ziv Epstein, Mike Cook, Simon Colton, and Andrew Lippman. 2023. Trash to Treasure: Using text-to-image models to inform the design of physical artefacts. In *The AAAI-23 Workshop on Creative AI Across Modalities*.
- [28] Ben Shneiderman. 2007. Creativity support tools: accelerating discovery and innovation. *Commun. ACM* 50, 12 (2007), 20–32.
- [29] Nikhil Singh, Guillermo Bernal, Daria Savchenko, and Elena L Glassman. 2023. Where to hide a stolen elephant: Leaps in creative writing with multimodal machine intelligence. *ACM Transactions on Computer-Human Interaction* 30, 5 (2023).
- [30] Amy Smith, Barrett R Anderson, Jasmine Tan Otto, Isaac Karth, Yuqian Sun, John Joon Young Chung, Melissa Roemmele, and Max Kreminski. 2025. Fuzzy Linkography: Automatic Graphical Summarization of Creative Activity Traces. In *Proceedings of the 17th Conference on Creativity & Cognition*.
- [31] Sarah Stermann, Molly Jane Nicholas, and Eric Paulos. 2022. Towards Creative Version Control. *Proceedings of the ACM on Human-Computer Interaction* 6, CSCW2 (2022).
- [32] Khonzoda Umarova, Talia Wise, Zhuoer Lyu, Mina Lee, and Qian Yang. 2025. How Problematic Writer-AI Interactions (Rather than Problematic AI) Hinder Writers' Idea Generation. *arXiv preprint arXiv:2503.11915* (2025).
- [33] David Zhou and Sarah Stermann. 2024. Ai.llude: Investigating Rewriting AI-Generated Text to Support Creative Expression. In *Proceedings of the 16th Conference on Creativity & Cognition*. 241–254.