

Towards Addressing Ableism on Social Media

Sharon Heung
ssh247@cornell.edu
Cornell Tech
New York, New York, USA

Abstract

Social media platforms have become breeding grounds for hate and harassment, impacting users' safety and well-being. Historically marginalized groups are common targets of online hate and harassment, including the disability community. While prior work has understood various forms of online discrimination (e.g. sexism and racism), ableism and disability-related discrimination have largely been overlooked. My work centers on understanding how disabled people experience ableism on social media and designing platform moderation to mitigate such harmful interactions. In this paper, I present my two-fold, research approach where I: 1) examine ableism on social media, and 2) explore the design of platform moderation to mitigate the harm of viewing ableist speech. Understanding ableism online magnifies the disability stereotypes and discrimination that exist today and designing tools to address ableist speech will further conversations on platform moderation and digital safety.

CCS Concepts

• **Social and professional topics** → **People with disabilities**;
• **Human-centered computing** → **Empirical studies in HCI**;
Social media; *Accessibility*.

Keywords

ableism, social media, platform moderation, online hate, accessibility, digital safety

ACM Reference Format:

Sharon Heung. 2024. Towards Addressing Ableism on Social Media. In *Companion of the 2024 Computer-Supported Cooperative Work and Social Computing (CSCW Companion '24)*, November 9–13, 2024, San Jose, Costa Rica. ACM, New York, NY, USA, 3 pages. <https://doi.org/10.1145/3678884.3682044>

1 Introduction

A growing body of work within Human-Computer Interactions has focused on understanding online hate and harassment towards historically marginalized communities. While disability identity is the largest minority group with an estimated population of 1.3 billion worldwide [1], ableism¹ has largely been overlooked. To address this gap in the literature, my research aims to tackle the

¹I refer to ableism as discrimination against disabled people [5]. More broadly, ableism is a "system that places value on people's bodies and minds based on societally constructed ideas of normalcy, intelligence, excellence, and productivity." [16]

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

CSCW Companion '24, November 9–13, 2024, San Jose, Costa Rica

© 2024 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-1114-5/24/11

<https://doi.org/10.1145/3678884.3682044>

following question: how do disabled people experience ableism and how might we design social media platforms to address the harm? More specifically, my research agenda is two-fold:

- (1) Examining ableism on social media (*completed and published*)
- (2) Designing moderation interventions to mitigate harm (*work in-progress and future work*)

While disabled people utilize social media to educate others about disability [3, 6], provide emotional and information peer support [7], share hobbies and stream games [19], and advocate for disability rights and accessibility accommodations [20], a few studies surfaced negative and hateful experiences online. For example, Sannon et al. [20] examined sociotechnical challenges of disability activists, finding that online visibility came at a cost of hate and harassment. While investigating the content moderation experiences of Blind TikTokkers, Lyu et al. [17] presented instances of Blind TikTokkers getting harassed and wrongfully reported. There ableist experiences are known to cause emotional harm, financial harm (especially towards content creators), and may exacerbate existing mental health conditions [11, 19]. Furthermore, online ableist speech perpetuates disability stereotypes and stigmatization that further marginalizes disabled people.

Although ableist speech manifests uniquely, how content moderation could mitigate the harm may be no different than addressing other forms of online hate and problematic content. Content moderation is the organized practice of screening and controlling unwanted content in online spaces whether that content is deemed as "irrelevant, obscene or illegal" [8, 21]. Platform and community moderation employ moderators and algorithms to filter and remove content [9]. Beyond top-down moderation, platforms provide capabilities that allow users to shape what they see (e.g. choosing who to follow or unfollow). Platforms also offer personal account moderation tools [13] to enable users to configure preferences on what accounts or types of content they want to avoid (e.g. blocking accounts and muting keywords). This form of personal content moderation or personalized moderation aligns with scholars' call to shift away from a one-size-fits-all approach given normative differences across cultures and communities [14, 22].

Currently, personalized moderation is being integrated with AI to create "smart" filters. For example, Intel's Bleep [18], a voice-based moderation tool, censors and filters out harmful content based on ableism, racism, and sexism. However, this begs an additional question: how should AI-based personalized moderation be designed to prevent the harm of viewing ableist text? For my future work, I propose designing and evaluating an ableism-specific personalized moderation tool with the online disability community that is rooted in my published work on understanding online ableism. In summary, my work uses a human-centered design and participatory approach by first understanding ableism on social

media and then designing platform interventions to mitigate the harm of viewing ableist speech.

2 Examining Ableism on Social Media

To date, my research [11, 12] has holistically examined how disabled people experience ableism on social media. My empirical work has surfaced disabled social media users' experiences with ableist speech from microaggressions to overt forms of hate.

Ableist Microaggressions. First, I investigated ableist microaggressions, subtle and nuanced forms of ableist speech. While several scholars have examined disabled people's experiences with in-person microaggressions [15], I wondered how disabled people experience online microaggressions and how might these experiences differ from offline settings. I conducted semi-structured interviews with 20 disabled social media users to understand how they experience and cope with ableist microaggressions. Across the interviews, I developed 12 types of ableist microaggressions, some of which are presented as archetypes. For example, *"Are you sure your husband loves you?"* is an archetype and direct quote that preserves the nuance and visceral impact.

This work also described the impact of ableist microaggressions on participants' wellbeing, self-worth, and social media use. For instance, one participant explained that they limit sharing photos of themselves *"to reduce discrimination and prevent unwanted questions on [their] vision impairment."* Although microaggressions are subtle, they left a lasting impact and changed the way participants engaged online (e.g. self-censorship and reduced social media use); hence the beginning of the paper title *"Nothing Micro About It."*

Ableist microaggressions may be difficult to detect by moderators or moderation algorithms. Furthermore, a removal of a microaggression may not be an appropriate platform response. Designs of platform moderation should carefully consider ways to address microaggressions, using both proactive (e.g. nudging perpetrators to not post microaggressions) and reactive (e.g. educating the perpetrator on why it's harmful) interventions.

Ableist Hate & Harassment. Prior research has surfaced instances of online hate towards specific types of disabled content creators (e.g. disability activists [20]) or certain disability groups (e.g. ADHD communities [7]). I extended this literature by holistically understanding various forms of ableism. To do this, I conducted a mixed-method study with varying types of creators with diverse disabilities. I recruited disabled creators because they are at-risk of experiencing online harassment due to their profiles' visibility. First, I disseminated a survey study (n=50) to understand the varying experiences with ableist hate and harassment and their self-reported frequency of ableist hate. Then I conducted follow-up interviews (n=20) to probe deeper on their discriminatory experiences and how it's impacted their content creation practices.

From the surveys and interviews, I developed a taxonomy of ableist hate and harassment with 5 main categories (e.g. Slurs and Derogatory Language) and 11 ableist types (e.g. Using Disability as an Insult). We also uncovered various forms of harassing tactics that escalated and intensified the hateful experience. For example, one participant's harasser found misinformation to support their accusation that the creator was faking their disability. Using statistical modeling, I found that LGBTQ disabled creators significantly

experience more and a disproportionate amount of ableist hate and harassment than non-LGBTQ disabled creators. This analysis contributes a descriptive model, signaling that intersectionality affects the amount of ableist hate received and as a result may lead to disproportionate impact on creators' with multiple marginalized identities.

Furthermore, I found multiple instances of platform moderation failing to support disabled content creators. For example, participants shared that the platforms oftentimes did not remove ableist hate even after the participant reported it. Due to these gaps in moderation, disabled creators experienced emotional harm (e.g. internalized ableism), and additional labor (e.g. crowdsourcing other users to report hate on their behalf). Over time, participants developed strategies to be resilient to these moderation flaws (e.g. avoiding reading the comments).

3 Designing Moderation Interventions to Mitigate Harm

After deeply understanding online ableism, I am now motivated to design moderation interventions to mitigate and reduce the harm of viewing ableist speech. Since my work has continuously centered the disability community, I hope to design user-facing tools that would empower them to reduce the likelihood of viewing triggering, ableist text. My previous participants acknowledged that the harmfulness of ableist speech is highly dependent on the individual; comments that are more ableist are not necessarily more harmful. Therefore, I was led towards imagining how to design personal content moderation tools that afford individuals to input their own preferences of what types of speech to be hidden. The main goal of this next research phase is to ideate, design, and evaluate a personalized moderation tool that mitigates the harm of viewing ableist text.

Designing Personalized Moderation to Address Ableist Text (Ongoing Work). I am currently exploring how personalized moderation on social media should be designed to mitigate the harm of viewing ableist text. This study examines two ends of a personalized moderation system: 1) how users configure their filters of ableist speech (e.g. intensity of ableism, types of ableism [11], or proportion of moderation) and 2) how users configure the presentation of the ableist speech (e.g. AI rewriting comment to be less toxic or obfuscating the original comment with an overlay categorizing the type of hate). By presenting design probes to disabled social media users, I plan to identify design recommendations of a personalized moderation tool. I will also ideate with participants on additional capabilities that will help support them in preventing the harm and managing their configured preferences. I anticipate this study will surface tensions of labor (e.g. the emotional labor of managing one's preferences), perceptions on trusting AI (e.g. to correctly identify and rewrite ableist speech), and fear of missing out on hateful comments (similar to what was found in [13]).

Building & Evaluating Personalized Moderation Tool (Future Work). I plan to build and evaluate an ableist-specific personalized moderation tool where disabled social media users can configure preferences to filter and change the presentation of ableist comments. This tool will integrate large language models (LLMs) to identify ableist comments and have additional capabilities to help

users' manage their filters; such as a timer (e.g. to only filter for the next 24 hours) and an analytics feature, detailing what types of hate is being detected and written by which users. By constructing the tool, I can evaluate LLMs' accuracy in identifying ableist text. Prior work has found LLMs to be biased with false positives, labelling sentences with disability-related words as toxic [10]. This proposed study can evaluate if current state of the art models can accurately identify implicit and explicit forms of ableist speech. Furthermore, it would be interesting to measure users' perception of the tool's accuracy and their overall trust of the system. While the onus of moderation should not fully be placed on the individual user, this study has the potential to contribute a working tool that reduces the harm of viewing triggering ableist text.

4 Contributions & Outcomes

The contributions of my research include: (1) theoretical framework of online ableism, (2) design considerations for personalized moderation, and (3) an evaluation of a system that identifies and obfuscates ableist speech, reducing the harm of viewing triggering text. Additionally, my work aligns with trauma-informed computing [2, 23], focused on alleviating online harm and actively designing systems to avoid re-traumatization. My research approach follows design justice practices [4] by directly and meaningfully involving the disability community in the participation of research and the design of technology. More broadly, my work engages with conversations on platform moderation and online safety. The outcome of my research advances our understanding of moderation techniques to mitigate content-based and social media-based harm, and could also inform tools to support researchers and moderators viewing traumatic content.

Acknowledgments

Thank you to all past participants for openly sharing their experiences. I'd like to acknowledge the endless support and mentorship from my advisors and committee, Shiri Azenkot, Aditya Vashistha, and Karen Levy. My work is supported by the National Science Foundation Graduate Research Fellowship Program.

References

- [1] 2023. <https://www.who.int/news-room/fact-sheets/detail/disability-and-health>
- [2] Janet X. Chen, Allison McDonald, Yixin Zou, Emily Tseng, Kevin A Roundy, Acar Tamersoy, Florian Schaub, Thomas Ristenpart, and Nicola Dell. 2022. Trauma-Informed Computing: Toward Safer Technology Experiences for All. In *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems (CHI '22)*. Association for Computing Machinery, New York, NY, USA, Article 544, 20 pages. <https://doi.org/10.1145/3491102.3517475>
- [3] Dasom Choi, Uichin Lee, and Hwajung Hong. 2022. "It's not wrong, but I'm quite disappointed": Toward an Inclusive Algorithmic Experience for Content Creators with Disabilities. In *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems (CHI '22)*. Association for Computing Machinery, New York, NY, USA, 1–19. <https://doi.org/10.1145/3491102.3517574>
- [4] Sasha Costanza-Chock. 2020. *Design Practices: "Nothing about Us without Us"* (1 ed.). <https://designjustice.mitpress.mit.edu/pub/cfohnud7>
- [5] Dana Dunn. 2021. Understanding ableism and negative reactions to disability. <https://www.apa.org/ed/precollege/psychology-teacher-network/introductory-psychology/ableism-negative-reactions-disability>
- [6] Jared Duval, Ferran Altarriba Bertran, Siying Chen, Melissa Chu, Divya Subramanian, Austin Wang, Geoffrey Xiang, Sri Kurniawan, and Katherine Isbister. 2021. Chasing Play on TikTok from Populations with Disabilities to Inspire Playful and Inclusive Technology Design. In *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems (CHI '21)*. Association for Computing Machinery, New York, NY, USA, 1–15. <https://doi.org/10.1145/3411764.3445303>
- [7] Tessa Eagle and Kathryn E. Ringland. 2023. "You Can't Possibly Have ADHD": Exploring Validation and Tensions around Diagnosis within Unbounded ADHD Social Media Communities. In *Proceedings of the 25th International ACM SIGACCESS Conference on Computers and Accessibility (ASSETS '23)*. Association for Computing Machinery, New York, NY, USA, Article 29, 17 pages. <https://doi.org/10.1145/3597638.3608400>
- [8] Casey Fiesler, Jialun Jiang, Joshua McCann, Kyle Frye, and Jed Brubaker. 2018. Reddit Rules! Characterizing an Ecosystem of Governance. *Proceedings of the International AAAI Conference on Web and Social Media* 12, 1 (Jun. 2018). <https://doi.org/10.1609/icwsm.v12i1.15033>
- [9] Tarleton Gillespie. 2018. *Custodians of the Internet: Platforms, Content Moderation, and the Hidden Decisions That Shape Social Media*. 1–288 pages. <https://doi.org/10.12987/9780300235029>
- [10] Saad Hassan, Matt Huenerfauth, and Cecilia Ovesdotter Alm. 2021. Unpacking the Interdependent Systems of Discrimination: Ableist Bias in NLP Systems through an Intersectional Lens. *ArXiv abs/2110.00521* (2021). <https://api.semanticscholar.org/CorpusID:238253456>
- [11] Sharon Heung, Lucy Jiang, Shiri Azenkot, and Aditya Vashistha. 2024. "Vulnerable, Victimized, and Objectified": Understanding Ableist Hate and Harassment Experienced by Disabled Content Creators on Social Media. In *Proceedings of the CHI Conference on Human Factors in Computing Systems (CHI '24)*. Association for Computing Machinery, New York, NY, USA. <https://doi.org/10.1145/3613904.3641949>
- [12] Sharon Heung, Mahika Phutane, Shiri Azenkot, Megh Marathe, and Aditya Vashistha. 2022. Nothing Micro About It: Examining Ableist Microaggressions on Social Media. In *Proceedings of the 24th International ACM SIGACCESS Conference on Computers and Accessibility (ASSETS '22)*. Association for Computing Machinery, New York, NY, USA, 1–14. <https://doi.org/10.1145/3517428.3544801>
- [13] Shagun Jhaver, Alice Qian Zhang, Quan Ze Chen, Nikhila Natarajan, Ruotong Wang, and Amy X. Zhang. 2023. Personalizing Content Moderation on Social Media: User Perspectives on Moderation Choices, Interface Design, and Labor. *Proc. ACM Hum.-Comput. Interact.* 7, CSCW2, Article 289 (oct 2023), 33 pages. <https://doi.org/10.1145/3610080>
- [14] Jialun Aaron Jiang, Morgan Klaus Scheuerman, Casey Fiesler, and Jed R Brubaker. 2021. Understanding international perceptions of the severity of harmful content online. *PLOS ONE* 16, 8 (August 2021), 1–22. <https://doi.org/10.1371/journal.pone.0256>
- [15] Richard Keller and Corinne Galgay. 2010. Microaggressive Experiences of People with Disabilities. *Microaggressions and Marginality: Manifestation, Dynamics, and Impact* (2010), 241–267.
- [16] Talia Lewis. 2021. Blog Archives — talialewis.com. <https://www.talialewis.com/blog/archives/01-2022>. [Accessed 24-05-2024].
- [17] Yao Lyu, Jie Cai, Anisa Callis, Kelley Cotter, and John M. Carroll. 2024. "I Got Flagged for Supposed Bullying, Even Though It Was in Response to Someone Harassing Me About My Disability": A Study of Blind TikTokers' Content Moderation Experiences. In *Proceedings of the CHI Conference on Human Factors in Computing Systems (CHI '24)*. Association for Computing Machinery, New York, NY, USA, Article 741, 15 pages. <https://doi.org/10.1145/3613904.3642148>
- [18] Jon Porter. 2021. Today I learned about Intel's AI sliders that filter online gaming abuse. <https://www.theverge.com/2021/4/8/22373290/intel-bleep-ai-powered-abuse-toxicity-gaming-filters>
- [19] Kathryn E. Ringland. 2019. "Autsume": Fostering an Autistic Identity in an Online Minecraft Community for Youth with Autism. *Information in Contemporary Society : 14th International Conference, iConference 2019, Washington, DC, USA, March 31-April 3, 2019, Proceedings. iConference (Conference) (14th : 2019 : Washington, D.C.)* 11420 (April 2019), 132–143. https://doi.org/10.1007/978-3-030-15742-5_12
- [20] Shruti Sannon, Jordyn Young, Erica Shusas, and Andrea Forte. 2023. Disability Activism on Social Media: Sociotechnical Challenges in the Pursuit of Visibility. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems (CHI '23)*. Association for Computing Machinery, New York, NY, USA, 1–15. <https://doi.org/10.1145/3544548.3581333>
- [21] Morgan Klaus Scheuerman, Jialun Aaron Jiang, Casey Fiesler, and Jed R. Brubaker. 2021. A Framework of Severity for Harmful Content Online. *Proc. ACM Hum.-Comput. Interact.* 5, CSCW2, Article 368 (oct 2021), 33 pages. <https://doi.org/10.1145/3479512>
- [22] Sarita Schoenebeck, Amna Batool, Giang Do, Sylvia Darling, Gabriel Grill, Daricia Wilkinson, Mehtab Khan, Kentaro Toyama, and Louise Ashwell. 2023. Online Harassment in Majority Contexts: Examining Harms and Remedies across Countries. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems (CHI '23)*. Association for Computing Machinery, New York, NY, USA, Article 485, 16 pages. <https://doi.org/10.1145/3544548.3581020>
- [23] Carol F Scott, Gabriela Marcu, Riana Elyse Anderson, Mark W Newman, and Sarita Schoenebeck. 2023. Trauma-Informed Social Media: Towards Solutions for Reducing and Healing Online Harm. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems (CHI '23)*. Association for Computing Machinery, New York, NY, USA, Article 341, 20 pages. <https://doi.org/10.1145/3544548.3581512>