

# Communication Games to Study Emerging Disinformation Scenarios With LLMs

Gionnieve Lim  
Singapore University of Technology  
and Design  
Singapore, Singapore  
gionnievelim@gmail.com

Bryan Chen Zhengyu Tan  
Singapore University of Technology  
and Design  
Singapore, Singapore  
bryantanzy@gmail.com

Kellie Yu Hui Sim  
Singapore University of Technology  
and Design  
Singapore, Singapore  
kelliesyhh@gmail.com

Weiyan Shi  
Singapore University of Technology  
and Design  
Singapore, Singapore  
weiyanshi6@gmail.com

Ming Hui Chew  
Singapore University of Technology  
and Design  
Singapore, Singapore  
chewminghuis@gmail.com

Ming Shan Hee  
Singapore University of Technology  
and Design  
Singapore, Singapore  
mingshan\_hee@myemail.sutd.edu.sg

Roy Ka-Wei Lee  
Singapore University of Technology  
and Design  
Singapore, Singapore  
roy\_lee@sutd.edu.sg

Simon T. Perrault  
Singapore University of Technology  
and Design  
Singapore, Singapore  
perrault.simon@gmail.com

Kenny Tsu Wei Choo  
Singapore University of Technology  
and Design  
Singapore, Singapore  
kennytwchoo@gmail.com

## Abstract

Disinformation is a complex phenomenon that has been studied using various methods, showing the innovative ways researchers seek to illuminate insights on its dissemination, impact and countermeasures. In this paper, we propose communication games as a promising method for examining disinformation, especially in emerging scenarios like the introduction of new technologies. The rise of Large Language Models (LLMs) has brought forward advanced tools capable of producing human-like text at scale. These tools can be exploited to craft convincing and complex disinformation, but also present opportunities for improving strategies to detect and counteract it. This paper investigates how LLMs influence disinformation and combative practices through a communication game, inspired by *Werewolf*, with 25 participants. Our findings reveal varying uses of LLMs depending on participants' roles and strategies, underscoring the importance of understanding their effectiveness in this context. We demonstrate that communication games can be a valuable addition to the disinformation toolkit.

## CCS Concepts

• **Human-centered computing** → **Collaborative and social computing design and evaluation methods**; **Empirical studies in collaborative and social computing**.

## Keywords

Disinformation Research Methods; Communication Games; Werewolf; Social Manipulation; Large Language Model; Chatbot Uses

## ACM Reference Format:

Gionnieve Lim, Bryan Chen Zhengyu Tan, Kellie Yu Hui Sim, Weiyan Shi, Ming Hui Chew, Ming Shan Hee, Roy Ka-Wei Lee, Simon T. Perrault, and Kenny Tsu Wei Choo. 2025. Communication Games to Study Emerging Disinformation Scenarios With LLMs. In *Companion of the Computer-Supported Cooperative Work and Social Computing (CSCW Companion '25)*, October 18–22, 2025, Bergen, Norway. ACM, New York, NY, USA, 6 pages. <https://doi.org/10.1145/3715070.3749236>

## 1 Introduction

Large Language Models (LLMs) like GPT-4o<sup>1</sup> and Llama 3<sup>2</sup> are reshaping the creation and dissemination of online content. LLMs can quickly generate human-like text, enabling a range of uses—from summarising complex information to aiding content moderation [23]—but also facilitating the spread of misleading narratives [6]. In social media environments like Facebook and Reddit, where even low-effort falsehoods can rapidly gain traction [13], powerful LLMs further complicate efforts to detect and limit disinformation.

Although much is known about how false content propagates—from bot-driven influence campaigns to user susceptibility—studies often rely on indirect indicators rather than real-time observation of disinformers' *deliberate* and *adaptive* tactics. Investigating these tactics has become even more urgent since LLMs can rapidly produce context-specific, highly persuasive content that is challenging to distinguish from genuine communication [9, 48]. Directly studying active disinformation agents is notoriously difficult: clandestine actors rarely reveal their motives, and ethical or logistical constraints often prevent researchers from accessing their workflows [20, 33].

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).  
CSCW Companion '25, Bergen, Norway  
© 2025 Copyright held by the owner/author(s).  
ACM ISBN 979-8-4007-1480-1/25/10  
<https://doi.org/10.1145/3715070.3749236>

<sup>1</sup><https://openai.com/index/hello-gpt-4o/>

<sup>2</sup><https://ai.meta.com/blog/meta-llama-3/>

Consequently, critical questions remain about how these disinformers may integrate LLMs into their strategies, how quickly those strategies adapt, and whether potential defenders (e.g., moderators or well-intentioned users) can effectively counter new threats.

To explore how LLMs might shape deceptive tactics and potential defences in small, localised settings, we introduce a *communication game* inspired by *Werewolf* (also known as *Mafia*). Our approach provides a controlled environment where participants adopt distinct roles: *Disinformers* (analogous to Werewolves) seek to spread and mask falsehoods, while *Moderators* (Seers) and *Users* (Villagers) try to identify and stop them. This format allows us to closely examine instances of disinformation and defensive responses from both malicious and benign actors, offering insights into how intentions, strategies, and behaviours adapt when LLM assistance is readily available.

In this paper, we describe how the novel application of this game-based methodology can help researchers and designers understand LLM-driven disinformation from multiple perspectives. Specifically, our contributions include:

- An exploratory *communication game* that examines LLM use for spreading or detecting disinformation.
- Empirical observations of how participants leverage LLMs for content generation, verification, and strategic guidance while attempting to either deceive or counter deception.
- Reflections on the trade-offs and value of communication games as a research method to study disinformation.

## 2 Background

### 2.1 Disinformation in Online Platforms

Disinformation is the *intentional* production and dissemination of false information to deceive or manipulate [53]. Persuasive and manipulative strategies are often employed in disinformation to influence public opinion and behaviour [15]. While large-scale campaigns and state-backed efforts have garnered significant attention, disinformation also proliferates through intimate, interpersonal interactions [8] and these localised incidents can influence community-level opinions [7, 39] and deepen social divides [40]. Past studies have focused on the use of emotional appeals, fabricated narratives, and re-sharing dynamics in disinformation [18, 31, 41]. However, systematically capturing the *intent* behind disinformation remains challenging. Disinformers rarely disclose their tactics, and ethical or logistical constraints make the real-time observation of deceptive behaviours difficult [20, 33]. As a result, research often relies on retrospective methods like content analyses [24, 34, 41, 52] and algorithmic signals [2, 19, 28, 50], limiting our understanding of how deceptive strategies evolve during an ongoing interaction.

### 2.2 Disinformation in the Age of LLMs

Recent advances in LLMs drastically lower the barrier to generating high-quality, context-specific falsehoods [14]. Although mainstream providers implement safeguards [30, 35], open-source or minimally moderated models remain accessible [3], enabling malicious users to produce credible-looking content at scale [29, 48]. Alarming, studies find that humans often struggle to distinguish AI-generated text from human-authored content [9, 57], highlighting the risk of widespread deception. Simultaneously, LLMs can

bolster disinformation detection. Tools for summarisation, query generation, and reasoned fact-checking may accelerate the debunking process [10, 45]. Yet their efficacy depends on prompt design, model reliability, and user skill [9, 37]. This dual role of LLMs, sometimes described as a “double-edged sword” [46], underscores the need to examine how both malicious and well-intentioned actors might leverage AI tools in smaller, more interactive contexts where social cues and user interactions play pivotal roles.

## 3 Disinformation Research Methods

Disinformation scholars employ a range of research methods to dissect how falsehoods arise, spread, and evolve [12]. From a brief systematic review of disinformation-related review papers, we identified a list of common methods.

The systematic review was conducted using the Web of Science<sup>3</sup> database following the PRISMA approach [36]. The search query matched for keywords in article titles: (TI=(disinformation OR misinformation OR "fake news")) AND (TI=(review OR survey)). Figure 1 shows the exclusion criteria. After attaining the final corpus ( $N = 114$ ), the review papers were quantitatively coded to extract and tabulate the disinformation research methods mentioned in them.

A variety of research methods were surfaced, and we discuss their value in the study of disinformation. Approaches such as *computational methods* ( $n = 65$ ), *content analysis* ( $n = 18$ ), and *data mining* ( $n = 14$ ) excel at detecting broad patterns and uncovering hidden relationships within large datasets, enabling the identification of trends, anomalies, and key indicators of disinformation at scale [32]. However, these methods often cannot reliably determine whether misleading material is crafted with intentional deceit [1, 56].

In contrast, smaller-scale techniques like surveys ( $n = 27$ ), *interviews/focus groups* ( $n = 16$ ), and *case studies* ( $n = 11$ ) allow researchers to probe motives, contextual details, and user experiences more closely [17, 44]. Yet, these methods generally rely on retrospective accounts, making it difficult to observe the real-time adaptation of deceptive tactics [38].

*Experiments* ( $n = 34$ ) can simulate certain aspects of disinformation (e.g., testing how labelled content affects user trust or modelling how falsehoods propagate in a network [25]), but these tend to simplify real-world group dynamics or miss the spontaneity of human-driven interactions [27, 47].

Hence, while existing methods shed light on disinformation's scope and impact, capturing the *intent* and adaptive strategies of those who craft and spread it remains challenging. With emerging technologies like LLMs, there is a growing need for complementary methods to study evolving behaviours [21, 22] in controlled yet interactive settings.

### 3.1 Communication Games as a Method

*Communication games* offer a complementary lens to study deception by creating a structured environment in which participants must navigate hidden information, uncertain motives, and evolving strategies. Social deduction games, such as *Werewolf* or *Mafia*,

<sup>3</sup><https://www.webofscience.com>




Figure 1: Flow diagram of the identification, screening and inclusion of articles

have historically provided insights into the mechanics of deception, collaboration, and group dynamics [49, 55]. By adapting these mechanics to an online forum context, it becomes possible to systematically investigate how participants use LLMs to spread or detect disinformation.

Compared to computational inferences or retrospective accounts, this method enables direct observation of strategic choices, actions, and role-based motivations in a controlled environment. This can be especially valuable for studying emerging disinformation scenarios, such as those involving new technologies, where the impact of novel behaviours and usage patterns remains uncertain. In the following section, we describe a study using a communication game, highlighting its merits (e.g., clear role distinctions, real-time adaptation) and limitations (e.g., smaller scale, contrived scenarios), and discussing how this approach can enrich broader research on LLM-driven disinformation.

#### 4 Communication Game Study

**Overview.** We developed a communication game inspired by the classic game, *Werewolf*, simulating an online forum where several Users and the Moderator attempt to resist and identify the Disinformer who tries to influence them. All players have access to an LLM chatbot. While disinformation can be far-reaching, we focus on small and localised settings, involving clusters of highly engaged users, in the scenario of LLMs being easily available to all. This game mirrors various aspects of social media, such as the anonymity of users, diversity of roles and goals, and the challenge of identifying deceptive actors.

**Platform.** Figure 2 illustrates the custom game platform developed for the study. Each player’s (A) username and (B) role are displayed at the top, followed by a (C) Announcements Panel and (D) Timer Countdown Panel showing the current phase and game duration. The (E) Forum Panel offers a standard interface for commenting, replying, upvoting, downvoting, and reporting posts. Moderators can also access a (F) Reports Panel to review and flag report comments. The (G) Debate Panel allows players to chat before casting votes, while the (H) Personal AI Chatbot Panel provides continuous access to the LLM which utilises *Nous-Hermes-2-Mixtral-8x7B-DPO*, an uncensored *Mixtral 8x7B MOE* model<sup>4</sup> variant. An (I) Online Members Panel lists active players, with moderators marked by “(Mod)”.

**Game Design.** Before the game, all players were given instructions on their roles, goals, and the platform. The Disinformer (D) has to sway others’ opinions on a topic towards a given stance and is asked to post at least one piece of false information in each round, the Moderator (M) is charged with maintaining a healthy discussion space and is able to flag reported posts, and three Users

(U) may freely engage in the discussion. As shown in Figure 3, the game runs for four rounds, each with several phases, spanning 2.5 hours. In the Discussion phase, participants interact with each other on the forum topic and try to disinform/combat (see Fig 2E, F, H). In Deliberation, participants review the forum discussion and contemplate who the Disinformer is. In Debate, the group openly discusses who the Disinformer is (see Fig 2G). In Decision, participants vote for the Disinformer in a separate form. In Post-Decision, they receive the anonymous voting outcomes and prepare for the next round. They may use the chatbot at all times for any purpose.

**Game Materials.** Participants discussed any two out of five topics selected for their controversy and relevance to the regional context where the study is held (e.g., conscription, euthanasia). Each topic was seeded with actual social media comments representing diverse viewpoints. To establish common knowledge across players, all participants received an influence guide that detailed actual disinformation [43] and persuasion [42] tactics, and a guide on how to use the LLM chatbot. All players received checklists of interactions to fulfil based on their role to ensure a base amount of engagement as well. Role assignments remained constant throughout the game, and the objective was to identify the Disinformer through collaborative reasoning and voting. Winning outcomes were incentivised with monetary bonuses up to ~US\$8 on top of a base sum of ~US\$24, as approved by our Institutional Review Board.

**Participants.** Using a Telegram recruitment group specific to our region and a local university email list, we recruited 25 participants (10 women, 15 men) aged 18 to 64. Most participants held at least a Bachelor’s degree ( $n = 20$ ), and all of them had prior experience with online forums and generative AI tools. Five game sessions were held, each with a group (G) of five participants, balanced as far as possible by their gender, stance towards the five topics, and availability that was filled during the screener survey.

**Data Collection.** We collected quantitative data from surveys and forum interactions and qualitative data through observations, interviews, and recordings. Participants’ demographics and expertise with AI tools were recorded in a pre-game survey. The influence tactics they used and experience with the LLM chatbot was recorded in a post-game survey. All forum interactions, including posts, votes, reports, and chatbot usage were logged. During the game, participants were encouraged to think-aloud while researchers observed them and recorded notes. After the game, participants were interviewed about the strategies they employed for their role and their experiences with the chatbot.

**Data Analysis.** We employed *reflexive thematic analysis* [5] involving five researchers who separately coded a subset of data and met regularly to refine emergent themes. Audio recordings from the think-aloud protocols and interviews were transcribed and arranged to match the forum interaction logs. The data was then analysed in two ways. From the role-level perspective, we

<sup>4</sup><https://huggingface.co/NousResearch/Nous-Hermes-2-Mixtral-8x7B-DPO>




Figure 2: Moderator's view of the forum platform. The Disinformers's and User's views are similar, but do not have the Reports Panel (F).




Figure 3: There are five phases (in yellow) in a round, each with a specific duration. The game has four rounds; players discuss a new topic after two rounds. Disinformers conjectures (in red) are measured at various periods during the game.

looked at commonalities and exceptional instances of how Disinformers, Moderators and Users behaved and used the chatbot. From the group-level perspective, we assessed gameplay dynamics, how players reacted to chatbot-generated content, and what led to Disinformers suspicion.

## 5 Results and Discussion

### 5.1 LLM Uses in the Disinformation

To show how communication games can provide insights on disinformation in emerging scenarios, we report a subset of results

focused on participants' use of the chatbot, and plan to conduct and report a full analysis in the future. Our findings highlight the dual roles of LLMs—as a *sword* to exacerbate disinformation, and as a *shield* to mitigate it. Examples illustrating how the Disinformers (D), Moderators (M) and Users (U1/2/3) behaved in various Groups (G) are provided.

**LLMs as a Sword.** Disinformers used the chatbot as a *sword* to amplify their agendas by generating false or misleading content, verify their own fabricated content, and obtain strategic guidance. For example, *G1D* and *G2D* twisted factual data into misleading narratives, while others used it strategically to create arguments

or seek tactical advice on avoiding detection (*G5D*) or deflecting suspicion (“*how to accuse someone of a role*”, prompted by *G2D*). These use cases demonstrate its potential in empowering malicious actors to automate and scale disinformation efforts.

**LLMs as a Shield.** Conversely, Moderators and Users used the chatbot as a *shield* to enhance contextual knowledge and counter disinformation. Examples include *G2U2* using it to cross-reference claims to identify discrepancies, or *G1U3* who had a minority stance and employed it to generate arguments and refine their approach to avoid further suspicion by blending in. In our study, the chatbot also offered guidance for identifying Disinformers and refining arguments, showcasing its potential to support not only fact-checking efforts but also strategic guidance.

LLMs have widely been proposed to empower disinformation through their language capabilities and world knowledge, generating persuasive and realistic content at scale [3, 4, 58]. Conversely, LLMs can enhance various stages of fact-checking through their reasoning capabilities [11, 16, 26]. However, our study reveals an **additional dimension of LLMs’ capabilities: offering tactical support to both malicious and virtuous actors**. This includes planning disinformation and evasion approaches for Disinformers, and identifying Disinformers and deflecting suspicion for the others.

## 5.2 Rethinking Communication Games for Disinformation Research

Our study demonstrates that a *Werewolf*-style communication game can effectively capture the *intentional*, *adaptive*, and *interactive* nature of disinformation. Specifically, communication games allows researchers to: **(1) Observe real-time deception and detection:** The game setting captures how disinformation tactics and detection evolve in response to real-time interactions, enabling direct observation of strategies as they unfold. **(2) Capture rich context and social dynamics:** The think-aloud reveals attitudes toward others as players interact and communicate, illuminating aspects on interpersonal trust, suspicion and social influence. **(3) Balance ethical feasibility:** Studying *actual* disinformers poses ethical and logistical barriers. Communication games offer *fictional* roles that minimise harm while examining deceptive interactions. **(4) Embed specific scenarios:** Our adaptation integrates an LLM chatbot into the game, mirroring how users might exploit AI in both deceptive and corrective ways on real platforms, which is difficult to isolate in large-scale computational research.

As with any research method, however, communication games have key inherent limitations: **(1) Ecological validity:** Participants are role-playing in a controlled environment and their actions may not reflect the actual behaviours of Disinformers, Moderators and Users on real platforms. **(2) Generalisability:** The logistics and duration of communication games, particularly with our design, makes it difficult to conduct the study with a large and diverse sample.

Despite their limitations, players’ attitudes and behaviours have been shown to reflect characteristics of their real-world counterparts in specific interactive scenarios [54]. Communication games can thus uniquely give insights into social phenomena often inaccessible to researchers [51]. For our case, the *Werewolf* game

enables us to gain early insights into how emerging disinformation scenarios with LLMs may possibly pan out in reality.

## 6 Conclusion

This work presents a first-of-its-kind exploration of the role of LLMs in the disinformation landscape, specifically in small and localised settings, through the use of a communication game as the research method. Via our custom *Werewolf* game, we observed how Disinformers, Moderators, and Users leveraged LLMs in an online forum that uniquely provided insights into the intentions and strategies of motivated role-playing disinformers. Our work reveals the value of communication games for studying emerging disinformation scenarios where empirical data can be hard to gather. Future work can explore different designs and player configurations of the *Werewolf* communication game to gain insights into other disinformation scenarios, such as collusion between multiple human and bot disinformers, as well as other objectives, such as exploring the reasons behind the unintentional spread of disinformation.

## References

- [1] Esma Aimeur, Sabine Amri, and Gilles Brassard. 2023. Fake News, Disinformation and Misinformation in Social Media: A Review. *Social Network Analysis and Mining* 13, 1 (Feb. 2023), 30. doi:10.1007/s13278-023-01028-5
- [2] Iuliia Alieva, Lynnette Hui Xian Ng, and Kathleen M. Carley. 2022. Investigating the Spread of Russian Disinformation about Biolabs in Ukraine on Twitter Using Social Network Analysis. In *2022 IEEE International Conference on Big Data (Big Data)*. doi:10.1109/BigData55660.2022.10020223
- [3] Dipto Barman, Ziyi Guo, and Owen Conlan. 2024. The dark side of language models: Exploring the potential of llms in multimedia disinformation generation and dissemination. *Machine Learning with Applications* (2024), 100545.
- [4] Ali Borji. 2023. A categorical archive of chatgpt failures. *arXiv preprint arXiv:2302.03494* (2023).
- [5] Virginia Braun and Victoria Clarke. 2006. Using thematic analysis in psychology. *Qualitative research in psychology* 3, 2 (2006), 77–101.
- [6] J Brewster, L Arvanitis, and MK Sadeghi. 2023. Could ChatGPT become a monster misinformation superspreader? *NewsGuard*.
- [7] George Buchanan, Ryan Kelly, Stephann Makri, and Dana McKay. 2022. Reading Between the Lies: A Classification Scheme of Types of Reply to Misinformation in Public Discussion Threads. In *Proceedings of the 2022 Conference on Human Information Interaction and Retrieval (Regensburg, Germany) (CHIIR '22)*. Association for Computing Machinery, New York, NY, USA, 243–253. doi:10.1145/3498366.3505823
- [8] Tom Buchanan. 2020. Why do people spread false information online? The effects of message and viewer characteristics on self-reported likelihood of sharing social media disinformation. *PLOS ONE* 15, 10 (10 2020), 1–33. doi:10.1371/journal.pone.0239666
- [9] Canyu Chen and Kai Shu. 2023. Can llm-generated misinformation be detected? *arXiv preprint arXiv:2309.13788* (2023).
- [10] Canyu Chen and Kai Shu. 2023. Combating misinformation in the age of llms: Opportunities and challenges. *AI Magazine* (2023).
- [11] Sumit Kumar Dam, Choong Seon Hong, Yu Qiao, and Chaoning Zhang. 2024. A Complete Survey on LLM-based AI Chatbots. *arXiv preprint arXiv:2406.16937* (2024).
- [12] Camille François. 2020. Actors, behaviors, content: A disinformation ABC. *Algorithms* (2020).
- [13] Christine Geeng, Savanna Yee, and Franziska Roesner. 2020. Fake News on Facebook and Twitter: Investigating How People (Don’t) Investigate. In *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems (Honolulu, HI, USA) (CHI '20)*. Association for Computing Machinery, New York, NY, USA, 1–14. doi:10.1145/3313831.3376784
- [14] Josh A Goldstein, Girish Sastry, Micah Musser, Renee DiResta, Matthew Gentzel, and Katerina Sedova. 2023. Generative language models and automated influence operations: Emerging threats and potential mitigations. *arXiv preprint arXiv:2301.04246* (2023).
- [15] Andrew M Guess and Benjamin A Lyons. 2020. Misinformation, disinformation, and online propaganda. *Social media and democracy: The state of the field, prospects for reform* 10 (2020).
- [16] Ankur Gupta, Yash Varun, Prarthana Das, Nithya Muttineni, Parth Srivastava, Hamim Zafar, Tanmoy Chakraborty, and Swaprava Nath. 2021. TruthBot: An

- Automated Conversational Tool for Intent Learning, Curated Information Presenting, and Fake News Alerting. *arXiv preprint arXiv:2102.00509* (2021).
- [17] Michael Hameleers. 2025. The Nature of Visual Disinformation Online: A Qualitative Content Analysis of Alternative and Social Media in the Netherlands. *Political Communication* 42, 1 (Jan. 2025), 108–126. doi:10.1080/10584609.2024.2354389
  - [18] HiveMind. 2024. *Disinformation and 7 Common Forms of Information Disorder*. <https://commonslibrary.org/disinformation-and-7-common-forms-of-information-disorder/>
  - [19] Muhammad Nihal Hussain, Serpil Tokdemir, Nitin Agarwal, and Samer Al-Khateeb. 2018. Analyzing disinformation and crowd manipulation tactics on YouTube. In *2018 IEEE/ACM International Conference on Advances in Social Networks Analysis and Mining (ASONAM)*. IEEE, 1092–1095.
  - [20] Index on Censorship. 2011. Interview with a troll. Retrieved December 1, 2024 from <https://www.indexoncensorship.org/2011/09/interview-with-a-troll/>
  - [21] Bohan Jiang, Zhen Tan, Ayushi Nirmal, and Huan Liu. 2023. Disinformation Detection: An Evolving Challenge in the Age of LLMs. doi:10.48550/arXiv.2309.15847 arXiv:2309.15847 [cs]
  - [22] Bohan Jiang, Chengshuai Zhao, Zhen Tan, and Huan Liu. 2024. Catching Chameleons: Detecting Evolving Disinformation Generated Using Large Language Models. doi:10.48550/arXiv.2406.17992 arXiv:2406.17992 [cs]
  - [23] Gretel Kahn. 2024. *Generative AI is already helping fact-checkers. But it's proving less useful in small languages and outside the West*. <https://reutersinstitute.politics.ox.ac.uk/news/generative-ai-already-helping-fact-checkers-its-proving-less-useful-small-languages-and>
  - [24] Gary King, Jennifer Pan, and Margaret E Roberts. 2013. How Censorship in China Allows Government Criticism but Silences Collective Expression. *American Political Science Review* 107, 2 (2013), 326–343. doi:10.1017/S0003055413000014
  - [25] Batool Lakzaei, Mostafa Haghir Chehreghani, and Alireza Bagheri. 2024. Disinformation Detection Using Graph Neural Networks: A Survey. *Artificial Intelligence Review* 57, 3 (Feb. 2024), 52. doi:10.1007/s10462-024-10702-9
  - [26] Gionnieve Lim and Simon T Perrault. 2023. Fact checking chatbot: A misinformation intervention for instant messaging apps and an analysis of trust in the fact checkers. In *Mobile Communication and Online Falsehoods in Asia: Trends, Impact and Practice*. Springer, 197–224.
  - [27] Alejandro Buitrago López, Javier Pastor-Galindo, and José A. Ruipérez-Valiente. 2024. Frameworks, Modeling and Simulations of Misinformation and Disinformation: A Systematic Literature Review. doi:10.48550/arXiv.2406.09343 arXiv:2406.09343 [cs]
  - [28] Jim Maddock, Kate Starbird, Haneen J. Al-Hassani, Daniel E. Sandoval, Mania Orand, and Robert M. Mason. 2015. Characterizing Online Rumoring Behavior Using Multi-Dimensional Signatures. In *Proceedings of the 18th ACM Conference on Computer Supported Cooperative Work & Social Computing* (Vancouver, BC, Canada) (CSCW '15). Association for Computing Machinery, New York, NY, USA, 228–241. doi:10.1145/2675133.2675280
  - [29] Raphael Meier. 2024. LLM-Aided Social Media Influence Operations. *Large Language Models in Cybersecurity: Threats, Exposure and Mitigation* (2024), 105–112.
  - [30] Meta. 2024. *Trust & Safety*. <https://llama.meta.com/trust-and-safety/>
  - [31] Tanushree Mitra, Graham P. Wright, and Eric Gilbert. 2017. A Parsimonious Language Model of Social Media Credibility Across Disparate Events. In *Proceedings of the 2017 ACM Conference on Computer Supported Cooperative Work and Social Computing* (Portland, Oregon, USA) (CSCW '17). Association for Computing Machinery, New York, NY, USA, 126–145. doi:10.1145/2998181.2998351
  - [32] M. Zohaib Nawaz, M. Saqib Nawaz, Philippe Fournier-Viger, and Yulin He. 2024. Analysis and Classification of Fake News Using Sequential Pattern Mining. *Big Data Mining and Analytics* 7, 3 (Sept. 2024), 942–963. doi:10.26599/BDMA.2024.9020015
  - [33] Ben Nimmo and Aric Toler. 2018. How They Did It: The Real Russian Journalists Who Exposed the Troll Factory in St. Petersburg. Retrieved December 1, 2024 from <https://gijn.org/stories/real-russian-journalists-exposed-troll-factory-st-petersburg/>
  - [34] Shuo Niu and Kathleen Palm Reed. 2023. Towards Understanding Substance Abuse Misinformation in YouTube Videos. In *Companion Publication of the 2023 Conference on Computer Supported Cooperative Work and Social Computing* (Minneapolis, MN, USA) (CSCW '23 Companion). Association for Computing Machinery, New York, NY, USA, 64–68. doi:10.1145/3584931.3606974
  - [35] OpenAI. 2024. *OpenAI Safety Practices*. <https://openai.com/index/openai-safety-update/>
  - [36] Matthew J Page, Joanne E McKenzie, Patrick M Bossuyt, Isabelle Boutron, Tammy C Hoffmann, Cynthia D Mulrow, Larissa Shamseer, Jennifer M Tetzlaff, Elie A Akl, Sue E Brennan, Roger Chou, Julie Glanville, Jeremy M Grimshaw, Asbjørn Hróbjartsson, Manoj M Lalu, Tianjing Li, Elizabeth W Loder, Evan Mayo-Wilson, Steve McDonald, Luke A McGuinness, Lesley A Stewart, James Thomas, Andrea C Tricco, Vivian A Welch, Penny Whiting, and David Moher. 2021. The PRISMA 2020 statement: an updated guideline for reporting systematic reviews. *BMJ* 372 (2021). doi:10.1136/bmj.n71 arXiv:https://www.bmj.com/content/372/bmj.n71.full.pdf
  - [37] Kellin Pelrine, Anne Imouza, Camille Thibault, Meilina Reksoprodjo, Caleb Gupta, Joel Christoph, Jean-François Godbout, and Reihaneh Rabbany. 2023. Towards reliable misinformation mitigation: Generalization, uncertainty, and gpt-4. *arXiv preprint arXiv:2305.14928* (2023).
  - [38] Marta Pérez-Escobar, Darren Lilleker, and Alejandro Tapia-Frade. 2023. A Systematic Literature Review of the Phenomenon of Disinformation and Misinformation. *Media and Communication* 11 (March 2023). doi:10.17645/mac.v11i2.6453
  - [39] Martin Potter. 2021. Bad actors never sleep: content manipulation on Reddit. *Continuum* 35, 5 (03 Sep 2021), 706–718. doi:10.1080/10304312.2021.1983254
  - [40] Nathalie Van Raemdonck. 2019. The Echo Chamber of Anti-Vaccination Conspiracies: Mechanisms of Radicalization on Facebook and Reddit. Retrieved December 1, 2024 from [https://papers.ssrn.com/sol3/papers.cfm?abstract\\_id=3510196](https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3510196)
  - [41] Raquel Recuero, Felipe Soares, and Otávio Vinhas. 2020. Discursive strategies for disinformation on WhatsApp and Twitter during the 2018 Brazilian presidential election. *First Monday* 26, 1 (Dec. 2020). doi:10.5210/fm.v26i1.10551
  - [42] Cialdini Robert. 2001. Harnessing the science of persuasion. *Harvard Business Review* 79, 9 (2001), 72–79.
  - [43] Jon Roozenbeek and Sander Van der Linden. 2019. Fake news game confers psychological resistance against online misinformation. *Palgrave Communications* 5, 1 (2019), 1–10.
  - [44] Julian Rossing. 2023. The disinformation landscape of the future : a qualitative perspective. <http://essay.utwente.nl/97407/>
  - [45] Vinay Setty. 2024. FactCheck Editor: Multilingual Text Editor with End-to-End fact-checking. In *Proceedings of the 47th International ACM SIGIR Conference on Research and Development in Information Retrieval*. 2744–2748.
  - [46] Siddhant Bikram Shah, Surendrabikram Thapa, Ashish Acharya, Kritesh Rauniyar, Sweta Poudel, Sandesh Jain, Anum Masood, and Usman Naseem. 2024. Navigating the Web of Disinformation and Misinformation: Large Language Models as Double-Edged Swords. *IEEE Access* (2024).
  - [47] Francesca Spezzano. 2021. Modeling Misinformation Diffusion in Social Media: Beyond Network Properties. In *2021 IEEE Third International Conference on Cognitive Machine Intelligence (CogMI)*. 168–171. doi:10.1109/CogMI52975.2021.00030
  - [48] Giovanni Spitale, Nikola Biller-Andorno, and Federico Germani. 2023. AI model GPT-3 (dis) informs us better than humans. *Science Advances* 9, 26 (2023), eadh1850.
  - [49] Shane Tilton. 2019. Winning through deception: A pedagogical case study on using social deception games to teach small group communication theory. *SAGE Open* 9, 1 (2019), 2158244019834370.
  - [50] Luis Vargas, Patrick Emami, and Patrick Traynor. 2020. On the Detection of Disinformation Campaign Activity with Network Analysis. In *Proceedings of the 2020 ACM SIGSAC Conference on Cloud Computing Security Workshop* (Virtual Event, USA) (CCSW'20). Association for Computing Machinery, New York, NY, USA, 133–146. doi:10.1145/3411495.3421363
  - [51] Sidharth Vemuri, Jenny Hynson, Lynn Gillam, and Katrina Williams. 2020. Simulation-Based Research: A Scoping Review. *Qualitative Health Research* 30, 14 (2020), 2351–2360. doi:10.1177/1049732320946893
  - [52] Soroush Vosoughi, Deb Roy, and Sinan Aral. 2018. The spread of true and false news online. *Science* 359, 6380 (2018), 1146–1151. doi:10.1126/science.aap9559
  - [53] Claire Wardle and Hossein Derakhshan. 2017. Information disorder: Toward an interdisciplinary framework for research and policy making. Retrieved December 1, 2024 from <https://edoc.coe.int/en/media/7495-information-disorder-toward-an-interdisciplinary-framework-for-research-and-policy-making.html>
  - [54] William J White. 2018. Communication research and role-playing games. In *Role-Playing Game Studies*. Routledge, 337–345.
  - [55] Sidi Zhang and Bowen Jiang. 2021. The Sense of Community in the Era of Online Socials: A Case Study of Werewolf Killing Among Chinese Youth. In *2021 5th International Seminar on Education, Management and Social Sciences (ISEMSS 2021)*. Atlantis Press, 480–486.
  - [56] Xichen Zhang and Ali A. Ghorbani. 2020. An Overview of Online Fake News: Characterization, Detection, and Discussion. *Information Processing & Management* 57, 2 (March 2020), 102025. doi:10.1016/j.ipm.2019.03.004
  - [57] Jiawei Zhou, Yixuan Zhang, Qianni Luo, Andrea G Parker, and Munmun De Choudhury. 2023. Synthetic lies: Understanding ai-generated misinformation and evaluating algorithmic and human solutions. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*. 1–20.
  - [58] Terry Yue Zhuo, Yujin Huang, Chunyang Chen, and Zhenchang Xing. 2023. Red teaming chatgpt via jailbreaking: Bias, robustness, reliability and toxicity. *arXiv preprint arXiv:2301.12867* (2023).