

Voices Lost in the Algorithm: Improving Instagram’s Moderation of Mental Health Communities

Eunice Mok*

University of Illinois Urbana-Champaign
Urbana, IL, USA
eymok2@illinois.edu

Sriya Gottiparthi*

University of Illinois Urbana-Champaign
Urbana, IL, USA
sriyag3@illinois.edu

Yoshee Jain*

University of Illinois Urbana-Champaign
Urbana, IL, USA
yoshee2@illinois.edu

Eshwar Chandrasekharan

University of Illinois Urbana-Champaign
Urbana, IL, USA
eshwar@illinois.edu

Abstract

As a method of content moderation, Instagram limits access to hashtags associated with content that may violate Meta’s Community Standards by making them unsearchable. To evaluate the precision of this blanket banning strategy, we manually labeled a dataset of 500 Instagram posts containing banned hashtags. We found that their current approach achieves only 6% precision, thereby misclassifying 94% of posts as problematic when they did not violate Meta’s Community Standards. We then developed a context-based classifier that slightly outperforms this method, showing promise for more nuanced approaches to flag norm-violating posts using contextual clues. Additionally, topic modeling and qualitative analysis of the banned posts revealed that approximately 20% of the overall dataset contained prosocial or recovery-related content, highlighting the limitations of keyword-based moderation in accurately distinguishing harmful content from supportive discourse. This work calls for more context-aware moderation systems that can reduce wrongful removals and better support users engaging with sensitive topics.

CCS Concepts

• **Human-centered computing** → **Empirical studies in collaborative and social computing**.

Keywords

social media; content moderation; banning; community standards

ACM Reference Format:

Eunice Mok, Yoshee Jain, Sriya Gottiparthi, and Eshwar Chandrasekharan. 2025. Voices Lost in the Algorithm: Improving Instagram’s Moderation of Mental Health Communities. In *Companion of the Computer-Supported Cooperative Work and Social Computing (CSCW Companion ’25)*, October 18–22, 2025, Bergen, Norway. ACM, New York, NY, USA, 5 pages. <https://doi.org/10.1145/3715070.3749259>

*These authors contributed equally to the paper.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

CSCW Companion ’25, Bergen, Norway

© 2025 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-1480-1/25/10

<https://doi.org/10.1145/3715070.3749259>

1 Introduction

Instagram is a social media platform that primarily uses photo and video sharing to enable communication between more than two billion users worldwide [33]. When sharing content, Instagram allows users to use hashtags, descriptive words preceded by a #, to describe their posts [39]. These hashtags not only allow users to summarize their post’s content, but also allow others to perform focused searches for related content on the platform [1]. These hashtags cover diverse content, including serious topics such as mental health, suicide, and body image disorders [6, 7, 29].

Instagram’s content moderation strategy restricts the use of certain hashtags (such as, #depression, #suicide, #anorexia) that can contain content that violates Meta’s Community Standards [32]. Specifically, while users are able to use banned hashtags in their posts, they are not able to look up content using the banned tags in the search fields and are instead referred to a Help Site or other resources. We refer to these tags as “banned” tags. These hashtags are subject to periodic review and their “banned” status may be revised if the associated content shifts away from norm-violating behavior [32]. While “banned” hashtags like #depression may be linked to promoting content that violates Instagram’s Community Guidelines, previous work also finds that these hashtags may be used on these platforms to share awareness-related content and recovery stories [14]. Wrongful penalization of such content can silence mental health conversations online, leading users to leave social media platforms due to chilling effects, reduced satisfaction, diminished motivation to engage [4], and growing concerns about the fairness of moderation practices [22]. To prevent unfair penalization from keyword-based banning, a context-based approach may be more effective in differentiating between harmful content and prosocial content [2] related to sensitive topics [19].

In this work, we evaluate the effectiveness of Instagram’s hashtag-based content moderation approach by analyzing the type of content being penalized and compare its performance against context-based moderation. We also examine whether prosocial content is commonly associated with “banned” tags. We address the following research questions:

- (1) **RQ1:** How effective is Instagram’s hashtag-based blanket “banning” approach compared to a context-based moderation approach?
- **RQ1a:** What proportion of content that uses banned hashtags violates Meta’s Community Standards?

- RQ1b: How does a context-based moderation strategy perform compared to Instagram’s blanket “banning” approach?
- (2) RQ2: What types of content are being penalized by Instagram’s current blanket “banning” approach?

Our contributions to CSCW are threefold: we shed light on the efficacy of Instagram’s “keyword-based” bans, present an alternative approach that relies on context-based classification, and highlight the types of content posted under “banned” tags and wrongly sanctioned on Instagram.

2 Related Work

Multiple studies on social media platforms have raised concerns regarding exposure to abusive behavior [18] and content that violate community norms [11] and platform rules/policies [5, 34]. Exposure to norm-violating content can potentially affect the individuals’ psychological well-being, contributing to increased levels of depression and distress [15, 36]. To combat norm-violating content and undesirable behavior, platforms employ a range of content moderation strategies [16], including teams of human moderators (e.g., commercial content moderators [35] or volunteer community moderators [23, 27]), automated techniques (e.g., keyword-based bans [21]), and human-in-the-loop approaches (e.g., AI-based triaging [8, 17, 40] and tools to guide manual review [12, 28]). Instagram uses a combination of these methods [30, 31]. The consequence of posting content that violates a platform’s Terms of Service can vary depending on the platform and offense committed, from content removal/warnings [13] and permanent account bans [20], to the quarantine [9] or removal of an entire community [10].

Previous research has found that when communities focused around harmful content on Reddit were banned, many users from these communities left, and users who migrated to other subreddits reduced their use of hate speech [10]. In addition, banning prominent Twitter users who spread conspiracy theories reduced the spread of these theories [20]. On the other hand, quarantining controversial Reddit communities did not significantly reduce the amount of hate speech that their users posted [9], indicating that the efficacy of content-moderation strategies may be platform-specific. A survey measuring user-moderator alignment on Reddit found that users may agree with moderation policies as stated, but disagree with the way they are carried out in practice [25]. In addition, moderation standards that are perceived as unclear or unfair can lead to user self-censorship [24]. Thus, to maintain a safe online environment that fosters meaningful conversations and allows users to express themselves freely, social media platforms must strike a balance between under-moderation and over-moderation.

Although there have been numerous studies on the efficacy of various content moderation methods on other platforms, this study evaluates the performance of Instagram’s current moderation methods in ensuring accurate penalization of harmful posts. We also conduct topical analysis on moderated content to highlight the importance of safeguarding prosocial mental health conversations.

3 Comparing Instagram’s Hashtag-Based Banning with a Context-Based Classifier

3.1 Methodology

3.1.1 Collecting Banned Posts from Instagram. We constructed our initial dataset by manually selecting a seed set of 14 banned hashtags relating to mental health, including “deadinside”, “depressed”, “endme”, “fucklife”, “hatemyself”, “killme”, “killmyself”, “noose”, “selfharm”, “selfhate”, “selfinjury”, “suicide”, “traumacore”, and “wannadie”. We collected these hashtags by referencing online resources [37] and prior work [7] related to suicide, self-injury, and eating disorders. We searched for these hashtags using the Instagram app on a mobile device. If these hashtags were unsearchable, we considered them “banned.” Although these hashtags were unsearchable on mobile devices, they were still accessible via the Instagram Web page on desktop, so we collected 500 posts that used these banned hashtags ($\mathcal{D}_{\text{Banned}}$) using the Google Chrome web browser. We collected the URL, images, captions, and hashtags for all posts. To respect anonymity, we did not collect identifiable information about the users who posted the content at any step of the pipeline.

3.1.2 Evaluating Instagram’s Hashtag Banning via Manual Review. To evaluate the effectiveness of Instagram’s hashtag banning strategy in enforcing Community Standards, we manually annotated the posts in $\mathcal{D}_{\text{Banned}}$ to determine if they contain norm-violating content. We referenced Meta’s terms of service and community standards, used by Instagram, detailing the types of content that are prohibited. We focused on the category regarding “Suicide, self-injury, and eating disorders”,¹ which aligned with our seed set of hashtags. All authors read the documentation and agreed on guidelines for labeling posts as: (1) containing norm-violating content, (2) containing prosocial or recovery-related content, (3) does not contain either of those options, and (4) unsure.

We first subsampled 100 posts from $\mathcal{D}_{\text{Banned}}$. Three authors then individually annotated these posts based on the guidelines agreed upon earlier. The authors ensured that any disagreements were due to subjective differences. Using Krippendorff’s α to measure Inter-Rater Reliability (IRR) [26] between the annotators, we obtained $\alpha = 0.92$, denoting strong agreement. After this, two annotators each labeled 200 posts from $\mathcal{D}_{\text{Banned}}$, in total labeling 500 posts from $\mathcal{D}_{\text{Banned}}$. During annotations, we ensured that the hashtag in consideration was included in the post. Any posts that did not contain the relevant hashtag were removed from the dataset. These results shed light on the effectiveness of Instagram’s blanket hashtag-based banning approach and are reported in Section 3.2.1.

3.1.3 Building the Context-Based Classifier. We adopted a two-stage pipeline to build our context-based classifier. Using images from all collected posts, we prompted CogVLM, a state-of-the-art open source vision model [38], to “describe the image in great detail.” For each post, we concatenated the descriptions of all images into a string to form {image, description} pairs. To mitigate privacy concerns, we used an open-source model that we could run locally. The anonymized CogVLM descriptions of the 100 posts that three authors initially labeled were used to train a BERT model. Since our

¹<https://transparency.meta.com/en-gb/policies/community-standards/suicide-self-injury/>

dataset was highly imbalanced in favor of posts that did not violate Community Standards, the descriptions from posts we classified as “banned” were oversampled to balance the dataset.

This model was then used to classify the remaining 400 posts. We used these 400 predictions to assess the accuracy of our moderator by comparing its outputs to the ground truth labels assigned by human annotators, and contrasted its performance with Instagram’s current hashtag-based banning approach. We report our classifier’s accuracy, precision, and recall in Section 3.2.2.

3.2 Results

3.2.1 Effectiveness of Instagram’s Hashtag-Based Banning. Among the initial 100 posts labeled by three authors, only six were found to contain norm-violating content. In the remaining 400 posts, just 22 violated Instagram’s Community Standards. In total, only 28 out of 500 banned posts violated Instagram’s guidelines. Based on our manual review, Instagram’s current approach yields an accuracy of 6%, a precision of 6%, and a recall of 100%. Our finding highlights that while the approach successfully flags all violating content (high recall), it does so at the cost of extremely low precision, penalizing several posts that do not violate the platform’s community standards.

3.2.2 Comparing Context-Based Moderation to the Hashtag-Based Banning Strategy. Table 1 shows the performance of our context-based classifier. With training on only 100 posts, of which just six were truly norm-violating, we could slightly outperform Instagram’s moderation approach in terms of precision (7.1%). Despite the minimal number of positive examples, our classifier correctly identified a small but meaningful portion of banned content from the 400 posts in the test dataset, demonstrating the potential of context-aware moderation. However, we also observed a significant drop in recall (27.3%), which can be attributed to the minimal sample of norm-violating posts. A larger training set would enable the model to better generalize the features of banned content and reduce false negatives, thereby improving recall.

4 Analyzing the Types of Content Penalized by Instagram’s Blanket Banning

4.1 Methodology

4.1.1 Filtering Prosocial Posts from $\mathcal{D}_{\text{Banned}}$. During labeling our initial dataset, the raters labeled posts containing pro-social content, defined as pro-recovery or pro-mental health; 109 out of 500 annotated posts received this label. However, by the time we started content analysis, some of our posts were removed, made unavailable, or found to be duplicates. These posts were removed from the prosocial dataset, resulting in 100 prosocial posts.

4.1.2 Performing Topic Modeling on Post Descriptions to Extract Themes. To gain a deeper understanding of the types of content posted using the “banned” hashtags, we performed topic modeling on post descriptions using Latent Dirichlet Allocation (LDA) [3]. We combined the image descriptions generated by CogVLM with the corresponding post captions, removing stop words and excluding phrases such as “the image shows” to account for the structured output typical of CogVLM. We then constructed bigrams and trigrams, filtering out those that appeared in fewer than five documents or

in more than 50% of the corpus. Using this cleaned dataset, we generated seven topics where we restricted the model to include a maximum of 20 words in each topic. We analyze the keywords in these topics and formulate overarching themes in Section 4.2.1.

4.1.3 Analyzing the Content in Posts Manually to Extract Coherent Themes. Our findings from topic modeling indicated the presence of common prosocial and recovery-related content (discussed in Section 4.2.1); however, due to noise in the dataset indicated by keywords in Topics 1 and 6 (refer to Section 4.2.1), we performed a manual qualitative analysis to understand these topics better. One rater annotated the 100 prosocial posts, re-evaluating and writing a short description of each post on how it contributed to the mental health discussion. After fully annotating the dataset, the posts were categorized into five overarching themes reported in Section 4.2.2.

4.2 Results

4.2.1 Overarching Themes Obtained from Topic Modeling. Topic 1 was composed primarily of vague descriptive terms such as “woman” and “hair,” while Topic 6 encompassed terms that lacked a coherent thematic focus. As a result, both topics were labeled irrelevant or insufficient for our analysis and were excluded from further interpretation. In contrast, Topic 2 included terms such as “help,” “message,” and “know,” which are commonly associated with outreach efforts and awareness-raising discourse. Topic 3 contained terms such as “selfharm,” “awareness,” and “mental health,” and was accordingly interpreted as relating to discussions of self-harm or, more broadly, mental health awareness. Topic 4 contained vocabulary related to data visualization, including “infographic,” “bar,” and “options,” and was therefore categorized as informative. Topic 5 was characterized by words indicating support and guidance, including terms like “overcoming” and “understanding,” and was labeled as advice-oriented. Lastly, Topic 7 included introspective and personal language, such as “statement,” “thoughts,” and “share,” which led us to categorize it as anecdotal.

4.2.2 Categories of Prosocial Content from Qualitative Analysis. Drawing from the topical analysis and the manual annotations, the posts were categorized into five themes: anecdotal, informative, advice-oriented, inspirational, and directive. Posts were assigned to a category based on the type of content that was reflected in their descriptions. Posts containing factual information and infographics were categorized as informative. Posts featuring artistic narratives, poems, or motivational quotes intended to uplift users were labeled as inspirational. Posts were also assigned to multiple categories if they reflected characteristics of more than one theme.

5 Discussion

In summary, our findings indicate that Instagram’s current moderation strategy, which relies on a blanket hashtag-based banning mechanism, over-penalizes posts that do not violate Community Standards, resulting in notably low precision. Such moderation strategies impede users’ free speech and discourage sharing natural, healthy emotional content on the platform. Our manual labeling indicated that while some posts used banned hashtags in an exaggerated, humorous, or “edgy” context, they did not typically violate the platform’s community guidelines. Further analysis of

Table 1: Confusion matrix for Instagram’s hashtag-based banning strategy and the context-based classifier using manual annotations as ground truth. Instagram correctly classified 28 posts that contained norm-violating content, however, it inaccurately flagged 472 posts that were indeed safe with its blanket ban strategy achieving a low precision of 5.6%. With training on 100 posts, out of which only six had contained violating Meta’s Community Standards, the context-based classifier correctly flagged 6 norm-violating posts and 300 safe posts, while misclassifying 16 harmful posts as safe and 78 safe posts as harmful in the remaining 400 posts that were used as the test dataset. Aggregating these results, the classifier had a 7.1% precision and 27.3% recall.

	Instagram			Context-Based Classifier	
Ground Truth	Banned	Not Banned	Ground Truth	Banned	Not Banned
Banned	28	0	Banned	6	16
Not Banned	472	0	Not Banned	78	300
Precision: 5.6% Recall: 100%			Precision: 7.1% Recall: 27.3%		

Table 2: Overview of themes identified in a qualitative analysis of Instagram posts labeled as containing “prosocial content”.

Theme	Description
Anecdotal	People sharing their struggles with mental health, encouraging others or seeking advice
Informative	Education on mental illness, self-harm, or self-care based on facts or statistics
Advisory	Casual or personal advice based on experience or feelings
Inspirational	Quotes or stories meant to be encouraging; words of encouragement
Directive	Refers users to mental health resources

the content within posts containing “banned” hashtags revealed that these tags were frequently used in prosocial contexts: to spread mental health awareness, offer support, and encourage recovery. Thus, Instagram’s current strategy not only wrongly penalizes non-violating posts but also hinders pro-recovery and mental wellness accounts from reaching a wider audience. We acknowledge that these hashtags can indeed be associated with toxic content. We suggest adopting a careful, nuanced, context-based moderation strategy as it may substantially increase precision, allowing platforms to uphold community standards while avoiding wrongful penalization of prosocial content.

6 Limitations and Future Work

We note some limitations in our methodology that affected the scope and reliability of our findings. Our limited dataset only presented very few examples of posts that contained “banned” hashtags and violated Instagram’s community standards. Future work should explore more extensive data collection strategies to ensure the generalizability of findings. A larger dataset would further enable a more comprehensive evaluation of the performance of our context-based classifier, allowing it to better learn the characteristics of “banned” posts and produce more accurate classifications, leading to higher recall. Moreover, our study focused exclusively on content violations related to mental health. This category was only one of the 25 content categories present in Meta’s community

standards. Performing similar studies at larger scales across all categories would lead to a more robust and generalized understanding of the performance of Instagram’s current banning approach. In conclusion, our findings call for future work that examines the efficacy of current moderation approaches on platforms like Instagram and highlights the need to develop context-aware moderation systems that can reduce wrongful removals and better support users engaging with sensitive topics.

References

- [1] Morgan Ames and Mor Naaman. 2007. Why we tag: motivations for annotation in mobile and online media. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems* (San Jose, California, USA) (CHI '07). Association for Computing Machinery, New York, NY, USA, 971–980. <https://doi.org/10.1145/1240624.1240772>
- [2] Jiajun Bao, Junjie Wu, Yiming Zhang, Eshwar Chandrasekharan, and David Jurgens. 2021. Conversations Gone Alright: Quantifying and Predicting Prosocial Outcomes in Online Conversations. In *Proceedings of the Web Conference 2021* (WWW '21). ACM, 1134–1145. <https://doi.org/10.1145/3442381.3450122>
- [3] David M Blei, Andrew Y Ng, and Michael I Jordan. 2003. Latent dirichlet allocation. *Journal of machine Learning research* 3, Jan (2003), 993–1022.
- [4] Jie Cai, Aashka Patel, Azadeh Naderi, and Donghee Yvette Wohn. 2024. Content Moderation Justice and Fairness on Social Media: Comparisons Across Different Contexts and Platforms. In *Extended Abstracts of the CHI Conference on Human Factors in Computing Systems* (CHI '24). ACM, 1–9. <https://doi.org/10.1145/3613905.3650882>
- [5] Stevie Chancellor, Yannis Kalantidis, Jessica A Pater, Munmun De Choudhury, and David A Shamma. 2017. Multimodal classification of moderated online pro-eating disorder content. In *Proceedings of the 2017 CHI Conference on Human Factors in Computing Systems*. 3213–3226.

- [6] Stevie Chancellor, Zhiyuan Lin, Erica L Goodman, Stephanie Zerwas, and Munmun De Choudhury. 2016. Quantifying and predicting mental illness severity in online pro-eating disorder communities. In *Proceedings of the 19th ACM conference on computer-supported cooperative work & social computing*. 1171–1184.
- [7] Stevie Chancellor, Jessica Annette Pater, Trustin Clear, Eric Gilbert, and Munmun De Choudhury. 2016. #thyghgapp: Instagram Content Moderation and Lexical Variation in Pro-Eating Disorder Communities. In *Proceedings of the 19th ACM Conference on Computer-Supported Cooperative Work & Social Computing* (San Francisco, California, USA) (CSCW '16). Association for Computing Machinery, New York, NY, USA, 1201–1213. <https://doi.org/10.1145/2818048.2819963>
- [8] Eshwar Chandrasekharan, Chaitrali Gandhi, Matthew Wortley Mustelier, and Eric Gilbert. 2019. Crossmod: A cross-community learning-based system to assist reddit moderators. *Proceedings of the ACM on human-computer interaction* 3, CSCW (2019), 1–30.
- [9] Eshwar Chandrasekharan, Shagun Jhaver, Amy Bruckman, and Eric Gilbert. 2022. Quarantined! Examining the Effects of a Community-Wide Moderation Intervention on Reddit. *ACM Trans. Comput.-Hum. Interact.* 29, 4, Article 29 (March 2022), 26 pages. <https://doi.org/10.1145/3490499>
- [10] Eshwar Chandrasekharan, Umashanthi Pavalanathan, Anirudh Srinivasan, Adam Glynn, Jacob Eisenstein, and Eric Gilbert. 2017. You Can't Stay Here: The Efficacy of Reddit's 2015 Ban Examined Through Hate Speech. *Proc. ACM Hum.-Comput. Interact.* 1, CSCW, Article 31 (Dec. 2017), 22 pages. <https://doi.org/10.1145/3134666>
- [11] Eshwar Chandrasekharan, Mattia Samory, Shagun Jhaver, Hunter Charvat, Amy Bruckman, Cliff Lampe, Jacob Eisenstein, and Eric Gilbert. 2018. The Internet's hidden rules: An empirical study of Reddit norm violations at micro, meso, and macro scales. *Proceedings of the ACM on Human-Computer Interaction* 2, CSCW (2018), 1–25.
- [12] Frederick Choi, Tanvi Bajpai, Sowmya Pratipati, and Eshwar Chandrasekharan. 2023. ConvEx: A Visual Conversation Exploration System for Discord Moderators. *Proceedings of the ACM on Human-Computer Interaction* 7, CSCW2 (2023), 1–30.
- [13] Conor Clune and Emma McDaid. 2023. Content moderation on social media: constructing accountability in the digital space. *Accounting, Auditing & Accountability Journal* 37 (05 2023). <https://doi.org/10.1108/AAAJ-11-2022-6119>
- [14] Elizabeth Eikey. 2017. Recovery and Maintenance: How Women with Eating Disorders Use Instagram. In *iConference 2017 Proceedings*.
- [15] Krauss Melissa J, Costello Shaina J, Floyd Glennon M, Wilfley Denise E, & Cavazos-Rehg Patricia A. FitzsimmonsCraft, Ellen E. 2020. Adolescents and young adults engaged with pro-eating disorder social media: eating disorder and comorbid psychopathology, health care utilization, treatment barriers, and opinions on harnessing technology for treatment. *Eating and weight disorders : EWD* 25, 6 (Dec. 2020), 1681–1692. <https://doi.org/10.1007/s40519-019-00808-3>
- [16] Tarleton Gillespie. 2018. *Custodians of the Internet: Platforms, content moderation, and the hidden decisions that shape social media*. Yale University Press.
- [17] Agam Goyal, Xianyang Zhan, Yilun Chen, Koustuv Saha, and Eshwar Chandrasekharan. 2025. MoMoE: Mixture of Moderation Experts Framework for AI-Assisted Online Governance. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*.
- [18] Evey Jiaxin Huang, Abhraneel Sarma, Sohyeon Hwang, Eshwar Chandrasekharan, and Stevie Chancellor. 2024. Opportunities, tensions, and challenges in computational approaches to addressing online harassment. In *Proceedings of the 2024 ACM Designing Interactive Systems Conference*. 1483–1498.
- [19] Shagun Jhaver, Darren Scott Appling, Eric Gilbert, and Amy Bruckman. 2019. "Did You Suspect the Post Would be Removed?": Understanding User Reactions to Content Removals on Reddit. *Proc. ACM Hum.-Comput. Interact.* 3, CSCW, Article 192 (Nov. 2019), 33 pages. <https://doi.org/10.1145/3359294>
- [20] Shagun Jhaver, Christian Boylston, Diyi Yang, and Amy Bruckman. 2021. Evaluating the Effectiveness of Deplatforming as a Moderation Strategy on Twitter. *Proc. ACM Hum.-Comput. Interact.* 5, CSCW2, Article 381 (Oct. 2021), 30 pages. <https://doi.org/10.1145/3479525>
- [21] Shagun Jhaver, Quan Ze Chen, Detlef Knauss, and Amy X. Zhang. 2022. Designing Word Filter Tools for Creator-led Comment Moderation. In *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems* (New Orleans, LA, USA) (CHI '22). Association for Computing Machinery, New York, NY, USA, Article 205, 21 pages. <https://doi.org/10.1145/3491102.3517505>
- [22] Shagun Jhaver, Alice Qian Zhang, Quanze Chen, Nikhila Natarajan, Ruotong Wang, and Amy Zhang. 2023. Personalizing Content Moderation on Social Media: User Perspectives on Moderation Choices, Interface Design, and Labor. arXiv:2305.10374 [cs.HC] <https://arxiv.org/abs/2305.10374>
- [23] Charles Kiene, Kate Grandprey-Shores, Eshwar Chandrasekharan, Shagun Jhaver, Jialun" Aaron" Jiang, Brianna Dym, Joseph Seering, Sarah Gilbert, Kat Lo, Donghee Yvette Wohn, et al. 2019. Volunteer work: Mapping the future of moderation research. In *Companion Publication of the 2019 Conference on Computer Supported Cooperative Work and Social Computing*. 492–497.
- [24] Daniel Klug, Ella Steen, and Kathryn Yurechko. 2023. How Algorithm Awareness Impacts Algospeak Use on TikTok. In *Companion Proceedings of the ACM Web Conference 2023* (Austin, TX, USA) (WWW '23 Companion). Association for Computing Machinery, New York, NY, USA, 234–237. <https://doi.org/10.1145/3543873.3587355>
- [25] Vinay Koshy, Tanvi Bajpai, Eshwar Chandrasekharan, Hari Sundaram, and Karrie Karahalios. 2023. Measuring User-Moderator Alignment on r/ChangeMyView. *Proc. ACM Hum.-Comput. Interact.* 7, CSCW2, Article 286 (Oct. 2023), 36 pages. <https://doi.org/10.1145/3610077>
- [26] Klaus Krippendorff. 2018. *Content analysis: An introduction to its methodology*. Sage publications.
- [27] Charlotte Lambert, Frederick Choi, and Eshwar Chandrasekharan. 2024. "Positive reinforcement helps breed positive behavior": Moderator Perspectives on Encouraging Desirable Behavior. *Proceedings of the ACM on Human-Computer Interaction* 8, CSCW2 (2024), 1–33.
- [28] Yijun Liu, Frederick Choi, and Eshwar Chandrasekharan. 2025. Needling Through the Threads: A Visualization Tool for Navigating Threaded Online Discussions. *arXiv preprint arXiv:2506.11276* (2025).
- [29] Samantha B Mackson, Paula M Brochu, and Barry A Schneider. 2019. Instagram: Friend or foe? The application's association with psychological well-being. *New Media & Society* 21, 10 (2019), 2160–2182. <https://doi.org/10.1177/1461444819840021> arXiv:https://doi.org/10.1177/1461444819840021
- [30] Inc Meta Platforms. [n. d.]. *How Instagram uses artificial intelligence to moderate content*. Retrieved May 13, 2025 from <https://help.instagram.com/423837189385631>
- [31] Inc Meta Platforms. [n. d.]. *How to Report Things*. Retrieved May 13, 2025 from <https://help.instagram.com/2922067214679225>
- [32] Inc Meta Platforms. [n. d.]. *Terms of Use*. Retrieved January 25, 2025 from <https://help.instagram.com/581066165581870>
- [33] Oskar Mortensen. 2024. How Many Users on Instagram? Statistics & Facts (2025). (December 2024). <https://seo.ai/blog/how-many-users-on-instagram> Accessed: 2025-01-24.
- [34] Harita Reddy and Eshwar Chandrasekharan. 2023. Evolution of Rules in Reddit Communities. In *Companion Publication of the 2023 Conference on Computer Supported Cooperative Work and Social Computing*. 278–282.
- [35] Sarah T Roberts. 2016. Commercial content moderation: Digital laborers' dirty work. (2016).
- [36] Koustuv Saha, Eshwar Chandrasekharan, and Munmun De Choudhury. 2019. Prevalence and psychological effects of hateful speech in online college communities. In *Proceedings of the 10th ACM conference on web science*. 255–264.
- [37] Natasha Tracy. [n. d.]. *Banned Instagram Mental Health Tags for Bipolar, Depression and More*. Retrieved May 3, 2025 from <https://natashatracy.com/mental-illness-issues/banned-instagram-mental-health-hashtags-bipolar-depression/>
- [38] Weihang Wang, Qingsong Lv, Wenmeng Yu, Wenyi Hong, Ji Qi, Yan Wang, Junhui Ji, Zhuoyi Yang, Lei Zhao, Xixuan Song, Jiazhen Xu, Bin Xu, Juanzi Li, Yuxiao Dong, Ming Ding, and Jie Tang. 2023. CogVLM: Visual Expert for Pretrained Language Models. (2023). arXiv:2311.03079 [cs.CV]
- [39] Michele Zappavigna. 2015. Searchable talk: The linguistic functions of hashtags. *Social semiotics* 25, 3 (2015), 274–291.
- [40] Xianyang Zhan, Agam Goyal, Yilun Chen, Eshwar Chandrasekharan, and Koustuv Saha. 2025. SLM-Mod: Small Language Models Surpass LLMs at Content Moderation. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*. Association for Computational Linguistics, Albuquerque, New Mexico, 8774–8790. <https://doi.org/10.18653/v1/2025-naacl-long.441>