

Jan Simson — Department of Statistics, LMU Munich; Munich Center for Machine Learning (MCML)

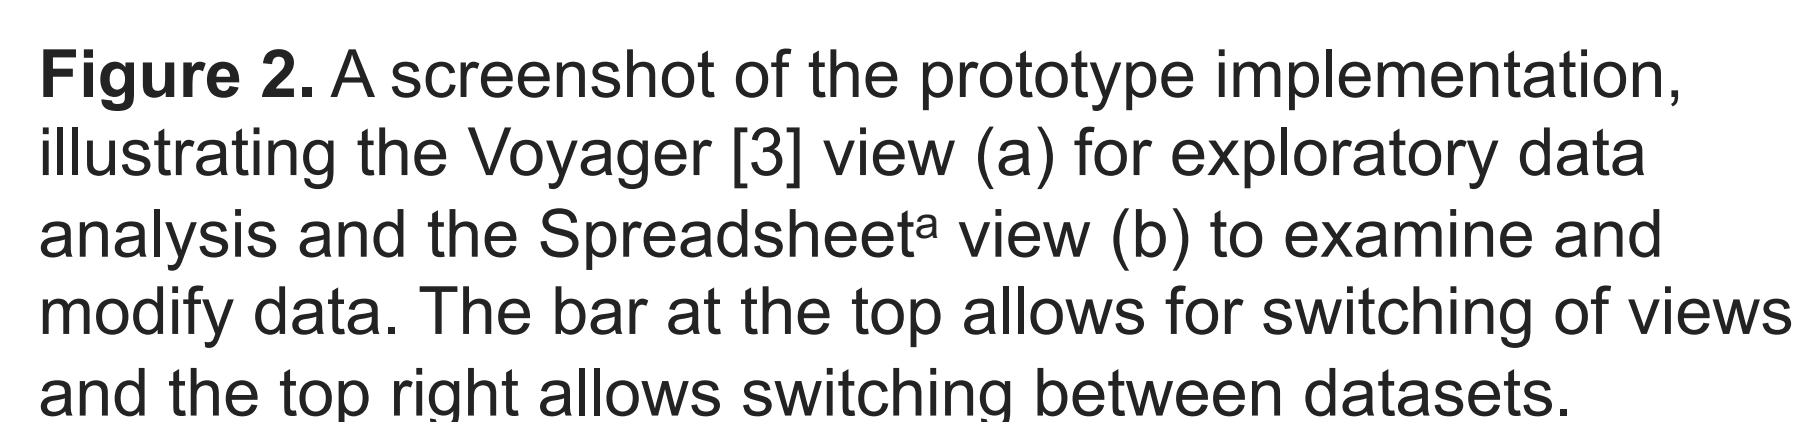
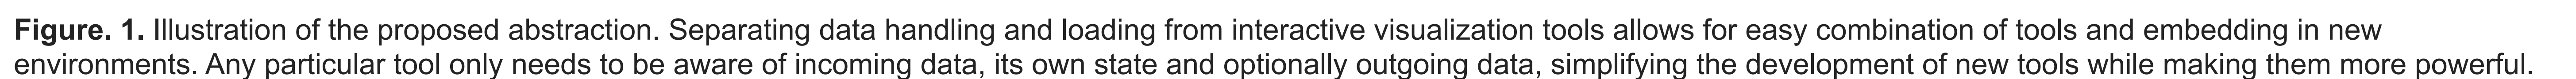
- There are many great interactive data visualization tools, often optimized for specific use cases
- Every tools needs to reimplement common steps such as data loading, processing and wrangling
- Data handling is usually not central to the design of a tool
- Different tools / visualizations often require data in specific formats, requiring external data wrangling
- This work proposes a new, modular approach separating *data handling* from *visualization* and allowing different tools to exist side-by-side

- Developed a prototype with different views: Data Loading, Spreadsheet^a, Data Wrangler^b, Sanddance [1], Voyager [3], Perspective^c, kepler.gl^d, plotly-react^e
- Relying on web technologies, allows for easy embedding

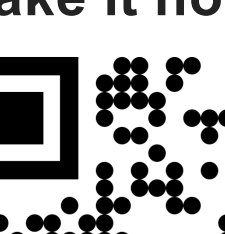
- Different build systems often require incompatible dependencies
→ separate builds; isolated iframes
- Different tools expect different data formats → internal processing layer and support for explicit processing

- Current prototype is exactly that:
A prototype (not yet finished)
- End-goal is to build a fully-featured version with plugin system for easy extendibility / addition of new tools
- Use duckdb [2] as a backend for flexible and fast data loading and wrangling + LLM based interaction
- Add integrations e.g. plugin for IDEs

- Better understand community needs and requirements
- See whether there is interest in a tool / framework like this



Take it home...



...or try it out

<https://data-studio.simson.io>

Anything else you
may want to add?



Write your answer on a post-it and add it onto (or next to the poster)