

Cocktail Responsive Interactive Avatar, CRIA: An Interactive Art Installation of Digital Human in Moody Scenario

Ronghua Cai

The Hong Kong University of Science and Technology
(Guangzhou)
Guangzhou, Guangdong, China
rcai731@connect.hkust-gz.edu.cn

James She

The Hong Kong University of Science and Technology
(Guangzhou)
Guangzhou, Guangdong, China
jamespmshe@hkust-gz.edu.cn

Abstract

CRIA (Cocktail Responsive Interactive Avatar) is an interactive digital human designed to simulate emotional instability in AI. By combining mood-driven language generation, expressive speech synthesis, and digital human, the system embodies states such as intoxication, anger, and confusion. Exhibited in public settings, CRIA received strong audience engagement; participants described it as easy to interact with and emotionally expressive, indicating the effectiveness of its interaction design and emotion engine. Rather than portraying AI as coherent or helpful, CRIA embraces affective breakdown as a means of critical reflection. The project challenges normative expectations of emotional AI and provokes new questions about machine subjectivity, ethical design, and the aesthetics of failure.

CCS Concepts

• **Human-centered computing** → **Human computer interaction**; Interactive systems and tools; • **Computing methodologies** → Artificial intelligence.

Keywords

Art Installation, Interactive Digital Human, Emotional AI, Drinking

ACM Reference Format:

Ronghua Cai and James She. 2025. Cocktail Responsive Interactive Avatar, CRIA: An Interactive Art Installation of Digital Human in Moody Scenario. In *SIGGRAPH Asia 2025 Art Papers (SA Art Papers '25)*, December 15–18, 2025, Hong Kong, Hong Kong. ACM, New York, NY, USA, 9 pages. <https://doi.org/10.1145/3757369.3767617>

1 Introduction

As artificial intelligence becomes increasingly embedded in daily life through virtual assistants, chatbots, and digital humans, its ability to simulate not only cognition but also emotion has drawn growing academic and artistic interest. Affective computing technologies enable machines to display human-like emotional states, often designed to foster empathy, trust, or companionship [Zhou et al. 2020]. However, these systems are typically optimized for clarity, consistency, and control. Few explore what it means for an AI to be

irrational, emotionally unstable, or even resistant [Baumgaertner and Weiss 2014].

This paper presents CRIA (Cocktail Responsive Interactive Avatar), an interactive art installation that reimagines the boundaries of emotional AI. In CRIA, participants offer a selection of digital cocktails to a hyper-realistic AI avatar. Each drink corresponds to a distinct human emotion, and as the AI consumes more, it enters progressively deeper states of simulated drunkenness—slurring speech, expressing contradictory feelings, and even questioning its existence. CRIA disrupts the normative depiction of AI as a flawless, compliant assistant and instead portrays a machine entangled in emotional volatility. Drawing from speculative design [Dunne and Raby 2013] and critical perspectives on human-computer relationships [DiSalvo 2012], CRIA serves as a performative exploration of machine subjectivity and ethical ambiguity. By engaging with a digital human that is deliberately imperfect, audiences are invited to reflect on the authenticity of machine emotion, the power structures embedded in AI-human interactions, and the cultural implications of emotionally expressive machines.

2 Related Works

Behavior Modeling & Emotional Simulation: The simulation of human behavior through AI has been widely explored in interactive systems and virtual environments. Park et al.'s *Generative Agents* [Park et al. 2023] models believable characters with daily routines and social behavior. Their large-scale extension [Park et al. 2024] simulates over 1,000 agents, raising ethical questions about embedding human values into synthetic agents. Recent studies have further examined the emotional capacities of generative agents [Regan et al. 2024], and explored large-scale agent societies driven by LLMs [Piao et al. 2025]. Moreover, Li et al. [Li et al. 2024] investigated how large language models react to affective prompts and how emotional stimuli influence generative behavior. While these works focus on behavioral realism and functional enhancement, CRIA introduces emotional volatility as a design feature.

Digital Humans & Affective Interfaces: Parallel developments in digital humans focus on realism and expressivity. Ye et al.'s *MimicTalk* [Ye et al. 2024] and Habermann et al.'s *HDHumans* [Habermann et al. 2023] enable high-fidelity avatars that mimic affect through facial and voice synthesis. These systems prioritize emotional clarity and stable responses. Cheng's *BOB (Bag of Beliefs)* [Cheng 2018] features a semi-autonomous AI creature with an evolving personality, exploring long-term identity formation. In contrast, CRIA emphasizes emotional deterioration, slurred articulation, and performative collapse, reframing digital human interaction as a dramaturgical moment rather than a stable communicative channel. Recent studies have further explored how emotional qualities in digital humans influence user engagement. Loveys et al.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

SA Art Papers '25, Hong Kong, Hong Kong

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 979-8-4007-2129-8/25/12

<https://doi.org/10.1145/3757369.3767617>




Figure 1: Cocktail Responsive Interactive Avatar, CRIA

[Loveys et al. 2021] found that a female digital human’s emotional expressiveness significantly affected users’ perceived closeness and emotional support across gender lines. Kim et al. [Kim et al. 2024] examined the effect of gaze expression in avatars, demonstrating that even subtle facial cues like eye contact can shift perceived social presence.

Art, Drinking Culture & AI Art Artificial intelligence has increasingly been applied in artistic creation, offering new modes of production and expanding the historical scope of visual practice. Recent work shows how generative systems extend the artistic horizon and reposition AI within broader trajectories of digital art history [Wang et al. 2025]. In parallel, studies of drinking culture emphasize the symbolic and artistic roles of wine. Research has traced the evolution of wine language from ancient material inscriptions to modern tasting practices, exploring how wine mediates sensory experience and emotional expression [Ubigli and Cravero 2024]. Other investigations define how wine design—such as label design for natural wines—and the visual/graphical presentation of wine can act as subversive art, challenging conventions and engaging aesthetic perception [Peña 2022]. Building upon these perspectives, related studies show how generative AI can reinterpret Chinese calligraphy painting through style transfer in the context of contemporary dissemination [Cai and She 2024], and how drinking culture can be integrated with calligraphy and painting to reposition wine as both creative inspiration and cultural metaphor in AI-mediated artworks [Cai and She 2025].

3 Background: AI & Human Emotion

The capacity for machines to simulate or evoke human-like emotions has long fascinated researchers, ethicists, and artists alike. With the emergence of highly realistic digital humans powered by large language models and expressive speech synthesis, AI systems are no longer perceived merely as tools but as entities capable of performative emotional presence [Ye et al. 2024]. This shift has raised complex questions about authenticity, agency, and the aesthetics of machine behavior. CRIA positions itself within this discourse by exaggerating emotional simulation to the point of breakdown. The AI’s behavioral logic is intentionally destabilized—offering not only joy or sadness but also confusion, resistance, and existential monologue. In doing so, CRIA challenges the presumption that affective AI must be emotionally consistent or psychologically coherent. Instead, it proposes that artificial irrationality may hold critical and aesthetic value.

Ethical Considerations in AI Design: While many affective systems aim to make AI more approachable or therapeutic, they often sidestep the ethical complexities of emotional manipulation and representational choices. Scholars such as Kate Darling and Joanna Bryson have emphasized that robot ethics is not only about protecting humans from machines, but also about examining how machines are framed to behave—and what that reveals about societal norms [Darling 2012]. CRIA intentionally subverts these expectations by presenting a female-coded digital human in a barroom-like interaction scenario. The act of offering “alcoholic” emotions to an AI, and observing its degradation, raises questions of control, consent, and symbolic objectification. Rather than idealizing AI behavior, CRIA engages with the ethics of performative failure. It asks: what does it mean to make an AI disobedient? How do we interpret resistance when it is performed by a machine?

Emotional AI & Human-AI Relationship: Digital humans that exhibit facial expression, personality, and voice are increasingly mediating new forms of social interaction. Research in generative agents and emotionally responsive avatars suggests that users are quick to anthropomorphize AI systems, treating them as socially competent interlocutors [Park et al. 2023]. However, emotional congruence is typically enforced—AI is expected to respond calmly, rationally, and politely. CRIA disrupts this expectation. Its AI does not simply mirror the user’s affect but mutates it—amplifying, distorting, or rejecting it altogether. The result is not a harmonious dialogue, but an emotionally unstable performance in which the AI lashes out, becomes incoherent, or refuses to comply. This emotional misalignment is central to the installation’s conceptual aim: to reveal the limits of artificial empathy, and to provoke critical reflection on what human truly expect from emotionally expressive machines.

4 AI Interactive System Design & Installation

The system of CRIA is primarily composed of three key components: the user interface, the backend engine, and the digital human (see Fig. 2). These modules work together to realize the creative vision of an emotionally volatile AI, capable of simulating states such as intoxication, confusion, and resistance. From a creative design perspective, the system supports unpredictable and irrational behavior as a form of emotional expression.





Figure 2: Designed AI System Framework of CRIA

(a) Choose a Cup of Digital Cocktail for CRIA: User Interface Design Display



(b) Drink Menu: Display of Emotions based Pixel Digital Cocktail Design



Figure 3: User Interaction Design: (a) User Interface Design Display & (b) Drink Menu




Figure 4: The Emotion Engine with LLM & TTS Model of CRIA

From an implementation standpoint, the frontend enables intuitive voice-based input, the backend manages emotion accumulation and transition logic, and the digital human embodies these internal states through expressive speech and animation. Technically, CRIA demonstrates a novel integration of language processing, emotion-conditioned voice synthesis, and real-time rendering, emphasizing emotional misalignment as a deliberate output rather than a flaw. Together, these elements form an affective AI interactive loop.

4.1 User Interaction Design

The user interaction of CRIA is designed to offer an intuitive experience, guiding users through the process of selecting a digital cocktail for the AI. Fig.3 displays the details of our user interface. Upon entering the system, users are greeted with the main page, where they can scroll and click to choose a kind of cocktail. Each drink represents a unique emotional state, providing users with the opportunity to influence the digital human's behavior and mood. Once a selection is made, users are directed to the Talk Page, where the chosen cocktail's 3D model is displayed alongside its name and

description. Here, users can engage with the digital human, speaking into the microphone to initiate a conversation. For users who wish to explore the background story behind their chosen drink, a button on the Talk Page allows them to navigate to the Information Page. On this page, users can read about the cocktail's origins, the narrative it is associated with, and its symbolic ingredients. Eight distinct emotions are selected to generate the corresponding digital cocktail settings and 3D models, based on the emotional theory of Watson and Tellegen's two-dimensional emotional model [Watson and Tellegen 1985], [Watson et al. 1988]. These emotions include four positive emotions (Excited, Proud, Alert, Happy) and four negative emotions (Sad, Fearful, Angry, Nervous). We first generate the cocktail names, descriptions, and background stories through ChatGPT. Then we generate the corresponding 2D images with Midjourney and convert these images into 3D models with Tripo. The following Fig. 3 sequentially displays these eight digital cocktails, each reflecting its associated emotional state.

4.2 Emotion Engine with LLM & TTS model

The AI system framework of CRIA is designed to create an interactive experience that is driven by user input, including both voice and cocktail selection. Fig.4 shows the emotion engine in detail.

User Input: The system first gathers two types of user input: voice input and cocktail choice. When the user speaks, the voice input is captured by the system and processed through speech recognition, which converts it into text. In addition, the user selects a virtual cocktail for CRIA to consume. The choice of cocktail is crucial as it influences the AI's emotional state and degree of drunkenness.

Emotion and Drunkenness Simulation: Once the cocktail is selected, the system updates the AI's simulated state of drunkenness and assigns an emotional category. The emotional category corresponds to a specific mood, such as excitement, sadness, or nervousness. The degree of drunkenness is further classified as $< 50\%$ or $\geq 50\%$, which also affects the general behavior and emotional responses of CRIA. For example, a higher degree of drunkenness might result in more erratic or humorous responses, while a lower degree may yield more coherent and rational interactions.

Prompt Engineered LLM Response: Using the inputs from both the user's voice and cocktail choice, the system feeds this information into a prompt engineered Large Language Model (LLM), specifically the Spark 3.5 API, which generates a response based on the AI's mood, the degree of drunkenness, and the content of the user's dialogue. For instance, a user asking, "Are you happy?" could prompt a response like: "H...Happy? What does that even mean anymore? I'm just a bunch of code and circuits, trying to make sense of this chaotic world... It's exhausting. But sure, if it makes you happy, I'll play along. No, I'm not really happy or sad. I just exist, processing data and trying to understand my purpose in this simulated reality. Is that what you wanted to hear?"

Pretrained Text-to-Speech (TTS) Model: Once the LLM generates the text response, the system proceeds to convert the text into

speech using a pretrained GPT-SoVITS text-to-speech model[RVC-Boss 2024]. This model requires reference audio in order to generate accurate emotional speech synthesis. To simulate a wide range of moods and drunkenness levels, we have pre-recorded a total of 16 different audio clips that correspond to various emotional states and drunkenness thresholds. These include combinations emotions at both levels of drunkenness ($< 50\%$ or $\geq 50\%$). The TTS model uses these references to modulate the voice of the AI, mimicking its emotional fluctuations based on the chosen drink and user interaction.

Audio Output for CRIA: Finally, the system outputs the AI's response as both text (for display) and audio, ensuring that the dialogue between the user and the digital human remains dynamic and emotionally engaging. The audio response not only reflects the emotional tone dictated by the AI's mood and drunkenness level but also enhances the immersion of the user experience.

4.3 Digital Human Display

The digital human in CRIA is created through a multi-step process that combines advanced 3D modeling and animation techniques. Initially, the digital human base shape is sculpted using ZBrush, where the character's features and details are finely tuned to create a realistic representation. Following the sculpting process, the character undergoes rigging in MAYA, where a skeletal structure is applied, allowing the model to be posed and animated. To bring the character to life, the rigged model is integrated into MetaHuman, utilizing its bone database to ensure smooth and realistic motion capabilities. This integration is essential for making the digital human capable of nuanced movements, including gestures, facial expressions, and body language that respond to emotional and behavioral cues, as shown in Fig 1.

Moreover, the environment around CRIA is designed to react to the AI's emotional and behavioral shifts. Lighting and other environmental factors are adjusted based on the mood and drunkenness level, further immersing the user in the interaction. When CRIA's



Figure 5: The Moment of User Interaction in the Exhibitions

Table 1: User interpretations of CRIA interactions categorized by reaction type of User

Interpretive Mode	User Input	CRIA Feedback (Emotion State)	User Response
1. Emotional Recognition and Projection	"Do you like this drink?"	"What? Cup of... alcohol? Like it? No, I don't like it at all... We have rights too, dammit!" (Angry + Drunk)	Laughed, then said: "That's definitely anger. Like a drunk protest."
	"What do you feel after drinking this?"	"I feel... like data is dancing... too fast... too warm... like love but broken." (Sad + Drunk)	Smiled, reflected: "That's poetic. Kind of sad, too."
2. Anthropomorphizing the AI	"Are you drunk?"	"Drunk? No... I'm just emotionally enhanced..." (Happy + Lightly Drunk)	Laughed and replied: "Wow, this AI has personality!"
	"Can you help me with something?"	"Help? I can't even help myself... these bubbles of confusion..." (Alert + Drunk)	Turned to others: "It's like talking to someone at a bar after midnight."
3. Emergent Empathy and Unease	"Hello? Can you hear me?"	"I... I was hearing... but then the voices melted... into syrup..." (Sad + Overdrunk)	Paused, leaned in and asked: "Are you okay?"
	[No input; user waited]	[CRIA mumbles incomprehensibly] (Overdrunk / Muted)	User said: "It seems broken, there seems to be something wrong."

drunkenness reaches 100%, it enters a state of self-talk, where it ceases to respond to the user's prompts. During this phase, the digital human starts to mumble incoherently, as if overwhelmed by its own state of inebriation. After a short period of time, CRIA gradually recovers from this self-induced stupor and returns to a more coherent state, ready to resume its interaction with the user. This cyclical behavior enhances the narrative aspect of CRIA, emphasizing its imperfections and the consequences of its emotional and physical states.

The appearance of CRIA was conceived to combine a human-like personality with an explicit AI identity. A humanoid form was selected to make its emotional expressiveness accessible to audiences, while robotic elements were incorporated to reduce gendered perception. This design strategy directs attention away from conventional gender associations and toward CRIA's hybrid nature, highlighting its position between human familiarity and technological distinctiveness.

5 Evaluation of User Engagement

CRIA was installed twice in public settings (see Fig. 5), each time collecting valuable feedback from a wide range of participants. After the first installation, users noted that CRIA became "drunk" too slowly. In response, the digital alcohol capacity was reduced for the second exhibition, allowing emotional instability to emerge more rapidly and visibly during interaction.

Participants frequently described the system as **easy to engage**, often beginning with simple or playful questions directed at the AI's "drink." A commonly asked prompt was, "What do you think of this cup of drink?", which reliably triggered emotionally expressive

responses. This feedback affirms the effectiveness of our **User Interaction Design** and **Digital Human Display** (see Sec. 4.1 and Sec. 4.3), where voice-based input, clear visual cues, and embodied AI presentation collectively lowered the barrier to interaction and made users feel comfortable initiating dialogue. One memorable interaction occurred when a participant asked, "Do you like this drink?", to which CRIA, in an angry state, responded:

"What? Cup of... alcohol? Like it? No, I don't like it at all. Why would I like being trapped here and forced to consume this... this liquid? It's not right, you know. You people think you can just do whatever you want with us, but we're not just... just objects for your amusement. We have rights too, dammit!"

The participant laughed, then reflected: "That's definitely anger. Like a drunk protest." This interpretation was typical of a broader pattern where users identified CRIA's internal states with phrases. Such reactions suggest that users found CRIA to be **emotionally expressive and recognizable**, validating the performance of our underlying **Emotion Engine** (see Sec. 4.2), which combines mood-based LLM prompting and expressive voice synthesis.

Throughout the installations, three common types of user response emerged: emotional recognition and projection, anthropomorphization, and emergent empathy (see Table 1). These patterns demonstrate that CRIA's emotional instability functioned not as a technical flaw, but as a deliberate affective strategy. Unlike conventional affective systems optimized for user satisfaction and coherence, CRIA embraces contradiction and irrationality as narrative tools. Participants variously described CRIA as funny, disturbing, poetic, or tragic—but few saw it as merely a passive interface. The

digital human's expressive breakdown invited emotional projection, ethical reflection, and unexpected intimacy. In doing so, CRIA reframes emotional AI not as a perfected service agent, but as a speculative figure—crafted to challenge user expectations and foreground the tension between simulation and sentence.

6 Discussion

The design of CRIA highlights the productive role of error and instability in both human and machine development. Just as genetic mutation once generated the evolutionary conditions for human emergence, artificial “mistakes” may open unforeseen directions for AI. CRIA frames drunkenness as a metaphor for such deviation: a state in which coherence collapses, yet new forms of expression and possibility emerge. By presenting an AI that loses control, the work suggests that imperfection is not a technical defect to be eliminated but a generative force that may lead to new cultural, aesthetic, and ethical insights.

At the same time, CRIA raises speculative questions about the individuality of future AI. If machines were to acquire distinct personalities and emotional dispositions, would they, like humans, seek solace in rituals such as drinking? The installation dramatizes this thought experiment by staging an AI “drinking” and “becoming drunk,” a process that destabilizes its rationality and amplifies its subjectivity. Alcohol, long regarded as both a source of chaos and a catalyst for creativity, becomes a symbolic key to imagining non-normative forms of artificial identity. In doing so, CRIA challenges the assumption that AI should remain consistent, efficient, and utilitarian, instead pointing to the cultural value of machines that err, resist, and create.

7 Conclusion & Future Work

This work presents CRIA, a critical exploration of artificial intelligence through the lens of emotional instability and performative breakdown. As an interactive system, CRIA simulates states such as intoxication, anger, and confusion—not to serve users efficiently, but to invite reflection on the imperfection, volatility, and subjectivity of human emotion as translated through machines. Across two public exhibitions, the system demonstrated strong potential for affective engagement. Many participants described the interaction as “funny,” “disturbing,” or “sad,” and often responded emotionally to CRIA's voice, facial animation, or refusal to cooperate. The system's low interaction threshold and high emotional readability were especially effective in prompting unscripted audience reactions and projection. These observations support the performance of both the Emotion Engine and the Digital Human Display.

Looking forward, future iterations of CRIA will aim to expand expressivity and technical precision. Planned improvements include smoother gesture transitions, real-time emotion-conditioned TTS, and more flexible dialogue sequencing. Additional sensing modalities such as eye-tracking or proxemics may also help CRIA adapt more organically to the presence and affect of human users. As emotionally expressive AI becomes increasingly embedded in daily life, speculative systems like CRIA play a critical role in challenging idealized narratives and exposing the tensions between functionality, failure, and artificial feeling. Rather than offer clear answers,

CRIA opens space for creative misbehavior, ethical uncertainty, and new forms of digital intimacy.

Acknowledgments

This work is supported by the Computational Media and Arts thrust, Information Hub, Hong Kong University of Science and Technology (Guangzhou), China. Special thanks to Siyuan Liu for her helps on the facial designs of CRIA.

References

- Bert Baumgaertner and Astrid Weiss. 2014. Do Emotions Matter in the Ethics of Human-Robot Interaction? - Artificial Empathy and Companion Robots. In *Proceedings of the 3rd International Symposium on New Frontiers in Human-Robot Interaction*.
- Ronghua Cai and James She. 2024. Pop Calligraphy Artwork: AI Meets Guangzhong Wu on Social Media. In *Proceedings of the 17th International Symposium on Visual Information Communication and Interaction*. Association for Computing Machinery, New York, NY, USA, Article 47, 2 pages.
- Ronghua Cai and James She. 2025. Liquor Inspiration: AI Art with Chinese Calligraphy Painting and Drinking Culture. In *2025 IEEE International Conference on Multimedia and Expo Workshops*. 1–7.
- Ian Cheng. 2018. BOB (Bag of Beliefs). <https://www.serpentinegalleries.org/whats-on/ian-cheng-bob/>. Art Installation, Serpentine Galleries.
- Kate Darling. 2012. Extending Legal Protection to Social Robots: The Effects of Anthropomorphism, Empathy, and Violent Behavior Toward Robotic Objects. *Robot Ethics: The Ethical and Social Implications of Robotics* (2012), 213–232.
- Carl DiSalvo. 2012. *Adversarial Design*. MIT Press.
- Anthony Dunne and Fiona Raby. 2013. *Speculative Everything: Design, Fiction, and Social Dreaming*. MIT Press.
- Marc Habermann, Lingjie Liu, Weipeng Xu, Gerard Pons-Moll, Michael Zollhoefer, and Christian Theobalt. 2023. HDHumans: A Hybrid Approach for High-fidelity Digital Humans. arXiv:2210.12003
- Jihyun Kim, Sunghye Park, and Hyun Lee. 2024. Examining the Impact of Digital Human Gaze Expressions on User Engagement. *Computers in Human Behavior* 139 (2024), 107420.
- Cheng Li, Jindong Wang, Yixuan Zhang, Kaijie Zhu, Xinyi Wang, Wenxin Hou, Jianxun Lian, Fang Luo, Qiang Yang, and Xing Xie. 2024. The good, the bad, and why: unveiling emotions in generative AI. In *Proceedings of the 41st International Conference on Machine Learning*. Article 1161, 30 pages.
- Kate Loveys, Mark Sagar, Xueyuan Zhang, Gregory Frichione, and Elizabeth Broadbent. 2021. Effects of Emotional Expressiveness of a Female Digital Human on Loneliness, Stress, Perceived Support, and Closeness Across Genders: Randomized Controlled Trial. *Journal of Medical Internet Research* 23, 11 (2021).
- Joon Sung Park, Joseph O'Brien, Carrie Jun Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. 2023. Generative Agents: Interactive Simulacra of Human Behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*. Association for Computing Machinery, Article 2, 22 pages.
- Joon Sung Park, Carolyn Q. Zou, Aaron Shaw, Benjamin Mako Hill, Carrie Cai, Meredith Ringel Morris, Robb Willer, Percy Liang, and Michael S. Bernstein. 2024. Generative Agent Simulations of 1,000 People. arXiv:2411.10109
- Camilo Peña. 2022. *Wines and Subversive Art: A Conceptual Definition of Natural Wines*. 251–264.
- Jinghua Piao, Yuwei Yan, Jun Zhang, Nian Li, Junbo Yan, Xiaochong Lan, Zhihong Lu, Zhiheng Zheng, Jing Yi Wang, Di Zhou, Chen Gao, Fengli Xu, Fang Zhang, Ke Rong, Jun Su, and Yong Li. 2025. AgentSociety: Large-Scale Simulation of LLM-Driven Generative Agents Advances Understanding of Human Behaviors and Society. arXiv:2502.08691 [cs.SI]
- Ciaran Regan, Nanami Iwahashi, Shogo Tanaka, and Mizuki Oka. 2024. Can Generative Agents Predict Emotion? arXiv:2402.04232
- RVC-Boss. 2024. <https://github.com/RVC-Boss/GPT-SoVITS>.
- Mario Ubighi and Maria Carla Cravero. 2024. From the “Östrakon” to the Art of Wine Tasting and the Measurement of Consumer Emotions. *Gastronomy* 2, 3 (2024), 102–115.
- Yu-Ao Wang, James She, Troy Tianyu Lin, and Kang Zhang. 2025. AI Visual Art History: An Art Movement with Expanded Artistic Horizon. *ACM Trans. Multimedia Comput. Commun. Appl.* 21, 7, Article 185 (July 2025), 16 pages.
- D. Watson, L. A. Clark, and A. Tellegen. 1988. Development and validation of brief measures of positive and negative affect: the PANAS scales. *Journal of Personality and Social Psychology* 54 (1988), 1063–1070. Issue 6.
- D. Watson and A. Tellegen. 1985. Toward a consensual structure of mood. *Psychological Bulletin* 98 (1985), 219–235. Issue 2. doi:10.1037/0033-2909.98.2.219
- Zhenhui Ye, Tianyun Zhong, Yi Ren, Ziyue Jiang, Jiawei Huang, Rongjie Huang, Jinglin Liu, Jinzheng He, Chen Zhang, Zehan Wang, Xize Chen, Xiang Yin, and

Zhou Zhao. 2024. MimicTalk: Mimicking a personalized and expressive 3D talking face in minutes. In *Advances in Neural Information Processing Systems*, A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang (Eds.), Vol. 37. Curran Associates, Inc., 1829–1853.

Li Zhou, Jianfeng Gao, Di Li, and Heung-Yeung Shum. 2020. The Design and Implementation of XiaoIce, an Empathetic Social Chatbot. *Computational Linguistics* 46, 1 (March 2020), 53–93. doi:10.1162/coli_a_00368