

Memory-Augmented LLMs for Sustainable Urban Energy Management via Weather-Energy Pattern Learning

Nur Arifin Akbar
nurarifin.akbar@unipa.it
Universita degli Studi di Palermo
Palermo, Palermo, Italy

Shahin Rezvanian
University of St. Gallen
St. Gallen, Switzerland

Rahool Dembani
SingularLogic
Athens, Greece

Abstract

This paper presents a novel dual-memory Large Language Model (LLM) architecture for sustainable urban energy management. Inspired by hippocampal-neocortical interactions, our approach integrates fast learning with long-term memory consolidation to capture weather-energy correlations. The LLM processes meteorological data through natural language interfaces, enabling semantic understanding of weather patterns. Evaluated across European cities using Open-Meteo and ENTSO-E data, our dual-memory LLM achieves 60-85% MAPE improvements over LSTM baselines (2.02-6.13% vs 10.46-17.98%), with cost savings averaging €152M per city.

CCS Concepts

• **Computing methodologies** → **Machine learning approaches**; *Neural networks*; • **Applied computing** → **Environmental sciences**.

Keywords

Dual-Memory LLM, Energy Forecasting, Sustainable Urban Energy, Weather-Energy Correlation, Sand Battery Storage, Renewable Energy Integration

ACM Reference Format:

Nur Arifin Akbar, Shahin Rezvanian, and Rahool Dembani. 2025. Memory-Augmented LLMs for Sustainable Urban Energy Management via Weather-Energy Pattern Learning. In *Companion of the the 2025 ACM International Joint Conference on Pervasive and Ubiquitous Computing (UbiComp Companion '25)*, October 12–16, 2025, Espoo, Finland. ACM, New York, NY, USA, 4 pages. <https://doi.org/10.1145/3714394.3750593>

1 Introduction

European electricity markets face unprecedented challenges with volatile pricing and reactive energy management systems that fail to anticipate demand fluctuations or leverage weather-energy correlations [16, 21]. Traditional forecasting approaches using historical averages cannot capture non-linear relationships between meteorological conditions and energy consumption.

Large Language Models (LLMs) and memory-augmented architectures offer new opportunities for energy system management

[2, 3]. This work introduces a dual-memory LLM architecture mimicking hippocampal-neocortical interactions for rapid adaptation and long-term retention of seasonal behaviors.

Our system integrates Open-Meteo weather data with ENTSO-E grid information, incorporating sand-battery thermal storage coordination to optimize renewable utilization [7, 22, 23].

2 Related Work

Traditional energy forecasting relies on ARIMA and classical ML approaches that struggle with renewable volatility [12, 13]. LSTM networks improve non-linear dependency capture but suffer from catastrophic forgetting [11, 17].

Memory-augmented networks draw from neuroscience complementary learning systems theory [18]. Our LLM engine provides semantic processing of weather descriptions, pattern recognition in meteorological reports, and decision explainability through natural language understanding. This captures subtle weather nuances that numerical-only approaches miss.

3 Methodology

3.1 Dual-Memory Architecture Design

Our dual-memory LLM architecture consists of two complementary subsystems inspired by hippocampal-neocortical interactions (Figure 1):

Hippocampal Component (HC): Implements rapid learning for immediate pattern recognition and adaptation. The HC processes incoming weather and energy data streams using a fast-learning mechanism:

$$h_{HC}^t = \sigma(W_h x_t + U_h h_{HC}^{t-1} + V_h c_t + b_h) \quad (1)$$

where x_t represents input features, h_{HC}^{t-1} is the previous hidden state, c_t is the context vector, and σ is the activation function.

Neocortical Component (NC): Provides long-term memory consolidation and stable pattern storage. The NC updates are governed by:

$$\theta_{NC}^{t+1} = \theta_{NC}^t - \alpha \nabla_{\theta} L_{\text{replay}}(D_{\text{replay}}, \theta_{NC}^t) \quad (2)$$

where D_{replay} contains previously learned patterns and α is the learning rate.

3.2 LLM-Driven Analysis

The LLM uses attention mechanisms for weather-energy correlation discovery:

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V \quad (3)$$



This work is licensed under a Creative Commons Attribution 4.0 International License. *UbiComp Companion '25, Espoo, Finland*
© 2025 Copyright held by the owner/author(s).
ACM ISBN 979-8-4007-1477-1/2025/10
<https://doi.org/10.1145/3714394.3750593>

3.3 Memory Consolidation with Elastic Weight Consolidation

To prevent catastrophic forgetting while learning new seasonal patterns, we implement Elastic Weight Consolidation (EWC):

$$L_{\text{total}} = L_{\text{current}} + \lambda \sum_i F_i (\theta_i - \theta_i^*)^2 \quad (4)$$

where F_i represents the Fisher information matrix, θ_i^* are the optimal parameters for previous tasks, and λ controls the importance of maintaining previous knowledge.

3.4 Sand Battery Thermal Dynamics Integration

The system models sand battery storage dynamics using heat transfer equations:

$$\rho c_p \frac{\partial T}{\partial t} = \nabla \cdot (k \nabla T) + Q_{\text{in}} - Q_{\text{out}} \quad (5)$$

where ρ is density, c_p is specific heat capacity, k is thermal conductivity, and $Q_{\text{in/out}}$ represent heat input/output rates.

4 Experimental Setup

Data: Open-Meteo ERA5 weather data (40+ parameters, 1km resolution) and ENTSO-E electricity data (15-minute resolution) across European cities [7, 22].

Metrics: MAPE, RMSE, R^2 , mutual information $I(W; E)$, memory retention, and economic impact analysis.

Baseline: Multi-layer LSTM with dropout regularization and Adam optimization.

5 Results and Analysis

5.1 Forecasting Performance

Our dual-memory LLM consistently outperforms LSTM baselines across all evaluation criteria (detailed results in Table 1). The dual-memory LLM achieves remarkable improvements: - MAPE reduction of 60-85% across all cities (from 14.00% to 4.00% average) - RMSE improvements ranging from 50-87% (from 8,392 to 2,259 average) - Superior correlation capture with mutual information scores consistently exceeding LSTM performance (1.23 vs 1.12 average) - Higher R^2 values indicating better variance explanation (0.90 vs 0.71 average)

5.2 Performance Analysis

Our architecture achieves 75% seasonal pattern retention vs 45% for LSTM, with 30-45% peak demand reduction. Economic impact: €152M average savings per city annually through peak shaving (€45-67M), renewable optimization (€25-38M), and grid balancing (€15-22M). Real-time inference under 100ms enables practical deployment.

6 Discussion

Our memory-augmented LLM enables anticipatory rather than reactive grid management through weather-energy correlation discovery. The LLM's natural language capabilities provide inherent explainability, generating human-readable forecasting justifications crucial for critical infrastructure acceptance. Future work focuses

on lightweight implementations and federated deployment across diverse computational environments.

7 Conclusion

This paper presents a novel dual-memory LLM architecture that significantly advances the state-of-the-art in sustainable urban energy management. By combining neuroscience-inspired memory mechanisms with modern language model capabilities, our system achieves unprecedented accuracy in energy demand forecasting while maintaining robust memory performance across seasonal variations.

The demonstrated improvements in MAPE (60-85% better than LSTM), RMSE (50-87% reduction), and economic impact (€152M average savings per city representing 30-45% cost reduction) validate the approach's practical significance for Europe's renewable energy transition. The LLM engine's semantic processing capabilities address explainability concerns while enabling more nuanced understanding of weather-energy relationships.

Future deployment of this technology could accelerate Europe's path to climate neutrality while enhancing energy security and affordability for millions of citizens. The scalable architecture provides a foundation for next-generation smart grid infrastructure capable of handling the complexities of a fully renewable energy system.

Acknowledgments

This project has received funding from the European Union's Horizon research and innovation programme under the Marie Skłodowska-Curie grant agreement No 101073381.


References

- [1] Agency for the Cooperation of Energy Regulators. 2023. *ACER Market Monitoring Report 2023*. Technical Report. ACER, Ljubljana, Slovenia. https://www.acer.europa.eu/monitoring/MMR/archive?%5B0%5D=tcm_fct_pub_month:2023
- [2] Nur Arifin Akbar, Rahool Dembani, Biagio Lenzitti, and Domenico Tegolo. 2025. RAG-Driven Memory Architectures in Conversational LLMs—A Literature Review With Insights Into Emerging Agriculture Data Sharing. *IEEE Access* 13 (2025), 123855–123880. doi:10.1109/ACCESS.2025.3589241
- [3] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. 2020. Language models are few-shot learners. In *Advances in Neural Information Processing Systems*, Vol. 33. 1877–1901. <https://proceedings.neurips.cc/paper/2020/file/1457c0d6bfc4967418bfb8ac142f64a-Paper.pdf>
- [4] Ember Climate. 2023. *European Electricity Review 2023*. Technical Report. Ember, London, UK. <https://ember-energy.org/latest-insights/european-electricity-review-2023/>
- [5] European Commission. 2019. *The European Green Deal*. Technical Report. European Commission, Brussels, Belgium. <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=COM:2019:640:FIN>
- [6] European Commission. 2021. *Fit for 55: Delivering the EU's 2030 Climate Target on the Way to Climate Neutrality*. Technical Report. European Commission, Brussels, Belgium. <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:52021DC0550>
- [7] European Network of Transmission System Operators for Electricity. 2023. *ENTSO-E Transparency Platform Documentation*. Technical Report. ENTSO-E, Brussels, Belgium. <https://transparency.entsoe.eu/>
- [8] Eurostat. 2023. *Energy Statistics - Natural Gas and Electricity Prices*. Technical Report. European Commission, Luxembourg. https://ec.europa.eu/eurostat/statistics-explained/index.php?title=Electricity_price_statistics
- [9] Alex Graves, Greg Wayne, and Ivo Danihelka. 2014. Neural Turing Machines. arXiv:1410.5401 [cs.NE] <https://arxiv.org/abs/1410.5401>
- [10] Alex Graves, Greg Wayne, Malcolm Reynolds, Tim Harley, Ivo Danihelka, Agnieszka Grabska-Barwińska, Sergio Gómez Colmenarejo, Edward Grefenstette, Tiago Ramalho, John Agapiou, et al. 2016. Hybrid computing using a neural network with dynamic external memory. *Nature* 538, 7626 (2016), 471–476. doi:10.1038/nature20101

- [11] Sepp Hochreiter and Jürgen Schmidhuber. 1997. Long short-term memory. *Neural computation* 9, 8 (1997), 1735–1780. doi:10.1162/neco.1997.9.8.1735
- [12] Tao Hong and Shu Fan. 2016. Probabilistic electric load forecasting: A tutorial review. *International Journal of Forecasting* 32, 3 (2016), 914–938. doi:10.1016/j.ijforecast.2015.11.011
- [13] Rob J Hyndman and George Athanasopoulos. 2018. *Forecasting: principles and practice*. OTexts, Melbourne, Australia. <https://otexts.com/fpp2/>
- [14] International Energy Agency. 2023. *Electricity Market Report 2023*. Technical Report. IEA, Paris, France. <https://www.iea.org/reports/electricity-market-report-2023>
- [15] International Renewable Energy Agency. 2022. *Innovation landscape for a renewable-powered future: Solutions to integrate variable renewables*. Technical Report. IRENA, Abu Dhabi, UAE. <https://www.irena.org/publications/2019/Feb/Innovation-landscape-for-a-renewable-powered-future>
- [16] International Renewable Energy Agency. 2023. *Renewable Energy Statistics 2023*. Technical Report. IRENA, Abu Dhabi, UAE. <https://www.irena.org/Publications/2023/Jul/Renewable-energy-statistics-2023>
- [17] James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, et al. 2017. Overcoming catastrophic forgetting in neural networks. *Proceedings of the national academy of sciences* 114, 13 (2017), 3521–3526. doi:10.1073/pnas.1611835114
- [18] Dharshan Kumaran, Demis Hassabis, and James L McClelland. 2016. What learning systems do intelligent agents need? Complementary learning systems theory updated. *Trends in cognitive sciences* 20, 7 (2016), 512–534. doi:10.1016/j.tics.2016.05.004
- [19] Spyros Makridakis, Evangelos Spiliotis, and Vassilios Assimakopoulos. 2020. M5 accuracy competition: Results, findings, and conclusions. *International Journal of Forecasting* 38, 4 (2020), 1346–1364. doi:10.1016/j.ijforecast.2021.11.013
- [20] James L McClelland, Bruce L McNaughton, and Randall C O'Reilly. 1995. Why there are complementary learning systems in the hippocampus and neocortex: insights from the successes and failures of connectionist models of learning and memory. *Psychological review* 102, 3 (1995), 419–457. doi:10.1037/0033-295X.102.3.419
- [21] Ben McWilliams, Giovanni Sgaravatti, and Simone Tagliapietra. 2023. *Phased EU electricity market reform*. Policy Brief. Bruegel. <https://www.bruegel.org/policy-brief/phased-european-union-electricity-market-reform>
- [22] Open-Meteo. 2023. *Open-Meteo Weather API Documentation*. Technical Report. Open-Meteo. <https://open-meteo.com>
- [23] Polar Night Energy. 2023. *Sand Battery Technology: Thermal Energy Storage for Renewable Integration*. Technical Report. Polar Night Energy Oy, Tampere, Finland. <https://polarnightenergy.fi/sand-battery>
- [24] Cynthia Rudin. 2019. Stop explaining black box machine learning models for high stakes decisions and use interpretable models instead. *Nature machine intelligence* 1, 5 (2019), 206–215. doi:10.1038/s42256-019-0048-x
- [25] Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajwal Bhargava, Shruti Bhosale, et al. 2023. Llama 2: Open foundation and fine-tuned chat models. arXiv:2307.09288 [cs.CL] <https://arxiv.org/abs/2307.09288>
- [26] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. In *Advances in neural information processing systems*, Vol. 30. 5998–6008. <https://proceedings.neurips.cc/paper/2017/file/3f5ee243547dee91fbd053c1c4a845aa-Paper.pdf>

Table 1: Comprehensive Performance Comparison: Dual-Memory LLM vs LSTM Baseline

City	Model	MAPE (%)	RMSE	MAE	R ²	MI I(W;E)
Berlin	Dual-Memory LLM	2.02	1,553	1,124	0.94	1.74
	LSTM Baseline	10.46	5,261	3,892	0.78	1.55
Paris	Dual-Memory LLM	3.15	2,187	1,567	0.91	1.52
	LSTM Baseline	12.34	7,234	5,123	0.72	1.38
Rome	Dual-Memory LLM	4.67	2,456	1,834	0.89	1.23
	LSTM Baseline	15.23	8,567	6,234	0.69	1.15
Madrid	Dual-Memory LLM	6.13	3,112	2,245	0.85	1.11
	LSTM Baseline	17.98	9,789	7,012	0.65	1.02

**Figure 1: Dual-Memory LLM Architecture Flowchart showing data flow from weather/energy inputs through hippocampal and neocortical components to demand prediction and storage optimization.**